

AI ADOPTION BEYOND THE SLOP

AI adoption for marketers.

How to recognize ideas whose time has come—and how to build markets in which use begins to reinforce itself.



READER TARGET

For entrepreneurs, investors, and marketing leaders who must bring AI solutions to market before the market is comfortable with them.

AI ADOPTION BEYOND THE SLOP

AI adoption for marketers.

How to recognize ideas whose time has come—and how to build markets in which use begins to reinforce itself.

For entrepreneurs, investors, and marketing leaders who must bring AI solutions to market before the market is comfortable with them.

Copyright © [2026] Axiologic Research

All rights reserved.

KDP publishing rights held by Outfinity SRL (Iași, Romania).

Available for free at www.axiologic.net.

Author's Note

This work was created under the umbrella of Axiologic Research as part of two interconnected research programs, one dedicated to Artificial Intelligence and the other to Outfinitism, both led by Sinică Alboaie. To explore the details of how these two programs intersect and build upon each other, we invite readers to visit the Outfinity Venture Studio page at www.outfinity.ch.

While Artificial Intelligence language models were used to assist with text generation, the foundational philosophy, thematic structure, and analytical approach derive exclusively from the author's decades of dedicated study of social phenomena, social technologies, and genuine hard technologies resulting from various European and self funded research projects. Recently, within the Achilles research project, we have been deeply studying AI generated literature and its automated verification using neuro symbolic approaches; all the free books made available to the public are part of the suite of works we use to test and popularise our latest ideas and technologies.

Map of the Book

The book follows the movement from the atypical user to a market that produces its own permission, infrastructure, and switching costs. The chapters are not a sequence of slogans about AI, but an analytical model for commercial decisions made before consensus.

Chapter	Central question
Preface	Why private utility appears before public legitimacy
1. Adoption Is Not an Opinion	How to separate attention, trial, use, and reorganization
2. The Abnormal User	Who crosses first through an incomplete product
3. The Slop Paradox	Why the frontier can rise while the visible median falls
4. The Visibility Inversion	Why failures are visible while assisted competence disappears
5. The Personality of Adoption	Which traits influence trial, use, and competent adoption
6. The Expertise Paradox	Who gains immediately and who extracts strategic value
7. The Permission Threshold	Why utility is not enough
8. Permission Cascades	How a practice becomes socially imitable
9. The Market That Teaches Itself	How users produce complementary assets
10. Apprenticeship Debt	How present productivity can mortgage future expertise

Chapter	Central question
11. Nested S-Curves	Why the ‘AI market’ is a collection of micro-markets
12. The Inevitability Radar	How to distinguish a cascade from a fashion
13. Permission Marketing	How to sell a behavior that is not yet legitimate
14. The Graveyard of Inevitability	How to avoid a self-confirming prophecy
15. What Comes Next, 2026–2031	Falsifiable predictions about adoption
Epilogue	When technology disappears into infrastructure

Methodological note. Research and statistics were verified through August 23, 2026. Original concepts are presented as working hypotheses and tested against the literature, not as laws already validated.

PREFACE

The Technology Used Before It Becomes Legitimate

An important technology rarely enters society in the form by which it will later be recognized. At first it appears incomplete, conspicuous, and poorly understood. It is used by people with unusually strong reasons to tolerate its defects, presented through demonstrations that are too perfect or through results so weak that they become caricatures, and debated by comparing the maximal promises of its advocates with the most visible failures of beginners. Between these extremes, however, the real market develops: a collection of private uses, improvised procedures, local advantages, and social justifications that can eventually turn an experiment into a norm.

Generative artificial intelligence entered precisely this ambiguous zone. Within three years, some aggregate estimates indicated a diffusion speed unprecedented for a consumer technology, while official business statistics still showed operational integration to be a minority practice. The Stanford AI Index 2026 estimates that generative tools reached approximately 53 percent population adoption in three years. The OECD reports that in 2025 more than one third of individuals in member countries had used such tools, while only 20.2 percent of firms reported using artificial intelligence; the gap between large and small firms was enormous, 52 percent versus 17.4 percent (Stanford HAI, 2026; OECD, 2026a). These figures are not contradictory. They measure different things and reveal the very subject of this book: trying a technology can become almost ordinary long before dependence on it, workflow integration, and institutional recognition become normal.

It is tempting to describe the present conflict through a simple moral opposition. On one side would stand courageous pioneers capable of seeing the future; on the other, conservatives who fear change and wait for group approval. This story is attractive marketing and weak explanation. Many AI critics have reasonable arguments about reliability, confidentiality, copyright, the externalization of verification, skill loss, and degradation of the information

environment. Many early users are not visionaries but opportunists, people overwhelmed by work, or consumers attracted by novelty. Some new products deserve rejection. Some social cascades turn an inferior technology into a standard. And many ideas whose time appears to have come disappear when cheap capital, novelty, and media attention recede.

The book nevertheless begins from a more precise intuition. Resistance to AI is not only an evaluation of technical performance. It can reflect a gap between private utility and social permission to use the technology. People try a tool when their pain is sufficiently intense, their autonomy sufficiently strong, or the potential advantage large enough to justify the risk. Then they observe what others do. They learn which results are acceptable, what must remain hidden, what can be acknowledged, and what becomes reputationally safe. When procedures, examples, professional roles, standards, complementary tools, and legitimizing managers emerge around use, each adoption reduces the cost of the next adoption. The market begins to teach itself.

This book proposes eight concepts that organize the phenomenon. The Abnormal User Principle explains why initial users are selected from unusual situations and do not represent the future market. The Slop Paradox explains why the frontier of quality can rise while the average quality of visible artifacts falls. The Visibility Inversion shows why AI failures are labelled and public while competent uses dissolve into the final result or remain hidden. The Permission Threshold separates technical utility from social legitimacy. The Permission Cascade explains the moment when each new user reduces perceived risk for the next. The Expertise Paradox distinguishes between those who receive the largest immediate gain and those able to extract the greatest strategic value. Apprenticeship Debt names the accumulated cost when organizations automate the work through which future experts used to become experts. The Inevitability Radar turns these observations into an analytical instrument for entrepreneurs, investors, and marketers.

Table 1 • Core Theoretical Contributions

Concept	Working hypothesis
Abnormal User Principle	Initial users are selected by atypical need, autonomy, or resources; they are not copies of the future majority.

Concept	Working hypothesis
Slop Paradox	The quality frontier can rise while the median quality of visible artifacts falls.
Visibility Inversion	Failures are public and labelled; competent uses are absorbed or concealed.
Permission Threshold	Deep adoption requires evidence, control, legitimacy, and identity safety.
Permission Cascade	Each relevant adoption lowers the social and operational cost of the next.
Expertise Paradox	Novices may gain most locally; experts may extract most strategically.
Apprenticeship Debt	Automating the present can borrow from the formation of future competence.
Inevitability Radar	Signals that use is beginning to produce its own conditions of continuation.

These concepts are theoretical contributions of the book, not findings already validated under these names by the scientific literature. They are built by synthesizing distinct fields: diffusion of innovations, technology acceptance, social influence, algorithm aversion, professional-identity psychology, and research on productivity, learning, and complementary assets. Each concept will be treated as a useful hypothesis, not as a label designed to end discussion. Where evidence supports the intuition, the contextual limits will be stated. Where it contradicts it, the contradiction will remain. A book about recognizing the moment of an idea cannot be credible if it turns every signal into confirmation.

The primary audience consists of people who must act before the evidence becomes comfortable: founders choosing a market, investors trying to distinguish

future infrastructure from the fashion of the quarter, marketers who must sell a new behavior, and leaders deciding whether introducing AI will produce capability or merely activity. For them, the relevant question is not whether “AI is good” or whether “AI will win.” The harder question is this: where, for whom, and under what conditions does a use begin to reinforce itself, allowing what looks eccentric today to become an implicit requirement tomorrow?

The answer is not found in the number of accounts created or the volume of content published. A crowd produces noise; a market produces complements. A technology begins to become inevitable when its repeated use creates processes, skills, data, standards, reputation, switching costs, and reasons for other actors to build around it. This is the book’s real object: not enthusiasm, but the mechanics through which a technical possibility acquires social infrastructure.

The methodological note matters. Research was verified through August 23, 2026. Preference was given to peer-reviewed articles, randomized experiments, field studies, meta-analyses, official statistics, and institutional reports with explicit methods. Recent AI studies are often limited to one profession, one model, one organization, or self-reported measures; generalization therefore requires caution. Adoption figures differ because some sources measure any use, others frequent use, organizational provision, integration into a function, or use of any form of AI. Rather than forcing these definitions into artificial uniformity, the book uses them to separate stages of adoption.

The book does not predict that every current form of AI will become indispensable. It predicts something narrower: legitimacy may reverse before the moral debate ends. Organizations will judge AI use by whether the process is verifiable, safe, and economically justified. AI then ceases to be an identity and becomes infrastructure.

CHAPTER 1

Adoption Is Not an Opinion

Debates about artificial intelligence frequently conflate four distinct phenomena: attention, trial, repeated use, and the reorganization of work. A person may discuss AI every day without using it. The same person may create an account, produce a few texts, and abandon the tool. Someone may use a general assistant every week without changing any important process. Or a person may reach the point where work, product, and competitive advantage are designed around AI capabilities. All these states appear in statistics under the label “adoption,” although their economic meanings are radically different.

This confusion fuels two symmetrical errors. Enthusiasts treat the large number of experiments as evidence that institutional transformation is already complete. Skeptics observe that aggregate productivity, profits, and work organization have not changed in proportion to the excitement and conclude that the technology is mainly spectacle. The evidence available in 2026 suggests a subtler interpretation. Individual use expanded rapidly, while deep organizational use remained far behind. In February 2026, Gallup found that half of US employees used AI in their role at least several times a year, 28 percent used it several times a week or more, and 13 percent used it daily. Yet only 41 percent said their organization had integrated AI tools to improve practices, and evidence of fundamental work redesign remained limited (Gallup, 2026a). In *State of the Global Workplace 2026*, 65 percent of employees in organizations that had implemented AI reported positive effects on personal productivity, but only 12 percent strongly agreed that AI had transformed how work was performed in the organization (Gallup, 2026b).

For an entrepreneur, the difference between use and transformation is essential. The first level can produce rapid user growth and temporary revenue. The second creates recurring budgets, integration, proprietary data, and switching costs. A product can be viral without becoming infrastructure; it can be infrastructure in a niche without being viral; or it can move from one to the other when individual behavior forces the organization to recognize a practice that is already widespread. Microsoft documented an early form of this transition in 2024: 78 percent of workplace AI users brought their own tools, and the behavior crossed generations. At the same time, 52 percent hesitated to disclose using AI for

important tasks, while 53 percent feared that doing so would make them appear replaceable (Microsoft, 2024). The market existed before the institution granted it permission.

Adoption should therefore be understood as a ladder of dependence and legitimacy, not as a binary choice. At the bottom lies exposure: the person knows the technology exists and can describe its promise. Trial follows, usually cheap and reversible. Higher up comes repeated use for peripheral tasks. The technology then enters important activities but may remain hidden, unofficial, or poorly governed. Genuine integration begins when processes, roles, and quality evaluation are redesigned. Only in the final stage does non-use become a deviation that requires justification. These rungs are not merely temporal stages; individuals, teams, and industries can remain trapped between them for years.

Table 2 · The Real Adoption Ladder

Stage	Economic meaning
Attention	The category enters the vocabulary without changing behavior.
Trial	The user tests cheaply and reversibly.
Repeated use	The tool recurrently solves peripheral tasks.
Important use	It enters consequential tasks, often still hidden or ungoverned.
Integration	Workflow, data, roles, and quality evaluation are redesigned.
Institutionalization	The practice receives policy, budget, ownership, and standards.
Normalization	Non-use begins to require justification.

In marketing, the first three rungs are relatively easy to purchase. An accessible interface, a free trial, and a spectacular demonstration can generate experiments.

Higher rungs require a different product and a different kind of evidence. The user must know when an output is good enough, who is accountable for an error, how data enter and where they go, how the system can be corrected, what happens when the model changes, and whether the advantage persists once competitors gain access to the same capability. Where verification costs exceed the time saved, “adoption” may merely move effort from the creator to the recipient. The phenomenon called workslop describes apparently finished content that transfers to colleagues the work of discovering its lack of substance. Existing commercial evidence remains too early for broad conclusions, but the strategic signal is clear: generated volume is not equivalent to created value.

A healthy AI market is measured not only by how quickly users generate something, but by how cheaply it can be established that the result deserves to be used. Production cost falls spectacularly; the costs of selection, verification, and accountability may remain constant or even rise. This asymmetry explains why some categories will be dominated by general models while others create value around provenance, control, audit, and local expertise. For founders, the opportunity is not always the model that produces more, but the system that turns abundance into decision.

General-purpose-technology research explains delayed organizational effects. Brynjolfsson, Rock, and Syverson call this the productivity J-curve: new technologies initially require hard-to-observe investments in process redesign, human capital, products, business models, and management. Measured productivity can lag until organizations learn to exploit these intangible assets. OECD evidence points similarly: AI-using firms tend to be more productive, but gains are associated with digital infrastructure, ICT skills, and other complementary assets rather than access to a model alone (Brynjolfsson, Rock, and Syverson, 2021; Calvino and Fontanelli, 2023; OECD, 2024).

Apparent AI-project failure must therefore be interpreted carefully. Sometimes the product solves no relevant problem; sometimes a model was installed without the social system that makes it useful. Organizations may buy licences while rewarding incompatible behavior, demand speed while measuring document volume, encourage experimentation while punishing the first error, or demand transparency while making users fear replacement. The result is not irrational resistance but incoherent institutional design.

Gallup offers a strong empirical indication of the role played by that design. In US organizations that had implemented AI, frequent use was much more likely

when tools were integrated into existing systems, 88 percent versus 55 percent, and when the manager actively supported use, 78 percent versus 44 percent. Yet fewer than one third of employees believed their manager provided such support (Gallup, 2026c). Access is a property of software; permission is a property of the organization. B2B marketing that confuses the two will inevitably overestimate the demonstration's capacity to produce change.

Among small and medium-sized enterprises, need appears to play a particular role. An OECD survey across seven countries found generative-AI use in 31 percent of SMEs. Among users, 65 percent reported improved employee performance, one third reported reduced workload, and 39 percent of firms facing skill shortages said AI had helped compensate. Yet 57 percent of non-users regarded the technology as unsuitable for their activity, about half cited legal or confidentiality risks or missing skills, and very few firms had prohibited use entirely (OECD, 2025a). Resistance is not primarily hatred of AI. It is a combination of insufficiently demonstrated relevance, unabsorbed risk, and inability to integrate.

This conclusion has a direct consequence for market strategy. Segmentation into “pro-AI” and “anti-AI” is almost useless. Someone who criticizes AI publicly may be an intensive private user. An enthusiast may produce many demonstrations and no economic dependence. An organization that buys thousands of licences may change nothing important. The market should be segmented by problem intensity, task frequency, error cost, outcome observability, verification competence, decision power, and the existence of complementary assets.

Adoption is reversible: users may return to earlier methods, restrict the tool to narrow tasks, or change suppliers without changing how they work. Mandated use can yield high penetration and fragile retention; a small voluntary community with verifiable outcomes may be a stronger base for a cascade. Measure whether use reorganizes work, creates durable assets, and survives failure. Relevant adoption is an investment the user no longer wants to lose.

CHAPTER 2

The Abnormal User Principle

The first user of a technology is not a miniature of its future market. Usually, the first user occupies a situation unusual enough to accept an incomplete product, an uncertain reputation, and a learning cost that most people do not consider justified. This selection creates a classic marketing error: founders look at their early users and believe they have discovered the mainstream customer profile, when they have actually discovered the extreme conditions under which the problem becomes stronger than the friction.

The Abnormal User Principle holds that early adoption appears when at least one of three variables is exceptional: pain, decision autonomy, or the potential value of advantage. A fourth variable, often ignored, is the capacity to absorb loss. An overwhelmed researcher may tolerate hallucinations because the alternative is missing the deadline. A programmer who can validate the code independently may experiment without colleagues' approval. A consultant may invest hours in the technology because a small productivity gain multiplies across dozens of projects. A corporation may test hundreds of solutions not because it sees the future more clearly, but because it can afford to be wrong ninety-nine times.

This idea is close to Eric von Hippel's theory of lead users. Lead users experience strong needs that will become widespread months or years later and have a powerful incentive to find solutions themselves. Precisely because they live ahead of the market, they can serve as laboratories for anticipating demand and as sources of product concepts (von Hippel, 1986). The difference introduced here is the emphasis on abnormality as a selection mechanism. Anticipatory need is only one form. Others include independence from professional norms, risk appetite, surplus resources, and extreme pressure.

In AI, the cornered user matters more than innovation rhetoric acknowledges. Microsoft reported that 68 percent of respondents struggled with the pace and volume of work, while 46 percent felt burned out. In the same study, employees were already bringing personal AI tools into the workplace before organizations provided adequate policy and infrastructure (Microsoft, 2024). The OECD found that for about one third of SMEs using generative AI, the technology reduced the

workload of employees or owners; firms also used it to compensate for skill or labor shortages (OECD, 2025a). Need does not automatically produce adoption, but it can lower the threshold at which risks become acceptable.

This category should not be romanticized. Desperation can trigger trial while obstructing learning. People under pressure have less time to verify, design a process, and distinguish a robust solution from a dangerous shortcut. For an AI product, the overwhelmed user is simultaneously an ideal customer and a failure environment. The value is perceived immediately, yet this user may produce exactly the results that sustain the reputation of slop. Responsible marketing should not exploit vulnerability; it should explicitly reduce verification cost and prevent assistance from becoming cognitive debt.

The second category is the sovereign user: the person who regards personal evaluation as sufficient reason to begin. This person is not necessarily more intelligent or more correct than the group. The user may be open to novelty, have a strong need for autonomy, or overestimate the quality of personal judgment. What distinguishes this user is a willingness to separate the instrumental truth of a technology from its current reputation. Such a person does not first ask whether the profession approves, but whether the tool helps in a concrete case. In environments with strong norms, the technology may be used clandestinely, creating an invisible market.

The third category is the arbitrage user. This person sees the interval between the general availability of a capability and the unequal distribution of the skill required to use it. Models may be accessible to everyone while the advantage remains undemocratized: some people know how to formulate the problem, construct a workflow, connect data, verify, and transform the result into a product. The arbitrage does not last forever. That temporality is precisely what attracts freelancers, agencies, consultants, and startups. They are not necessarily betting on the permanent superiority of a technology, but on the market learning more slowly than they can execute.

The fourth category is the experimenter endowed with optionality. Wealthy organizations, laboratories, funded founders, and professionals with discretionary time can test technologies without urgent need. Paradoxically, they may become pioneers precisely because the problem does not constrain them. They have permission to fail. This category corrects the hypothesis that early adopters are mainly desperate or rebellious. Sometimes privilege precedes courage: what looks like vision is the ability to absorb the cost of trying.

Table 3 · Archetypes of the Early Adopter

Archetype	Adoption mechanism
Cornered adopter	Adopts because volume, complexity, or lack of resources has made the old process untenable.
Sovereign adopter	Experiments without group permission and accepts the risk of personal judgment.
Arbitrage adopter	Seeks advantage in the interval between a technology becoming available and the market learning to use it competently.
Resourced experimenter	Can afford time, money, and failure before the product matures.
Legitimitor	Is not necessarily first; is the credible actor who makes the practice socially imitable.

These types are not stable personalities. The same person may be cornered in one project, sovereign in a domain they know, conformist where reputational risk is high, and a well-funded experimenter inside an organization. In market segmentation, context beats demographics. A founder who identifies a community of professionals overwhelmed by documentation may find faster adoption than one who searches generically for “people open to AI.” A marketer who notices a high-frequency activity with rapid feedback and a verifiable outcome may find the nucleus of a cascade even in an apparently conservative industry.

The principle also generates an important prediction: what attracts initial users may repel the mass market. A configurable, powerful, and unpredictable product gives the sovereign user control; for the majority, the same flexibility means responsibility and risk. A message about transcending professional limits excites the arbitrageur; to the ordinary employee, it may sound like a threat to the job. The informal community in which pioneers exchange prompts and improvisations creates energy, but the enterprise buyer demands governance,

traceability, and repeatable results. Crossing the chasm is not merely moving from a small segment to a large one; it is changing the reason the product is purchased.

Rogers distinguishes innovators, early adopters, the early majority, the late majority, and laggards. A nuance often lost in startup discourse is that innovators and early adopters do not play the same role. Innovators tolerate risk and complexity. Respected early adopters act as interpreters and legitimizers within a community: they reduce uncertainty by showing that the technology can be integrated into recognized norms and objectives (Rogers, 2003). In this book's terms, abnormal users discover possibility; legitimizers make it socially imitable.

For AI, this separation explains why technical communities can appear years ahead of a profession without transforming it. An engineer can demonstrate a capability but cannot decide what constitutes an acceptable error in a hospital, law firm, or public institution. The transition requires people with local authority who understand both the tool and the domain's norms. The most important early customers are therefore not always those who use the product most, but those whose results can be imitated without loss of status.

This is where the concept of permission capital appears. A respected physician, senior partner, admired professor, or recognized institution can turn use from deviation into prudent practice. Their reputational capital absorbs part of the risk perceived by others. Conversely, a highly visible influencer without authority in the target community can increase awareness while reinforcing the impression that the product belongs to a superficial subculture. Marketing a controversial technology requires choosing not only an audience but a legitimate carrier of the message.

The Abnormal User Principle can be used as a method of market research. Instead of asking the average market what it wants from a category it does not yet understand, we seek people for whom the future is already a present problem. We study not only which tool they use, but which costs they accept, what they improvise, conceal, verify, and where they give up. We then separate the components likely to become general from adaptations produced by their extreme situation. Their need may anticipate the market; their tolerance for friction almost certainly does not.

This method also protects against a trap. Initial users may demand sophisticated features that make the product harder for others to adopt. They may value granular control, technical language, and freedom of configuration because mastery of complexity is part of their identity. If the company mistakes pioneer

competence for the future standard, it builds a product that is admired and rarely used. The goal is not to generalize the abnormal user, but to generalize the advantage that user discovered while progressively removing the abnormality required to obtain it.

The first signal that an idea's time is arriving is therefore not that ordinary people approve it. It is that different groups of abnormal users independently reach the same need and begin constructing similar solutions. When the overwhelmed, the autonomous, the arbitrageurs, and well-funded experimenters converge on the same capability, we have more than enthusiasm. We have pressure from several directions. The market does not yet exist in full, but the problem has begun searching for its own infrastructure.

CHAPTER 3

The Slop Paradox

The most effective popular criticism of generative artificial intelligence is not that models can produce nothing of value, but that they make the production of things unworthy of attention too cheap. Slop has become the informal name for texts, images, videos, and documents that look finished but are poor in intention, verification, or substance. The term captures a real experience: users do not encounter the frontier of capability; they encounter the distribution of published results. When the cost of publication approaches zero, that distribution is flooded by attempts that previously would not have crossed the effort threshold required to exist.

The Slop Paradox begins here: a technology can simultaneously raise the highest accessible quality and lower the average quality of the visible artifacts associated with it. This is not a contradiction. The frontier is set by people who combine the model with knowledge, selection, iteration, and editing. The visible average is shaped by the total number of outputs and the threshold for publication. If a technology enables an expert to produce a result that is 20 percent better but enables one million people to publish ten unverified outputs instead of one, public perception will be dominated by the second change.

An experiment published in *Science Advances* offers a controlled version of this tension. Doshi and Hauser found that access to AI-generated ideas increased ratings of creativity, writing quality, and enjoyment for short stories, with larger benefits for less creative participants. At the same time, AI-assisted stories became more similar to one another, reducing collective diversity (Doshi and Hauser, 2024). At the individual level, use was rational; at the ecosystem level, convergence could reduce the value of variety. This does not prove that all AI content becomes uniform, but it shows that improvement in the individual average and degradation of collective diversity can coexist.

Slop is not defined only by low quality. It is also a problem in the distribution of effort. A document generated quickly may look good enough to send, yet force the recipient to identify its assumptions, repeat verification, and discover whether a real idea exists behind the language. The author saves time by transferring cost to the reader. In organizations, this externalization can turn local productivity into

systemic unproductivity. Commercial research by BetterUp Labs and the Stanford Social Media Lab popularized the term workslop for such outputs; because it was a survey study rather than an independent measurement of productivity, its figures should be treated cautiously. The concept remains useful: technology can move work rather than eliminate it.

For marketing, this transfer changes the product promise. If the message is “create ten times more,” the company may stimulate exactly the behavior that degrades trust in the category. A user produces more presentations, emails, specifications, and alternatives, but colleagues possess finite evaluation capacity. When volume grows faster than the quality of selection, every new artifact competes not only with rival products but with the recipient’s patience. Abundance becomes a tax on attention.

That tax creates a distinct opportunity. As synthetic production becomes cheaper, value migrates toward mechanisms that demonstrate provenance, intention, responsibility, and quality. Editing, evaluation, reputation, audit, selection, and controlled-access communities are not merely accessories. They become the infrastructure that makes abundance tolerable. An anti-slop market does not oppose generation; it imposes admission costs. Valuable products will reduce not only creation time but also the time others need to trust the result.

The preprint “Measuring AI Slop in Text” attempts to turn the term into an evaluable taxonomy and begins from the absence of an agreed definition. The authors identify interpretable dimensions of slop and the problem of recognizing structural or lexical clichés associated with generated text (Shaib et al., 2025). The fact that the field is only beginning to build measurement instruments matters. Public criticism preceded scientific operationalization. For entrepreneurs, “slop” is therefore both a market signal and an unstable category: customers feel the cost before they can specify it in procurement requirements.

The paradox has a second component. As tools improve, slop may become harder to detect rather than less common. A fluent text without substance consumes more trust than one that is obviously poor. An error in an awkward sentence activates vigilance; the same error in an elegant report can enter a decision. The stylistic frontier advances faster than the accountability frontier. Surface quality can therefore no longer serve as a sufficient proxy for effort, truth, or competence.

Aversion to AI can consequently be a rational response to the degradation of signals. Professionals learned to evaluate work through indirect cues: coherence,

structure, vocabulary, detail, and the apparent time invested. Artificial generation weakens the connection between these cues and the substance behind them. When anyone can produce a convincing form, the recipient must demand different evidence. The critic is not necessarily defending the romance of manual labor; the critic may be defending a trust system whose signals have been compromised.

This observation limits the thesis that accusations of slop are merely a pretext used by people waiting for consensus. Sometimes they are a defensive mechanism against real costs. Slop can also delay high-quality adoption by contaminating the category's reputation. A manager who receives three superficial AI-generated reports will not automatically distinguish incompetent use from a technological limit. The manager will raise the permission threshold for the whole team. Every user who externalizes verification consumes the reputational capital of users who employ the tool carefully.

Criticism also has a limit. A decline in the average quality of public outputs does not mean the frontier lacks economic importance. Markets do not always pay for the average. A studio can use AI to explore thousands of alternatives and publish one. A researcher can generate ten hypotheses and retain only what survives verification. A law firm can use a model for search and structure without delivering the raw answer to a client. In such cases, abundance remains internal and selection protects the exterior. Slop becomes dangerous when the filter disappears, not when alternatives appear.

This leads to a product distinction between generators and editorial systems. A generator optimizes the ability to produce. An editorial system optimizes the probability that what leaves the process deserves to be seen. The first category dominated the initial phase because the demonstration is immediate. The second will become more important as organizations discover that output is not outcome and local speed is not end-to-end performance. In markets with a high cost of error, a system that refuses, requests evidence, and preserves traceability may be more valuable than the model that answers most fluently.

Table 4 · Where Value Migrates in an Abundant Market

From → to	Product consequence
Generation → selection	Advantage moves from producing alternatives to choosing the result worth seeing.

From → to	Product consequence
Volume → attention	The scarce cost becomes the recipient's time, not the generator's.
Fluency → evidence	Persuasive form must be replaced by sources, tests, and traceability.
Answer → audit trail	Enterprise relevance depends on who approved the result and how it can be verified.
Claimed originality → provenance	The market will pay for origin and accountability.
Local speed → end-to-end outcome	Savings matter only when they do not transfer work downstream.

The Slop Paradox also provides an instrument for competitive analysis. When a category is flooded with similar outputs, differentiation through production alone becomes unsustainable. Value may move toward proprietary data, distribution access, editorial authority, integration with real processes, or the ability to demonstrate effect. An agency promising “AI content” enters a race toward zero marginal cost. A platform able to show why a claim is supported, who approved it, and how it performed in the market builds an asset that is harder to copy.

In this sense, slop can accelerate market maturation. It provokes the institutional response that was missing: quality standards, attribution rules, filters, certifications, and new reputational conventions. Like spam, it will not disappear through indignation but through infrastructure. Infrastructure appears when the cost of pollution becomes high enough for selection to become a product.

An idea's time is signaled not by the disappearance of slop, but by an economy that learns to filter it.

CHAPTER 4

The Visibility Inversion

The public does not observe AI use uniformly. Poor results are often visible, labelled, and easily attributed to the technology. Good results are edited, integrated into a process, or presented as ordinary work. This asymmetry produces a visibility inversion: the more mature the use, the less obvious its traces may become, while incompetent use becomes the public representation of the category.

A generic text published without revision reveals its origin through repeated formulas and predictable structure. An analysis in which AI assisted search, proposed objections, and accelerated rewriting does not necessarily carry a distinct signature. Code accepted after testing, a document reviewed by a lawyer, and a hypothesis abandoned before an experiment disappear into the final product. The technology's contribution was real but is not counted by the public. By contrast, the absurd image, invented citation, and meaningless email circulate precisely because they are memorable failures.

Microsoft documented an additional reason for concealment. In 2024, more than half of workplace AI users hesitated to acknowledge using the tool for important tasks and feared that disclosure could make them appear replaceable (Microsoft, 2024). When acknowledgement carries a reputational cost, the real market moves underground. Managers see less competent use than actually exists and more problems than good practices because people report errors that affect them, not private conversations that helped them.

The inversion complicates market research. The question “do you use AI?” may receive answers distorted by shame, ambiguity, and differences in definition. An employee may regard autocorrection, semantic search, or summarization embedded in software as normal functions rather than AI. Another may call every chatbot interaction AI. A professional may deny use because the final product was verified and rewritten personally. Someone else may exaggerate use to signal modernity. Self-reported data measure both behavior and the identity respondents wish to project.

For a startup, the hidden market is both opportunity and risk. Opportunity appears when people improvise with general tools, use personal accounts, or construct clandestine procedures for a need the organization has not recognized.

This behavior indicates pain and willingness to adopt. Risk appears when the founder confuses clandestine use with enterprise demand. What is individually useful may be legally unacceptable, impossible to audit, or incompatible with institutional responsibility. A commercial product must turn hidden practice into legitimate practice, not merely make concealment more convenient.

This is an important source of category creation. The company sells not only capability but a translation between two regimes: from private use to governed use; from answer to evidence; from a personal account to an audited process; from “it helped me” to “we can demonstrate why it is acceptable.” In many domains, the base model is already available. Value moves to the product that removes the reasons for concealment.

The Visibility Inversion also explains why public perception can change abruptly. For a long time, competent use accumulates without an explicit signal. People develop practices, firms test internally, managers observe advantages, and suppliers build integrations. When a respected institution formally recognizes the behavior, what appeared marginal suddenly becomes ubiquitous. The change did not occur on the day of the announcement; the announcement made a prior change visible.

This mechanism resembles a social phase transition. Below the threshold, individual users believe they are exceptions and hide their behavior. Above it, each discovers that many others were doing the same, and the cost of acknowledgement falls. For marketing, the moment is critical. Before it, the message should provide safety, control, and discreet communities. After it, the message shifts toward standardization, interoperability, and the advantage of not being left behind.

There is also a moral inversion. If valuable use is invisible, advocates can always invoke unobservable benefits to dismiss visible criticism. “The best cases are secret” is not evidence. This book treats the inversion as a hypothesis about selection among observations, not as a universal excuse. It must be tested through usage data, field studies, and measurable outcomes. Without them, invisibility may conceal only the absence of value.

A stronger signal is the distance between personal use and official infrastructure. When employees bring their own tools at scale, informal procedures emerge, and managers cannot explain organizational policy, a latent market exists for integration and governance. When use remains occasional, produces no retention, and creates no improvisation, the market may be only

curiosity. The distinction appears in the sacrifices users make: time invested in learning, money paid personally, risks accepted, and processes rebuilt.

The inversion has consequences for positioning. An AI product that boasts excessively about its technological identity may increase the customer's reputational cost. For a mature segment, AI should be visible as a control function and invisible as spectacle. The customer is buying a professional outcome, not participation in a cultural movement. In other segments, pioneer identity remains part of the value. The choice depends on the technology's social status within that community.

A marketing rule follows: observe which part of the value users want to display and which part they want to hide. A designer may display the speed of exploration but does not want the final work perceived as automatic. A lawyer may value accelerated research but must demonstrate human accountability. An entrepreneur may promote team efficiency but does not want to suggest that the product is a superficial wrapper around a common model. Messaging that ignores this economy of attribution creates resistance even when utility is recognized.

AI will gradually become inseparable from software as generative functions enter editors, search, analytics, collaboration tools, and automated processes. The label will lose explanatory power; markets will divide by outcome, risk, and integration. A technology wins completely when it no longer needs to be announced and the category disappears into the infrastructure it changed.

CHAPTER 5

The Personality of Adoption

When a new product encounters resistance, marketers often look for a simple psychological explanation. Some people appear unusually willing to explore uncertain tools, while others demand concrete evidence, stable routines, and clear guarantees before changing how they work. The distinction is commercially relevant, but it should not be turned into a morality play between imaginative pioneers and timid conservatives. Personality can influence the first approach to AI; it does not, by itself, determine who will become a capable, loyal, or profitable user.

The most useful studies for this book are those that connect measurable traits directly to AI acceptance, use, engagement, or performance. Even within that narrower literature, four outcomes must be kept separate: willingness to try, declared intention to use, repeated behavior, and competent use. A person may enthusiastically test every new model and never integrate one into serious work. Another may resist for months, then adopt a single verified system and reorganize an entire workflow around it. For market analysis, personality is therefore a modifier of the adoption path, not a substitute for observing behavior.

A 2026 cross-cultural meta-analysis synthesized 44 studies, 187 effect sizes, and 14,594 students using AI-assisted writing tools. Openness showed the strongest positive association with acceptance, followed by conscientiousness and extraversion; agreeableness had a smaller positive association, while neuroticism was negatively associated with acceptance. The effects also varied with cultural context (Li, 2026). The result is directionally useful, but its boundaries matter. Most samples came from higher education, many outcomes concerned acceptance constructs rather than durable behavior, and only a minority of the included studies used behavioral indicators. A founder should read the evidence as a map of tendencies, not as a licence to profile customers as fixed psychological types.

A study of 1,016 university students produced a similar but not identical pattern. Openness, conscientiousness, and extraversion were positively related to educational use of generative AI; neuroticism was negatively related, agreeableness was not significant, and openness emerged as the strongest predictor in the authors' neural-network analysis (Arpaci, Kuşci, and Gibreel, 2025). This

strengthens the case that curiosity toward novelty matters at the trial stage. It also suggests that disciplined, goal-oriented users may adopt once AI is presented not as entertainment but as a structured instrument for completing work.

Workplace-oriented evidence is thinner and less orderly. In a pilot study of 62 participants performing marketing-like tasks, Kovbasiuk and colleagues found interactions involving conscientiousness and agreeableness, task quality, and future-use intentions, but not the simple pattern in which openness predicts every form of early adoption (Kovbasiuk et al., 2025). The sample was too small for broad segmentation claims, yet that limitation is itself instructive. Traits may affect different moments of the journey differently: exploration, persistence, collaboration with the system, evaluation of output, and willingness to return are not the same behavior.

A cross-cultural experiment with 639 participants in the United Kingdom and Arab contexts tested the Big Five together with need for cognition and locus of control across scenarios that varied AI autonomy and task criticality. Personality mattered, but system design, task context, and culture often mattered more; the effects of need for cognition, intuition, and locus of control were narrower and scenario-dependent (Rahman et al., 2026). For marketers, this is a strategically important finding. Acceptance does not depend only on who the person is. It depends on the relationship the product constructs between that person, the task, and the machine.

The practical interpretations should remain probabilistic. Openness may increase experimentation, but experimentation is not retention. Conscientiousness may favor systems that provide structure, reliability, and visible progress, yet the same trait may produce rejection when outputs are difficult to verify. Extraversion may support social learning and peer-led diffusion, although it is not required for sustained individual use. Higher neuroticism or anxiety may increase avoidance of uncertain systems, but anxiety can also produce reliance when AI is marketed as reassurance. Agreeableness may facilitate cooperative interaction with an assistant in one context and preference for human help in another. The same trait can support opposite behaviors when product design changes the perceived risk.

Adjacent constructs may be closer to the behavior that a market actually needs to predict. Personal innovativeness in information technology describes an individual's willingness to try new technologies within a particular domain (Agarwal and Prasad, 1998). AI self-efficacy concerns whether users believe they

can operate, judge, and recover from errors in the system. Need for cognition affects whether people want to inspect reasoning, assumptions, and evidence rather than merely receive an answer. Locus of control influences whether automation feels like an extension of agency or a transfer of authority. These constructs do not divide the market into permanent tribes; they change the probability of trial, persistence, and delegation under specific conditions.

Table 5 · Personality Traits and Plausible Market Responses

Variable	Marketing and design implication
Openness and curiosity	Exploration, novelty, and possibility resonate; the risk is trial without retention.
Need for cognition	Calls for sources, comparisons, hypotheses, and useful transparency—not merely fluent answers.
Locus of control	Acceptance rises when AI feels like an extension of agency rather than a takeover of decision-making.
Self-efficacy	Onboarding must create a legitimate sense of competence and control.
Tolerance for ambiguity	Shapes patience with the jagged frontier and uneven results.
Conscientiousness and risk sensitivity	Favor predictability, verification, and explicit limits.
Sensitivity to social proof	Calls for confirmation from relevant peers and clear local norms.

The unit of adoption is often the person–task pair, not the person in the abstract. A physician may experiment aggressively with diagnostic literature searches and remain conservative about clinical documentation. A lawyer may adopt AI for

discovery while refusing generated legal argument. An engineer may use it daily for code and avoid it in people management. Even personal innovativeness is domain-specific: the same customer can be an early adopter in one workflow and a late majority user in another. Market research that asks only whether someone is “pro-AI” or “anti-AI” discards the very variation that product strategy needs.

This book therefore proposes the concept of a permission profile. The profile does not attempt to diagnose personality. It describes the conditions under which a segment grants itself permission to use the product. For one group, permission appears when there is freedom to explore and obtain surprising results. For another, it appears when a respected expert validates the process. For another, when the system preserves human approval, inspectability, and a decision trail. For another, when non-use begins to threaten professional status. A permission profile is dynamic, contextual, and directly actionable through design, demonstrations, distribution, training, and positioning.

This corrects a common error in technology marketing: trying to persuade the wrong person with more information. Resistance may not come from ignorance of the benefit. It may come from a threat to autonomy, identity, competence, or professional duty. A more spectacular demonstration can intensify the threat. For a control-oriented segment, the right message is not “watch what it can do alone,” but “see how precisely you can define what it may do.” For a status-sensitive segment, the message is not “you no longer need to know,” but “your expertise becomes the mechanism that makes the system reliable.” For an exploration-oriented segment, excessive governance can destroy the initial attraction.

The commercial implication is encouraging: personality changes slowly, whereas control, competence, safety, and experience can be designed. Products can offer different modes of agency, from suggestion to execution; provide visible sources and reversible actions; let users begin with bounded tasks; and use credible peers to supply social proof. The goal is not to manipulate a trait. It is to remove a barrier that matters disproportionately to a given segment. A product that makes uncertain users feel capable expands its addressable market. One that repeatedly exposes users to their own lack of skill merely confirms its natural enthusiasts.

Personality matters, but it does not close the sale. Traits shape the first move; experience, task fit, organizational permission, and product architecture shape the trajectory. Mature marketing does not classify customers as brave or fearful. It

discovers what evidence, control, and context make adoption rational for each market.

CHAPTER 6

The Expertise and Seniority Paradox

One of the most important intuitions about the AI market appears contradictory at first. Experienced people may extract more value from a model because they know what to ask, recognize errors, and can integrate the output into a real decision. The same people may be more reluctant because they have established procedures, reputations to protect, and a higher cost of change. Beginners sometimes adopt more readily, yet lack the criteria needed to distinguish a plausible answer from a correct one. Who, then, benefits most?

The evidence does not support a single verdict because “benefit” can mean different things. Experiments on well-bounded tasks often show a compression of the performance distribution: weaker or less experienced workers gain the most in percentage terms. In a field study of 5,172 customer-support agents, access to generative assistance increased productivity by roughly 15 percent, with the largest effects among less experienced and lower-performing employees. Their language became more similar to that of the best agents, and retention among new workers improved. For highly experienced agents, speed gains were small, and following the system’s suggestions could marginally reduce conversational quality (Brynjolfsson, Li, and Raymond, 2025).

An experiment involving 453 professionals performing writing tasks found an average 40 percent reduction in time and an 18 percent increase in quality; participants with weaker baseline performance improved more (Noy and Zhang, 2023). These findings contradict the simple hypothesis that AI primarily helps senior people. For repeatable tasks with relatively clear feedback and strong examples in the model’s training data, AI can act as a compressor of experience: it gives the novice an approximation of practices that would otherwise take months or years to acquire.

Compression, however, is not the same as expertise. A study of 758 BCG consultants shows why. Across eighteen tasks inside the model’s capabilities, AI users completed more tasks, faster and at higher quality. On a complex managerial task outside the frontier, AI users were 19 percent less likely to reach the correct answer. Their prose could be more persuasive even when the recommendation was

wrong (Dell’Acqua et al., 2026). Along a jagged technological frontier, value comes not only from using the tool but from recognizing where it ceases to be trustworthy.

A preregistered meta-analysis of 106 experiments and 370 effect sizes also shows why “human plus AI” should not be treated as automatic superiority. On average, combinations performed below the better component working alone; losses were more common in decision tasks, while gains were more common in creative work. When the human was better than the AI, the combination was more likely to help; when the AI was already better, human intervention could degrade the result (Vaccaro, Almaatouq, and Malone, 2024). The market implication is not “human in the loop” as a slogan, but explicit loop design: who generates, who evaluates, and where each party has a genuine informational advantage.

This is the precise form of the expertise paradox proposed by this book. The novice may receive the largest immediate performance gain because the starting point is lower and codified routines can be borrowed. The expert may receive the greatest strategic leverage because the expert can formulate better problems, connect domains, identify exceptions, and accept or reject the output. Between them lies a dangerous zone: the user competent enough to produce elaborate work but not competent enough to notice when the system has crossed its frontier.

Seniority adds a different dimension. It is not synonymous with age, expertise, or performance. A senior manager may have decision power, contextual experience, and information access but little competence in a technical task. An individual contributor may possess deep expertise without organizational authority. An older employee may be a technological experimenter, while a young graduate may avoid AI for ethical reasons or career anxiety. Market research that treats “senior people” as a homogeneous category mixes four variables: experience, age, status, and power to redesign the process.

Usage data nevertheless suggest an age effect. A representative US study of workplace generative-AI adoption found in 2024 that roughly 34 percent of workers under forty used it, compared with 17 percent of workers aged fifty and above (Bick, Blandin, and Deming, 2024). OECD data on individual use in 2025 show an even wider difference across age groups, although those statistics also include people outside the labor market (OECD, 2026a). Such figures confirm an adoption gap, not its cause. Access to digital roles, task composition,

organizational culture, time for learning, and prior exposure may explain a substantial part of it.

Experienced professionals' resistance should not be dismissed as rigidity. An established way of working is capital. It contains routines, relationships, criteria, and responsibilities that do not appear in a tutorial. Change temporarily devalues that capital. The cost is not merely the time required to learn a tool, but the return from unconscious competence to visible awkwardness. For someone whose status rests on fluency and control, this local regression is reputationally expensive. The new product asks that person to become a beginner again, often in front of the people whose work they lead.

Research on professional-identity threat supports this explanation but also corrects it. In a study of medical students and professionals, threats to professional capabilities and recognition predicted resistance to AI. Yet students reported stronger threats and resistance than professionals, not the reverse (Jussupow, Spohrer, and Heinzl, 2022). People still building their professional identity may experience AI as cancelling an investment before it has consolidated. Senior people have more capital to protect, but may also have greater security, autonomy, and capacity to redefine their role.

A 2025 study of 635 Chinese employees over the age of forty-five adds a relevant nuance. AI adoption was associated both with identity threat and with greater intergenerational knowledge transfer, while digital self-efficacy reduced the threat. The evidence was cross-sectional and self-reported within a specific cultural setting, so it does not establish causality. Even so, it points to an important commercial possibility: the experienced user need not be treated only as a beneficiary or an obstacle, but as a carrier of knowledge that the product can help circulate (Guo and Wei, 2025).

This is a crucial commercial result. The message "AI frees experts from routine work" may appeal to the senior employee while threatening the junior if routine work is precisely the route into the profession. The message "anyone can perform like an expert" attracts users who lack access to expertise, but directly attacks expert status and promises more than the system can guarantee. The same slogan can accelerate one category and block another.

Table 6 · The Expertise Paradox by Segment

Segment	Likely gain and dominant risk
Junior or novice	Large gain on standardized tasks; high risk of overconfidence and shallow learning.
Individual expert	Strategic gain through framing, verification, and exception handling; requires granular control.
Senior manager	Can redesign the process and distribute permission even without deep technical expertise.
Older employee	Not synonymous with seniority; adoption depends on self-efficacy, role, task, and identity.
Organization	Can distribute expertise while simultaneously eroding the pipeline that produces new expertise.

For market analysis, separating productivity gains from authority gains is more useful than dividing customers by age. Juniors often buy access to a standard, rapid feedback, and relief from blank-page anxiety. Experts buy coverage, exploration, information compression, and the ability to test more alternatives. Managers buy consistency, visibility, and process control. The organization buys practice transfer and reduced dependence on a few people, but risks eroding the very source of future expertise.

Product design should reflect these differences. A beginner mode needs guidance, explanations, exercises, and mechanisms that discourage automatic acceptance. An expert mode needs granular control, citations, comparison among alternatives, integration with proprietary data, and the ability to impose criteria. A managerial mode needs policies, metrics, roles, and approval trails. Putting the same chat box in front of everyone and calling it transformation confuses access with adoption.

The expertise paradox also creates a measurement problem. If evaluation considers only time saved, beginners may appear to be the absolute winners. If it considers decision value, the avoidance of rare errors, novelty, and effects on other people, experts may have far greater leverage. A CEO who uses AI to formulate a better question may save only thirty minutes and still create millions in value. A junior employee who generates ten times as many pages may create volume rather than outcomes. Adoption metrics must be connected to the economics of the task, not the frequency of interaction.

There is also a systemic risk. If the best professionals increasingly follow system suggestions, practices may become more homogeneous and the supply of genuinely new examples may shrink. The support-agent study explicitly raises the possibility that expert adherence to AI recommendations could reduce the original contributions from which future models might learn (Brynjolfsson, Li, and Raymond, 2025). The market can distribute existing expertise while slowing the production of new expertise. This tension returns in the chapter on apprenticeship debt.

The initial intuition can therefore be retained in a more defensible form. Senior people and experts are not, in general, either the principal beneficiaries or the principal opponents. They face a different structure of benefits and costs. When AI offers a shortcut through a standardized task, the novice gains more. When value depends on problem definition, contextual judgment, and verification, the expert may turn the same technology into a superior multiplier. And when a product demands the surrender of identity, autonomy, or routine without clear compensation, resistance is rational regardless of age.

CHAPTER 7

The Permission Threshold

A product can be useful and still remain unused. The statement sounds obvious, yet much AI marketing behaves as though demonstrating capability should end the discussion. The model writes, classifies, summarizes, or proposes faster than a person; therefore the market should adopt. When adoption is slow, the explanation shifts to customer ignorance, fear, or inertia. In reality, technical utility is only one condition of permission.

The permission threshold is the cumulative level of evidence, control, legitimacy, and safety that an individual or institution requires before turning a possibility into normal behavior. It is not identical to intention to use in classical technology-acceptance models, although it builds on them. The Technology Acceptance Model established the importance of perceived usefulness and ease of use, while UTAUT integrated performance expectancy, effort expectancy, social influence, and facilitating conditions (Davis, 1989; Venkatesh et al., 2003). AI adds a sharper problem: the product can participate in producing judgment, not merely in executing a command. The user must therefore grant permission to delegate part of the competence by which the user is evaluated.

Permission has at least four sources. The first is epistemic: can I trust the output to be good enough for this task? The second is operational: does the system integrate without merely transferring cost to verification, security, or repair? The third is institutional: am I allowed to use it, and can compliance be demonstrated? The fourth is identity-based: is its use compatible with my understanding of what it means to be a professional? A product may pass the first three tests and fail the fourth. It may pass the identity test in a startup community and fail the legal test in a hospital.

Research on algorithm aversion shows that people may penalize system errors more severely than human errors. Dietvorst, Simmons, and Massey found that participants avoided algorithmic forecasts after observing the algorithm make mistakes, even when they had also seen it outperform a person (2015). The literature does not, however, support universal aversion. Across several experiments, Logg, Minson, and Moore found that people sometimes gave greater weight to advice believed to come from an algorithm, especially when they did not

regard themselves as experts (2019). Markets move between appreciation and aversion according to the task, experience, stakes, control, and visibility of error.

This oscillation explains why transparency alone does not automatically solve the problem. Users do not always demand an understanding of the internal mechanism; they demand an ability to manage the consequence. A classic result shows that allowing people to modify even a small part of an algorithmic forecast increases willingness to use it, satisfaction, and performance (Dietvorst, Simmons, and Massey, 2018). A preregistered replication in 2025 confirmed the effect of adjustability, while additional transparency did not produce an equally clear effect in that sample (Bohlen et al., 2025). Symbolic control can matter more than an extensive technical explanation.

A product rule follows: permission is often constructed through agency. The approval button, the ability to require sources, the right of veto, comparison with a baseline, and a clear allocation of responsibility are not accessories. They are components of value. A system promising maximal autonomy may be technically impressive and commercially premature. For many markets, the optimal product is not the one that removes the person but the one that shows precisely where that person remains indispensable.

Professional identity raises the threshold in a less visible way. AI can threaten not only a job but also the story by which a person justifies status. If a lawyer is valuable because the lawyer writes, a designer because the designer imagines, a programmer because the programmer produces code, and a teacher because the teacher explains, a technology that imitates these outputs raises an existential question before it produces an ROI calculation. Studies of professional-identity threat indicate that perceived losses of capability and recognition can fuel resistance (Jussupow, Spohrer, and Heinzl, 2022; Mirbabaie et al., 2022).

Marketing that ridicules this response makes the problem worse. The message ‘AI will not replace you, but someone using AI will’ sounds conciliatory, yet also functions as a threat. It demands adoption under the pressure of status degradation and binds the product to anxiety. It may generate trials, but also hidden use, passive sabotage, or ceremonial adoption. A more mature strategy describes the new role in positive and verifiable terms: the expert defines criteria, inspects exceptions, expands coverage, and becomes the owner of the assisted process.

Table 7 · The Architecture of Permission

Form of permission	Question the product must resolve
Epistemic	Is there evidence that the result is good enough for this task?
Operational	Does integration avoid transferring cost to verification, security, or repair?
Institutional	Are there policies, accountability, audit, and a clear right of use?
Identity-based	Is the practice compatible with professional status and standards?
Social	Do relevant colleagues and managers treat use as legitimate?

Permission also has a collective component. An employee may regard the tool as useful and still refuse to be the first to apply it to an important task. The employee does not know whether the manager will reward initiative or conclude that the role is easily automated. Microsoft's 2024 data capture this ambiguity: 52 percent of workplace AI users hesitated to disclose use on their most important tasks, while 53 percent feared appearing replaceable (Microsoft, 2024). Personal permission existed; reputational permission did not.

The manager is therefore a distributor of legitimacy. Gallup reported in 2026 that workflow integration and managerial support were strongly associated with frequent use. Employees who said their manager actively supported AI use were about 1.7 times as likely to use it several times a week or more (Gallup, 2026e). This result alone does not establish causality, but it suggests that local social licence matters. Training without managerial clarity can produce knowledge of the tool without permission to apply it.

In enterprise sales, the purchasing unit is not only the user but a permission coalition. The economic sponsor wants value, security wants limits, legal wants accountability, the manager wants control, and the user wants protection from exposure. A demonstration that persuades one function can increase opposition

from the others. Founders often interpret user enthusiasm as proof that procurement is irrational. In reality, the two parties are buying different products: the user buys capability; the institution buys a change in its risk regime.

The threshold can be lowered through three mechanisms. Local proof shows that the product works on data, tasks, and constraints close to those of the customer. Local control lets the organization define limits, approvals, and responsibility. Distributed legitimacy provides examples from comparable and credible actors rather than only famous testimonials. None of these mechanisms requires the user to believe in AI as a general idea. It asks only for acceptance of a bounded process.

This changes how the category should be built. Broad language about intelligence and the future is useful for attention but weak for crossing the threshold. The higher the stakes of the task, the narrower the promise should become. ‘We write everything’ becomes ‘we prepare a first draft with sources and controls.’ ‘We automate research’ becomes ‘we reduce screening time while preserving the evidence trail.’ ‘We replace support’ becomes ‘we give the agent candidate responses and escalate exceptions.’ Narrowing does not reduce ambition; it builds the permission from which ambition can grow.

The threshold does not move only through product improvement. It can rise after a scandal, a public failure, or a regulatory change. A rapidly growing category can lose legitimacy faster than it loses capability. Reputation is therefore a shared asset among suppliers. Companies that externalize verification costs and publish exaggerated claims do not harm only their own customers. They raise the threshold for the entire market.

An early sign of maturity is a change in the questions customers ask. At first, the customer asks whether AI can do something. Then whether it can do it well enough. Later the customer asks how it integrates, who is accountable, what data it uses, and how the outcome is measured. The later questions sound more skeptical, but they are actually signs of proximity to purchase. The market is no longer debating whether the capability exists; it is negotiating the conditions of permission.

CHAPTER 8

Permission Cascades and Dense Islands

An idea does not spread merely because people hear about it. Information can cross the world in an afternoon while behavior remains unchanged for years. To use a new tool on an important task, a person must absorb costs, accept risk, and believe that someone else's experience is relevant to their own situation. Adoption is therefore closer to a complex contagion than to the transmission of a rumor.

Mark Granovetter's threshold model describes behaviors in which each actor's decision depends on the number or proportion of people who have already acted. One person may have a threshold of zero and experiment without confirmation; another enters after seeing a colleague; another waits for the majority; another almost never adopts (Granovetter, 1978). Different distributions of these thresholds can produce radically different collective outcomes in groups that look similar on the surface. A few low-threshold individuals do not guarantee a cascade. They must be connected in a way that lets their behavior reduce the thresholds of those who follow.

Damon Centola's experiments show that costly or uncertain behaviors often require reinforcement from multiple sources. In an online health community, adoption was more likely when participants received social reinforcement from several neighbors, and behavior spread faster and farther in clustered networks than in random networks with long-range ties (Centola, 2010). This reverses the simplest intuition about virality. Weak ties are excellent for carrying information; redundant ties may be better at legitimizing action.

AI fits this model because results are difficult to evaluate from the outside. A post may display a spectacular demonstration without revealing preparation time, discarded attempts, or verification work. Users need to see several examples from people they regard as comparable. A physician does not adopt because a technology influencer declares a model intelligent. Adoption occurs when respected colleagues demonstrate a bounded use, the hospital defines policy, and the risk becomes intelligible. A lawyer does not need another viral video; the lawyer needs evidence that other lawyers can preserve confidentiality, accountability, and quality.

This book calls a permission cascade the process in which each relevant adoption produces not only a new user but also a reduction in the social, cognitive, or institutional cost faced by the next one. The user publishes an example, writes a procedure, trains a colleague, persuades a manager, requests an integration, or creates a standard. Behavior leaves infrastructure behind. Once enough traces accumulate, the group no longer has to be persuaded from the beginning.

A permission cascade is not the same as a network effect. A product has a network effect when its value rises directly with the number of users. An AI editor may become more legitimate as adoption grows even when its functional value to one user does not directly depend on the presence of others. This is an effect of legitimation and market learning, not necessarily a technical network effect. The distinction matters to investors because ‘network effects’ rhetoric is often applied to any product that is growing.

Dense islands are communities in which users can observe one another’s behavior, compare outcomes, and transfer practices. They may be departments, professions, research groups, startup ecosystems, open-source communities, or creator networks. Their social density enables rapid feedback. A successful case is replicated, a failure is discussed, a workflow is refined, and a common language emerges. Such islands are more valuable market laboratories than large, heterogeneous audiences.

For a founder, the question is not only where many potential users exist, but where users can learn from one another. A market of one million isolated individuals may adopt more slowly than a community of ten thousand in which reputation circulates and examples are comparable. This is the power of apparently small professional niches. They can produce density of evidence, and density can lower thresholds faster than advertising reach.

Table 8 · Signals of a Permission Cascade

Signal	What it indicates in the market
Independent confirmations	Several relevant actors arrive separately at the same use.

Signal	What it indicates in the market
Transferred templates	Procedures circulate among colleagues without supplier intervention.
Managerial legitimation	Experimentation time and controlled errors become permitted.
Voluntary sacrifice	Users invest money, data, reputation, or redesign.
Replication across teams	The practice moves from a local island into adjacent contexts.
Institutional language	Roles, policies, budgets, standards, and responsibilities appear.

Marketing for a complex contagion must behave differently from marketing for a nearly frictionless application. Optimizing the individual click is not enough. It must create independent and observable confirmations: case studies within the same profession, communities of practice, local benchmarks, peer demonstrations, and champion roles. This does not mean manufacturing consensus. Confirmations work precisely when they appear autonomous and can include limits, failures, and conditions of use.

Managers function as nodes with unusual power. They influence not only opinion but also the allocation of time, access, protection, and rewards. Gallup reports that managerial support and workflow integration are associated with much more frequent AI use (Gallup, 2026c). A manager can lower the threshold with a clear statement: this use is permitted, experimentation time is legitimate, and controlled failures will not be punished. A manager can also block the market through ambiguity: ‘innovate, but take no risks; use AI, but own every error completely; save time, but keep every old obligation.’

Cascades can be beneficial or perverse. A community may legitimize a weak practice because no one wants to appear left behind. Demonstrations may select one another, and metrics may be changed to validate a prior decision. An idea can become standard not because it is superior but because it accumulated enough

coordination, capital, and complements. This book does not use inevitability as a synonym for truth or goodness. A cascade describes a dynamic, not a moral justification.

This warning also changes investor due diligence. Growth generated by executive pressure, subsidies, and promotional contracts can imitate a permission cascade. A stronger signal is voluntary behavior that creates complementary assets. Users change their processes, create templates, recommend the product in communities where their reputation is at stake, and continue after incentives end. A market that permanently requires the same subsidy has not lowered its threshold; it has merely paid it temporarily.

Another error is to search for one ‘chasm’ across society as a whole. AI does not cross a single divide. Every community has its own threshold distribution, risks, and legitimate actors. Programmers, teachers, lawyers, and civil servants may use the same family of models while occupying different social stages. A campaign that appears late in one ecosystem may be premature in another.

The best strategy is often sequential. The company enters an island where pain is intense and the result observable, helps the community produce practices and evidence, then uses those assets to enter an adjacent segment. Expansion occurs not only through new features but through the transfer of legitimacy. The relevant reference is not ‘thousands of companies use the product,’ but ‘organizations with risks and processes like yours have established how to use it.’

As a cascade approaches its threshold, the market’s language changes. The question is no longer ‘who has tried it?’ but ‘how do you do it?’ Conventions, benchmarks, consultants, jobs, and auxiliary tools appear. Critics may remain vocal, but they begin discussing the proper mode of use rather than whether use is possible. This is one of the best indications that an idea’s time is arriving: opposition is absorbed into institutional design.

An idea does not become difficult to stop because it has persuaded everyone. It becomes difficult to stop when enough islands have built around it and each island reduces the crossing cost for the next. What spreads is no longer only the technology. Permission spreads with it.

CHAPTER 9

The Market That Teaches Itself

New products are often described as finished objects entering a pre-existing market. The team builds, marketing communicates, and the customer decides. For general-purpose technologies, this picture is false. Product and market construct each other. Every implementation discovers a task, an exception, an integration, a rule, and a form of value that the supplier could not fully know in advance. Adoption does not merely distribute the product; it produces the knowledge required to make the product adoptable.

Eric von Hippel called lead users those who experience intense needs before those needs become widespread and therefore have an incentive to improvise solutions. They can function as laboratories for future demand because they live today with the problem the wider market will encounter later (von Hippel, 1986). The abnormal user principle extends this idea to the social conditions of AI. Early users do more than supply product requirements. They tolerate ambiguity, invent verification processes, negotiate acceptability, and create the vocabulary through which others will understand the category.

These contributions are complementary assets. A powerful model is not enough to reorganize work. Clean data, integrations, roles, policies, metrics, training, and organizational redesign are also required. OECD research finds that productivity advantages associated with AI adoption are linked to complements such as ICT skills, infrastructure, and use of other digital technologies (Calvino and Fontanelli, 2023; OECD, 2025b). The OECD Skills Outlook 2026 reports that skill shortages are a central barrier and that most workers do not need advanced AI expertise so much as digital, data, managerial, and problem-solving capabilities (OECD, 2026b).

This is why a technology can look revolutionary in demonstrations and disappointing in aggregate productivity at the same time. The productivity J-curve theory argues that general-purpose technologies require costly, difficult-to-measure investments in complementary intangible capital. Productivity may appear stagnant or even weaker while organizations rebuild processes, with gains becoming visible only after that capital accumulates

(Brynjolfsson, Rock, and Syverson, 2021). For AI, prompts, evaluations, task taxonomies, proprietary data, and new allocations of responsibility are forms of organizational capital, not incidental accessories.

A self-teaching market emerges when these investments do not remain isolated. Suppliers convert exceptions into features, consultants codify implementation, communities distribute practices, and customers demand common standards. The cost of the next implementation falls. The product improves, but more importantly the environment becomes more capable of using it. A technology can advance without a spectacular leap in the underlying model if the ecosystem reduces enough surrounding friction.

The process creates positive feedback loops. More users generate more cases; cases attract integration and capital; integration reduces risk; lower risk attracts more conservative customers; those customers demand governance; governance opens high-stakes industries. Each stage changes the nature of the market. What begins as an individual tool can become organizational infrastructure, while the buying criterion shifts from apparent intelligence to process reliability.

Table 9 • The Assets Through Which a Market Learns

Asset	Accumulated value
Exception map	Real problems, edge cases, and moments when the system must refuse.
Evaluation set	Examples and metrics that define ‘good enough’ for the customer.
Workflow template	Steps, approvals, and escalations that turn capability into process.
Permission scheme	Who may view, generate, approve, and execute.
Shared vocabulary	Language through which vendor and customer discuss risk and value.
Integration library	Connections that reduce the marginal cost of the next deployment.

For startups, the danger is confusing lessons from lead users with requirements of the majority. The abnormal user accepts configuration, errors, and difficult interfaces because the need is exceptional. The wider market will not accept the same costs. A lead user reveals where value may exist, not the final form of the product. The team must extract the valuable mechanism from the improvisation and remove the sacrifice that made it impossible for everyone else.

Another danger is capture by the most vocal users. Early communities may demand complexity, freedom, and features that increase their expert status while reducing accessibility. The product becomes perfect for the people who helped it emerge and unsuitable for those needed to turn it into a large market. Moving toward the mainstream sometimes requires a careful break with pioneer taste: fewer options, stronger product opinion, more predictable results, and integration with processes that early enthusiasts find boring.

Investors should observe what kind of knowledge is accumulating. Click and conversation data may improve the model, but they are not the only asset. Harder to copy may be the map of exceptions, the relationship with institutions, the evaluation scheme, or the way the product connects roles with different interests. A startup using the same base model as its rivals can still hold a durable advantage if it learns faster what an acceptable implementation looks like in a difficult niche.

This yields the concept of market-learning density: the amount of reusable knowledge produced by each customer. A large contract may deliver revenue and almost no generalizable learning if it is built as a one-off project. Ten smaller customers in the same segment may produce a repeatable pattern, a control set, and commercial evidence of much greater value. Before product–market fit, learning density may be a more important metric than the gross value of the pipeline.

A market also learns through failure. A pilot that never reaches production may reveal that the problem is not the model but the absence of a process owner or the impossibility of measuring the outcome. Abandonment after enthusiasm may show that the product saves time in one stage while creating work in another. Legal resistance may reveal the real requirement for provenance. An organization that refuses can be a better teacher than one that buys from fear of missing out.

Marketing therefore becomes a research function, not only an amplification function. It must discover which formulation of the problem produces recognition, which evidence moves the threshold, which objection conceals an identity conflict, and which use spreads among colleagues. Good content does not

repeat the promise; it codifies what the market has learned. A good case study does not merely state that a customer adopted. It explains what had to become true before adoption was possible.

At some point the loop reverses. Initially the supplier educates the market. Later the market educates the supplier and penalizes deviation. Standard expectations emerge around security, citation, interoperability, and control. Customers compare products using a vocabulary they did not possess two years earlier. At that point, the category has matured enough that a missing feature is no longer interpreted as a design choice but as a deficiency.

The signal that an idea's time is arriving is therefore deeper than rising demand. The idea begins to produce external memory. People no longer need to reinvent every use; they can inherit procedures, examples, and institutions. Each new participant enters a market slightly more intelligent than the one encountered by the previous participant. This is the real acceleration: not only do models learn from data; society learns how to use them.

CHAPTER 10

Productivity Without Learning and the Apprenticeship Debt

AI can make a beginner produce better before the beginner understands better. For an organization under deadline pressure, the distinction may appear academic. If the result is acceptable and the cost falls, productivity has risen. But a company does not merely consume competence; it must reproduce it. The simple, slow, and sometimes frustrating work assigned to juniors is often the infrastructure through which tomorrow's experts emerge. When that work is automated, the organization may save in the present while accumulating an invisible debt.

This book calls that cost apprenticeship debt. It arises when a system removes the opportunities through which people built mental representations, learned exceptions, and received feedback without creating an alternative path of development. Like any debt, it may be rational. Not every repetitive activity deserves preservation for its pedagogical value. The problem begins when immediate benefit is measured while the lost capacity to produce competence is excluded from the balance sheet.

Psychology uses the term cognitive offloading for delegating cognitive demands to the body or the environment. Offloading is often adaptive: external memory and tools release capacity for other activities. Yet supported performance and learning are distinct outcomes; a system can be an excellent offloading device and a poor teacher (Risko and Gilbert, 2016). Apprenticeship debt does not condemn delegation. It asks the organization to identify which cognitive function has been externalized and where the capacity for unaided judgment will be rebuilt.

Research offers reasons for both hope and concern. In the contact center, generative assistance appeared to transfer practices from strong agents to less experienced ones, and some gains persisted, suggesting learning rather than mere dependence on suggestions (Brynjolfsson, Li, and Raymond, 2025). AI can be an always-available teacher, explain the same issue in several ways, and offer feedback without the social cost of asking a 'stupid' question. For people without access to mentors, this is a genuine democratization.

An experiment with nearly one thousand secondary-school students shows how easily assisted performance can separate from actual learning. During mathematics practice, access to a general GPT-4-based system substantially improved performance. When access was removed for the examination, students in that condition performed 17 percent worse than the control group. A version designed as a tutor, with guardrails and pedagogical guidance, produced large gains during practice without the same harm in unaided assessment (Bastani et al., 2025). AI did not have one uniform effect; interaction architecture determined whether the tool shortened thinking or supported it.

Another experiment, involving 194 physics students, found that a carefully designed AI tutor produced learning gains more than twice those of an active-learning lesson, in less time and with higher engagement. The system was not a generic chat interface. It followed pedagogical principles, sequenced problems, used verified solutions, and delivered adaptive feedback (Kestin et al., 2025). The result rejects the claim that AI inevitably degrades learning. It degrades learning when the product optimizes the answer instead of the formation of capability.

The distinction is equally critical at work. An assistant that writes the report can increase output. A system that asks the user to formulate the hypothesis, compare alternatives, and justify acceptance can increase competence. The first minimizes cognitive friction; the second preserves productive friction. Early markets usually reward the first product because its demonstration is more spectacular and its time savings easier to quantify. Mature organizations will eventually have to purchase the second effect as well.

A 2025 CHI study of 319 knowledge workers and 936 examples of use found that greater confidence in AI was associated with less self-reported critical effort, while confidence in one's own ability was associated with more critical thinking. The authors also observed a shift of effort from producing information toward verification, integration, and supervision (Lee et al., 2025). The result does not prove long-term cognitive decline, but it describes a role transition. The question is not whether the person thinks less at every stage; it is whether the person knows how to think in the new stage.

Table 10 • What Must Be Measured Beyond Output

Dimension	Evaluation question
Assisted output	What does the person produce when the system is available?
Unaided transfer	What can the person do after assistance is removed?
Error detection	Can the person recognize an answer that is fluent but wrong?
Explanation	Can the person reconstruct the reasoning and justify acceptance?
Adaptation	Can the competence transfer to a new task or exception?
Mentoring pipeline	Can the organization still turn juniors into experts?

Apprenticeship debt can arise through three mechanisms. The first is removal of entry tasks. If summarization, preliminary research, documentation, and simple coding are automated, juniors lose opportunities to encounter raw material and understand how it becomes professional work. The second is shortened feedback. Immediate answers may replace the struggle required to build one's own representation. The third is homogenization. If everyone begins from the same suggestion, the variation in approaches from which the organization might learn is reduced.

The labor market offers an early signal, although not a verdict. US administrative data analyzed by the Stanford Digital Economy Lab show, in the August 2026 update, that employment among workers aged twenty-two to twenty-five in highly AI-exposed occupations remained substantially below the counterfactual trend, while no comparable decline appeared among experienced workers. The effect was concentrated where AI automates rather than augments work (Brynjolfsson, Chandar, and Chen, 2026). The study cannot attribute every change to AI and does not prove that apprenticeship is disappearing, but it raises the exact strategic question: if the entry route narrows, where will future experts come from?

For entrepreneurs and investors, this problem creates a product category, not merely a warning. Markets will need systems that measure unaided progress, introduce adaptive difficulty, record reasoning, generate practice, and turn real work into curriculum. Companies selling acceleration without development may retain customers until those customers discover that they have become dependent on people they no longer train. Companies able to demonstrate both output and capacity growth will have a more defensible enterprise story.

The debt also affects employer brand. A junior may interpret the tool as a mentor or as evidence that the organization no longer intends to invest in development. A senior employee may be freed from repetitive review or become the only person capable of verifying an expanding volume of output. If AI compresses execution without reallocating time to mentoring and harder problems, the organization does not liberate talent; it turns talent into quality control for a machine.

The solution is not artificial preservation of every old task. It would be absurd to retain inefficient work merely because earlier generations learned through it. What is needed is an accounting of formation. For every automated stage, leaders should ask what tacit knowledge was acquired there, how its absence will be detected, and which alternative experience can produce it. Sometimes the answer is simulation, sometimes rotation, comparative review, or assisted mentoring.

Marketing to education and HR should avoid promising that learning will become effortless. Frictionless learning is often consumption. A valuable product does not eliminate effort; it moves effort from mechanical activity toward model-building, feedback, and transfer. The difference should be measured through performance after assistance is removed, ability to explain, transfer to new tasks, and error detection.

Apprenticeship debt may become one of AI's largest secondary markets. As systems absorb entry-level work, organizations will need deliberate ways to form beginners. The winners will raise today's productivity without mortgaging tomorrow's competence.

CHAPTER 11

Nested S-Curves and the Anatomy of Segments

The phrase ‘AI adoption’ suggests a single curve that begins slowly, accelerates, and eventually stabilizes. For market decisions, this is almost useless. AI is not a homogeneous product, and society is not a single population. There are different curves for trial, repeated use, integration, and dependence; different curves across professions, companies, countries, and tasks; different curves for general tools and vertical systems. These curves are nested, they reinforce one another, and they can also block one another.

Data from 2026 make this fragmentation clear. Stanford HAI estimates extraordinarily rapid population adoption of generative tools, at roughly 53 percent within three years. The OECD reports that in 2025 more than one third of individuals had used generative AI, while only 20.2 percent of firms reported using any form of AI. Among firms, adoption in large enterprises was roughly three times that of small firms, while ICT and professional-services sectors were far ahead of others (Stanford HAI, 2026; OECD, 2026a). These figures are not contradictory; they measure different curves at different points.

An individual may adopt a chatbot before the company does. A company may purchase licences before teams adopt. A team may use the tool frequently without transforming the process. An industry may contain a few advanced firms that conceal a passive majority. A country may exhibit high personal use and low enterprise integration. Each layer influences the next without determining it.

This book proposes analyzing markets along four axes. The first is problem intensity: how costly, frequent, and urgent is the task? The second is verifiability: how quickly can an error be detected, and at what cost? The third is social density: how effectively do users communicate, and how relevant are their examples to one another? The fourth is permissibility: how much legal, reputational, and professional control must be constructed before use? Their combination explains why the same capability moves rapidly through one segment and remains blocked in another.

A frequent, painful, and easily verifiable task is fertile ground. Code completion, structuring an internal text, or preliminary classification can produce rapid feedback. A rare task with delayed feedback and severe consequences is much harder. Strategy, diagnosis, or scientific evaluation cannot be sold on speed alone because correctness may become known only later and fluent answers can conceal uncertainty.

Table 11 · Useful Axes for Market Segmentation

Axis	Commercial question
Pain intensity	How costly is the old process, and how urgent is the alternative?
Verifiability	How quickly and cheaply can output quality be established?
Social density	How often do users communicate, and how visible are the results?
Permission	What level of evidence, control, and legitimacy does the segment require?
Recurrence	How often does the task return, and how quickly can a habit form?
Complementarities	What data, integrations, skills, and institutions must be built around the product?

Nested curves also explain why an apparently saturated market can contain enormous open space. General chatbots may have broad adoption, while an enterprise capability with audit, proprietary data, and accountability may still be at the beginning. Conversely, a vertical category can look small in aggregate statistics while already being mature within its community. An investor who looks only at total addressable market misses the order in which a market can actually be penetrated.

That order matters. A startup should not necessarily begin with the largest segment, but with the one in which a local curve can accelerate and produce transferable assets. The ideal initial segment has a severe problem, users who communicate, observable outcomes, and enough autonomy to experiment. If it also generates standards or data that generalize to neighboring segments, it becomes a beachhead.

Rogers distinguished innovators, early adopters, the early majority, the late majority, and laggards, while Geoffrey Moore popularized the difficulty of moving from early adopters to the early majority in technology markets (Rogers, 2003; Moore, 1991). These frameworks remain useful, but AI complicates category membership. The same person can be an innovator for one task and a laggard for another. A department may belong to the early majority for summarization and remain before the chasm for autonomous agents. Segmentation by a general ‘attitude toward AI’ misses this granularity.

The minimum unit of analysis is therefore the workflow. A workflow has inputs, criteria, accountability, frequency, and an error cost. It permits an observable curve: how many have tried, how many use repeatedly, how many have integrated, how many have redesigned roles, and how many would suffer if the product disappeared? The last question is more valuable than declared satisfaction. Benign dependence indicates infrastructure.

Markets with nested curves also produce statistical illusions. An average may show growth while the target segment stagnates. Declared use may rise because AI features are embedded by default in software, not because people are buying specialized products. Revenue may increase through enterprise contracts before end users create value. Founders must track indicators at every layer rather than using one metric as universal proof.

During trial, activation matters. During repeated use, frequency and return to a real problem matter. During integration, the number of connected workflows and data sources matters. During institutionalization, policies, recurring budgets, roles, and audit matter. During dependence, withdrawal cost and the counterfactual outcome matter. Growth may look healthy at one level and fragile at the next.

Nested curves also change the message. The innovator buys possibility and access. The early adopter buys advantage within a community. The early majority buys a proven outcome and a whole product. The late majority buys compatibility

with the norm and protection against falling behind. A message celebrating pioneering can repel the very segment the company needs at the next stage.

A signal that the market has entered a new phase is the appearance of a different buyer. Initially the individual user decides. Then the functional manager. Later IT, security, legal, and procurement enter. This ‘complication’ is not necessarily useless friction; it is evidence that the product is attempting to move from private utility to institutional responsibility. A company that continues to sell only individual excitement in this phase will mistake maturation for slowdown.

In this context, an idea whose time has come is not one that rises simultaneously across every segment. It is one that creates a credible sequence of local curves. One segment produces evidence for another, an integration opens a new function, a standard reduces industry risk, and the product becomes easier to buy as the market learns. Irreversibility does not emerge from a global chart. It emerges when nested curves begin pulling one another upward.

CHAPTER 12

The Inevitability Radar

Saying after the fact that an idea's 'time had come' is not analysis. It is a retrospective story that erases all similar ideas that disappeared. To be useful to entrepreneurs and investors, the concept must generate predictions before the outcome is secure and permit the unwelcome conclusion that excitement will not become a market. The Inevitability Radar is this book's operational proposal for that task.

The radar does not measure how impressive a technology is. It measures whether adoption is beginning to generate the conditions of its own continuation. An idea becomes difficult to stop when repeated use creates complementary assets, lowers the cost of the next use, changes expectations, and increases the cost of non-use. Each signal can exist separately without producing inevitability. Strength comes from convergence.

The first signal is persistent pain. People must do more than declare the problem important; they must already spend money, time, risk, or reputation to solve it. Early users who tolerate an incomplete product because the alternative is worse provide stronger evidence than an audience impressed by a demonstration. The second signal is recurrence. Repeated use in a workflow matters more than account volume. The third is sacrifice: personal payment, learning time, improvised integration, and accepted risk reveal private value before public legitimacy.

The fourth signal is social density. Relevant users communicate, compare, and influence one another. The fifth is the appearance of complements: services, data, evaluations, training, roles, policies, and standards. The sixth is a falling verification cost. If every output requires expertise more expensive than the work saved, the market remains fragile. The seventh is legitimacy: credible actors transform experiment into permitted practice. The eighth is integration: the product enters organizational systems and responsibilities. The ninth is cost reversal: at some point non-use has to be justified.

Table 12 • The Inevitability Radar

Signal	Interpretation
Persistent pain	The problem survives fashion and the innovation budget.
Recurrence	Use returns without campaigns or mandates.
Sacrifice	Customers invest resources that are difficult to recover.
Social density	The practice circulates in credible communities.
Complementarities	Processes, data, evaluations, roles, and adjacent suppliers appear.
Falling verification cost	Quality becomes cheaper to demonstrate.
Legitimacy	Respected professionals and managers make the practice imitable.
Integration	The product enters existing systems and metrics.
Cost of non-use	Absence of the product begins to impair performance or status.

These signals should not be mechanically reduced to a pseudo-scientific score. A single number creates an illusion of precision and hides dependencies. Severe pain without verifiability may generate desperate experiments rather than a stable market. Social density without value may create fashion. Mandated integration without voluntary use may create contracts and resentment. The radar should be read as a configuration, not a sum.

For founders, the most important comparison is between demand signals and infrastructure signals. Demand can rise rapidly through attention, fear, and novelty. Infrastructure grows more slowly: procedures, budgets, roles, and control systems appear. When demand grows without infrastructure, the market is

vulnerable to scandal, cost, or fashion reversal. When infrastructure continues to grow during periods of reduced attention, the idea acquires real inertia.

For investors, the radar helps distinguish a company temporarily capturing a wave from one helping to build the sea. One startup may have strong revenue from manual services disguised as a product. Another may have less revenue but be building the evaluation standard or integration on which an entire category will depend. Due diligence should ask what becomes easier for the market after each customer and what remains proprietary to the company.

Another strong signal is a change in the type of objection. In the earliest phase, people say the technology cannot work. In the intermediate phase, they concede capability but challenge quality, ethics, or economics. During institutionalization, they demand policy, accountability, and interoperability. Objections become more concrete and more expensive, yet this is precisely what shows that the idea has entered reality. A product whose responsibility no one discusses has not yet acquired enough power.

The radar also includes counter-signals. Use stops when an incentive disappears. Customers cannot identify an owner of the outcome. Local savings create downstream costs. Every successful case depends on exceptional consultants. Data cannot move from pilot to production. The product does not survive comparison with a simplified non-AI process. Users enjoy the experience but would not be affected if it vanished. These signals do not prove final failure, but they indicate that inevitability remains a marketing claim.

The radar must be applied at the level of the segment and workflow. ‘AI in health care’ is too broad. Document triage, note drafting, image analysis, and clinical recommendation have entirely different pain, verifiability, and thresholds. An investor can reasonably conclude that one function’s time is arriving while another still lacks a decade of institutional development.

A practical use is to update a hypothesis map quarterly. The team records which signals it has observed, what evidence would falsify them, and which change would alter the thesis. This discipline prevents every event from being reinterpreted as confirmation. If usage rises while integration stalls, the thesis must be revised. If regulation raises cost while compliance providers and shared standards emerge, the market may become slower and more defensible at the same time.

The radar does not predict the exact date of inflection. Social systems are sensitive to sequence, threshold distributions, and external shocks. A

demonstration, price change, competitor, or policy can trigger a cascade that prior conditions made possible rather than inevitable. The tool is better at comparing directions and identifying missing conditions than at forecasting a calendar date.

The most interesting question is not ‘will AI win?’ but ‘what must become true for this behavior to continue without evangelists?’ When the answer includes recurrent use, complements, legitimacy, integration, and withdrawal cost, the market is approaching autonomy. When the answer is only ‘models will get better,’ the company has outsourced its strategy to technical progress.

The Inevitability Radar does not offer certainty. It offers a way to be explicitly wrong. That is its principal value. Important ideas deserve courage, but capital and time require hypotheses that can be revised. The best founder is not the one who declares inevitability with the greatest conviction, but the one who sees before others which mechanisms would make it real and begins to build them.

CHAPTER 13

Marketing a Technology That Does Not Yet Have Permission

Traditional marketing assumes that a product offers a known benefit within an intelligible category. AI marketing often has to sell the capability, the new behavior, the right to practice it, and the criterion by which the result should be judged—all at once. The customer must do more than choose a supplier. The customer must accept that part of the work can be reorganized and that this reorganization does not compromise identity, accountability, or position. Marketing adoption is therefore closer to designing an institution than to writing a slogan.

The first rule is not to sell AI when the customer is buying something else. During the novelty phase, the label attracts attention. As the market matures, it may increase perceived risk and make the product appear interchangeable with hundreds of wrappers. The customer is buying response time, fewer errors, volume coverage, continuity, audit, or access to expertise. AI explains the mechanism, not necessarily the category of value.

The second rule is to choose the segment by the structure of its threshold, not only by market size. A good early segment has a problem painful enough to tolerate imperfection, observable results, and a community capable of distributing evidence. Lead users should be sought where present need anticipates future need, not among people who love every new technology. An enthusiast without a problem provides feedback about the demonstration. A user pushed to the limit provides feedback about value.

The third rule is to sell a ladder of delegation. Products often leap in their language from ‘assistant’ to ‘autonomous agent,’ although markets need intermediate steps. At the first level, the system searches and proposes. At the second, it prepares a draft with sources. At the third, it executes within limits and requests approval for exceptions. At the fourth, it operates autonomously under supervision. Each step has different criteria and responsibilities. Customers can adopt value without accepting the entire ideology of autonomy.

The fourth rule is to turn objections into features or processes. If the customer fears a breach of confidentiality, the answer is not a post about the future but architecture, contract, and data control. If the customer fears hallucination, the answer is evaluation, citation, refusal, and an approval trail. If the customer fears status loss, the answer is a stronger professional role, not ridicule of resistance. A recurring objection is often an unformulated product requirement.

Table 13 · Barriers and Market Interventions

Barrier	Credible intervention
Confidentiality	Data control, isolation, contracts, permissions, and explicit retention rules.
Hallucination	Citations, evaluations, refusal, testing, and approval of exceptions.
Professional status	The expert becomes the owner of the criteria and the assisted process.
Workflow incompatibility	Integration into the current system and end-to-end redesign.
No outcome owner	A role, SLA, metric, and operational responsibility.
Learning risk	Practice mode, explanations, and measurement of unaided performance.
Price and ROI	Total cost, a local benchmark, and verified production savings.

The fifth rule is to demonstrate in proximity. Cases from distant industries create attention, not necessarily permission. Evidence should resemble the customer's task, data, risk, and institution. A general benchmark can show model capability; a local pilot shows that the organization can live with it. For complex contagions, several modest confirmations from peers may be more effective than one famous testimonial.

The sixth rule is to measure the outcome end to end. ‘Generation time’ is a convenient and often misleading metric. Preparation, verification, correction, escalation, and the consequence of error must also be included. The concept of workslop describes output that looks finished but transfers clarification and repair work to colleagues. A product can make the local user faster and the organization slower. Credible marketing shows where cost disappears, not merely where it moves.

The seventh rule is to protect user dignity. Adopting a technology that imitates competence can generate shame. Interfaces and messages that infantilize experts or celebrate the elimination of effort raise the threshold. The product should attribute judgment clearly, make human contribution visible, and let users explain the process without appearing to conceal their work. In many professions, attribution is part of the product.

The eighth rule is to create a community of practice before building an audience. An audience consumes promises; a community produces procedures. Early customers should be connected so that they can compare implementations, publish limits, and develop a shared language. This reduces support cost, accelerates learning, and creates the multiple confirmations a cascade requires. The community need not be enormous. It must be relevant and reputationally serious.

The ninth rule is to change the story as the market changes. At first, the message can celebrate possibility and difference. Once early adopters appear, it should show advantage and methods. For the early majority, emphasis moves to the whole product, references, integration, and risk. For the late majority, the product becomes standard and the message becomes compatibility with the norm. A company that remains permanently ‘the rebel from the future’ may become unable to sell to institutions that have finally decided to adopt.

The tenth rule is not to confuse fear of missing out with retention. FOMO can accelerate purchase and destroy trust when the result fails to appear. Maximalist promises create an impossible psychological contract. It is healthier to define a verifiable wedge and show the path to expansion. A category with a jagged frontier needs marketing that says where the system should not be used. Explicit limitation can build more trust than a catalogue of capabilities.

Marketing must also produce signals for the product team. Objections, pilot-to-production conversion, time to first value, usage distribution across roles, and reasons for hidden use are data about the permission threshold. The commercial team should be rewarded not only for closing the contract but also for

the quality of its hypothesis about adoption. A contract that never enters the workflow can be revenue and a market failure at the same time.

For investors, a strong AI marketing company is not the one that spends the most on category creation, but the one that can show how each marketing act removes structural friction. Content educates a control function, the community produces references, the pilot creates a reusable template, and an integration opens a new segment. Customer-acquisition cost is then partly an investment in category infrastructure.

Ultimately, permission marketing must answer an uncomfortable question: what will customers have to believe about themselves after they adopt? That they are less competent and easier to replace? That they cheated? That they abandoned their craft? Or that they became the owner of greater capability with clearer responsibility? Products do not enter markets through utility alone. They enter through identity stories that people can inhabit.

A technology whose time is arriving no longer needs to be presented as magic. It needs a social form in which people can use it without denying their standards. Marketing does not invent that form, but it can discover, name, and accelerate it. In AI markets, this is often the competitive advantage the model alone cannot supply.

CHAPTER 14

The Graveyard of Apparently Inevitable Ideas

Every technological era produces objects that intelligent people declare will inevitably change the world. Some do. Others remain memorable demonstrations, stable niches, or warnings about excessive confidence. If this book interpreted every burst of enthusiasm, criticism, and investment as evidence that an idea's time had come, the theory would be impossible to falsify. The Inevitability Radar must therefore be tested in the graveyard.

Google Glass combined the capital of a dominant company, fascination with augmented reality, and an easily understood demonstration. Yet the consumer edition was withdrawn from sale in 2015, while Google ended sales of the Enterprise Edition in March 2023 and support in September of the same year (Google, 2023). Reasons included price, battery life, unclear utility, and social discomfort with cameras worn in public space. The capability was real, but the permission threshold was not merely individual. Other people in the room also had to accept the object on the user's face.

Three-dimensional television offers a different lesson. The industry had spectacular cinema content, major manufacturers, and apparently ready distribution infrastructure. At home, however, the experience required glasses, compatible content, and extra cost for a benefit that was not frequent enough. By 2017, the last major manufacturers had abandoned the feature (Frankel, 2017). The technology was not impossible; the complements and viewing habit did not justify the friction. An impressive event failed to become a recurrent workflow.

Segway was introduced as a transformation of urban mobility but remained far closer to commercial, tourist, and recreational niches than to the promised revolution. Production of the Segway PT ended in 2020 after adoption remained far below what had been imagined at launch (Wilson, 2020; Thilmany, 2021). The product required cities, sidewalks, regulation, and social norms to adapt at the same time. When value depends on changing the environment, a technology competes not only with products. It competes with institutions, infrastructure, and accumulated reflexes.

These cases do not show that augmented reality, immersion, or micromobility were false ideas. Sometimes a product fails and the direction returns in a more suitable form. Smartphones absorbed functions imagined for wearables; screens and games continued to evolve; electric scooters found another combination of cost and convenience. The ‘time of the idea’ may belong to the underlying mechanism rather than to its first incarnation. For investors, the distinction between the inevitability of a capability and the success of a company is fundamental.

Apparently inevitable failures share several illusions. The first is confusing a demonstration with a habit. A result people want to see once is not necessarily one around which they will reorganize an entire day. The second is confusing attention with permission. Public curiosity can coexist with social rejection of the behavior. The third is underestimating complements. The product looks complete in the laboratory but requires content, infrastructure, standards, and organizational change that no actor has an incentive to finance.

The fourth illusion is assuming that performance will eliminate every non-technical friction. A better model does not automatically solve accountability, integration, identity, or economics. The fifth is survivorship selection. Entrepreneurs study companies that persisted until the market understood them and ignore those that persisted in a mistaken thesis. Persistence is a virtue only when feedback can change direction.

Table 14 • False Signals of Inevitability

Confusion	Why it can mislead
Demo ≠ habit	One-time amazement does not guarantee recurrent use.
Attention ≠ permission	Curiosity can coexist with social rejection.
Contract ≠ use	Purchase does not prove integration into work.
Pilot ≠ production	The protected environment removes precisely the relevant frictions.

Confusion	Why it can mislead
Model progress $\neq$ defensible advantage	Base capability does not automatically validate the application layer.
Criticism $\neq$ proof of adoption	Opposition can signal either high stakes or a real acceptability problem.
Subsidy $\neq$ economics	Capital-dependent growth can vanish when the full cost reaches the customer.

In AI, false signals can be unusually strong because demonstrations are cheap. A prototype can produce an astonishing result before the team knows how it will be evaluated at scale. Contracts can be purchased from innovation budgets and never renewed by operations. Users can experiment from fear of falling behind and then return to the old process. Investors can mistake progress in foundation models for automatic validation of every application layer.

The pilot may be the most dangerous artifact. It confirms that the technology can be used in a protected environment, with a motivated team and exceptional attention. Production requires an owner, continuous data, integration, support, metrics, and acceptance of error. Many products die between the two not because the demonstration was false, but because the demonstration removed precisely the real conditions of the market.

Another false signal is growth without economic autonomy. Discounts, cloud credits, manual services, and top-down contracts can produce revenue. The question is whether the product survives when customers compare its full cost with the alternative. A market whose price is subsidized by capital can look inevitable until the real cost reaches the user.

Criticism can also be misread. Intense opposition sometimes indicates that the product touches an important source of value or identity. At other times it reveals a genuine acceptability problem that familiarity will not remove. Google Glass was not blocked only by people unable to see the future; a permanent camera changed social relations and imposed costs on non-users. When a product imposes costs on people who did not choose it, the market cannot be reduced to buyer preference.

To avoid a self-confirming prophecy, every inevitability thesis should be accompanied by a failure thesis. What observation would show that the pain is not strong enough? How long without retention would change the view? At what verification cost does the workflow become uneconomic? Which critical complement has no owner? Which institutional response could constrain the market? A thesis that cannot answer these questions is not an investment thesis but a belief.

The graveyard should not produce cynicism. Many ideas fail because they appear before complementary assets exist. Some can return when cost, form, and norms change. But ‘too early’ is a dangerous explanation because it can protect every failure from criticism. An idea was too early only if we can identify the missing condition and later observe its formation. Otherwise, it may simply have offered insufficient value.

Predicting which ideas have reached their time also means accepting that some good ideas do not yet deserve a company and that an inevitable technology does not make any particular startup inevitable. The market may absorb the capability into incumbents, turn the product into a feature, or create a different category. Investors do not buy the future in the abstract. They buy the right of a concrete organization to capture part of it.

The theory becomes useful only when it can say ‘not yet,’ ‘not in this segment,’ or ‘not in this form.’ The graveyard of inevitable ideas is not the opposite of optimism. It is the intellectual price of an optimism that intends to survive reality.

What We Think Comes Next, 2026–2031

Predictions about AI have two recurring defects. The boldest are memorable and poorly calibrated; the cautious become true through language so broad that they guide no decision. This book will risk a middle ground. The following claims are not prophecies about model capability but hypotheses about adoption mechanisms. They should be evaluated as evidence accumulates.

The first prediction is that the AI label will gradually lose commercial power in mature categories. Generative functions will enter ordinary software, and buyers will compare outcomes, integration, risk, and total cost. Companies defined only by access to a model will be pressured by commoditization. Value will migrate toward data, distribution, workflow, evaluation, and sector authority. AI will not disappear; the need to announce it at every step will.

The second prediction is that the anti-slop market will grow faster than the term itself suggests. Provenance, selection, verification, citation control, reputation, and editorial accountability will become central functions. Content abundance does not eliminate demand for quality; it increases the cost of finding it. Platforms and organizations will pay for systems that reduce how much must be inspected and explain why something deserves acceptance.

The third prediction is that enterprise adoption will move from licences toward redesign. Gallup's 2026 data already show that access and declared adoption do not guarantee transformation, while workflow integration and managerial support are associated with more frequent use (Gallup, 2026a; 2026c). Organizations will discover that adding a chatbot on top of old processes produces activity rather than necessarily productivity. Budgets will shift toward integration, evaluation, role redesign, and adoption management.

The fourth prediction is that interfaces with calibrated agency will win in high-stakes markets. Fully autonomous systems will advance in bounded processes with rapid feedback and low error costs. In regulated professions and contextual decisions, winning products will provide delegation levels, veto rights, escalation,

and traceability. Control will not be treated as a concession to the fearful but as a productive capability.

The fifth prediction is that the expertise paradox will divide the market into two classes of value. Horizontal products will continue to lift baseline performance quickly and compress differences on standardized tasks. High-value products will amplify experts, providing coverage and exploration without eliminating judgment. Companies promising both that ‘everyone is an expert’ and that ‘experts become far more powerful’ will eventually have to choose which economy they serve.

The sixth prediction is that apprenticeship will become a board-level issue, not only an educational one. Early labor-market signals show pressure on entry roles in exposed occupations, particularly where AI automates work (Brynjolfsson, Chandar, and Chen, 2026). As organizations remove junior work, they will notice the absence of a competence pipeline. Products and policies will emerge that separate production mode from learning mode and measure performance without assistance.

The seventh prediction is that adoption differences will not follow generations in a simple way. Younger people will use general tools more frequently but may experience the threat to professional entry more intensely. Senior people may experiment more slowly but capture value faster once they control the process and see a clear case. Useful segmentation will combine expertise, status, task, and permission profile rather than year of birth.

The eighth prediction is that the line manager will become the principal node of diffusion. Strategy may be set at the top and tools purchased centrally, but daily permission is granted locally. The manager decides whether experimentation time is legitimate, whether errors are learning or misconduct, and whether saved time becomes better work or additional volume. Companies that do not redesign this role will demand transformation from people who are already overloaded.

The ninth prediction is that regulation will mature some markets rather than stop them. Rules raise cost and slow experimentation, but they can also create shared language, standards, and trust. Categories able to turn compliance into a product will enter organizations that remain inaccessible to general tools. Regulation will destroy business models built on ambiguity and strengthen the infrastructure of those able to demonstrate control.

The tenth prediction is that model differentiation will matter less than the ability to learn from implementation. Frontier models will continue to change,

while raw advantage at the application layer will be difficult to preserve. Companies with market memory—exceptions, evaluations, integrations, and relationships—will survive a change of model provider. Those whose product has been codified around one specific demonstration will be reabsorbed.

Table 15 • Twelve Falsifiable Predictions for 2031

Prediction	What would count as confirmation
The AI label becomes ordinary	Mature categories sell outcomes and integration rather than model access.
The anti-slop market grows	Provenance, verification, and selection become explicit budget lines.
Enterprise moves from licences to redesign	Budgets shift toward workflow, evaluation, and role change.
Calibrated agency wins	Products with delegation levels dominate high-stakes domains.
The market bifurcates by expertise	Horizontal tools lift the baseline; vertical systems amplify experts.
Apprenticeship reaches the board	Organizations measure the competence pipeline and unaided performance.
Age yields to the person–task pair	Segmentation combines role, expertise, control, and risk.
The line manager becomes a diffusion node	Use correlates increasingly with local legitimacy.
Regulation matures some markets	Compliance becomes a product and a barrier to entry.
Implementation memory beats model differentiation	Exceptions, evaluations, and relationships persist across model providers.

Prediction	What would count as confirmation
Hidden use migrates	It falls in mature segments and reappears at the capability frontier.
Critics build the institution	Valid objections become requirements, standards, and companies.

The eleventh prediction is that hidden use will decline in mature segments and temporarily increase in newly affected ones. As policies become clear, personal practices will be brought into governed systems. But each new capability will open other areas in which employees experiment ahead of the institution. The visibility inversion does not disappear; it moves along the frontier.

The twelfth prediction is that AI's opponents will contribute to its institutionalization. Criticism concerning slop, rights, bias, confidentiality, and competence loss will force markets to build filters and accountability. A technology becomes mature not when it eliminates critics but when it can convert valid objections into operational requirements. Some of the most valuable companies will exist precisely because critics were right about the defects of the first generation.

These predictions can fail in two directions. Technical progress may reduce verification cost faster than expected and accelerate autonomy. Or incidents, energy costs, regulation, and weak ROI may slow integration. In both cases, the relevant mechanism remains the same: capability must be translated into permission, complements, and processes.

The most important thing to watch through 2031 is not parameter count or account count. It is the change in the norm. In which professions does non-use begin to require explanation? Where do organizations include verification and provenance in the standard workflow? Where do juniors learn through new systems rather than being removed? Where do budgets move from experimentation to operations? These signs will show not only that the technology is present but that the market has reorganized around it.

When an idea's time arrives, it does not become immune to criticism. It becomes capable of turning criticism into infrastructure. This is the book's

fundamental prediction for AI: the next phase will be less spectacular than the first and much harder to reverse.

EPILOGUE

After Consensus, the Technology Disappears

At first, using AI is a declaration. The user appears optimistic, reckless, modern, lazy, brave, or suspect. The result is judged together with the method, and the method becomes a character trait. This period produces conflict precisely because society has not yet decided where the tool ends and the author begins.

Later, the distinction fades. Functions enter software, procedures incorporate verification, organizations assign accountability, and people learn to use different systems without turning each interaction into a philosophical position. ‘AI user’ no longer remains a meaningful identity, just as there is no powerful cultural category for people who use databases, spellcheckers, or search engines. What remain are strong and weak professionals, responsible and negligent processes, useful and useless products.

This does not mean that every problem has been solved. Invisible technologies can concentrate power, standardize thought, and distribute costs that users do not see. That is precisely why the disappearance of the label must be accompanied by the appearance of institutions. When AI is no longer a spectacle, its governance becomes more important, not less.

The book began from a provocative hypothesis: part of the aversion to AI may resemble the hesitation of people unwilling to cross the adoption gap first. Research made the hypothesis less comfortable and more valuable. Resistance is sometimes conformity, sometimes status defense, and sometimes a correct assessment of a weak product. Early users may be visionary, desperate, independent, privileged, or simply curious. Senior people may extract superior leverage and adopt more slowly; beginners may gain more in performance and lose more in formation. Personality matters, but competence, task, and environment may matter more.

The book’s contributions are tools for preserving these tensions without abandoning decision. The Abnormal User Principle tells us where to look for the future while warning us not to confuse the pioneer with the majority. The Slop Paradox separates the frontier from the median. The Visibility Inversion explains

why strong uses may be under-observed. The Permission Threshold and Permission Cascade show how utility becomes a norm. The Expertise Paradox and Apprenticeship Debt force markets to measure more than speed. The Inevitability Radar asks whether use is producing the conditions of its own continuation.

An idea does not become inevitable when everyone agrees. Complete consensus is rare and may arrive too late to matter. An idea becomes difficult to stop when enough people have built processes, skills, roles, rules, and expectations around it. At that point, abandonment no longer means merely changing an opinion; it means dismantling infrastructure.

For the entrepreneur, this is an invitation to build more than a demonstration. For the investor, to look for learning loops and complements rather than growth alone. For the marketer, to understand that people also buy permission to become different. For the leader, to protect both current productivity and the formation of future capability.

For the critic, it is an invitation to a position stronger than reflexive refusal. Effective criticism does not remain outside adoption, merely repeating that slop exists. It designs the standards that make technology tolerable, contestable, and useful. Future institutions will be built not only by those who believed early but also by those who saw early what could go wrong.

In the end, the surest sign of success will not be the ubiquity of the word AI. It will be its disappearance from important sentences. We will not ask whether a company ‘uses AI’ but whether it understands its work, verifies its results, and develops its people. We will not ask whether a book, analysis, or decision is ‘AI-generated’ but who is accountable for it and by what standard it was accepted.

Before everyone agrees, a few people enter, fail, learn, and construct permission. After almost everyone has entered, the technology ceases to be the main subject. What remains is the world it made possible—and our responsibility for its form.

SOURCES AND LITERATURE

References

The references include literature used to validate, refine, or challenge the book's hypotheses. Web sources and institutional reports were consulted in the versions available on August 23, 2026.

Agarwal, R., & Prasad, J. (1998). A conceptual and operational definition of personal innovativeness in the domain of information technology. *Information Systems Research*, 9(2), 204–215. doi:10.1287/isre.9.2.204 [source](#)

Arpaci, I., Kuşci, I., & Gibreel, O. (2025). The role of personality traits in predicting educational use of generative AI in higher education. *Scientific Reports*, 15, 30440. doi:10.1038/s41598-025-16339-0 [source](#)

Bastani, H., Bastani, O., Sungu, A., Ge, H., Kabakcı, Ö., & Mariman, R. (2025). Generative AI without guardrails can harm learning: Evidence from high school mathematics. *Proceedings of the National Academy of Sciences*, 122(26), e2422633122. doi:10.1073/pnas.2422633122 [source](#)

BetterUp Labs & Stanford Social Media Lab. (2025). Workslop: The hidden cost of AI-generated busywork. Research report. [source](#)

Bick, A., Blandin, A., & Deming, D. J. (2024, revised 2025). The rapid adoption of generative AI. NBER Working Paper 32966. doi:10.3386/w32966 [source](#)

Bohlen, L., Kruschel, S., Rosenberger, J., Zschech, P., & Kraus, M. (2025). Overcoming algorithm aversion with transparency: Can transparent predictions change user behavior? arXiv:2508.03168. doi:10.48550/arXiv.2508.03168 [source](#)

Brynjolfsson, E., Chandar, B., & Chen, R. (2025, revised August 2026). Canaries in the coal mine? Six facts about the recent employment effects of artificial intelligence. Stanford Digital Economy Lab. [source](#)

Brynjolfsson, E., Li, D., & Raymond, L. R. (2025). Generative AI at work. *Quarterly Journal of Economics*, 140(2), 889–942. doi:10.1093/qje/qjae044 [source](#)

Brynjolfsson, E., Rock, D., & Syverson, C. (2021). The productivity J-curve: How intangibles complement general purpose technologies. *American Economic Journal: Macroeconomics*, 13(1), 333–372. doi:10.1257/mac.20180386 [source](#)

Calvino, F., & Fontanelli, L. (2023). A portrait of AI adopters across countries: Firm characteristics, assets' complementarities and productivity. *OECD Science, Technology and Industry Working Papers* 2023/02. doi:10.1787/0fb79bb9-en [source](#)

Centola, D. (2010). The spread of behavior in an online social network experiment. *Science*, 329(5996), 1194–1197. doi:10.1126/science.1185231 [source](#)

Centola, D., & Macy, M. (2007). Complex contagions and the weakness of long ties. *American Journal of Sociology*, 113(3), 702–734. doi:10.1086/521848 [source](#)

Davis, F. D. (1989). Perceived usefulness, perceived ease of use, and user acceptance of information technology. *MIS Quarterly*, 13(3), 319–340. doi:10.2307/249008 [source](#)

- Dell'Acqua, F., McFowland, E. III, Mollick, E., Lifshitz, H., Kellogg, K. C., Rajendran, S., Kraye, L., Candelon, F., & Lakhani, K. R. (2026). Navigating the jagged technological frontier: Field experimental evidence of the effects of artificial intelligence on knowledge worker productivity and quality. *Organization Science*, 37(2), 403–423. doi:10.1287/orsc.2025.21838 [source](#)
- Dietvorst, B. J., Simmons, J. P., & Massey, C. (2015). Algorithm aversion: People erroneously avoid algorithms after seeing them err. *Journal of Experimental Psychology: General*, 144(1), 114–126. doi:10.1037/xge0000033 [source](#)
- Dietvorst, B. J., Simmons, J. P., & Massey, C. (2018). Overcoming algorithm aversion: People will use imperfect algorithms if they can (even slightly) modify them. *Management Science*, 64(3), 1155–1170. doi:10.1287/mnsc.2016.2643 [source](#)
- Doshi, A. R., & Hauser, O. P. (2024). Generative AI enhances individual creativity but reduces the collective diversity of novel content. *Science Advances*, 10, eadn5290. doi:10.1126/sciadv.adn5290 [source](#)
- Frankel, D. (2017, January 23). 3D TV is officially dead as Sony and LG stop making sets. *StreamTV Insider*. [source](#)
- Gallup. (2026a, April 12). Rising AI adoption spurs workforce changes. [source](#)
- Gallup. (2026b). State of the Global Workplace 2026. [source](#)
- Gallup. (2026c). AI in the workplace: What separates adopters and holdouts. [source](#)
- Gallup. (2026e, updated May 2026). Global indicator: Artificial intelligence. [source](#)
- Google. (2023, March 15). Glass Enterprise Edition announcement FAQ. [source](#)
- Granovetter, M. (1978). Threshold models of collective behavior. *American Journal of Sociology*, 83(6), 1420–1443. doi:10.1086/226707 [source](#)
- Guo, Y., & Wei, L. (2025). AI technology adoption and intergenerational knowledge transfer among older employees. *Frontiers in Psychology*, 16, 1673730. doi:10.3389/fpsyg.2025.1673730 [source](#)
- Hoff, K. A., & Bashir, M. (2015). Trust in automation: Integrating empirical evidence on factors that influence trust. *Human Factors*, 57(3), 407–434. doi:10.1177/0018720814547570 [source](#)
- Jussupow, E., Spohrer, K., & Heinzl, A. (2022). Identity threats as a reason for resistance to artificial intelligence: Survey study with medical students and professionals. *JMIR Formative Research*, 6(3), e28750. doi:10.2196/28750 [source](#)
- Kestin, G., Miller, K., Klales, A., Milbourne, T., & Ponti, G. (2025). AI tutoring outperforms in-class active learning: An RCT introducing a novel research-based design in an authentic educational setting. *Scientific Reports*, 15, 17458. doi:10.1038/s41598-025-97652-6 [source](#)
- Kovbasiuk, A., Triantoro, T., Przeglasińska, A., Sowa, K., Ciechanowski, L., & Gloor, P. (2025). The personality profile of early generative AI adopters: A Big Five perspective. *Central European Management Journal*, 33(2), 252–264. doi:10.1108/CEMJ-02-2024-0067 [source](#)
- Lee, H.-P., Sarkar, A., Tankelevitch, L., Drosos, I., Rintel, S., Banks, R., & Wilson, N. (2025). The impact of generative AI on critical thinking: Self-reported reductions in cognitive effort and confidence effects from a survey of knowledge workers. *Proceedings of CHI 2025*, 1–22. doi:10.1145/3706598.3713778 [source](#)
- Li, L. (2026). Personality, culture, and AI writing tool acceptance: A cross-cultural meta-analysis of Big Five traits and technology acceptance among university students. *Frontiers in Psychology*, 17, 1860730. doi:10.3389/fpsyg.2026.1860730 [source](#)
- Logg, J. M., Minson, J. A., & Moore, D. A. (2019). Algorithm appreciation: People prefer algorithmic to human judgment. *Organizational Behavior and Human Decision Processes*, 151, 90–103. doi:10.1016/j.obhdp.2018.12.005 [source](#)

- Microsoft. (2024). 2024 Work Trend Index Annual Report: AI at work is here. Now comes the hard part. [source](#)
- Mirbabaie, M., Brünker, F., Möllmann Frick, N. R. J., & Stieglitz, S. (2022). The rise of artificial intelligence: Understanding the AI identity threat at the workplace. *Electronic Markets*, 32, 73–99. doi:10.1007/s12525-021-00496-x [source](#)
- Moore, G. A. (1991). *Crossing the Chasm*. HarperBusiness. [source](#)
- Noy, S., & Zhang, W. (2023). Experimental evidence on the productivity effects of generative artificial intelligence. *Science*, 381(6654), 187–192. doi:10.1126/science.adh2586 [source](#)
- OECD, Boston Consulting Group, & INSEAD. (2025b). The adoption of artificial intelligence in firms: New evidence for policymaking. *OECD Artificial Intelligence Papers*. doi:10.1787/f9ef33c3-en [source](#)
- OECD. (2024). *OECD Digital Economy Outlook 2024, Volume 2: Strengthening connectivity, innovation and trust*. OECD Publishing. doi:10.1787/3adf705b-en [source](#)
- OECD. (2025a). Generative AI and the SME workforce: New survey evidence. *OECD Artificial Intelligence Papers*. doi:10.1787/2d08b99d-en [source](#)
- OECD. (2026a). Artificial intelligence use by individuals and firms continues to rise. Statistical release and analysis. [source](#)
- OECD. (2026b). AI and skills: What we know so far. *OECD Artificial Intelligence Papers*. doi:10.1787/f843b352-en [source](#)
- Rahman, M. M., Yankouskaya, A., AlShakhsi, S., Babiker, A., Liebherr, M., & Ali, R. (2026). How personal and cognitive traits influence AI attitudes, ethics, and well-being: The role of AI autonomy and criticality in design. *SN Computer Science*, 7, 686. doi:10.1007/s42979-026-05214-y [source](#)
- Risko, E. F., & Gilbert, S. J. (2016). Cognitive offloading. *Trends in Cognitive Sciences*, 20(9), 676–688. doi:10.1016/j.tics.2016.07.002 [source](#)
- Rogers, E. M. (2003). *Diffusion of Innovations* (5th ed.). Free Press. [source](#)
- Shaib, C., Chakraborty, T., Garcia-Olano, D., & Wallace, B. C. (2025). Measuring AI ‘slop’ in text. arXiv:2509.19163. doi:10.48550/arXiv.2509.19163 [source](#)
- Stanford Institute for Human-Centered Artificial Intelligence. (2026). *AI Index Report 2026*. Stanford University. [source](#)
- Thilmany, J. (2021, November 23). How the Segway evolved autonomous transportation. *ASME*. [source](#)
- Vaccaro, M., Almaatouq, A., & Malone, T. W. (2024). When combinations of humans and AI are useful: A systematic review and meta-analysis. *Nature Human Behaviour*, 8, 2293–2303. doi:10.1038/s41562-024-02024-1 [source](#)
- Venkatesh, V., Morris, M. G., Davis, G. B., & Davis, F. D. (2003). User acceptance of information technology: Toward a unified view. *MIS Quarterly*, 27(3), 425–478. doi:10.2307/30036540 [source](#)
- von Hippel, E. (1986). Lead users: A source of novel product concepts. *Management Science*, 32(7), 791–805. doi:10.1287/mnsc.32.7.791 [source](#)
- Wilson, M. (2020, June 23). Segway, the most hyped invention since the Macintosh, ends production. *Fast Company*. [source](#)