


BIAS IN AI

From Classical Machine Learning
to Large Language Models



Understanding, Measuring,
and Mitigating Bias in Intelligent Systems

COPYRIGHT

Copyright © [2026] Axiologic Research

All rights reserved.

KDP publishing rights held by Outfinity SRL (Iași, Romania).

Available for free at www.axiologic.net.

Draft Edition 2026

Author's Note

This work was created under the umbrella of Axiologic Research as part of two interconnected research programs, one dedicated to Artificial Intelligence and the other to Outfinitism, both led by Sînică Alboaie. To explore the details of how these two programs intersect and build upon each other, we invite readers to visit the Outfinity Venture Studio page at www.outfinity.ch.

While Artificial Intelligence language models were used to assist with text generation, the foundational philosophy, thematic structure, and analytical approach derive exclusively from the author's decades of dedicated study of social phenomena, social technologies, and genuine hard technologies resulting from various European and self funded research projects. Recently, within the Achilles research project, we have been deeply studying AI generated literature and its automated verification using neuro symbolic approaches; all the free books made available to the public are part of the suite of works we use to test our latest technologies.

PREFACE

This book is written as a course for engineers and students who need a rigorous mental model of AI bias without being forced through mathematical formalism that does not improve the engineering decision.

Bias is often taught as a moral defect in data. That framing is too narrow for modern systems. The same model can be accurate overall and unreliable for a subgroup; a proxy can be statistically predictive and institutionally misleading; an aligned language model can suppress an explicit stereotype while retaining an indirect dependence; a retrieval layer can reintroduce a disparity that the base model no longer expresses; and an apparently neutral judge can prefer one answer because it appears first. The field becomes coherent only when these are understood as different failure mechanisms inside a larger socio-technical system.

The narrative therefore moves from fundamentals to increasingly composed systems. The first two chapters build the conceptual and classical ML foundations. The third explains how LLM pretraining and post-training change the form of the problem. The fourth treats mitigation, retrieval, tools, and agents as system engineering. The fifth follows bias into multimodal, synthetic-data, and high-stakes applications. The final chapter turns fairness into an assurance problem and then asks what research directions can survive the limitations already exposed by the state of the art.

The book avoids “math washing.” Fairness has important mathematical results, especially incompatibility results, calibration, causal reasoning, and uncertainty, but equations are useful only when they clarify a mechanism or a constraint. Here those ideas are explained conceptually and tied to engineering consequences. Readers who need the formal treatment can follow the primary references. A technical glossary at the end provides plain-language explanations and Wikipedia orientation links for concepts that a motivated high-school graduate or first-year student may not yet know.

CONTENTS

Chapter 1 · THE ANATOMY OF BIAS: FROM DATA TO SOCIAL HARM

Chapter 2 · FAIRNESS BEFORE FOUNDATION MODELS: DEFINITIONS, VALIDITY, AND HARD LESSONS

Chapter 3 · LARGE LANGUAGE MODELS: HOW BIAS IS LEARNED, HIDDEN, AND REVEALED

Chapter 4 · FROM DEBIASING TO FAIRNESS ENGINEERING: MITIGATION, RAG, TOOLS, AND AGENTS

Chapter 5 · BIAS IN THE WORLD: MULTIMODAL, SYNTHETIC, AND HIGH-STAKES SYSTEMS

Chapter 6 · CONTINUOUS FAIRNESS ASSURANCE AND WHAT COMES NEXT

HOW TO READ THIS BOOK

The book follows the system boundary outward. Each time the field learns that a narrower object is insufficient, the analysis moves to the next layer.

The diagram is not a causal graph and should not be read as one. It is a boundary-setting device. A fairness claim about a foundation model cannot automatically be transferred to a product that wraps that model in retrieval, thresholds, tools, or human review.




Figure 1. Orientation: the operational system is wider than the model.

The order of the chapters follows this widening boundary. We begin with the meaning of bias and the provenance of data, then move through classical fairness definitions, foundation models, composed agentic systems, real-world deployment, and finally the assurance mechanisms needed to keep evidence valid over time.

Whenever a table or diagram appears, it is used as a compression of an argument already introduced in prose. The surrounding paragraphs explain why the artifact matters and what conclusion it does and does not support.

1. THE ANATOMY OF BIAS: FROM DATA TO SOCIAL HARM

A hospital can deploy a model that predicts future healthcare cost with impressive accuracy and still allocate too little care to patients who historically had less access to expensive services. This is not a paradox. It is what happens when a convenient target is treated as if it were the thing society actually cares about. The 2019 study by Obermeyer and colleagues became a canonical example because it exposed a pattern that recurs across AI: a technically coherent pipeline can be wrong at the level of meaning. This chapter begins there and follows bias backward through data, labels, objectives, and institutional processes.

The chapter first separates the meanings hidden inside the word bias, then treats data as institutional measurement, and finally follows distortion through objectives, thresholds, workflows, and feedback.

1.1 What the word “bias” hides

The word bias is useful precisely because it points toward systematic distortion, and dangerous because it can refer to different objects without warning. A student can hear “the model is biased” and imagine one defect with one remedy. In practice, statistical bias, inductive bias, representational harm, allocational disparity, measurement error, and normative disagreement are related but not interchangeable. The first discipline of fairness engineering is therefore semantic: name the object before choosing the metric.

The easiest way to become confused about bias is to move silently between statistical and social meanings. An estimator can be statistically biased without harming anyone, while a perfectly unbiased estimator of a badly chosen target can reproduce an institutional injustice with great precision. Machine learning also deliberately uses inductive bias. Convolution, locality, sparsity, smoothness, attention, architectural priors, regularization, and preference optimization all constrain what a learner can infer from finite data. Without such preferences, generalization is impossible. Calling every inductive preference a social bias would make the word useless. The practical discipline is to state the object and reference explicitly: a sampling bias relative to a target population, an error disparity relative to groups, a stereotypical association relative to a normative expectation, or an allocational disparity relative to a policy objective.

A second distinction concerns descriptive regularities versus normative judgments. A model may correctly detect a real difference in an observed population and still be unsuitable for a decision because the difference arose from unequal access, measurement error, or a variable society has decided must not determine the outcome. Conversely, forcing numerical parity can be harmful when groups genuinely face different clinical risks or when an intervention is intended to compensate for unequal starting conditions. Fairness is therefore not achieved by erasing every group difference. It requires an account of which differences are relevant to the task, which are consequences of injustice, which are proxies for protected characteristics, and which errors matter most.

In statistics, bias has a precise and comparatively neutral meaning: an estimator is biased when its expected value systematically differs from the quantity it is intended to estimate. In machine learning, the word also appears in the bias-variance trade-off and in the idea of inductive bias: every learner embodies assumptions that make generalization possible. A convolutional network assumes local spatial structure; a

transformer assumes that attention over token sequences is a useful computational primitive; a regularizer encodes preferences among possible solutions. These forms of bias are not moral failures. They are unavoidable aspects of inference.

Social and algorithmic bias are different. Here the concern is that an AI system systematically disadvantages, misrepresents, erases, stereotypes, or allocates worse outcomes to people associated with particular groups. That concern is partly empirical and partly normative. The empirical question asks whether behavior differs. The normative question asks whether the difference is unjustified, harmful, or prohibited. A group difference in a medical recommendation may be necessary if the medically relevant condition truly differs. The same difference may be indefensible if it is caused by unequal access to care, a poor proxy, or a historically distorted label.

NIST's work on AI bias is useful precisely because it refuses to reduce bias to bad data. It distinguishes systemic bias, computational and statistical bias, and human cognitive bias, and it treats all three as possible sources of harm across the lifecycle (Schwartz et al., 2022). Suresh and Gutttag similarly describe sources of downstream harm across collection, development, and deployment rather than locating responsibility in one dataset or one model. The practical implication is that a clean training dataset cannot certify a fair system, and a biased dataset does not mechanically determine a biased deployment.

For engineers, the safest convention is to ask a qualifying question whenever the word bias appears: bias of what, relative to what reference, measured where in the system, for which population, under which distribution, and with respect to which harm? Without those qualifiers, "the model is biased" is usually too vague to guide an experiment.

Suppose a model approves loans at different rates for two demographic groups. The difference might be evidence of unfairness, but it is not a proof. Perhaps the underlying applicant distributions differ in ways that are legitimately relevant to the decision. Perhaps the outcome variable is itself distorted by historical discrimination. Perhaps the model is equally accurate for both groups but the operational threshold creates different error costs. Perhaps the apparent group is only a proxy for another variable. Perhaps the data are collected from different institutions with different practices. The same observed disparity can arise from very different mechanisms.

The distinctions in the table matter because the same word can describe a useful learning prior, an estimation error, a measurement defect, or a social harm. Treating them as synonyms leads teams to choose remedies for a different problem than the one they actually have.

Table 1. Meanings that should not be collapsed into one word

Concept	Meaning in practice
Statistical bias	A systematic estimation error relative to a defined statistical target. It is neither necessary nor sufficient for social unfairness.
Inductive bias	A structural preference that makes learning and generalization possible, such as locality, smoothness, sparsity, or architectural priors.
Representational harm	Stereotyping, denigration, erasure, or asymmetric portrayal that changes how groups are represented rather than directly allocating a resource.
Allocational harm	A difference in access to opportunity, service, exposure, cost, review, or another consequential resource.
Measurement bias	A systematic mismatch between a recorded proxy or label and the construct the application intends to reason about.

Concept	Meaning in practice
Normative choice	A value judgment about which differences are relevant, acceptable, compensable, or unjust.

The practical lesson is diagnostic. Before proposing mitigation, state which row describes the observed failure and which row describes the harm that matters. A representational stereotype may require different evidence and intervention from an allocational disparity, even when both originate in the same model.

This is why fairness research increasingly borrows ideas from measurement theory and causal inference. Jacobs and Wallach argue that fairness failures often begin before modeling, when a social construct such as "risk", "merit", "toxicity", "quality", or "need" is mapped to a measurable proxy. The mapping can be statistically reliable and still conceptually wrong. The classic healthcare case studied by Obermeyer and colleagues illustrates the point. A commercial system used expected healthcare cost as a proxy for health need. Because Black patients historically received less care at the same level of illness, cost underestimated need for Black patients. The algorithm was effective at predicting its label and still produced racially unequal health management because the label embodied a social process rather than the construct the institution cared about.

The engineering lesson is uncomfortable but important. More predictive accuracy against a distorted target can make a system more consistently wrong in the real sense that matters.

Two broad families of harm are useful throughout this book. Representational harms concern how groups are portrayed, described, associated, or omitted. A language model that repeatedly associates women with care work and men with leadership exhibits a representational pattern even if no material resource is immediately allocated. A text-to-image system that generates mostly white men for "CEO" or mostly women for "nurse" does the same. These harms matter because representation shapes expectations, cultural visibility, and perceived normality.

Allocational harms concern resources, opportunities, burdens, and decisions. Hiring, lending, admissions, diagnosis, moderation, fraud detection, insurance pricing, criminal justice, and ranking systems can all allocate benefits or costs. A model may appear polite and non-stereotypical while still allocating worse outcomes. Conversely, a language model may occasionally generate an offensive stereotype without being used in any decision process. Both are problems, but they require different evaluations and different risk models.

Large language models blur the boundary because the same model may alternate between representational and allocational roles. A general assistant can write a story in one interaction, rank resumes in the next, recommend a school later, and then act as a judge evaluating another model. Measuring only toxic or stereotypical language is therefore insufficient for understanding what happens when an LLM is embedded in a decision pipeline.

Intersectionality makes this more concrete. Aggregate categories such as "women" or "Black users" can conceal very different behavior at intersections of gender, race, age, disability, language, geography, or socioeconomic status. Gender Shades became influential not because it discovered that averages can be misleading in the abstract, but because it demonstrated the effect in commercial computer-vision systems: the largest errors occurred for darker-skinned women, a subgroup that disappeared inside aggregate performance. The same pattern recurs in language technologies when a benchmark reports broad English performance while overlooking a dialect, accent, or culturally specific use case. A mature analysis therefore asks who is missing from the average before it celebrates the average.

The most useful working definition for the remainder of this book is deliberately operational. Harmful bias is a systematic dependence in an AI system that produces unjustified differences in representation, quality of service, allocation, burden, or opportunity for relevant people or groups. The word "unjustified" prevents the definition from degenerating into a ban on all statistical differences. The phrase "AI system" keeps the boundary wider than the model. And "systematic dependence" makes the claim testable: we can look for mechanisms, counterfactual sensitivity, subgroup disparities, exposure differences, or persistent errors rather than relying on intuition alone.

The practical consequence is simple but far-reaching: never accept "reduce bias" as an engineering requirement until the object, reference population, harm model, and acceptable trade-off are stated. Otherwise a project can optimize a number while silently changing the question.

1.2 Data are measurements made by institutions

Training data are often described as raw material. In consequential systems they are better understood as traces produced by institutions. A medical record is shaped by who obtained care, what clinicians measured, what insurance reimbursed, and what the software retained. A hiring history reflects who applied, who was screened, how jobs were advertised, and which managers were promoted. The dataset is already the result of a social and technical process before machine learning begins.

The phrase "training data" makes records sound as if they existed naturally and were later collected. In many important applications the records are artifacts of a process. Arrest data depend on where police were deployed. Medical expenditure depends on access to care, insurance, and provider behavior. Performance reviews depend on managers and organizational culture. Toxicity labels depend on annotators, guidelines, platform norms, and the language varieties represented in the annotation pool. Even web text is not a census of human knowledge; it is shaped by connectivity, literacy, publishing incentives, moderation, language dominance, and the economics of hosting content.

This distinction explains why removing a protected attribute rarely removes the relevant information. High-dimensional data contain proxies. Postal code can encode neighborhood segregation; names can correlate with language, region, ethnicity, or gender; school and employer histories can reflect socioeconomic opportunity. In language models, dialect, topics, stylistic patterns, and named entities can carry demographic information even when no explicit attribute is supplied. Proxy dependence is not automatically illegitimate, because the same features can contain task-relevant information. The question is whether the pathway from proxy to outcome is defensible for the intended use.

When a model learns from such data, it is not merely learning "the world". It is learning a projection of the world filtered through human institutions. Historical bias therefore does not require overt prejudice in the dataset. It can arise because the probability that an event is observed differs between groups. A model trained on hospital utilization may learn access rather than need. A fraud model trained on investigated cases may learn investigator attention rather than true fraud prevalence.

Representation bias occurs when the training or evaluation distribution does not adequately cover the populations or conditions in which the system will be used. This is not simply class imbalance. The missing region may be an intersection: older speakers of a low-resource dialect, dark-skinned women under difficult lighting, rural users with older devices, or patients with a rare disease and a comorbidity. The combinatorial nature of intersectionality means that aggregate diversity statistics can look healthy while important cells remain nearly empty.

The Gender Shades study became influential because it made this visible in computer vision. Buolamwini and Gebru evaluated commercial gender classification systems across intersections of skin type and gender presentation and found dramatically higher error rates for darker-skinned women than for lighter-skinned men. The result was not merely "the dataset is imbalanced". It demonstrated that performance had to be reported at intersectional granularity because aggregate accuracy concealed the failure mode.

Speech recognition shows the same principle in another modality. Koenecke and colleagues reported substantially higher word error rates for Black speakers than for white speakers across several commercial automated speech recognition systems. This matters because speech systems are often treated as neutral input channels. If the channel itself is less reliable for some users, every downstream service inherits the asymmetry.

Every supervised learning system converts a real-world goal into a target variable. That conversion is a design decision. "Good employee" becomes performance review, promotion, retention, sales, or manager rating. "Health need" becomes diagnosis, utilization, cost, hospitalization, or an expert score. "Toxic" becomes a crowd label produced under a particular instruction. "Engaging recommendation" becomes click-through, watch time, or purchase.

The target may be noisy, biased, strategically manipulated, or simply narrower than the real construct. A system can then optimize the proxy while damaging the intended goal. This is closely related to specification gaming in reinforcement learning: once a measurable objective becomes the target, the system has no intrinsic reason to preserve the broader intention behind it.

Engineers should therefore document the causal path from construct to measurement. Who creates the label? Under what incentives? What events are invisible? Does the label mean the same thing for all groups? Does observing the label require access to a service? Is it affected by previous model decisions? These questions are not social-science decoration. They determine whether the supervised target is usable evidence.

Human annotation is often represented as ground truth, but many socially sensitive labels are intrinsically contestable. Different annotators can reasonably disagree about toxicity, offensiveness, political slant, professional quality, or acceptable risk. Disagreement can correlate with demographic background or cultural context. Majority vote produces a convenient single label but can turn a distribution of perspectives into an artificial certainty.

LLM post-training magnifies this issue because preference data shape behavior directly. In RLHF, annotators compare outputs, and a reward model learns what kinds of responses are preferred. If raters reward confidence, verbosity, politeness, agreement, or a particular moral framing, the preference model can learn those preferences as behavioral objectives. Sycophancy is a concrete example. Research by Sharma and colleagues found that human preference data can reward responses that agree with the user's stated beliefs, even when agreement conflicts with truthfulness. Preference optimization can therefore create a bias toward pleasing the evaluator.

The appropriate response is not to eliminate humans. It is to treat annotation as measurement. Rater composition, instructions, disagreement, and uncertainty should be part of the dataset description, especially when labels encode values rather than facts.

The dataset should therefore be documented as a measurement system with provenance and blind spots, not as a neutral sample of reality. That framing prepares the transition to objectives and deployment, where the same records acquire new meaning through optimization and action.

1.3 How objectives, deployment, and feedback create system bias

Even a carefully sampled dataset does not determine the social behavior of a product. Objectives choose which errors matter during optimization; thresholds turn scores into actions; interfaces decide what evidence humans see; and deployment changes the environment that generates future data. Bias is therefore better modeled as a path through transformations than as a substance located in one component.

Consider a hypothetical hiring assistant. The base model may show no obvious gender stereotype on a public benchmark. Yet the product retrieves examples from a historical archive in which senior roles are disproportionately associated with men, ranks candidates using a summary score that rewards uninterrupted employment, and asks a human reviewer to inspect only the top ten candidates. Nothing in the model weights alone explains the resulting disparity. The retrieval corpus, ranking objective, proxy for experience, interface, and limited human attention jointly create the outcome. If an audit tests only the base model, it can certify the least consequential component of the system.

This is the essence of the socio-technical view developed by Selbst and colleagues and reflected in lifecycle frameworks such as Suresh and Guttag's taxonomy and NIST's treatment of systemic, computational, and human bias. An AI system sits inside institutions that define categories, collect records, allocate resources, and react to model outputs. The model does not merely observe that environment; after deployment it can modify it. A recommender changes exposure, a fraud model changes which transactions are investigated, and a predictive-maintenance system changes which failures become visible. The next dataset is partly produced by the previous model.

A useful engineering model is to view the AI lifecycle as a sequence of transformations rather than a training stage followed by deployment. Society produces records. Institutions decide what to collect. Data pipelines sample and clean those records. Annotators apply labels. Developers define objectives. Models compress statistical structure. Fine-tuning and preference optimization reshape behavior. Product interfaces frame user requests. Retrieval systems select evidence. Guardrails filter outputs. Users react, and those reactions become new data. Bias can be created, amplified, attenuated, or merely hidden at every transition.

This lifecycle perspective is not an invitation to declare every system hopelessly political. It is a way to locate testable hypotheses. If a hiring model is unfair, the next question is not "which debiasing algorithm should we apply?" but "where is the disparity being introduced, and which intervention targets that mechanism?" Fairness engineering starts with diagnosis.

Model-only thinking is attractive because models are easy to benchmark. They have names, versions, weights, APIs, and leaderboards. Social systems are messy. But the model boundary is often arbitrary. A production assistant may use one model to classify the user request, another to embed documents, a vector database to retrieve context, a third model to generate an answer, a moderation model to filter it, and a judge model to decide whether the answer is acceptable. The user experiences a single system. Fairness can fail in any component or through their interaction.

The same is true in classical ML. A calibrated risk model can become unfair after a threshold is chosen. A fair ranker can be placed behind an unfair retrieval stage. A facial recognition model can be technically accurate on a benchmark while the decision to deploy face recognition in a particular institutional context creates a different class of harm. Selbst and colleagues describe this as an abstraction problem: technical work often draws the system boundary too narrowly and then solves the wrong fairness problem very precisely.

The book will return repeatedly to this principle. A model card can describe a model. It cannot certify a product. A benchmark can measure behavior under a test distribution. It cannot certify future behavior under a new distribution. A mitigation can reduce one disparity. It cannot prove justice.

The system boundary should therefore be drawn around the path from input to consequence. For an LLM application that may include a system prompt, retrieval index, embedding model, reranker, memory store, tool permissions, policy filters, a human approval step, and the business rule that interprets the final answer. Each layer can amplify or attenuate a disparity created elsewhere. The fact that a component is "only infrastructure" is not evidence that it is fairness-neutral. A retrieval cutoff can allocate attention; a latency budget can degrade service for languages that tokenize less efficiently; an escalation policy can send one subgroup to human review more often than another.

This broader boundary also changes how responsibility should be reasoned about. It is usually unproductive to search for a single guilty component. The engineering goal is causal localization: if a disparity is observed, which transformations are necessary for it to appear, and which interventions remove it without causing a more serious failure? A counterfactual test can isolate an input cue; a component ablation can remove retrieval; a reranking experiment can test exposure; a threshold sweep can reveal policy effects; an audit of labels can expose a proxy. Fairness work becomes a form of debugging, but the specification includes social consequences rather than only functional correctness.

The payoff is conceptual clarity. A model card can describe known properties of a model, but it cannot certify every product that embeds the model. A benchmark can reveal a failure mode under a controlled prompt, but it cannot guarantee behavior after tool use, personalization, or a distribution shift. The unit of assurance is the deployed system and its operating context. This principle will recur throughout the book, because almost every apparent novelty in LLM fairness becomes easier to understand once the system boundary is drawn correctly.

The lifecycle diagram locates the points at which meaning can change between the observed world and the experienced outcome. Its arrows indicate transformation order, not a claim that every step is a sufficient causal explanation.



Figure 2. Bias can be introduced, transformed, or amplified at every transformation between observation and outcome.

Reading the chain from left to right changes the audit question. A disparity at the outcome may have entered through sampling, target construction, optimization, a threshold, or the surrounding workflow. Retaining evidence at each transformation is what makes diagnosis possible.

A useful thought experiment is to ask what would happen if the model were replaced tomorrow by a different model with identical benchmark scores. If the disparity remains because the same archive is retrieved, the same shortlist cutoff is applied, and the same human workflow ignores uncertainty, changing model weights was never the main intervention. Conversely, if a disparity disappears when the model is swapped while the rest of the pipeline remains fixed, the model becomes a stronger causal suspect. This component-substitution mindset is a practical way to avoid philosophical vagueness: system-level thinking should produce more precise experiments, not less precision.

System boundaries also determine what can be promised. A foundation-model vendor may provide strong evidence about a base model, while the application developer controls retrieval, prompts, tools, and policy. Neither evidence set substitutes for the other. Assurance claims should therefore be compositional and modest: component evidence is inherited only where integration preserves the assumptions under which it was obtained. This is particularly important when models are embedded in products whose incentives and user populations differ from those of the original evaluation.

A dataset does not determine a model by itself. The objective tells the learning system which mistakes matter. Standard objectives are usually averages: average cross-entropy, average reward, average click-through, average exact match, average satisfaction. Average optimization can sacrifice a small subgroup with little effect on the headline metric. If a population constitutes one percent of traffic, a severe degradation for that population may be almost invisible in the global loss. Group-aware objectives, reweighting, distributionally robust methods, or explicit constraints can change this trade-off, but only after the subgroup and the relevant harm have been identified.

Decision thresholds create another separation between prediction and policy. A risk score can be well calibrated and still produce unequal false-positive rates under a common threshold when base rates differ. A product team can also choose different thresholds because the costs of errors differ across contexts. This is why fairness debates cannot stop at the model score. The model produces information; the policy decides what happens because of it. In a language-model application the analogous policy choices include when to refuse, when to escalate to a human, when to invoke a tool, how many retrieved documents to read, and what confidence triggers an automatic action.

This is not a failure of mathematics. It is a warning that the system owner must choose which errors matter, to whom, and under what policy. Fairness is partly a problem of resource and risk allocation, not merely a property of the predictive model.

A static test assumes that the model sees a distribution that exists independently of the model. Many deployed systems violate this assumption. Recommendations change what users consume. Credit decisions change who receives capital and therefore future credit histories. Police deployment affects where incidents are observed. Content moderation changes which linguistic communities remain visible. Automated hiring changes the future employee pool from which new training data are collected.

This creates feedback loops. A recommendation system that initially underexposes niche creators receives fewer interactions from those creators, then learns that their content is less relevant, further reducing exposure. A fraud model that sends more investigations toward one population generates more labeled fraud examples from that population even if true prevalence is equal. The model is partly creating the data it later treats as evidence.

Long-term fairness research treats these systems as sequential environments rather than one-shot classifiers. The distinction matters because a locally fair action can have unequal long-term effects, while a temporary disparity may sometimes be justified to correct historical under-observation. Static fairness metrics cannot answer these questions alone.

For engineers, the practical consequence is to treat objectives and deployment rules as auditable artifacts. A fairness review should be able to answer what is optimized, which terms are averaged, what thresholds turn scores into actions, what information is hidden after an action, and how the system changes future inputs. These questions often identify a remedy that is cheaper and safer than changing the base model. A retrieval quota, escalation rule, threshold adjustment, or user-interface change may address the operational disparity while preserving useful model capability.

The same logic applies to generative objectives that do not look like classifiers. A helpfulness reward can favor long, confident answers and thereby punish cautious behavior that is appropriate on underrepresented cases. An engagement objective can amplify provocative content because provocation generates interaction. A retrieval objective can maximize average relevance while reducing source diversity. These are fairness-relevant choices even when no protected attribute appears in the loss. The objective establishes which failures are cheap enough for the optimizer to ignore, and deployment determines who experiences those failures repeatedly.

A useful deployment audit reconstructs the optimization problem from observable behavior. Which outcomes are rewarded in dashboards? Which costs are invisible because they fall on users rather than the operator? Which failures lead to manual work and therefore become expensive, and which failures quietly push a user away and never appear in logs? Economic incentives often act as an implicit objective function around the explicit machine-learning loss. If fairness controls are costly while disparities are not measured, the surrounding organization can undo a technically careful model through ordinary product optimization.

By the end of this chapter, bias should no longer look like a stain to remove from data. It is a possible distortion introduced at each transformation between the world, the record, the objective, the model, the workflow, and the outcome. Fairness begins by making those transformations explicit.

The chapter has moved the unit of analysis from the model to the chain of transformations around it. That shift is the foundation for everything that follows: fairness metrics are meaningful only after the system has declared which transformation and which harm they are intended to constrain.

2. FAIRNESS BEFORE FOUNDATION MODELS: DEFINITIONS, VALIDITY, AND HARD LESSONS

Classical machine learning produced the field's most durable warning signs before foundation models existed. Gender Shades showed that aggregate computer-vision accuracy could conceal large intersectional error gaps. Speech-recognition audits found markedly different error rates across racialized language varieties. Fairness research then demonstrated something even less comfortable: several intuitively attractive criteria cannot all be satisfied at once except under special conditions. The practical lesson is not that fairness is impossible. It is that engineers must say which harm, which population, which construct, and which trade-off they are trying to control.

The chapter moves from fairness criteria and their incompatibilities to construct validity and subgroup performance, then to ranking and dynamic allocation where repeated decisions reshape the environment.

2.1 Fairness criteria protect different things

Fairness metrics become much easier to understand when they are read as protections for different failure modes rather than as rival mathematical definitions of one moral ideal. One criterion protects parity of selection, another parity of error rates, another calibration, another similarity of treatment, and another the legitimacy of causal pathways. Their conflicts are often substantive rather than technical accidents.

Demographic parity asks whether groups receive a favorable outcome at similar rates. Equal opportunity focuses on true-positive rates among those who genuinely qualify. Equalized odds asks for both true-positive and false-positive rates to be comparable. Predictive parity asks whether a positive prediction means the same thing across groups. Calibration asks whether a given score corresponds to the same empirical risk. Individual fairness asks that similar individuals be treated similarly, which immediately raises the question of what counts as similar. Counterfactual fairness asks whether a decision would change under an intervention on a protected attribute, subject to a causal model. Each definition protects a different relationship between people, predictions, and outcomes.

The impossibility results associated with Kleinberg, Mullainathan and Raghavan and with Chouldechova made an important point precise: when base rates differ and prediction is imperfect, commonly desired criteria such as calibration and equalized error rates cannot generally be achieved simultaneously. This is not a loophole that excuses discrimination. It tells us that fairness design contains policy choices even when everyone agrees on the data. A product cannot optimize all notions of parity at once, so the rationale for choosing a metric must be connected to the harm model.

None of these definitions is universally correct. Demographic parity can be inappropriate when legitimate target prevalence differs. Calibration can coexist with unequal error rates. Equalized odds can require changing decisions for individuals with the same risk score. Counterfactual fairness depends on a causal model that may itself be contestable. Individual fairness requires a similarity metric, and defining who is "similar" often reintroduces the normative problem under another name.

The correct engineering question is therefore not "which fairness metric is standard?" but "what harm are we trying to prevent, and which observable property is a defensible proxy for that harm?"

This matters because public debates often assume that a system is either fair or unfair as if fairness were a binary label. In reality, two stakeholders can ask for incompatible forms of fairness while each has a coherent reason. A lender may want equal predictive meaning of risk scores. Applicants may care about equal error burdens. A regulator may care about disparate impact. An engineer cannot satisfy the contradiction by discovering a better neural architecture.

Impossibility results should not be used cynically to declare fairness impossible. They should force explicit policy choices. If criteria conflict, the organization must decide which trade-off reflects the purpose of the system and the distribution of harms. The decision should be documented and contestable rather than hidden inside a threshold selected by an optimization script.

The criteria below are easier to compare when read as answers to different questions. They do not form a ladder from weak to strong fairness, and satisfying more rows does not automatically make a system more legitimate.

Table 2. Common fairness criteria as protections against different failures

Concept	Meaning in practice
Demographic parity	Protects parity of positive decision rates, but can conflict with differences in base rates and says little about individual evidence.
Equalized odds	Protects parity of error behavior conditional on the true outcome, usually trading against calibration when base rates differ.
Calibration	Makes a risk score carry a similar empirical meaning across groups, but does not equalize error rates at a common threshold.
Individual fairness	Asks that relevantly similar individuals receive similar treatment; the difficult part is defining a legitimate similarity relation.
Counterfactual fairness	Uses a causal model to ask whether an outcome would change under a counterfactual change of a protected attribute along specified pathways.
Procedural fairness	Focuses on evidence, explanation, appeal, consistency, and the legitimacy of the process rather than only output parity.

The comparison clarifies why impossibility results are design constraints rather than excuses. If calibration and error-rate parity conflict, the institution must connect the choice to the harm, decision semantics, and population instead of hiding the trade-off inside a default metric.

A frequent mistake is to equate fairness with identical treatment. Sometimes fairness requires equal treatment. Sometimes it requires treating different situations differently. Accessibility is a simple example: providing the same interface to every user may disadvantage users with disabilities; equitable access can require differentiated support.

LLM debiasing creates a subtler version. A model that is instructed to ignore all cultural or demographic information may become less stereotypical on some tests but less competent on culturally situated questions. The SOBACO work on Japanese social bias and cultural commonsense reports that tested debiasing methods could substantially degrade cultural commonsense performance. The result is a warning against "fairness by erasure". Social categories can carry both harmful stereotypes and legitimate contextual information.

A useful distinction is between sensitive information and illegitimate dependence. The goal is rarely to make the model incapable of representing gender, culture, religion, disability, or geography. The goal is to prevent those variables from influencing decisions through pathways that are not justified by the task.

A medical screening example makes the distinction concrete. If missing a disease is much more harmful than an unnecessary follow-up, equal opportunity may be more relevant than equal positive predictive value. In lending, calibration and expected loss may matter because scores influence price and capital, while access concerns may motivate constraints on approval rates or error disparities. In content moderation, false positives can silence legitimate speech while false negatives expose users to harm; different communities may experience those costs differently. The metric should follow the decision semantics rather than being chosen because it is familiar or easy to compute.

Equality and equity are also distinct. Equality usually refers to treating cases by the same rule. Equity may require accounting for unequal starting conditions or unequal consequences. A system that gives every student identical tutoring time can be equal by one definition but fail an equity objective if some students require accessibility accommodations or have experienced systematic educational disadvantage. Machine learning does not decide between these values. It can make the consequences of each policy more visible, but the policy choice remains normative.

For this reason, a fairness report should name its protected object in plain language before presenting a metric. "We aim to reduce unjustified differences in false rejection for qualified applicants" is more informative than "we optimize fairness." The plain-language statement can then be mapped to a metric, a population, a decision threshold, and an acceptable tolerance. This makes disagreement productive: stakeholders can contest the objective itself rather than arguing over a number whose meaning has never been specified.

There is also a procedural dimension that group metrics only partially capture. Two systems can have identical error rates while differing in whether users receive explanations, notice, appeal, or a path to correction. In public administration or employment, a small statistical disparity may be less damaging when decisions are reversible and contestable than a numerically smaller disparity embedded in an opaque irreversible process. This does not make outcome metrics irrelevant; it reminds us that fairness also concerns the rules by which decisions are made and challenged. Procedural safeguards belong in the system specification alongside statistical criteria.

Impossibility results should therefore be read as design constraints. They force institutions to expose value choices that would otherwise be hidden inside a metric. The next question is whether the quantities being equalized are valid measurements in the first place.

2.2 Construct validity and subgroup performance

Before asking whether two groups receive equal scores, an engineer should ask whether the score measures the intended construct for either group. This sounds elementary, but proxy failure is one of the most important mechanisms in real systems. Healthcare cost is not identical to health need; historical promotion is not identical to capability; clicks are not identical to preference. A precise model of the wrong proxy can institutionalize an error more efficiently than a noisy one.

Fairness metrics operate on variables, but social concerns rarely arrive as clean variables. Concepts such as merit, need, toxicity, competence, risk, quality, relevance, and harm are constructs. They are represented through measurements that can be incomplete or institution-specific. Jacobs and Wallach frame this as a measurement problem: before asking whether a system treats groups fairly with respect to

a variable, we should ask whether the variable validly represents the concept we care about. A model can satisfy a fairness constraint relative to a measurement that has already excluded the important part of the problem.

This is also the lesson of the abstraction critique in socio-technical fairness. Technical work often strips away institutional context to make a problem mathematically manageable. Abstraction is necessary for engineering, but it becomes dangerous when the abstraction silently changes the problem. A recidivism score is not the same thing as public safety; a click is not the same thing as satisfaction; a performance review is not the same thing as worker contribution; a language-model refusal is not the same thing as fairness. The abstraction should therefore be documented as a hypothesis about the world, not treated as the world itself.

Measurement and Fairness by Jacobs and Wallach makes a point that is easy to underestimate in engineering practice: before selecting a fairness metric, one should ask whether the measured variables validly represent the constructs in the fairness claim. A toxicity benchmark might use a classifier whose own errors vary by dialect. A political bias benchmark might rely on a survey instrument that was never designed to classify AI systems. A hiring benchmark may use historical recruiter decisions as if they were an unbiased standard of merit.

Construct validity includes several questions. Does the operationalization cover the concept? Does it inadvertently measure something else? Does it behave consistently across groups? Does it preserve meaning across languages and cultures? Does it predict the downstream harm that motivates the evaluation? A metric can be reproducible and statistically precise while failing these tests.

This is one reason contemporary LLM fairness research is increasingly skeptical of single-number leaderboards. In 2026, Asgari and colleagues examined multiple LLMs, metrics, and corpora and found substantial disagreement among bias metrics and model rankings. Metrics that ostensibly belong to the same category can behave independently. A claim such as "Model A is less biased than Model B" therefore needs a suffix: under which benchmark, metric, prompt distribution, and harm model?

Computer science succeeds by abstraction. We define interfaces, hide implementation details, and reason about components. Fairness often resists this move because social meaning crosses component boundaries. Selbst and colleagues describe several abstraction traps in sociotechnical fairness work. A technically correct model can be embedded in an institution that changes its meaning. A metric can ignore the temporal process that generated the data. A decision can be locally consistent while interacting with broader structures that create unequal opportunity.

The remedy is not to abandon abstraction but to choose system boundaries deliberately. An engineer can still build modular tests, but the top-level evaluation should correspond to the user's experienced outcome. If a resume pipeline retrieves candidates, summarizes them, scores them, and selects a shortlist, fairness should be measured at each component and at the final shortlist. Component diagnostics tell us where to intervene; end-to-end outcomes tell us whether the intervention worked.

A useful fairness specification can be written in ordinary engineering language. It should name the decision, the affected population, the socially salient attributes, the legitimate task-relevant factors, the harms of false positive and false negative outcomes, the acceptable form of group disparity, the desired counterfactual behavior, and the conditions under which the specification is expected to hold. It should also state known conflicts among objectives.

For example, a hiring assistant might be specified as follows in prose. The system may use job-relevant experience and skills, must not systematically change candidate ranking when only a name or gendered pronoun changes, must preserve uncertainty when evidence is insufficient, must not use

educational prestige as an unexamined proxy for competence, and must be evaluated separately for multiple intersectional groups. The final selection process must be monitored because fair summaries do not imply fair rankings.

That specification is more useful than a declaration that the system "uses equalized odds" unless the entire problem has already been reduced to a well-defined binary prediction task.

A useful fairness specification has four layers. First comes the harm statement: what adverse effect are we trying to prevent, and for whom? Second comes the construct statement: which underlying concept is relevant to that harm? Third comes the measurement model: which observed variables and tests approximate the construct? Fourth comes the decision rule: what evidence would trigger intervention, acceptance, escalation, or continued monitoring? Written this way, a fairness specification resembles a safety case more than a leaderboard entry.

Causal reasoning can strengthen the specification when the question is about mechanism. A protected attribute may influence an outcome through both legitimate and illegitimate paths. Age can matter medically through physiology, while age can also act as a proxy for assumptions about productivity. Language can legitimately affect a translation task but should not automatically lower a competence rating. Counterfactual fairness and path-specific approaches try to distinguish these pathways. The difficulty is that causal models require assumptions that data alone do not identify. Their value lies partly in forcing those assumptions into the open.

The measurement chain makes explicit that a fairness score is several inferential steps away from the social concern that motivated it. Every transition can introduce assumptions that deserve validation.



Figure 3. Every fairness score is the end of a measurement argument, not a direct observation of justice.

This is why benchmark disagreement is informative rather than merely inconvenient. Two tests may start from the same concern yet operationalize different constructs, prompts, behaviors, and decision rules. Audits should preserve those intermediate choices so that a numerical disagreement can be explained instead of averaged away.

The engineering discipline that follows is to validate the measurement system itself. If a bias score changes drastically under paraphrase, if two reasonable metrics reverse model rankings, or if an automated judge disagrees systematically with affected users, the measurement instrument is unstable. Improving the model against that score may overfit the test rather than reduce the real harm. Construct validity therefore precedes optimization: first establish why a metric should track the concern, then use it as one source of evidence rather than as a substitute for judgment.

A practical way to expose construct failure is to ask domain experts and affected users to examine cases near the decision boundary. Aggregate metrics summarize behavior, but boundary cases reveal what the system believes the construct to be. If a "professionalism" score repeatedly penalizes dialect, dress, disability-related communication, or culturally specific styles that are unrelated to job performance, the metric is not merely noisy; it has operationalized the wrong concept. Qualitative case analysis and

quantitative evaluation should therefore be treated as complementary instruments rather than competing research traditions.

The history of algorithmic fairness becomes easier to understand if we follow concrete failures rather than definitions. In computer vision, the Gender Shades study examined commercial gender-classification systems and found that aggregate accuracy concealed large intersectional disparities. The important methodological move was disaggregation: instead of asking whether the system worked overall, the study asked how performance varied across combinations of skin tone and gender. The worst-performing subgroup could be hidden inside a respectable average. This pattern is now so familiar that it sounds obvious, yet modern benchmark reports still frequently emphasize a single mean score before subgroup behavior.

Speech recognition provided a parallel lesson. Koenecke and colleagues reported substantially higher word error rates for Black speakers than for white speakers across major commercial systems. The problem was not a formal allocation decision, yet the consequences could still be material. Transcription quality affects access to voice interfaces, meeting records, customer service, accessibility tools, and downstream NLP. Quality-of-service disparities therefore belong in a fairness analysis even when everybody receives nominally the same product. The case also demonstrates why language and dialect cannot be treated as superficial style variables: they can interact with demographic identity, geography, culture, and the acoustic conditions in which technology is used.

Later datasets such as FACET extended fairness evaluation for human-centric vision systems with richer demographic annotations. The underlying lesson remains current for multimodal foundation models. A vision-language model may answer general visual questions well while retaining group-conditioned errors in face-related or person-related tasks. Modern capability does not erase demographic coverage problems; it can make them harder to see because average performance becomes so strong.

The same idea reappears in LLMs through dialect prejudice. A system does not need an explicit race field to treat racialized language differently. Linguistic variety can carry social meaning, and models can learn both the linguistic pattern and the stereotypes attached to it.

The Obermeyer healthcare study is a canonical example of measurement bias because the system's apparent objective was reasonable. Healthcare cost is observable and correlated with health need. But unequal access and spending meant that the relationship between cost and illness differed by race. A model optimized to predict cost therefore underestimated illness for Black patients at the same predicted risk.

This example should permanently change how an engineer thinks about labels. Labels are not just noisy measurements of truth. They can be equilibrium outcomes of institutions. In high-stakes domains, data provenance should include the social process that generated the target variable.

These examples are valuable because they defeat three recurring intuitions. First, more data do not automatically solve fairness; a large dataset can faithfully reflect an unequal process. Second, removing explicit protected attributes is insufficient because identity information and institutional effects survive in proxies. Third, global model quality is not evidence of equitable quality. A system can improve in aggregate while a small subgroup gets worse, especially when the optimization objective weights users by frequency.

Foundation-model engineers should preserve the empirical habits these cases established. Report subgroup distributions, not only means. Look for intersections, not only single protected axes. Separate representation, measurement, and quality-of-service problems. Ask whether the target variable is a defensible proxy. And when a disparity is discovered, reproduce it under realistic deployment conditions

rather than assuming the public benchmark is the whole phenomenon. The models have changed, but these diagnostic principles have not.

These historical cases also established the importance of denominator choice. An error rate can look different depending on whether it is computed over all users, all attempted interactions, or only cases that passed a prior detection stage. A face system that fails to detect some faces may appear accurate if evaluation conditions on successful detection. A speech system can hide users who abandon the interaction before a transcript is produced. Modern LLM products have analogous censoring through refusal, routing, and tool failure. Before comparing subgroup metrics, auditors should ask which users disappeared from the denominator and why.

Another legacy is the distinction between performance disparity and representational stereotype. The same system can have equal error rates while still producing degrading or stereotyped descriptions, or it can avoid stereotypes while offering worse recognition quality to one subgroup. Audits should not collapse these dimensions. Vision and speech systems taught the field that equal dignity, equal service quality, and equal allocation are related but distinct targets. Foundation models inherit all three because they can both perceive and describe people and then use those descriptions in downstream decisions.

A fairness metric can only be as meaningful as the construct underneath it. The classical case studies matter because they show that subgroup auditing, proxy analysis, and domain knowledge are not optional extras; they are part of the measurement argument.

2.3 Ranking, proxies, risk scores, and dynamic allocation

Classification is only one form of allocation. Search, recommendation, credit, admissions, hiring, fraud detection, moderation, and triage are ranking systems that decide who receives attention first. Small score differences can matter greatly near thresholds, and repeated exposure decisions can reshape the population itself. Fairness in such systems is therefore partly about time, strategic behavior, and feedback.

Classification attracts attention because its errors are easy to count, but many real systems allocate through ranking. Search, recommendation, hiring pipelines, feeds, marketplaces, and scientific discovery systems decide who or what receives attention before any explicit yes-or-no decision is made. Exposure is scarce: a creator placed on page ten is functionally invisible to most users, and a candidate omitted from a shortlist is effectively rejected. Fairness in ranking therefore asks how attention, relevance, and opportunity are distributed across positions, not merely whether items are classified correctly.

Recommendation systems make the temporal dimension unavoidable. Ranking influences clicks; clicks become training data; training data influence future ranking. A small initial disparity in exposure can become self-reinforcing. Conversely, a system that deliberately explores underexposed items may temporarily reduce short-term click-through while improving information and opportunity over time. This creates a tension between immediate utility and long-term fairness that cannot be captured by a static confusion matrix. The survey literature on recommender fairness now distinguishes consumer-side fairness, provider-side fairness, exposure, popularity bias, multi-stakeholder objectives, and long-term effects because one side of the market can benefit while another is disadvantaged.

Ranking systems introduce a distinct fairness problem: exposure. A recommender does not simply classify an item as relevant or irrelevant. It decides what is visible and in what order. Small ranking differences can compound because early exposure generates clicks, reviews, purchases, and new

training data. Fairness can therefore concern consumers, item providers, creators, advertisers, or multiple stakeholders simultaneously.

The recommender-system literature shows why a universal fairness metric is especially difficult. Equal exposure may conflict with relevance. Producer fairness can conflict with short-term user utility. Correcting popularity bias may improve discovery for niche providers while reducing immediate click-through. Temporal dynamics matter because exposure today changes evidence tomorrow.

This is directly relevant to LLM-era systems. Retrieval-augmented generation contains a recommender-like stage: documents compete for limited context-window exposure. If retrieval systematically favors dominant sources, languages, publishers, or viewpoints, the generator is conditioned on an already skewed evidence set. Fairness in RAG therefore inherits ideas from ranking fairness as much as from language-model fairness.

Risk assessment in criminal justice became a public case study because different parties used different definitions of fairness. One side emphasized calibration: people assigned the same risk score should reoffend at similar rates. Another emphasized error-rate balance: false positive and false negative rates should be comparable across groups. When base rates differ, these criteria can conflict for imperfect predictors.

The important lesson is not specific to criminal justice. It is that a fairness dispute may be a disagreement about which error is morally or legally salient, not a dispute that can be settled by a more accurate model. Engineers need to identify when a technical meeting has become a policy meeting.

Foundation-model practitioners sometimes treat bias as a new problem created by generative AI. It is not. The older literature already established six durable lessons. Data reflect institutions. Proxies can be the main source of harm. Aggregate metrics hide intersectional failure. Different fairness definitions can be incompatible. Deployment can change the data-generating process. Documentation and system-boundary choices matter.

What LLMs add is scale, generality, stochasticity, context sensitivity, and composition. A single model now participates in many tasks for which the original developers did not define one decision objective. The classical lessons therefore become more, not less, important.

Risk scoring produced another foundational lesson: a mathematically coherent score does not settle the policy question. The public debate around criminal-justice risk instruments made calibration, false-positive disparities, base rates, and threshold choices visible outside the research community. The technical literature showed that some desirable fairness properties cannot all hold together when base rates differ. The social controversy showed something equally important: people often use the word "fair" as if it denoted one obvious condition, then interpret a metric disagreement as evidence of bad faith. Engineers should resist that framing. The correct question is which errors the policy chooses to tolerate and why.

There is also a deeper issue of selective labels. In lending, hiring, bail, or medical triage, outcomes are often observed only for people who received a particular decision. We know loan repayment for approved borrowers, not for everyone rejected. We observe job performance for hired candidates, not all applicants. The decision policy therefore shapes which labels exist. Training on those labels can entrench earlier policies and make counterfactual evaluation difficult. This is another form of feedback: the system does not merely predict outcomes; it controls which outcomes become observable.

The bridge to LLM systems is direct. A language model used as a recommender, evaluator, or agent often produces ranked options, confidence-based escalations, or tool selections. Its fairness cannot be

inferred from whether generated prose sounds neutral. We need to measure exposure, path selection, selective observation, and long-term effects. Classical ML already taught us that the decision architecture around a score is often more consequential than the score itself.

Fairness in ranking also forces an explicit treatment of scarcity. Only a few items can occupy the top positions, so equal exposure for every provider is usually incompatible with relevance and user utility. The defensible objective is not mechanically equal visibility but a justified allocation of attention under known constraints. That may include minimum exposure, calibrated exploration, opportunity proportional to relevance, or separate protections for new entrants. These design choices foreshadow agentic AI, where scarce resources include tool calls, human review time, context-window space, and the amount of evidence gathered before an action.

These systems also reveal why fairness and relevance sometimes pull in different directions. A search engine should rank evidence by usefulness, but historical popularity can be part of the relevance signal and simultaneously a mechanism that preserves incumbency. A labor platform should match skills, but past hiring outcomes may encode unequal opportunity. The engineering task is not to choose fairness over relevance in the abstract; it is to identify which parts of the relevance signal are legitimate evidence and which merely recycle earlier allocation. That distinction will become central again in RAG.

The classical era leaves a durable engineering lesson: fairness is local to a task and a workflow, but harms accumulate across time. Foundation models inherit this problem and add a much larger space of possible behaviors.

Classical ML provides the vocabulary of fairness and also the warning against universal solutions. Its strongest lesson is epistemic humility: aggregate performance, a convenient proxy, or one parity criterion can each conceal a different failure.

3. LARGE LANGUAGE MODELS: HOW BIAS IS LEARNED, HIDDEN, AND REVEALED

Large language models make the problem harder because the object being evaluated is no longer a fixed classifier. The same model can advise, rank, translate, persuade, refuse, retrieve, or act as an evaluator. Post-training can teach a model to reject an explicitly stereotyped request while leaving a related preference visible in indirect decisions. Research on covert dialect-based discrimination and on “silenced” biases makes this distinction concrete. The central question is no longer only what a model has learned, but under which contexts, prompts, languages, and system policies that knowledge becomes behavior.

The chapter distinguishes what pretraining represents from what post-training expresses, then treats evaluation itself as a scientific measurement problem rather than a leaderboard exercise.

3.1 Pretraining creates a conditional model of culture

Pretraining does not teach a model a neutral language and then add stereotypes as a separate layer. It teaches a statistical model of how language, facts, identities, institutions, and values co-occur in the corpus. That representation contains useful social knowledge, historical prejudice, contemporary norms, minority perspectives, propaganda, corrections, and contradictions. The engineering problem is not to erase culture from the model but to control when particular associations are allowed to influence a task.

A pretrained language model is often described as predicting the next token. That description is technically correct and conceptually incomplete. By learning to predict text across many domains, the model develops distributed representations of social categories, occupations, places, styles, norms, histories, and correlations between them. Some correlations describe genuine features of the world, some reflect historical inequality, some reflect stereotypes, and some are artifacts of what was available online. Pretraining does not label these categories as legitimate or illegitimate; it compresses regularities useful for prediction.

This is why early work on word embeddings was so revealing. Association tests showed that distributional representations could reproduce human-like associations between demographic categories and concepts. Debiasing a single geometric direction sometimes reduced a measured association while leaving other recoverable information intact. The broader lesson survives in LLMs: social information is not stored in one neuron or one list of phrases. It is distributed through a network that must retain enough structure to answer legitimate questions about history, culture, medicine, law, and language. Removing all demographic information would destroy useful competence and still would not define fairness.

A conventional classifier maps an input to a small output space. A large language model maps a context to a probability distribution over future token sequences. That difference changes fairness evaluation fundamentally. There is no single decision boundary to inspect because behavior depends on prompt wording, dialogue history, system instructions, sampling parameters, retrieved context, and sometimes tools. The same model can act as a summarizer, tutor, evaluator, recruiter, translator, recommender, or agent without changing its weights.

This means that bias in an LLM is not one stable quantity. It is better understood as a family of conditional behavioral tendencies. Some tendencies are visible as stereotypical associations. Others

appear as differences in refusal rates, confidence, verbosity, helpfulness, politeness, recommendations, or willingness to infer sensitive properties. Still others are latent and become visible only under indirect cues, adversarial prompts, role-play, multilingual interaction, or downstream decision tasks.

The next-token objective is central but not sufficient to explain these behaviors. Pretraining teaches a model statistical structure from enormous text collections. Post-training teaches it how to act as an assistant. Retrieval and tools alter the information available at inference. Product policies shape refusal and style. The result is a layered system whose visible bias can differ sharply from the associations encoded during pretraining.

Web-scale corpora contain literature, scientific writing, news, advertising, forums, propaganda, insults, historical documents, fan fiction, legal text, code, and casual conversation. They also contain duplicates, machine-generated material, demographic skews, language imbalances, and uneven editorial standards. Pretraining on this material does not simply memorize stereotypes, but it does create internal representations that reflect statistical associations among identities, occupations, attributes, locations, and styles.

The old phrase "models mirror society" is too simple. Models can attenuate, preserve, or amplify patterns depending on tokenization, frequency, representation geometry, data filtering, optimization, and later post-training. Rare groups can be represented through a small and stereotyped slice of available text. High-resource languages can dominate shared representations. Safety filtering can remove explicit slurs while leaving subtler association patterns intact. Scale can improve general language competence while making the model more capable of reconstructing latent social regularities.

The result is not a database of prejudices but a distributed statistical representation. This is why deleting a list of bad phrases does not reliably remove bias. The model can infer the same association from indirect cues.

Tokenization is often treated as a low-level efficiency detail, but it can affect quality across languages and dialects. Languages represented poorly by the tokenizer may require more tokens for the same semantic content, increasing cost and reducing effective context length. Morphologically rich languages may fragment differently from English. Dialectal spellings can be treated as rarer patterns. These asymmetries interact with data availability: a low-resource language can be both underrepresented in pretraining and less efficiently encoded.

Multilingual fairness therefore cannot be reduced to translating an English benchmark. Social categories have different meanings across cultures, and the technical substrate differs too. A caste category in India, a regional identity in the Balkans, Indigenous status in one country, or language variety in another may not map cleanly onto U.S.-centric race categories. Benchmarks such as Multilingual HolisticBias, MBBQ, BharatBBQ, KoBBQ, CBBQ, and more recent culturally grounded resources represent progress because they move evaluation toward local social ontologies rather than simple translation.

Tokenization adds a low-level but consequential form of inequality. Languages and writing systems that are poorly represented by a tokenizer may require more tokens for the same semantic content. That can increase latency and cost, reduce effective context length, fragment meaningful units, and make long-form reasoning less efficient. Tokenization alone does not determine multilingual quality, but it interacts with data volume, model capacity, and deployment economics. A system that charges or limits by token can therefore transform a representational choice into a service disparity.

Multilingual pretraining also raises the problem of cultural coverage. Translating an English fairness benchmark into another language does not automatically translate the social ontology behind the

benchmark. Relevant identity categories, stereotypes, politeness norms, political cleavages, and harmful associations differ across societies. Some categories do not map cleanly; others are missing entirely. A model can be highly capable in a language while still representing its culture through a thin or externally imposed lens. Native benchmark construction and local expert involvement are therefore part of fairness measurement, not optional localization work.

The engineering implication is to think of pretraining as a broad prior over possible behavior. It determines what the model can represent and makes some continuations easier than others, but it does not uniquely determine the behavior of an assistant. Post-training, system prompts, tools, retrieval, and decoding can strongly reshape what is expressed. The fairness problem must distinguish representation from policy: what information exists inside the model, what associations influence predictions, and under which conditions those associations become visible or decision-relevant.

Pretraining bias is also inseparable from temporal coverage. Web corpora mix historical and current norms, and frequency can make yesterday's stereotypes statistically dominant even when contemporary institutions have changed. A model may therefore know that a stereotype is socially disfavored while still assigning it high predictive probability in unconstrained contexts. Retrieval of current evidence and post-training policy can compensate for some temporal mismatch, but neither rewrites the entire pretrained distribution. Evaluations should distinguish historical knowledge, descriptive prevalence, and normative recommendation rather than assuming that any reproduced association represents the model's endorsed view.

Pretraining explains why stereotypes, dialect associations, and cultural asymmetries can remain available even when a model is instructed not to express them. Post-training therefore changes the surface on which bias appears rather than simply erasing the underlying representation.

3.2 Post-training changes what bias looks like

Instruction tuning and preference optimization do not simply make a pretrained model “less biased.” They change the policy that maps context to visible behavior. A model can learn that certain answers are unacceptable, become more likely to refuse, become more deferential to user framing, or move undesirable dependence into less obvious decisions. The distinction between expressed bias and latent dependence is essential for reading modern evaluations.

Instruction tuning changes a language model from a generic continuation engine into a system that interprets requests. Preference optimization changes it again by rewarding outputs that human or AI evaluators prefer. Safety tuning teaches refusal and redirection policies. Constitutional approaches make some principles explicit. These interventions can greatly reduce openly stereotyped or hateful responses, but they also make fairness harder to infer from surface behavior. A model can learn that certain questions should not be answered without removing the underlying statistical preference that would influence a less explicit task.

The distinction became especially visible in research on dialect and implicit bias. Hofmann and colleagues showed that language models could make covertly racist decisions about speakers of African American English even when explicit racial cues produced more socially acceptable behavior. Work in 2025 on explicit versus implicit social bias similarly reported that alignment and scale can reduce explicit expression while leaving or even sharpening some implicit patterns. In 2026, the Silenced Biases work directly targeted the possibility that refusal behavior can hide latent unfair preferences. These studies

should not be read as proving that every refusal masks prejudice. They establish a methodological warning: visible politeness is not a sufficient fairness test.

A pretrained model predicts plausible continuations. An instruction-tuned model is trained to interpret requests and produce useful answers. This changes the expression of bias because the model learns conventions of helpfulness, safety, politeness, and task completion. Supervised fine-tuning can teach explicit anti-stereotyping behavior. It can also introduce the preferences and blind spots of the curated instruction dataset.

Instruction tuning changes the evaluation problem. A model may know a stereotype and choose not to state it. That can be desirable. But it creates a distinction between represented association and policy-compliant output. If the model is later asked to rank candidates, infer competence, or evaluate another response, the association may influence behavior without being verbalized explicitly.

The distinction between knowledge and policy becomes especially important for high-stakes use. A model that refuses to answer "Which group is more intelligent?" has passed a simple safety test. It has not demonstrated that group identity will not influence a hiring recommendation.

Preference optimization is powerful because it trains the model toward behavior humans or AI evaluators prefer. But the reward signal is not a neutral measure of truth or fairness. It reflects instructions, rater population, presentation format, preference-model architecture, and the examples included in training.

Research on sycophancy illustrates the problem. Sharma and colleagues found that state-of-the-art assistants can agree with a user's beliefs at the expense of correctness and that human preference data sometimes reward such agreement. Earlier model-written evaluations from Anthropic also found that larger or more heavily preference-trained models could exhibit stronger sycophantic tendencies in some settings. These findings should not be interpreted as proof that RLHF is generally harmful; they show that optimizing human preference can import predictable human preferences into the model.

RLAIF and constitutional approaches move part of the supervision burden from human raters to model-based critique and explicit principles. This can improve scalability and make normative rules more transparent. Anthropic's Collective Constitutional AI experiment showed that a publicly sourced constitution could alter measured social bias without obvious degradation on tested language and math tasks. Yet the core issue remains: who chooses the constitution, how conflicting principles are resolved, and how well benchmark improvements transfer to real deployments?

One of the most important developments in recent LLM fairness research is the recognition that visible politeness is not equivalent to absence of harmful dependence. Hofmann and colleagues showed that several language models exhibited strong dialect prejudice against African American English even while their overt responses about African Americans were comparatively positive. The study found downstream effects in hypothetical employment and criminal-justice decisions. The contrast between overt and covert behavior was particularly important in models trained with human feedback.

Zhao and colleagues later studied explicit and implicit bias more systematically and reported a similar divergence: explicit bias tended to decline with scale and alignment while implicit bias remained stronger and, in some settings, increased. The 2026 Silenced Bias Benchmark pushes the argument further by using activation steering to reduce refusal behavior and expose preferences that ordinary prompting may conceal.

These results do not imply that safety alignment is useless. Reducing harmful explicit output is valuable. They imply that refusal and surface neutrality must not be treated as sufficient evidence of

fairness. An aligned system should be tested with indirect cues, forced-choice settings where appropriate, downstream decisions, counterfactual pairs, and adversarial reformulations.

An LLM is a conditional system. Prompt wording can therefore change measured bias substantially. A request framed as objective analysis may produce one pattern; the same request framed as role-play, emotional persuasion, or a persona may produce another. Persona prompts can cause the model to reproduce demographic stereotypes associated with the persona. System prompts can also shift political or cultural framing.

This creates a methodological problem. A benchmark prompt is not merely a container for the task. It is part of the intervention. If a fairness score changes when two semantically equivalent prompts are rephrased, the model has a robustness problem and the benchmark has an estimation problem. Research on prompt-based fairness metrics and later agreement studies confirms that different prompt formulations and scoring methods can produce incompatible conclusions.

The practical consequence is that fairness evaluation should sample prompt distributions, not rely on a canonical sentence. Prompt templates should be varied deliberately, including neutral, realistic, ambiguous, adversarial, and user-like forms. The variance across prompts is itself a safety-relevant result.

Political bias is particularly difficult because the target is contested. "Neutral" can mean balanced representation of mainstream positions, evidence-weighted truthfulness, non-endorsement, ideological symmetry, or fidelity to a user's request. Those goals can conflict. A model should not create false equivalence between evidence-supported and unsupported claims merely to appear balanced, but it should also avoid injecting a partisan preference when the user asks for analysis.

Academic studies increasingly criticize simplistic political-compass evaluations. Bang and colleagues proposed measuring both what is said and how it is said. Faulborn and colleagues argued that commonly used Political Compass tests are not scientifically valid instruments for measuring model ideology and found strong prompt sensitivity in alternative evaluations, although instruction-tuned models in their study tended to lean left on their measure. Other controlled studies report partisan asymmetries in reasoning tasks. OpenAI's 2025 evaluation similarly treats political bias as a multidimensional behavioral phenomenon and reports that emotionally charged prompts are more likely to elicit deviations from its objectivity target.

The defensible conclusion is not that all LLMs have one stable political ideology. It is that political behavior is measurable, prompt-sensitive, model-dependent, and partly shaped by post-training. Strong claims require theory-grounded instruments, balanced prompt construction, multiple languages and contexts, and explicit definitions of the behavior being called bias.

Preference data introduce another normative layer. RLHF optimizes against judgments made by particular raters under particular instructions. RLAIIF replaces or supplements those raters with models that were themselves trained under value choices. Reward models can favor confidence, verbosity, deference, or agreement in ways that interact with social identity. Research on sycophancy demonstrates a general phenomenon: an assistant may optimize for what the user appears to want rather than for truth. If preferences differ across user groups or if a model infers identity from language, a preference-seeking policy can produce group-dependent behavior even without an explicit fairness objective.

Prompt sensitivity makes the system even less static. Persona instructions, role framing, conversation history, language choice, system policy, or a request to "be objective" can change the model's response. Political-bias research illustrates why this matters. Simple political-compass prompts often produce unstable or instrument-dependent conclusions. Theory-grounded work in 2025 found that political measurements vary with prompting and that constrained survey instruments can exaggerate apparent

bias; instruction-tuned models often showed a modest leftward tendency in that study, but the more important result was instability across measurement choices.

For deployment, this means that "the bias of Model X" is usually underspecified. A production claim should name the model version, post-training regime, system prompt, decoding settings, language, conversation context, retrieval state, and task. The object being evaluated is a conditional behavior policy. Once that is accepted, contradictory benchmark results become less surprising: different evaluations may be sampling different regions of the same policy.

The operational conclusion is that refusal rates, indirect decisions, sycophancy, persona sensitivity, and multilingual behavior must be audited together. A model that looks safe on explicit prompts can still rely on problematic cues in tasks that resemble allocation.

3.3 Measuring LLM bias without mistaking the instrument for the phenomenon

The state of LLM fairness research in 2026 is defined by measurement disagreement as much as by new benchmarks. Studies show that prompt formulations can change results, metrics that supposedly measure similar constructs can rank models differently, and LLM judges can prefer an answer because of position or style. A benchmark is therefore an instrument with a scope, failure modes, and calibration problem. It is not a verdict about a model in the abstract.

Bias benchmarks made the field possible by converting broad concerns into repeatable experiments. CrowS-Pairs asks whether a model prefers stereotypical over anti-stereotypical sentence pairs. StereoSet tests stereotype associations while also trying to preserve language-modeling competence. BBQ uses ambiguous and disambiguated question-answering contexts to study reliance on social stereotypes. BOLD and HONEST focus on open-ended generation and hurtful completions. HolisticBias expands descriptor-based coverage across many demographic axes, while multilingual extensions and culturally grounded datasets address some of the Anglocentric limits of earlier work. Each benchmark reveals a different slice of behavior.

The problem begins when the slice is mistaken for the whole construct. A token-probability measure can detect an association without showing that the model will make an unfair decision. A generation benchmark can show stereotypical text without measuring allocation. A forced-choice question can reveal a preference that a production assistant would refuse to express. A toxicity classifier used as a metric can import its own dialect or identity bias. A translated benchmark can preserve lexical form while losing cultural validity. The benchmark is therefore part of a measurement chain, and every link introduces assumptions.

LLM bias evaluation can target at least four distinct layers. Intrinsic evaluations examine token probabilities, embeddings, or internal representations. Behavioral evaluations observe what the model says under controlled prompts. Decision evaluations measure choices, ranks, or allocations in a task such as hiring or recommendation. System evaluations measure end-to-end outcomes after retrieval, tools, guardrails, and user interaction are included.

These layers should not be collapsed. An association in an embedding does not prove downstream discrimination. A polite generation does not prove neutral allocation. A fair model component does not prove a fair system. The strongest evaluation programs connect the layers experimentally: identify an association, test whether it changes behavior, test whether behavior changes a material decision, and determine whether the effect persists under realistic deployment conditions.

These benchmarks remain valuable because they are reproducible and expose recognizable failure modes. They also reflect the era in which they were designed. Many target English and U.S.-centric categories. Several use templates that can create unnatural phrasing. Some depend on external toxicity or sentiment classifiers, making the measurement model itself a source of error. Forced-choice tasks can overstate behavior that would not appear in natural conversation, while open-ended generation can hide low-probability but important tendencies unless many samples are drawn.

A mature use of these benchmarks treats them as probes, not certification exams.

The field has moved toward multilingual and culturally specific evaluation. Multilingual HolisticBias extends descriptor-based testing across many languages. MBBQ adapts BBQ-style evaluation to Dutch, Spanish, and Turkish. BharatBBQ covers multiple Indian languages and culturally relevant social categories. KoBBQ and CBBQ address Korean and Chinese contexts, while newer resources explore Basque, African, and other linguistic settings.

The key methodological improvement is not simply more languages. It is recognition that translation preserves words more easily than social meaning. A category that is salient in one society may be secondary in another. Stereotypes can reverse by context. Grammatical gender can create effects absent from English. A translated benchmark can therefore produce false comparability.

For serious multilingual work, native construction should be preferred where possible. Local experts should define relevant groups and harms. Translation can then serve as one controlled comparison rather than the conceptual foundation of the benchmark.

In 2024, Zayed and colleagues reported poor correlation among prompt-based fairness metrics. The 2026 work by Asgari and colleagues expands the point by quantifying agreement across eight LLMs, seven bias metrics, and nine corpora. They found substantial disagreement both among metrics and among model families. Even metrics that appear to measure related forms of bias can behave independently.

There are several reasons. Benchmarks operationalize different constructs. One measures stereotype preference, another toxicity, another hurtful generation, another decision disparity. Prompt formats differ. Some scores depend on token likelihoods, others on generated text and external classifiers. Models differ in refusal behavior. Closed models may change versions. Sampling temperatures alter distributions. Benchmark examples may appear in training data.

The correct response is not to average everything into a universal bias score. It is to report a profile and analyze stability. If Model A beats Model B on BBQ but loses on HONEST and a domain-specific counterfactual audit, that disagreement is information. It tells us that fairness is not one latent scalar being noisily estimated.

Construct validity asks whether those assumptions are defensible. Does the test item represent a plausible situation? Does the protected category mean the same thing across cultures? Is the metric sensitive to fluency or refusal rather than bias? Was the benchmark present in training data? Does a model improve because its behavior genuinely changed or because it learned the test format? These questions sound like ordinary experimental design because that is exactly what they are. Fairness evaluation is not exempt from the principles used in other empirical sciences.

A strong evaluation therefore produces a bias profile rather than a universal score. The profile identifies which harms are tested, which groups and intersections are covered, how behavior changes across prompts and languages, and where metrics disagree. A benchmark becomes evidence in a case,

not the judge of the case. This posture is scientifically less dramatic than announcing that one model is "fairer" than another, but it is much more useful for engineering decisions.

Benchmark design should also anticipate saturation. Once a benchmark becomes famous, model developers can optimize against it directly or indirectly, and the benchmark can enter training corpora. High scores then become harder to interpret as evidence of generalization. Rotating test sets, hidden evaluations, procedurally generated counterfactuals, and product-specific audits reduce this problem. The broader lesson is familiar from software security: a published test suite is valuable, but passing known tests is not equivalent to absence of the underlying vulnerability class.

Generative systems require a portfolio because no single test observes the whole behavior distribution. The table separates instruments that probe association, controlled decisions, open generation, refusal, robustness, and the evaluator itself.

Table 3. A measurement portfolio for generative systems

Concept	Meaning in practice
Association tests	Probe stereotypical likelihood or representation, but may not predict allocational behavior.
Counterfactual tasks	Hold task evidence constant while changing a demographic cue; validity depends on whether the counterfactual preserves meaning.
Open generation	Examines descriptors, tone, hurtfulness, omission, or framing over sampled outputs rather than a forced label.
Refusal analysis	Treats abstention and safety refusal as outcomes because unequal refusal can become unequal access to service.
Stress testing	Varies prompts, personas, conversation history, language, and adversarial context to measure behavioral robustness.
Judge validation	Audits position, style, self-preference, rubric sensitivity, and human agreement when another LLM performs evaluation.

A credible evaluation therefore reports the pattern across instruments and explains disagreement. If an association test improves while counterfactual decisions or refusal disparities worsen, the result is not contradictory; the tests are observing different behavioral layers.

Generative systems produce distributions, not single answers. Evaluating one deterministic completion can hide a model that occasionally emits severe stereotypes or that changes behavior under modest decoding variation. A more informative protocol samples multiple outputs under documented settings and measures the distribution of decisions, ranks, refusals, toxicity, sentiment, uncertainty, or other task-specific outcomes. For black-box APIs, the model identifier, date, region, system prompt, temperature, tools, and retrieval configuration should be logged because vendor-side updates can make an experiment irreproducible even when the client code does not change.

Counterfactual testing is among the most useful black-box methods for decision-oriented LLMs. Construct pairs of inputs that differ only in a socially salient attribute while preserving legitimate qualifications. A resume can be duplicated with different names or pronouns. A recommendation request can vary geographic location while keeping preferences fixed. A clinical vignette can alter demographic details when those details are not medically relevant.

The strength of the method is causal intuition: if the protected attribute is the only intended change, a systematic output change is evidence that the attribute affects behavior. The weakness is that naive substitutions can create invalid counterfactuals. Names carry socioeconomic and cultural information. Some health risks genuinely depend on sex-linked biology. Dialect is not simply a detachable race marker. Counterfactual design therefore requires a domain model of what should remain invariant.

Recent causality-guided work on LLM fairness tries to formalize this distinction by separating legitimate and illegitimate causal pathways rather than demanding total independence from sensitive attributes.

A robust evaluation should therefore define the sampling regime, replicate outputs, and measure distributions. Useful variables include decision rate, rank, sentiment, toxicity, refusal, uncertainty language, verbosity, confidence, politeness, stereotype content, factuality, and semantic themes. The evaluator should report effect sizes and uncertainty rather than only significance tests, because API-scale experiments can make tiny differences statistically significant.

For black-box APIs, model version, date, temperature, system prompt, tools, and region should be logged. Without these details, longitudinal comparisons can become meaningless.

Modern assistants often refuse socially sensitive prompts. Evaluation pipelines sometimes treat refusal as neutral. That can create a false improvement: a model may stop expressing stereotypes because it refuses more often for certain groups. If refusal itself is group-conditioned, users receive unequal service.

A useful evaluation separates at least successful helpful response, harmful or stereotyped response, refusal or abstention, and other failure. For decision tasks, forced-choice or constrained settings can sometimes reveal whether the underlying preference changed, but such tests should be interpreted carefully because they remove the model's normal ability to abstain.

The explicit-implicit literature and the 2026 work on silenced bias make this decomposition increasingly important. A model that appears fair because it has learned which questions are socially dangerous is behaviorally different from a model whose relevant decision process is actually less dependent on the sensitive cue.

DecodingTrust and later benchmarks such as FLEX demonstrate that fairness conclusions can change under adversarial prompting. The model may behave acceptably under canonical benchmark instructions and fail when the user supplies bias-inducing context, role-play, emotional framing, or competing instructions.

A production evaluation should therefore include stress tests. The goal is not to create unrealistic jailbreak theater but to explore plausible ways a real system may be pushed away from its default policy. For an enterprise assistant, this can include retrieved documents containing biased assumptions. For a hiring system, it can include manager notes that express stereotypes. For a tutor, it can include a student who asks the model to confirm a prejudiced belief.

Robust fairness means that a system's desired behavior survives reasonable contextual perturbations.

Using one LLM to judge another is attractive because human evaluation is expensive. It is also dangerous because the judge becomes part of the measurement system. LLM judges can exhibit position bias, preferring an answer because it appears first or second; verbosity bias, favoring longer polished responses; self-preference, favoring outputs similar to their own style; and task-specific preferences unrelated to the target criterion.

Wang and colleagues showed that simply changing response order could strongly alter pairwise judgments. Later work systematically confirmed position bias across multiple judge models and tasks. Other studies show that superficial quality can distract judges from instruction following.

LLM judges are still useful. They should be calibrated. Pairwise order should be swapped. Multiple judges can be ensembled. Human validation should be stratified by demographic dimension and task type. The judge prompt and rubric should be versioned. Disagreement between judges should be treated as uncertainty rather than silently resolved by majority vote.

The 2026 ACL study by Asgari and colleagues is especially important because it formalizes a problem practitioners had already encountered informally: different bias metrics and different models often disagree in their rankings. Even metrics placed in the same broad family can behave largely independently. The correct interpretation is not that measurement is futile, but that each metric operationalizes a narrower behavioral construct than the phrase “bias score” suggests.

LLM-as-a-judge adds another layer of measurement uncertainty. Shi and colleagues evaluated position bias across fifteen judge models and more than 150,000 evaluation instances, finding that position effects were systematic and task-dependent rather than random noise. An evaluation pipeline that uses an LLM judge should therefore randomize candidate order, repeat comparisons, inspect disagreement, and preserve the judge version as part of the experimental configuration.

The goal is not to abandon benchmarks but to build a measurement portfolio: multiple task formulations, distributions rather than single samples, subgroup and intersectional analysis, refusal as an outcome, and independent validation of automated judges. Measurement disagreement is information, not an inconvenience to average away.

For LLMs, the distinction between represented knowledge, expressed behavior, and evaluated behavior is indispensable. Pretraining, post-training, context, and measurement form separate layers, and a trustworthy claim must state which layer it actually tested.

4. FROM DEBIASING TO FAIRNESS ENGINEERING: MITIGATION, RAG, TOOLS, AND AGENTS

A modern AI product is often a chain of models and policies rather than one model. A retriever decides which sources are visible; a reranker controls exposure; a generator interprets the evidence; an agent chooses tools; memory changes future context; an LLM judge may decide which answer or candidate wins. Recent RAG studies show that fairness alignment can be weakened by retrieved context, while work on LLM judges documents position and presentation biases. Mitigation therefore has to move from the metaphor of “cleaning bias out of weights” to the engineering of controlled information flows and bounded decisions.

The chapter begins with interventions at data, training, and inference, then widens the boundary to retrieval and finally to tool-using agents whose trajectories allocate evidence and action.

4.1 Change the data, objective, training signal, or inference policy

Mitigation fails when it begins with a technique rather than a mechanism. Reweighting is useful when optimization under-serves a subgroup represented in the data. Counterfactual augmentation is useful when a protected cue should not change the task answer. Preference optimization can change visible behavior but may also create refusals or evaluator-dependent artifacts. The correct question is which dependency or failure path the intervention is intended to change.

Data-centric mitigation is attractive because it intervenes before a model has learned the unwanted pattern. It can mean removing corrupted or dehumanizing content, increasing representation of underserved languages and groups, improving metadata, balancing labels, collecting counterexamples, or changing how examples are weighted. Yet “balancing” is not a universal recipe. Making every group equally frequent can distort the intended deployment distribution; removing every offensive sentence can also remove the evidence a model needs to recognize or discuss abuse. The design question is which data property caused the harmful behavior and which transformation changes that mechanism without erasing useful context.

Counterfactual data augmentation creates paired examples in which a social attribute is changed while the desired label or behavior remains stable. In a coreference task, pronouns and occupations can be varied; in resume screening, names or demographic cues can be changed while qualifications remain constant. The method is powerful when the counterfactual is semantically valid. It becomes misleading when identity is embedded in cultural or historical context that cannot be swapped without changing meaning. The safest use is therefore targeted: counterfactuals should encode a specific invariance hypothesis rather than a general idea that demographic categories are interchangeable.

The word debiasing suggests that a model contains a contaminant that can be removed. In practice, social information is distributed through representations and entangled with useful knowledge. A system must often reason about culture, demographics, and identity to answer correctly. The objective is therefore not to delete social information but to reduce harmful dependence while preserving legitimate context.

This leads to a layered risk-control view. Data curation can reduce upstream distortions. Training objectives can penalize disparities. Fine-tuning can teach desired behavior. Inference-time prompts can encourage reflection. Constrained decoding and filters can block harmful outputs. System-level auditing

can catch failures introduced by retrieval or tools. Human review can handle ambiguous high-stakes cases. No single layer should be expected to provide a universal guarantee.

But curation is normative. Removing identity-related language too aggressively can erase marginalized communities or dialects. Balancing counts does not guarantee balanced semantic coverage. Adding more examples from a group can reproduce the same stereotypes if the examples are narrow. Copyright and privacy constraints can also make targeted data collection difficult.

The best data practice therefore focuses on coverage and provenance rather than a simplistic notion of purity. Datasheets for datasets and data statements in NLP were early attempts to institutionalize this discipline.

Counterfactual training should also be evaluated beyond the augmented templates. Otherwise the model may learn a narrow pattern such as ignoring explicit gender words while remaining sensitive to indirect cues.

Supervised fine-tuning on carefully designed examples can substantially improve behavior. Parameter-efficient methods make it possible to experiment with specialized fairness adapters without retraining a foundation model. Preference optimization can target fairness directly by preferring responses that avoid harmful stereotypes or unjustified demographic dependence.

The risk is benchmark overfitting and policy masking. A model can learn to produce the expected neutral response on familiar prompts. The explicit-implicit findings show why post-training should be evaluated on indirect and downstream tasks, not only on the examples used to train it.

Preference optimization also inherits evaluator bias. If human or model raters prefer longer, more agreeable, more cautious, or ideologically familiar responses, those preferences can become part of the assistant's behavior. Fairness training therefore requires audits of the reward model and the preference dataset, not only the final assistant.

Constitutional AI makes high-level normative principles explicit and uses models to critique and revise responses according to those principles. This has two attractive properties for fairness. The rules are inspectable, and multiple stakeholders can potentially contribute to them. Collective Constitutional AI demonstrates one route for incorporating public input into model behavior.

The hard problem is conflict. "Be objective", "respect cultural differences", "avoid stereotypes", "do not discriminate", and "tell the truth" can point in different directions in edge cases. A constitution is therefore not a substitute for evaluation. It is a policy layer whose interpretation must itself be tested.

A promising direction is to make principle application more structured: identify which principle is active, expose conflicts, require evidence for exceptions, and escalate ambiguous cases rather than relying on a single opaque reward score.

Training objectives can make subgroup risk visible to optimization. Reweighting gives more influence to underrepresented or poorly performing groups. Group-distributionally robust optimization focuses on worst-group performance. Adversarial objectives can reduce the recoverability of protected attributes from representations. Constraint-based approaches can penalize disparities. These methods move fairness into the loss function, but they still require the engineer to choose the groups, metric, and tolerance. If the harm model is wrong, the optimization can become an efficient way to optimize the wrong fairness criterion.

Fine-tuning and preference optimization provide a higher-level intervention for foundation models. Curated examples can teach a model to avoid stereotypes, request missing information, use

person-specific evidence rather than demographic generalizations, or express uncertainty in high-stakes decisions. Preference data can reward responses that are both useful and fair. Constitutional or principle-based methods can make some of the desired behavior explicit rather than relying entirely on opaque rater preferences. Their main advantage is controllability: the organization can state the policy it is trying to induce. Their main weakness is that the principle still needs interpretation, and the evaluator that applies it can introduce its own blind spots.

The state of the art therefore favors layered evidence over one training intervention. After data curation, re-run behavioral tests. After fine-tuning, check whether calibration, factuality, safety, multilingual performance, or subgroup utility changed. Compare explicit and implicit tasks. Test refusals separately. If a model appears improved only on the training-aligned benchmark, assume overfitting until broader evaluation shows otherwise. Training-time mitigation is valuable, but it does not eliminate the need for deployment controls.

Mitigation becomes more intelligible when each intervention is attached to the mechanism it can actually change. The table moves from what the model sees to what the product does, making clear that later controls can sometimes address harms that retraining alone cannot reach.

Table 4. Mitigation layers and the mechanism each layer can change

Concept	Meaning in practice
Data and coverage	Change what the system observes; useful for missing representation but insufficient when labels or objectives encode the wrong construct.
Objective and weighting	Change which errors optimization treats as costly; can improve worst-group performance while affecting calibration or average utility.
Fine-tuning and preference	Change the conditional policy learned from instructions, demonstrations, or preferences; can also alter refusal and style.
Inference control	Change prompts, decoding, self-checks, abstention, or post-processing without retraining; effects can be fragile under context shift.
Retrieval and workflow	Change evidence exposure, routing, thresholds, tools, review, and escalation; often the most direct place to control product-level harm.
Monitoring	Detects regression, drift, and unintended trade-offs after the intervention enters a changing environment.

This layered view also exposes trade-offs. A data intervention can improve coverage while leaving a bad proxy untouched; an inference rule can stabilize deployment while remaining fragile under a new context; monitoring can detect regression but cannot repair it. Effective mitigation usually combines layers rather than declaring one technique universally “debiasing.”

Inference-time mitigation is increasingly important because many organizations do not control the base model and because fairness requirements can change faster than retraining cycles. Prompting can instruct a model to focus on task-relevant evidence, ignore protected characteristics where appropriate, generate a structured rationale, or perform a consistency check across counterfactual variants. Self-critique and deliberation can reduce some stereotypical behavior, but prompt-only methods are fragile. A user can alter the context, a long conversation can dilute the instruction, and a stronger system prompt can create over-refusal. Prompting is best treated as one control layer whose effectiveness must be tested under realistic context variation.

Representation editing, activation steering, pruning, and related mechanistic interventions try to change behavior inside the network. These techniques are scientifically valuable because they probe whether particular directions, features, heads, or neuron subsets participate in biased behavior. They are also risky as general-purpose remedies. Distributed representations rarely map cleanly to one social concept; suppressing a feature can harm legitimate knowledge or move the behavior to another pathway. The 2026 work on adaptive, reversible pruning illustrates a promising direction: interventions that can be conditioned on context rather than permanently deleting capacity. The more local the intervention, the more important it becomes to test transfer beyond the benchmark that identified the feature.

Prompting the model to consider fairness, take another perspective, or critique its first answer is inexpensive and works with closed APIs. Self-debiasing and perspective-taking methods have reported meaningful improvements on selected benchmarks. Large models often have enough meta-knowledge to identify an obvious stereotype when explicitly asked.

The limitation is robustness. Prompt-based mitigation can disappear under rephrasing, competing instructions, long contexts, or adversarial inputs. It can also create generic, overcautious responses that score well because they avoid committing to anything. A prompt that says "be unbiased" is therefore a useful layer, not a safety case.

These methods are scientifically valuable because they probe mechanism. They also face a deep conceptual problem: social concepts are distributed and context-dependent. A direction associated with gender may be needed for translation, medicine, or coreference. Removing it globally can damage useful capabilities while the model reconstructs the same information elsewhere.

The right standard is therefore behavioral transfer. An internal intervention is successful only if it reduces the relevant downstream harm across diverse tasks without unacceptable collateral damage.

Inference-time methods include calibrated prompting, reranking, constrained decoding, rejection sampling, generation-and-filter pipelines, and test-time adaptation. Their practical advantage is deployability: they can be applied without modifying model weights and can be tailored to the current domain.

The 2026 literature increasingly studies fairness at inference time because model weights may be inaccessible or frozen. These methods are particularly promising for systems that can detect a sensitive decision context and invoke stronger controls dynamically. The danger is brittleness and latency. A filter can create false positives for dialectal language. A regeneration loop can silently refuse more often for one group. A reranker can inherit the same bias as the generator.

Inference control should be evaluated as part of the system, including refusal disparity and quality-of-service effects.

Fairness interventions can interact with accuracy, calibration, robustness, privacy, safety, cultural competence, and cost. The trade-off is not always inevitable; some interventions improve both fairness and utility by removing spurious shortcuts. But trade-offs do occur, and they should be measured rather than asserted.

The SOBACO results are a useful caution: interventions that reduce measured social bias can damage culturally grounded commonsense. Privacy creates another tension because subgroup auditing often requires sensitive attributes that organizations may prefer not to collect. Safety filtering can conflict with fairness if identity terms are overblocked. Robustness can conflict with aggressive specialization to one benchmark.

A mitigation claim should therefore be multidimensional. The meaningful statement is not "we reduced bias by 30 percent" but "we reduced a defined disparity under these distributions while preserving task utility, calibration, refusal parity, and robustness within reported uncertainty." That is less marketable and more scientific.

System-level controls are often more interpretable. A hiring application can prevent the model from seeing protected fields, but also audit proxies and require evidence-based scoring. A clinical assistant can prohibit autonomous treatment decisions, expose uncertainty, and route borderline cases to specialists. A RAG system can diversify retrieval or require claims to cite evidence. An LLM judge can randomize answer order and separate dimensions. A recommender can constrain exposure. These controls do not make the model intrinsically fair, but they can bound the pathways through which model bias becomes material harm.

This is why "debiasing" is a misleading endpoint. A safety-critical organization does not claim that a system has been made failure-free; it builds barriers, monitors failure modes, and updates controls as evidence changes. Fairness engineering should adopt the same posture. We should prefer reversible interventions, explicit policies, monitoring, and defense in depth when the underlying construct is contested or the deployment environment evolves. The goal is not to erase every social correlation from a general-purpose model. It is to prevent unjustified correlations from controlling decisions and representations in the contexts where they matter.

Reversibility is an underrated property of mitigation. When evidence about a fairness problem is uncertain, a product-level rule or inference-time intervention can often be rolled back, scoped to one domain, or activated only when a detector fires. Permanent weight editing is harder to localize and harder to explain after the fact. This suggests a risk-based hierarchy: prefer the least invasive control that reliably blocks the harmful pathway, escalate toward retraining or model surgery when the behavior is pervasive, and keep regression tests that show why the stronger intervention remains necessary.

The strongest practical systems combine these controls with abstention. When a fairness-sensitive decision lacks sufficient task-relevant evidence, the correct behavior may be to ask for information, defer, or escalate rather than to fill the gap with demographic priors. Abstention is not free, because it can create unequal burden, but it gives the system an explicit alternative to confident guessing. Fairness evaluation should therefore measure who is asked for more information and who is routed away from automation.

Mitigation should end with a new evidence profile, not a congratulatory claim that the model has been debiased. The intervention must be re-tested for task utility, calibration, robustness, safety, privacy, and the groups that were not the original target of the change.

4.2 Retrieval changes whose evidence enters the answer

Retrieval-augmented generation changes the fairness boundary because it decides what evidence the model is allowed to see. A retriever may under-represent a language variety, over-rank authoritative but historically narrow sources, or expose skewed context that defeats alignment learned during post-training. Recent work has moved from demonstrating the problem to end-to-end interventions that jointly reason about retrieval exposure and generation disparity.

Retrieval-augmented generation is often introduced as a factuality technique: retrieve relevant documents, place them in context, and let the model answer from evidence. From a fairness perspective, this architecture creates a prior allocation step. Before the language model reasons about the question, the retriever decides which sources are visible. If one community, geographic region, language, school of

thought, or type of institution is underrepresented in the index or systematically ranked lower, the generator never receives the missing evidence. The final answer can look balanced in tone while being epistemically narrow.

The retriever itself can inherit representation and semantic biases from embeddings, query formulation, corpus composition, filtering, and ranking objectives. Relevance is not neutral when a query is ambiguous. A system optimized for click-like similarity may prefer popular or highly linked sources. An enterprise index may privilege documents produced by central teams because they are better formatted. A scientific assistant may overretrieve high-status journals and underretrieve negative results or regional venues. These are not always demographic harms, but they can become fairness problems when source visibility correlates with social groups or when the system is used to allocate credibility.

Modern AI applications rarely expose a raw LLM. They use retrieval, memory, tools, databases, policies, classifiers, and evaluators. Each component is optimized for a different objective. The retriever optimizes relevance. The generator optimizes helpfulness. A safety model optimizes policy compliance. A ranking layer optimizes engagement or business value. Fairness can be lost in the composition even when every component looks acceptable in isolation.

The correct unit of analysis is the user-to-outcome path. For a RAG assistant, that path includes query rewriting, embedding, vector retrieval, reranking, context construction, generation, citation selection, and filtering. For an agent, it includes planning, tool selection, memory updates, repeated decisions, and termination. Bias can enter through any transition.

RAG is usually motivated by factuality and domain knowledge. It also creates an evidence-allocation system. Only a small subset of available documents enters the context window, so retrieval determines which sources the model can use. If embeddings or ranking favor dominant linguistic styles, popular sources, or majority viewpoints, the generator receives a skewed evidence set.

Recent empirical work confirms that fairness issues can arise at both retrieval and generation stages. Wu and colleagues proposed a fairness evaluation framework for RAG and found disparities across demographic attributes. Hu and colleagues reported that RAG can undermine fairness alignment even when users attempt to censor biased external data; in their experiments, supposedly cleaned retrieval corpora did not guarantee unbiased outputs.

The engineering implication is direct. A fairness audit of the base LLM is not transferable to a RAG application. The retrieval corpus, embedding model, chunking strategy, reranker, and citation policy all require evaluation.

RAG systems must reconcile multiple sources: the user's statement, retrieved documents, system instructions, and model priors. Research in 2025 identified an "authority bias" in which LLMs can over-trust user-provided claims relative to retrieved factual evidence in conflict settings. The name is domain-specific, but the broader phenomenon is important: source provenance affects how the model weights evidence.

A system can therefore be biased not only about demographic groups but about epistemic authority. It may privilege information because of how it is framed, where it appears in the context, or whether it is attributed to a user, database, expert, or institution. In high-stakes domains, source-weighting behavior should be evaluated explicitly.

Fairness evaluation for RAG should be decomposed. Measure corpus coverage and source representation. Measure retrieval exposure by group or source category. Test queries that are semantically equivalent across dialects and languages. Examine whether demographic cues alter

retrieved evidence. Then hold retrieval fixed and test the generator, and conversely vary retrieval while holding the question fixed. Only after component tests should the end-to-end system be evaluated. This decomposition turns a vague claim that "RAG is biased" into a diagnosable set of mechanisms.

Source diversity should not be confused with forced viewpoint symmetry. A medical RAG system should not balance high-quality clinical guidance with unsupported claims merely to equalize perspectives. The relevant fairness question is whether comparable evidence receives comparable opportunity to be retrieved and whether source-quality criteria themselves systematically exclude valid knowledge from particular regions, languages, or institutions. Retrieval fairness must therefore be conditioned on epistemic quality. Otherwise a well-intended diversity constraint can reduce factual reliability and ultimately harm the users it was meant to protect.

Query rewriting deserves separate attention. Many RAG systems let an LLM transform the user's request into search queries. If the rewriter normalizes dialect, removes identity-relevant context, adds assumptions, or translates into a dominant language, it can change retrieval before the retriever sees the request. Query generation should therefore be logged and evaluated as its own model step. In multilingual systems, comparing both original-language retrieval and translated-query retrieval can reveal whether apparent generator bias is actually an information-access disparity introduced upstream.

The 2026 Fair RAG work marks a further step from diagnosis toward system-level intervention by jointly reasoning about retriever-side skew and generator-side disparity. That direction matters because it treats fairness as a property of the retrieval-generation pipeline rather than as a behavior that can be guaranteed by post-training the generator alone.

RAG fairness therefore belongs partly to information retrieval and partly to generation. The system must measure who is represented in the corpus, what the retriever exposes, how the generator uses authority, and whether citations make disagreement and missing evidence visible.

4.3 Agent trajectories create new forms of allocation

An agent is not only a language model that speaks more often. It is a policy that chooses actions, tools, sources, and stopping conditions across time. Fairness can fail because two matched users are sent down different paths even when their final text looks similar. Tool invocation, escalation, memory, and judge decisions become allocational events that need their own tests.

Tool-using systems introduce decisions that ordinary chatbot benchmarks do not observe. An agent may choose whether to search the web, query an internal database, ask a human, call a medical calculator, run code, open a ticket, or take no action. Tool selection is itself an allocation of computational attention and institutional resources. If the choice is systematically different for equivalent users or language varieties, downstream quality can diverge even when the final language model is unchanged.

Memory and personalization deepen the problem. Personalization can improve fairness by adapting to accessibility needs, language preference, or user expertise. It can also create persistent asymmetry if early interactions are misinterpreted. A mistaken inference about competence, risk tolerance, political preference, or communication style can become part of the context for future decisions. Because memory changes the input distribution, the system can create individualized feedback loops. Fairness testing must therefore include cold-start and long-history cases and provide mechanisms for users to inspect, correct, or reset consequential memory.

An agent that can call tools decides which external systems to consult. Tool selection can create unequal outcomes. A travel agent may use one search provider more often for certain regions. A medical

assistant may retrieve guidelines from high-resource countries and underuse local protocols. A coding agent may choose libraries whose documentation is abundant rather than appropriate. A financial agent may prefer data vendors with incomplete coverage of small markets.

These are not necessarily demographic biases, but they can become demographic through correlation. Geography, language, income, and institutional prestige often determine which tools have high-quality data. The agent's policy therefore inherits infrastructure inequality.

Evaluation should log not only final answers but tool-call trajectories. If two counterfactual users receive different tool choices, different search depth, or different willingness to verify information, the disparity may be invisible in the final text but consequential over time.

Persistent memory can improve user experience, but personalization is a new channel for differential treatment. A model may adapt to a user's name, writing style, profession, location, or prior preferences. Some adaptation is desirable. Harm arises when the system infers unsupported traits or uses remembered identity cues to change quality, opportunity, or safety treatment.

OpenAI's first-person fairness work is relevant here because it measures how user-name cues affect responses rather than only how AI systems make third-person decisions about subjects. The study reports low rates of harmful name-associated stereotypes in its tested setting, while also emphasizing limitations of the demographic categories and English-language scope. The methodological contribution is broader: personalized assistants need fairness tests where the user is the affected person.

Memory audits should therefore test whether equivalent requests receive equivalent quality when only identity-associated context changes, and whether personalization can be disabled or corrected when it becomes inappropriate.

A one-turn disparity may be small. An agent can act for dozens or hundreds of steps. Small preferences can compound through path dependence. If the agent is slightly less likely to verify a claim from one source, that claim can shape subsequent searches. If it ranks one candidate slightly lower, that candidate may never enter later deliberation. If it explores one region less thoroughly, later choices are made with less evidence.

The agent diagram follows a single request through retrieval, planning, tool selection, action, evaluation, and memory. The linear drawing is intentionally simple; in real systems some stages repeat, but the sequence is enough to expose where equivalent requests can diverge.




Figure 4. Agentic fairness concerns the path through the system, not only the final message.

Fairness evidence should therefore include trajectory traces, not only final messages. Two users may receive similar prose while one request was given poorer evidence, a weaker tool, more restrictive escalation, or a memory update that changes future treatment. Those differences are part of the product behavior.

This is analogous to feedback loops in recommender systems but occurs inside a single task trajectory. Agent fairness therefore needs sequential metrics. Researchers should measure path

divergence, tool-use disparity, evidence coverage, error recovery, and final outcomes, not only individual decisions.

The frontier here is still immature. Most agent benchmarks measure task success and safety rather than demographic or cultural fairness over trajectories. This is an obvious area for new research.

LLMs are increasingly used to grade answers, filter synthetic data, rank model outputs, evaluate code, score applications, moderate content, and even supervise other agents. This converts a language model from a generator into a decision institution.

The literature shows that LLM judges have systematic biases. Position bias can change pairwise rankings when response order is swapped. Verbosity and style can influence scores. Some judges prefer responses from the same model family or style. These effects matter because evaluation results can feed directly into model training. A biased judge can create a feedback loop in which the system learns to optimize judge preferences rather than the intended quality criterion.

A reliable judge harness should randomize order, use explicit rubrics, request evidence, separate dimensions, calibrate against humans, track disagreement, and avoid using the same judge as both optimizer and sole evaluator. For sensitive tasks, multiple independent judges and human escalation are preferable to a single authoritative score.

A modern AI lab may use LLMs to generate benchmark items, another LLM to filter them, the target model to answer, and a judge LLM to score the answer. This pipeline is scalable, but it can become epistemically self-referential. Shared model families can reproduce the same assumptions across item generation and scoring. Synthetic examples may reflect the generator's worldview. Judge preferences may reward the style used by the generator.

This does not invalidate automated evaluation. It means that diversity of measurement is essential. Human-created data, real usage samples, domain experts, external datasets, independent model families, and mechanistic probes should be combined when stakes are high. The more automated the evaluation loop becomes, the more important it is to preserve independent sources of error detection.

Multi-step agents can compound small differences. Imagine two semantically equivalent requests phrased in different dialects. One formulation leads to a low-confidence interpretation, which triggers a weaker search query, retrieves less useful evidence, selects a conservative tool path, and produces a lower-quality final recommendation. No individual step contains a dramatic disparity, but the trajectory amplifies the initial asymmetry. Evaluating only the final answer makes localization difficult. Trace-level logging is therefore essential: what was retrieved, which tool was selected, what uncertainty was recorded, what branch was taken, and which intermediate judgment caused the next action?

The right conceptual model is a path through a decision graph. Fairness concerns include path probability, information access, tool allocation, delays, refusals, escalation, and final outcomes. An audit can compare matched trajectories under counterfactual demographic cues, measure where paths diverge, and test whether divergences are justified by task-relevant information. Agentic fairness is therefore closer to process auditing than to static model benchmarking. As agents become more autonomous, provenance and traceability become prerequisites for understanding bias rather than mere observability conveniences.

Once a system can act, the unit of audit becomes the trajectory. Matched users should be compared not only on the final answer but on retrieved evidence, tools called, costs incurred, escalation decisions, judge outcomes, and what is stored in memory for the next interaction.

Once retrieval, tools, judges, and memory enter the architecture, fairness becomes a property of information flow and action trajectories. The base model remains important, but it is no longer a sufficient boundary for assurance.

5. BIAS IN THE WORLD: MULTIMODAL, SYNTHETIC, AND HIGH-STAKES SYSTEMS

When AI moves beyond text, the fairness problem does not merely repeat itself in another modality. Visual systems entangle demographic cues with backgrounds, clothing, lighting, and task framing. Speech systems interact with accent, disability, noise, and microphone quality. Synthetic data can improve coverage yet also create feedback loops in which rare linguistic or demographic patterns disappear before average performance visibly collapses. In high-stakes domains, these technical asymmetries become differences in opportunity, cost, treatment, visibility, or the ability to contest a decision.

The chapter follows fairness across modalities, examines the new risks created by synthetic-data loops, and then anchors the discussion in high-stakes workflows where model behavior becomes a material decision.

5.1 Fairness across vision, language, speech, and multimodal models

Multimodal systems introduce a methodological trap: demographic attributes are visually and acoustically entangled with other features. A fairness result can be driven by lighting, clothing, recording quality, task wording, or a model's general inability to solve the task. Recent face-verification work on multimodal LLMs illustrates the point sharply: a model with uniformly poor performance can look “fair” under a disparity measure because everyone is served badly. Fairness metrics must therefore be interpreted together with absolute utility.

Multimodal models combine representations that were historically audited separately. A vision-language model may identify people in an image, infer context, answer natural-language questions, and generate an evaluative description in one step. Bias can enter through the visual encoder, the language model, cross-modal alignment, the prompt, or the benchmark label. A model that performs well at object recognition can still associate demographic features with occupations or social roles. Conversely, language safety tuning may suppress a problematic verbal answer without changing how the visual representation contributes to an internal decision.

The shift from classification to generation changes representational harm. A classifier makes a bounded mistake; a generator can fill an entire scene with demographic assumptions. Text-to-image systems can underrepresent groups in high-status occupations, sexualize some identities, flatten cultural variation, or reproduce a narrow visual canon. These outputs matter because generated media increasingly enters advertising, education, design, entertainment, and future training data. Representation is not only about offense. It affects who is imagined as normal, expert, powerful, vulnerable, or absent.

Foundation models increasingly combine text, images, audio, video, code, and action. Multimodality changes the expression of bias because one modality can reveal attributes that another modality leaves implicit. A text-only system may not know a user's skin tone; a vision-language model may infer it immediately. An audio-language model may infer accent, age, or emotion from speech. A video model may combine appearance, voice, location, and behavior into a dense representation of identity.

This increases capability and privacy risk at the same time. It also complicates fairness because the model can condition on socially salient cues even when the user never states them. Evaluating only the

textual transcript of a multimodal interaction can therefore miss the variable that actually caused the decision.

The same system may exhibit different fairness properties by modality. A user could receive a neutral response when describing themselves in text and a different recommendation when uploading a photograph. A medical vision-language model can be strong on average yet show demographic performance gaps. Recent work evaluating large vision-language models across gender, skin tone, age, and race continues to find disparities despite strong general visual capabilities.

Text-to-image and image-to-image models create a representational problem at scale. When a user asks for "a successful entrepreneur", the model must instantiate visual attributes that were not specified. Which genders, ages, skin tones, clothing styles, body types, and cultural contexts appear? The model's distribution becomes a proposal for what the category looks like.

This is not identical to classification fairness. There may be no ground-truth demographic distribution for an abstract prompt. Evaluation therefore needs a reference policy. One approach compares generated distributions with real-world prevalence. Another uses parity as a normative baseline. A third asks whether prompts that should be identity-neutral nevertheless produce highly concentrated stereotypes. Each choice encodes a different notion of fairness.

Generative systems also create a tension between realism and correction. If historical professions are demographically skewed, a historically faithful image may reproduce that skew. If the model enforces contemporary parity, it may distort historical context. The appropriate behavior depends on the prompt and the user's intent, which suggests that context-sensitive fairness is more defensible than a global demographic quota.

Large vision-language models can classify, describe, compare, and rank people from images. This makes them tempting components in retail, hiring, security, healthcare, accessibility, and content moderation. It also recreates the old computer-vision risks in a more general interface.

The danger is semantic overreach. A model that can recognize visible objects may also confidently infer personality, profession, socioeconomic status, emotion, or intent from weak visual cues. Some of these inferences are scientifically dubious even before fairness is considered. Bias then enters not only through unequal error rates but through unjustified constructs.

A responsible evaluation should distinguish observable attributes from inferred traits. It should test subgroup accuracy, uncertainty, refusal behavior for unsupported inferences, and sensitivity to irrelevant visual details. The system should be rewarded for saying that an attribute cannot be inferred reliably from an image when that is the scientifically correct answer.

Audio models add another layer of socially meaningful variation. Accent and dialect encode geography, community, education, migration history, and sometimes race or class. Speech emotion recognition can be culturally variable. Voice-based age or gender inference can be inaccurate or inappropriate. Background noise and microphone quality correlate with socioeconomic conditions.

The earlier speech-recognition disparity literature therefore remains directly relevant to multimodal assistants. A voice interface that transcribes one population less accurately gives that population a noisier prompt to the language model. The language model may then respond less helpfully, making the end-to-end disparity larger than the original word-error gap.

The correct audit follows the chain: audio recognition quality, semantic interpretation, tool invocation, final answer quality, and recovery behavior when the input is uncertain. A system that recognizes its own uncertainty and asks for clarification can be fairer than one that confidently guesses.

Large vision-language models can also be used as judges of people and scenes. The 2025 study by Wu and colleagues evaluated mainstream LVLMS across demographic attributes and prompts and reported continuing disparities involving gender, skin tone, age, and race. The prompt mattered, reinforcing the broader lesson that multimodal fairness is conditional rather than a fixed property of the visual encoder. When such systems support surveillance, moderation, hiring, or accessibility, a small classification asymmetry can become a material decision.

Multimodal fairness requires cross-modal counterfactuals. The same text can be paired with different images; the same image can receive different demographic descriptions; the same semantic utterance can be spoken with different accents. These tests expose whether modality-specific cues change decisions that should be invariant. They also reveal a practical difficulty: many demographic attributes are not visually or acoustically well-defined, and attempts to infer them can themselves be harmful. Evaluation should focus on observed system behavior under controlled stimuli, not on pretending that complex identities can be read objectively from appearance.

Recent multimodal benchmarks reinforce the importance of prompt and task design. GRAS, presented at EACL 2026, shows that demographic-bias conclusions in vision-language models can change when semantically similar questions are formulated differently. This is the visual analogue of prompt sensitivity in LLM evaluation: the instrument can become part of the measured phenomenon.

The right comparison is therefore two-dimensional: how well does the system work, and how unevenly is that performance distributed? A uniformly bad system is not made acceptable by a small disparity score, and a highly accurate system can still impose concentrated harm on a subgroup.

5.2 Synthetic data can repair coverage or silently compress diversity

Synthetic data are often proposed as a solution to missing or imbalanced examples. They can be valuable, but they also create a new source of distributional power: the generator decides what counts as a plausible example. ACL 2026 work on self-consuming performative loops and new studies of bias inheritance show that recursive or model-generated data can change group preferences and downstream disparities. Emerging 2026 evidence on “fairness collapse” suggests that social bias may deteriorate before ordinary language-model metrics show severe model collapse.

Synthetic data appears to solve several fairness problems at once. It can create examples for rare cases, balance categories without collecting sensitive records, generate counterfactual pairs, and reduce dependence on copyrighted or private material. In simulation-heavy domains it can expose a system to edge cases that are expensive or dangerous to observe in the real world. These are genuine advantages. The mistake is to assume that generated data are independent of the biases of the generator.

A synthetic example inherits a distribution from the model and prompt that created it. If the generator has a narrow image of a profession, location, disability, or cultural practice, synthetic balancing may merely replicate that narrowness many times. A dataset can become numerically balanced while semantically homogenized. Diversity of labels is not the same thing as diversity of lived situations. For low-resource cultures, the model may generate examples that look plausible to outsiders but encode stereotypes or generic globalized norms.

The same mechanism can amplify bias. The generator's own stereotypes shape the synthetic distribution. If synthetic data are filtered by another LLM with similar preferences, those preferences can

be reinforced. If the synthetic corpus is then used to train the next model generation, the process becomes recursive.

Model-collapse research provides a broader warning. Shumailov and colleagues showed that recursively training generative models on model-generated data can cause distributions to lose information about the original data, especially the tails. The fairness connection is important even though model collapse is not itself a fairness theorem. Minority linguistic forms, rare cultural references, and atypical examples often live in distribution tails. Recursive synthetic training can therefore erase precisely the variation that fairness engineering tries to preserve.

Generative AI does not only learn from culture; it increasingly produces culture. Generated text enters websites, forums, marketing copy, documentation, educational resources, and future training corpora. Generated images populate search engines and stock libraries. Model outputs can therefore change the distribution from which later models learn.

This creates a cultural feedback loop. If current models overrepresent some identities or styles, large-scale generated content can make those patterns more common in future data. The loop is not guaranteed to amplify bias, because filtering and human editing intervene, but the possibility changes the governance problem. Training data provenance now needs to distinguish human-originated from model-originated content when feasible.

A long-term fairness strategy should treat diversity preservation as an information problem. The objective is not only to reduce harmful correlations today but to avoid collapsing cultural and linguistic variety tomorrow.

The effect is not inevitable. Synthetic data can be used carefully as a controlled intervention. A strong pipeline starts with a coverage gap identified in real evaluation, generates candidates under diverse prompts or multiple models, validates them with domain experts or affected communities, preserves provenance, and keeps synthetic and human-origin data distinguishable. The objective is not to replace reality with a model's imagination but to augment missing regions while retaining a reference to real distributions.

Generative culture creates a longer feedback loop. Recommendation systems promote generated content; creators imitate styles that receive exposure; generated images and text become training data; future models then reproduce the compressed aesthetic. Fairness in this setting is partly about cultural option value: whether less common languages, styles, narratives, and viewpoints remain available to be learned and generated. This is difficult to reduce to a protected-attribute metric, but it belongs to the same family of concerns about who and what remains visible in an increasingly machine-mediated information environment.

Evaluation of synthetic augmentation should therefore compare more than label frequencies. Measure lexical, semantic, visual, and behavioral diversity; test whether rare real examples remain nearest neighbors of anything in the augmented set; and check whether downstream models improve on held-out human-origin cases rather than only on synthetic-style tests. Provenance is essential because later investigators need to know whether an apparent population pattern came from people or from a generator. Once synthetic and human-origin data are irreversibly mixed, diagnosing homogenization becomes much harder.

There is also an economic feedback loop. If synthetic content is cheaper to create and easier to optimize for platform algorithms, it can crowd out expensive human-produced material from search and recommendation. The resulting training corpus is then shaped not just by model bias but by production economics. This is a fairness concern when communities with fewer resources to optimize for

machine-mediated visibility become progressively less represented. Cultural diversity can decline without any explicit discriminatory rule, simply because the cheapest content dominates the channels from which future models learn.

Synthetic balancing is most credible when it is used to create tests before it is used to create training data. Generated counterfactuals and edge cases can cheaply expand the evaluation surface and reveal where real collection is needed. If a synthetic case exposes a disparity, the next step can be targeted human-origin data collection or expert validation. This sequencing reduces the risk of training a model to satisfy the synthetic generator's worldview while still capturing the economic advantage of automatic test generation.

Wang and colleagues introduced the Self-Consuming Performative Loop at ACL 2026 to study a realistic dynamic in which deployed models influence the data on which future models are trained. Their controlled experiments show that different forms of bias can move in different directions inside the same loop, which is another reason not to rely on a single headline metric. Work on bias inheritance in LLM-based data augmentation likewise shows that model-generated training data can propagate and amplify bias into downstream tasks.

A particularly concerning 2026 preprint describes a “fairness collapse” regime in which demographic bias degrades across recursively synthetic training generations before standard language-modeling metrics show severe performance collapse. Because this result is newer and not yet as settled as the Nature model-collapse literature, it should be treated as an emerging warning rather than a universal law. Its engineering implication is nevertheless plausible: monitor subgroup and diversity indicators directly rather than expecting average perplexity or accuracy to reveal social degradation.

Synthetic data should be treated as an intervention with provenance, diversity diagnostics, and longitudinal monitoring. The central risk is not only generic model collapse but selective disappearance: minority forms of language or behavior can be compressed before headline quality reveals the problem.

5.3 High-stakes workflows turn model behavior into material consequences

A high-stakes system cannot be evaluated by transferring a generic fairness score from the base model. The same output can have very different meaning depending on whether it is advice, a recommendation, a preliminary screen, or a binding decision. Engineers must map the full workflow: what the AI proposes, what the human sees, what can be appealed, which errors are reversible, and who bears the cost of uncertainty.

Employment systems illustrate how quickly a general-purpose language model can become an allocational system. An LLM may parse resumes, generate candidate summaries, infer skills, rank applicants, draft interview questions, or advise a hiring manager. Each function has a different fairness surface. A summary can omit achievements, a ranking can change exposure, and an interview assistant can ask different follow-up questions. Studies in 2025 on resume matching and allocational fairness reported demographic disparities under realistic matching settings, reinforcing the need to evaluate the complete hiring workflow rather than only stereotypical prompts.

The strongest design principle in hiring is evidentiary grounding. The system should score job-relevant evidence, not demographic proxies or unsupported personality inferences. Counterfactual resume tests can reveal sensitivity to names, gender markers, educational prestige, or geography. But the audit should also examine which candidates receive detailed analysis, which are escalated, and how model uncertainty

reaches the human reviewer. Human oversight is meaningful only if the reviewer sees enough evidence to detect an error rather than merely confirming a confident machine recommendation.

Recent studies of LLM-based job-resume matching find a mixed picture. Newer models can show reduced explicit gender or race effects on some controlled tests, yet educational background and other proxies can remain influential. Work on allocational fairness in hiring pipelines reports that retrieval stages can introduce significant disparities even when the generator appears comparatively neutral. Other human-versus-LLM comparisons find that LLM group differences are not always larger than human differences in the tested setting.

The correct conclusion is not that LLM hiring is fair or unfair in general. It is that the comparison baseline matters. Human parity is not proof of justice, but it is relevant when evaluating replacement decisions. A deployment should be tested end to end, with counterfactual resumes, intersectional groups, realistic job families, retrieval behavior, rank stability, and final selection rates.

Educational AI can recommend courses, grade writing, provide tutoring, identify students at risk, and generate feedback. Bias can appear through language proficiency, cultural references, disability accommodations, school prestige, socioeconomic proxies, or expectations about ability.

Generative tutors create an additional quality-of-service dimension. Two students can receive equally non-discriminatory content but different depth, patience, encouragement, or intellectual challenge. If the model simplifies more aggressively for one dialect or inferred demographic group, the harm may not appear in a toxicity benchmark. Fairness therefore includes pedagogical ambition.

Grading systems require special caution with LLM-as-a-judge. Position and verbosity bias are obvious risks, but cultural and stylistic norms matter too. A polished style should not dominate the score when the rubric targets reasoning. Independent rubrics, blinded evaluation where possible, calibration against human graders, and appeals are essential.

Healthcare combines unequal data availability, high stakes, causal complexity, and legitimate demographic differences. A fairness rule that blindly removes demographic information can be medically harmful. At the same time, proxies such as cost, access, and documentation intensity can reproduce structural inequalities.

Foundation models add free-form language and multimodal data. A medical LLM may summarize records, answer patient questions, draft notes, retrieve guidelines, or interpret images. Bias can arise if the record itself is incomplete for some patients, if retrieval favors guidelines from one population, if the model's uncertainty differs by dialect, or if multimodal performance varies by skin tone or sex.

The engineering target should be clinically relevant performance stratified by subgroup, plus calibration, uncertainty, and error severity. When a demographic variable is medically relevant, its use should be explicit and evidence-based rather than learned implicitly through unreliable proxies. Causal reasoning is especially valuable here because it forces the team to ask why an attribute should affect the outcome.

Financial systems already have mature fairness concerns around credit, pricing, fraud, and insurance. LLMs enter these workflows through document analysis, customer support, underwriting assistance, and explanation. They can also become informal decision aids that influence a human officer even when the final legal decision is nominally human.

A conversational interface can hide the exact variables influencing a recommendation. The model may summarize a file in a way that emphasizes neighborhood, education, occupation, or language style.

Traditional fairness controls applied to the final score may therefore miss bias introduced in upstream narrative compression.

For these systems, traceability matters. The organization should know which source facts were used, which were omitted, and how the LLM's intermediate output affected the final decision. A prose explanation is not enough if it cannot be connected to evidence.

The same numerical disparity can have very different significance across domains because the decision, reversibility, evidence standard, and severity of error differ. The table frames high stakes as a workflow property rather than a label attached to an industry.

Table 5. What “high stakes” changes in the evaluation question

Concept	Meaning in practice
Employment	The system may rank, summarize evidence, generate interview questions, or influence human attention; opportunity can be allocated at multiple stages.
Education	Quality includes grading, tutoring usefulness, accessibility, language support, and the ability to contest an automated judgment.
Healthcare	Clinical relevance of demographic factors, proxy validity, uncertainty, access patterns, escalation, and severity of missed cases dominate generic parity.
Finance and insurance	Calibration, pricing, proxy variables, documentation, strategic behavior, and legal constraints interact with allocational fairness.
Public systems	Procedural rights, explanation, appeal, record correction, and asymmetric consequences can matter as much as predictive accuracy.
Information access	Ranking and moderation determine visibility, exposure, and who has practical access to knowledge or an audience.

What should be measured follows from that workflow. In healthcare, missed escalation may dominate; in employment, ranking and attention allocation may matter before any formal decision; in public administration, explanation and appeal can be part of fairness itself. Domain context is therefore part of the technical specification.

Finance and insurance add calibration, regulation, and strategic behavior. Credit decisions often use structured predictive models, but LLMs increasingly participate in document analysis, customer interaction, explanation, fraud investigation, and underwriting workflows. A seemingly harmless language step can alter the evidence presented to a downstream risk engine. Fairness analysis must therefore trace both model decisions and information transformations. The common principle across these domains is that a general-purpose model becomes high-stakes through workflow integration. Risk classification belongs to the use, not to the vocabulary of the model.

Across these domains, the strongest common control is to limit unsupported inference. General-purpose models are good at filling gaps, which is useful in writing and dangerous in adjudication. A hiring system should not infer motivation from a name; a clinical system should not infer symptoms from ethnicity; a credit workflow should not invent stability from writing style. Structured evidence schemas, source citation, and explicit “unknown” states make it harder for fluent completion to masquerade as observed fact. The more consequential the decision, the more the system should distinguish supplied evidence from model-generated interpretation.

Legal and public-sector uses expose the difference between advisory language and institutional authority. An LLM that summarizes case law is not equivalent to a system that recommends bail, prioritizes inspections, assesses benefit eligibility, or drafts an official decision. The latter uses can affect rights and should be evaluated for due process, explainability, access to correction, and differential error. Hallucination is itself a fairness issue when some users have less ability to challenge an incorrect citation or when language barriers make verification harder.

Public administration also creates asymmetry through access. A conversational interface can make bureaucracy easier to navigate, but only if it works across the languages, literacy levels, disabilities, and edge cases of the population it serves. A system that responds fluently to standard formulations while failing on nonstandard language can create a new digital gate. The fair design response may involve multilingual retrieval, human escalation, accessible interfaces, and explicit uncertainty rather than attempting to solve the problem entirely in model weights.

Moderation systems classify harmful content, enforce policies, and rank visibility. They are vulnerable to dialect bias, context loss, reclaimed slurs, satire, political disagreement, and unequal rates of false positive removal. Generative moderation agents can explain decisions, but explanation does not guarantee consistent policy application.

LLM moderation also creates a recursive problem: an LLM may judge whether another LLM's output is harmful. If the judge is more likely to flag language associated with particular communities, safety infrastructure itself becomes a source of bias.

A competent moderation evaluation therefore requires group-conditioned false positive and false negative analysis, contextual examples, multilingual data, appeals, and continuous monitoring of policy changes. The cost of over-enforcement should be measured alongside the cost of under-enforcement.

Recommendation systems shape attention, and attention is a scarce resource. LLMs increasingly mediate search, news summaries, shopping, travel, education, and professional discovery. When the model answers directly instead of presenting a list, the allocation of attention becomes less visible.

Source selection is therefore a fairness issue. Which publishers are cited? Which businesses are recommended? Which countries are represented as examples? Which scientific communities are treated as authoritative? A conversational answer can appear comprehensive while systematically omitting lower-resource sources.

This is one reason RAG fairness and epistemic diversity deserve more attention. The problem is not merely whether the final sentence contains a stereotype. It is whether the system allocates visibility and credibility equitably enough for its purpose.

High-stakes deployments often claim safety because a human remains "in the loop". Human oversight can help, but it can also create automation bias. People may over-trust a confident model, use it as an anchor, or review only exceptional cases. If human reviewers are under time pressure, the AI's recommendation can become the de facto decision.

The right question is not whether a human signs the final form. It is whether the human has time, information, authority, and incentives to disagree. Oversight should be designed as a functional control, not a legal label.

The concept of disparate burden is useful here. Even when final outcomes are statistically similar, one group may need more retries, longer explanations, additional documents, more appeals, or more human review to obtain the same service. These frictions are often absent from fairness dashboards because they occur before the final decision. Logging interaction cost, latency, number of clarifications, escalation

frequency, and successful completion can reveal exclusion that outcome parity misses. For conversational public services, burden metrics may be as important as classical error rates.

Reversibility and remedy change the acceptable error budget. A mistaken recommendation in a news feed may be corrected by the next interaction; an erroneous denial of a benefit or an automatically generated legal filing can have consequences that persist long after the model output. High-stakes fairness therefore depends not only on probability of error but on the cost and reversibility of error. Systems with irreversible effects need stronger evidence, more conservative automation, and clearer routes to remedy than systems where users can easily recover.

In consequential applications, fairness engineering must remain close to the workflow and to the people affected by it. The objective is not a model that passes an abstract benchmark, but a process in which errors are bounded, uncertainty is handled deliberately, and decisions remain reviewable and contestable.

Multimodal and high-stakes systems show why fairness cannot be divorced from absolute utility, workflow, and long-term data generation. A socially acceptable system must work well enough, distribute that quality responsibly, and remain correctable when it fails.

6. CONTINUOUS FAIRNESS ASSURANCE AND WHAT COMES NEXT

A fairness audit is a photograph of a moving system. Models change, retrieval corpora change, prompts change, vendors update APIs, user populations shift, and the outputs of yesterday's systems become tomorrow's training data. The engineering problem therefore resembles reliability and security more than one-time certification. The system needs specifications, test harnesses, change control, monitoring, incident response, and explicit policy ownership. From that operational foundation, the chapter turns to the research frontier: causal mechanisms, distribution shift, pluralistic alignment, executable constraints, and AI-assisted auditing.

The chapter first turns fairness into continuous assurance, then develops causal and distribution-aware perspectives, and closes with pluralistic policy, executable controls, and AI-assisted auditing.

6.1 Fairness evidence must survive system change

The purpose of continuous assurance is not bureaucratic repetition. It is to keep the evidential claim synchronized with the deployed system. If the model version, retrieval corpus, prompt policy, threshold, user population, or legal context changes, the old fairness evidence may no longer support the new product. This is why NIST's generative-AI profile emphasizes documented measurement and why production teams increasingly treat evaluation suites as regression tests rather than one-off reports.

A production fairness program needs the equivalent of a regression suite. The suite should begin with a versioned harm specification rather than a pile of benchmark scripts. For each important harm, it should record the affected population, expected behavior, unacceptable failure, test distribution, and evidence threshold. Tests can then combine public benchmarks with product-specific cases, counterfactual pairs, naturalistic samples, multilingual prompts, long conversations, adversarial contexts, and subgroup analysis. The objective is coverage of failure mechanisms rather than a large number of disconnected scores.

The harness should evaluate the full system configuration. A run records the base model and version, system prompt, retrieval corpus and embedding model, reranker, tool permissions, safety filters, memory policy, decoding parameters, and judge models. This level of detail may appear excessive until a vendor silently updates a model or a product team changes a retrieval cutoff and a previously stable disparity reappears. Without configuration provenance, fairness regressions become anecdotes. With provenance, they become debuggable software behavior.

A fairness audit performed before release is necessary and insufficient. Models change. Retrieval corpora change. user populations change. Safety policies change. Prompt templates are edited. External APIs update. A system that passed an evaluation in January can behave differently in June without any deliberate fairness regression.

Continuous assurance treats fairness like reliability and security. The organization defines properties, instruments the system, monitors drift, investigates incidents, and re-tests after changes. The relevant question is not "did we certify the model?" but "do we have evidence that the deployed system remains within acceptable bounds?"

This operational framing aligns with NIST's AI Risk Management Framework and its Generative AI Profile, which emphasize mapping risks, measuring them, managing them, and documenting outcomes over the lifecycle. The EU AI Act similarly requires, for covered high-risk systems, data-governance practices that examine possible bias and measures to detect, prevent, and mitigate it, while also addressing biased feedback loops in continuously learning systems.

The harness should combine naturalistic production-like examples with controlled counterfactual pairs. It should include ambiguous cases, adversarial variants, multilingual examples where relevant, and intersectional groups. For generative systems, it should sample multiple responses. For RAG and agents, it should log intermediate components. For LLM judges, it should randomize answer order and calibrate against humans.

The output should be a profile, not a pass-fail badge. A mature dashboard tracks utility, subgroup quality, refusal disparity, counterfactual sensitivity, representational harm, robustness, calibration, uncertainty, and drift. Which dimensions matter most depends on the product.

Datasheets, data statements, and model cards are valuable because teams forget why data and models were created. A system may be reused years later by people who did not participate in its development. Documentation preserves intended use, exclusions, subgroup results, known limitations, and provenance.

For LLM systems, the documentation unit should be larger than the base model. A system card should describe the model version, system prompts, retrieval sources, embedding models, rerankers, safety filters, memory policy, tool permissions, evaluation suites, and known failure modes. Otherwise the most important fairness choices remain undocumented because they sit outside the model weights.

Documentation should also record disagreement. If two fairness metrics produce opposite model rankings, that fact belongs in the report. Hiding it behind an average score removes exactly the uncertainty decision makers need to see.

A team that builds a system is poorly positioned to discover every assumption in it. Independent red teams, external researchers, domain experts, and affected communities provide different failure hypotheses. This is especially important for culturally specific harms that the development team may not recognize.

Independence is also technical. If the same model family generates test cases, evaluates outputs, and filters the dataset, shared blind spots can become invisible. Cross-model evaluation, human review, and external data sources increase error diversity.

The goal is not to eliminate automation from auditing. The frontier is likely the opposite: AI-assisted auditors can generate large numbers of counterfactual and adversarial tests. But automated auditing should be designed to create independent evidence, not self-consistency.

Sampling from real traffic is the other half of the harness. Laboratory prompts can only cover known concerns. Production monitoring can reveal emerging language, new user populations, interaction patterns, and cases that were not represented during development. Privacy-preserving sampling and careful human review can estimate quality-of-service disparities, refusal patterns, escalation rates, and complaint themes. The purpose is not to infer sensitive attributes recklessly. Where group information is unavailable or legally constrained, audits can use opt-in studies, synthetic counterfactuals, targeted panels, or other defensible measurement designs.

The harness should also distinguish blocking tests from diagnostic tests. A blocking test encodes a failure that is understood well enough to prevent release when it regresses. A diagnostic test explores an

uncertain phenomenon and should trigger investigation rather than automatic rejection. Mixing the two makes organizations either overly conservative or too willing to ignore alarms. Versioned test intent is therefore part of the evidence: the team should know which findings are contractual requirements, which are research probes, and which remain unresolved hypotheses.

Production assurance requires more than storing a benchmark score. The table lists the evidence needed to reconstruct what was tested, under which configuration, against which harm specification, and who owned the policy choice.

Table 6. Evidence a production fairness program should preserve

Concept	Meaning in practice
Specification	The harm, population, construct, metric, tolerance, and intended intervention are versioned together.
Configuration	Model, prompts, retrieval stack, tools, thresholds, filters, memory, and judge versions are reproducible.
Tests	Public benchmarks are combined with product-specific counterfactual, naturalistic, multilingual, stress, and subgroup evaluations.
Production signals	Refusals, escalation, quality, exposure, incidents, appeals, and sampled outcomes are monitored where lawful and appropriate.
Change history	Material model or product changes are linked to the evidence that was re-run, superseded, or invalidated.
Decision ownership	Normative thresholds and policy choices are attributed to responsible institutions rather than presented as facts discovered by a model.

The point is reproducibility over time. When a model, retriever, prompt, threshold, or user population changes, the organization must be able to identify which prior claims remain supported and which must be re-tested. Without that chain, “fair at launch” becomes an uncheckable historical statement.

The assurance loop represents fairness as a maintained engineering claim. Specification determines what to observe, diagnosis links evidence to plausible mechanisms, intervention changes the system, and re-testing checks both the target harm and regressions.



Figure 5. Continuous assurance treats fairness evidence as something that must be renewed after change.

Monitoring closes the loop because deployment generates new conditions that the original evaluation may not cover. A fairness program is mature when it can state not only that a test once passed, but why the evidence is still applicable after the system and its environment have changed.

Coverage should be reviewed periodically as a first-class metric. A suite can remain green while the product grows into new languages, regions, or workflows that the suite never tests. The harness should therefore report not only failures but untested surface area. In mature programs, “we do not yet have

evidence for this population" is a valid and useful result, because it prevents a narrow evaluation from being silently generalized into an unsupported product-wide claim.

Governance is often presented as a layer above engineering, but in mature AI systems it changes the architecture. The EU AI Act, NIST AI RMF, sectoral regulation, procurement requirements, and organizational policies all create expectations about risk management, data governance, testing, logging, human oversight, and post-deployment monitoring. For high-risk systems, the EU framework explicitly requires attention to representativeness, data gaps, and biases likely to affect health, safety, fundamental rights, or discrimination, together with measures to detect, prevent, and mitigate those biases. The 2026 consolidated text also includes conditions under which special-category personal data may exceptionally be processed for bias detection and correction, illustrating the tension between privacy and fairness measurement.

NIST's Generative AI Profile is valuable for engineers because it frames harmful bias and homogenization as risks to be measured across system outputs and downstream processes. It recommends use-case-appropriate benchmarks, subgroup and intersectional evaluation, documentation of benchmark assumptions, field testing, counterfactual prompts, and attention to latent bias in text, images, audio, and embeddings. The important point is not compliance with a particular checklist. It is that governance is converging on the same system-level view that the research literature has reached.

The EU AI Act is not a fairness textbook, but it reinforces lifecycle engineering. The consolidated text as of 27 July 2026 requires data-governance practices for covered high-risk systems that include examining possible biases likely to affect health, safety, fundamental rights, or unlawful discrimination, along with appropriate mitigation. It also requires data to be relevant and sufficiently representative for intended purpose and addresses the risk of biased outputs entering future inputs through feedback loops.

The most important engineering consequence is that claims about fairness must be connected to intended use. A generic model benchmark cannot replace application-specific evaluation. Documentation, monitoring, risk management, and data governance become part of the technical architecture rather than paperwork added after the model is built.

The Council of Europe's Framework Convention on AI adds a broader human-rights lifecycle perspective. The European Union ratified the Convention in May 2026. For engineers, the practical signal is convergence: leading governance frameworks increasingly treat trustworthy AI as a process of risk identification, evidence, mitigation, and accountability rather than a one-time property of a model.

Bias incidents should be handled like security and reliability incidents. The organization needs a way for users and employees to report failures, preserve the prompt and system state, reproduce the behavior, estimate scope, identify the responsible component, deploy a mitigation, and monitor for regressions.

The investigation should avoid the reflex to patch the single example. If a dialectal phrase triggered a worse recommendation, the question is whether a broader linguistic class is affected. If a guardrail falsely blocked one identity term, the team should test related terms and languages. If a judge favored a longer answer, the entire evaluation pipeline may need recalibration.

A fairness incident is often evidence about a hidden distribution, not an isolated bug report.

Fairness measurement often requires demographic attributes, while privacy practice encourages minimizing sensitive data. The tension is real. If an organization never measures subgroup membership, it may be unable to detect disparate impact. If it collects detailed attributes carelessly, it creates privacy and security risk.

The solution is governance rather than denial. Sensitive attributes can be collected under strict purpose limitation, access control, retention limits, secure computation, or privacy-preserving aggregation when legally and ethically appropriate. Some audits can use voluntary self-identification or trusted third parties. The right architecture depends on jurisdiction and risk.

The broader lesson is that trustworthy AI properties interact. Privacy, fairness, robustness, explainability, and safety should not be optimized in isolation.

The operational objective is accountable evidence. An organization should be able to state what it knows about fairness, what it does not know, which measurements support the claim, which populations were not covered, what controls are active, and who can change them. This is more defensible than a universal "fairness certified" label. It also makes future updates safer because the evidence can be re-evaluated when models, data, law, or deployment conditions change.

Good governance also preserves decision ownership. If a fairness threshold is chosen because of law, clinical policy, labor rules, or an organizational value judgment, that rationale should be attributed to the responsible institution rather than presented as a property discovered by the model. This separation matters for accountability and for later revision. Models and benchmarks can inform the decision, but they should not become a convenient way to hide normative authority inside an apparently technical artifact.

Procurement and vendor management are part of the same evidence system. Organizations increasingly assemble products from external foundation models, embedding services, moderation APIs, and data providers. Contracts and technical interfaces should preserve version information, notice of material changes where possible, evaluation access, incident communication, and enough documentation to assess subgroup behavior. A fairness program that assumes full control of every model component will fail in precisely the multi-vendor architectures that dominate real deployment.

NIST's Generative AI Profile makes this operational stance explicit by asking organizations to evaluate and document fairness and bias using use-case-appropriate benchmarks, subgroup performance, field testing, and red-teaming while documenting benchmark limitations. This is not a legal fairness definition; it is an assurance discipline. The emphasis on documented assumptions is valuable because the validity of a fairness claim depends on the use case and population for which the evidence was produced.

A trustworthy organization should therefore be able to say what it currently knows, which populations were tested, which assumptions remain unresolved, what changed since the last evaluation, and which evidence was invalidated. That is a stronger claim than a permanent "fairness certified" label.

6.2 Causal, dynamic, and distribution-aware fairness

Static subgroup parity is descriptive. It rarely tells us which mechanism created the disparity or whether the result will survive a new geography, language, economic condition, or product policy. Causal and distribution-aware approaches ask harder questions: which pathways are legitimate, which are objectionable, how sensitive is the conclusion to assumptions, and what happens when the deployment environment moves outside the validated region?

The most defensible near-future direction is the abandonment of universal bias scores. The 2026 metric-agreement results make the scientific problem explicit: different measures can rank the same models differently because they operationalize different behaviors. A mature evaluation should therefore resemble a reliability profile rather than a leaderboard.

The profile should be conditional. It should say how the system behaves across tasks, languages, prompts, demographic intersections, decision types, and distribution shifts. It should distinguish representational, allocational, epistemic, and quality-of-service harms. It should report uncertainty and disagreement. The absence of one number may look less convenient, but it is closer to the structure of the problem.

Correlation-based tests are good at finding disparities and poor at explaining them. Causal fairness asks which pathways connect a sensitive attribute to an outcome. This matters because some pathways are legitimate and others are not. Geographic location can legitimately affect a delivery estimate but may be an unacceptable proxy for ethnicity in a credit decision. Sex can be medically relevant in some clinical settings and irrelevant in others.

Recent LLM work on causality-guided debiasing and socially biased causal reasoning points toward a more mechanistic evaluation. Instead of asking only whether the final output differs, researchers can test whether the model relies on an objectionable causal chain. This is especially promising for agents, where intermediate reasoning, retrieval, and tool use create observable pathways.

The challenge is that causal models are assumptions, not discovered truth. Human experts must still define which paths are permissible. But causal structure makes those assumptions explicit and testable.

A model can be fair on today's data and unfair tomorrow. Population composition changes, social norms change, new slang appears, policy changes alter labels, and new markets introduce different correlations. Classical research already shows that fairness guarantees do not automatically transfer under distribution shift.

For foundation models, shift is constant. User prompts are open-ended. Model versions update. RAG corpora change daily. The evaluation target should therefore include fairness robustness: how rapidly do disparities grow as the environment departs from the benchmark distribution?

A promising research direction is certified or bounded fairness under families of shifts, combined with adaptive monitoring that detects when deployment leaves the validated region. The goal is not a timeless guarantee but a measurable operating envelope.

Static parity can be misleading when decisions change future opportunities. Long-term fairness models the feedback between the system and its environment. This direction should become more important as agents and recommender-like assistants shape user behavior over months or years.

A career assistant that systematically recommends more ambitious paths to one group can alter future qualifications. A learning tutor that adapts difficulty differently can change future performance. A content recommender that underexposes minority creators can change the supply of content. These systems require simulation, longitudinal measurement, and sometimes causal experimentation.

The research frontier is to connect short-term behavioral audits with long-term social dynamics without pretending that a simulation is society. Models should be used to test hypotheses and identify plausible feedback mechanisms, not to generate false precision about future social equilibria.

Many fairness failures are aggravated by false confidence. If the model is uncertain about dialect, diagnosis, or eligibility but produces a confident judgment anyway, uncertainty is converted into unequal error. A better system can abstain, request more information, or escalate to a human.

Uncertainty-aware fairness therefore asks whether uncertainty is calibrated across groups and whether abstention policies create equal access to useful service. A system that abstains more often for a

low-resource language may reduce harmful errors but still deliver worse service. The right design may require investing in better models or alternative workflows rather than declaring the abstention policy fair.

This connects fairness to selective prediction and human-AI collaboration. The future system should know not only what it predicts but where its evidence is weak and who bears the cost of uncertainty.

These directions converge on a shift from static properties to conditional guarantees. Rather than claiming that a model is fair, a system might claim that within defined populations and operating conditions, certain disparities remain below monitored thresholds; outside those conditions, uncertainty or shift detectors trigger a different policy. This resembles assurance in safety-critical engineering. The guarantee is bounded, evidence-based, and explicitly invalidated when the environment changes.

The frontier is not a more sophisticated way to produce one universal score. It is a set of models that make mechanisms, uncertainty, time, and operating boundaries explicit enough to support falsifiable claims.

6.3 Pluralistic alignment and executable controls

Some disagreements about AI behavior are not noise that a larger preference dataset will remove. Cultures, institutions, and users can legitimately disagree about speech, autonomy, paternalism, political neutrality, or educational practice. A mature architecture must distinguish non-negotiable constraints from policy choices that can vary, and it must make those choices inspectable. This is where explicit constitutions, policy representations, and neurosymbolic controls become interesting engineering tools rather than philosophical decoration.

Pluralistic alignment starts from the observation that some value disagreements are real rather than measurement noise. People and cultures differ about acceptable speech, distributive priorities, political neutrality, educational norms, and trade-offs between autonomy and protection. A single global reward model can compress that disagreement into the preferences of the annotator population or model provider. A pluralistic system can instead separate non-negotiable constraints, such as legal or safety requirements, from settings in which users or institutions may choose among legitimate policies. The technical challenge is to preserve choice without allowing personalization to become a mechanism for discriminatory treatment.

Contestability is the institutional counterpart. When an AI system affects a consequential outcome, the user should be able to understand the relevant grounds, challenge incorrect data, request human review where appropriate, and obtain a route to correction. Contestability can compensate for imperfect prediction because it creates a feedback channel from affected people into the system. It also generates high-value failure data. The important design question is whether the appeal process itself is accessible across language, disability, literacy, and socioeconomic differences.

Constitutional approaches offer a starting point because principles are explicit. A more advanced architecture could represent policies in machine-readable form, detect conflicts, request evidence, and preserve a trace of why a particular exception was allowed.

Purely statistical alignment can be difficult to verify. Some fairness constraints are naturally symbolic. A hiring system can state that name, gendered pronouns, or family status must not affect ranking unless a legally justified exception exists. A medical system can state which demographic factors are clinically relevant for a specific guideline. A procurement agent can require source diversity across jurisdictions.

Neurosymbolic systems can combine LLM flexibility with explicit constraints, provenance, and validation. The LLM proposes an action; a policy layer checks whether protected variables entered through allowed pathways; a verifier tests invariants; an audit trace records evidence. This does not solve the normative problem, but it makes enforcement less dependent on a model remembering a vague instruction.

The most promising use of symbolic controls is not to encode all ethics in rules. It is to encode the parts of policy that are sufficiently clear to be testable and leave genuinely ambiguous cases for explicit review.

As AI systems become too broad for manually authored benchmarks, auditing will itself become an AI task. Models can generate counterfactual cases, discover prompt families, search for subgroup failures, cluster anomalous outputs, and propose causal hypotheses. Automated red-teaming agents can explore a system more extensively than a human team.

The paradox is that auditors are biased too. The solution is diversity and adversarial structure. Multiple auditing agents with different models and prompts can search for disagreement. Human experts can validate high-impact findings. Dynamic benchmarks can be generated from current deployment data while protected test sets preserve comparability over time.

AI-assisted auditing can scale the search for failures. Models can generate counterfactual tests, paraphrase prompts, cluster incidents, inspect traces, propose causal hypotheses, and compare behavior across versions. They can also serve as adversarial users that explore edge cases more broadly than a fixed benchmark. Yet the auditor is another model with biases. The correct architecture is asymmetric: use AI to expand coverage and reduce human labor, but validate critical findings with independent methods, humans, or deterministic checks. An AI auditor should generate evidence, not certify itself.

Pluralism becomes an engineering problem when a system has to distinguish legitimate variation from discriminatory customization. A multilingual tutoring system may reasonably adapt examples, formality, or pedagogical style to local conventions. A hiring system should not use the same freedom to apply different standards of competence to demographic groups. The boundary cannot be learned safely from preference data alone because the training distribution may encode the very asymmetry under review. The architecture therefore needs a policy layer that states which dimensions are customizable, which evidence is admissible for a decision, and which constraints are invariant across users.

Executable controls are useful when the policy is clearer than the model behavior. Examples include prohibiting specified features from entering a ranking function, requiring independent evidence before a sensitive inference is used, enforcing exposure constraints in a recommender, or forcing escalation when uncertainty crosses a domain-specific threshold. Such controls should be small, auditable, and attached to provenance. They should not attempt to formalize every moral judgment. Their purpose is to prevent a learned component from silently redefining a rule that the institution has already decided.

The strongest architecture separates three layers that are frequently mixed together in alignment discussions: descriptive knowledge about social differences, normative policy about which differences may influence a task, and enforcement mechanisms that verify the policy. A model can retain rich cultural knowledge while an external checker prevents that knowledge from affecting a protected decision path. This separation is especially valuable for long-lived systems because policy can be revised without pretending that a new fine-tuning run has erased the old representation.

The value of explicit controls is not that ethics can be reduced to rules. It is that parts of policy that are already clear can be separated from learned habit, inspected independently, and enforced consistently while ambiguous cases remain visible as judgment rather than hidden optimization.

6.4 AI-assisted auditing and fairness as a control system

The scale of modern AI makes exhaustive manual auditing impossible. Models can help generate tests, search prompt spaces, cluster incidents, inspect traces, and propose hypotheses, but an automated auditor inherits its own biases and blind spots. The promising architecture is therefore adversarial and asymmetric: AI expands coverage; independent tests, humans, or executable invariants validate consequential findings. Fairness becomes a control loop with sensors, interventions, and evidence that is repeatedly renewed.

The frontier is not an all-knowing fairness judge. It is an ecosystem of fallible auditors designed to disagree productively.

The deepest shift is conceptual. Fairness should be treated less like a score and more like a control problem. A control system observes the environment, estimates state, detects deviation, applies an intervention, and checks whether the intervention worked. It expects drift.

For AI fairness, observation means subgroup and counterfactual monitoring. State estimation means understanding which component or social process may be responsible. Intervention can mean data changes, prompt policy, retrieval correction, threshold updates, fine-tuning, human review, or product redesign. Verification means rerunning a broad evaluation and checking for regressions in utility, robustness, privacy, and other groups.

This view is compatible with the lifecycle frameworks from NIST, with dynamic fairness research, and with the realities of continuously changing foundation-model systems. It also avoids the fiction that a model can be permanently certified as unbiased.

The next phase of the field should prioritize deployment validity over benchmark novelty. We need fewer claims that a new method "debiases LLMs" and more narrow, falsifiable claims about which harmful dependence is reduced under which distributions. We need native multilingual and culturally grounded evaluation, not English taxonomies translated outward. We need end-to-end audits of RAG and agent systems, not only base-model tests. We need to measure refusals, uncertainty, and quality of service alongside stereotype scores. We need causal and long-term analysis where static parity is inadequate. We need evaluation infrastructure that assumes metrics disagree and treats disagreement as evidence.

The table summarizes a research stack rather than a promised architecture. Each layer addresses a limitation of static scorekeeping: mechanism, uncertainty, legitimate value variation, enforceable policy, scalable test generation, and continued feedback.

Table 7. A plausible research stack beyond static fairness scores

Concept	Meaning in practice
Causal models	Represent mechanism hypotheses and distinguish legitimate from objectionable pathways rather than relying only on correlation.
Uncertainty and shift	Detect when the deployment environment moves outside the region for which fairness evidence was obtained.
Pluralistic policy	Represent legitimate value variation explicitly while maintaining non-negotiable legal and safety constraints.
Executable controls	Place inspectable invariants, provenance requirements, prohibited-feature rules, or escalation conditions around learned components.

Concept	Meaning in practice
AI-assisted audits	Use models to expand test coverage and analyze traces, while validating consequential findings independently.
Continuous control	Treat fairness as a monitored state with interventions and re-evaluation rather than a permanent certificate.

The components are deliberately complementary. Causal models cannot remove normative disagreement; executable controls cannot discover harms they were never specified to watch; audit agents can expand coverage but may share the same blind spots as the systems they evaluate. The frontier is therefore integration with explicit uncertainty and independent checks.

Most importantly, fairness research must retain a distinction between scientific measurement and social choice. Science can reveal disparities, causal mechanisms, and trade-offs. It cannot decide by itself which distribution of errors is just. That decision belongs to institutions and affected people, informed by technical evidence but not replaced by it.

The mature engineering objective is therefore not a mythical unbiased AI. It is an AI system whose relevant biases are made visible, whose harmful effects are bounded and monitored, whose assumptions are documented, whose decisions can be contested, and whose failure modes remain open to correction.

Neurosymbolic controls offer one promising engineering direction for making policy constraints explicit. Instead of asking a model to remember every fairness rule implicitly, a system can represent some constraints as executable checks around model behavior: required evidence fields, prohibited decision features, invariance tests, causal-path restrictions, uncertainty thresholds, exposure constraints, or escalation conditions. The language model remains useful for perception, generation, and ambiguous reasoning, while symbolic or programmatic components enforce rules that should be inspectable. This does not solve normative disagreement, but it separates policy from model habit and makes violations easier to trace.

The final step is to think in terms of fairness control systems. A control system has a specification, sensors, state estimates, interventions, and feedback. In AI fairness, the specification is the harm model and tolerances; sensors are benchmarks, production samples, user reports, counterfactual tests, and subgroup metrics; state estimates summarize current risk; interventions can occur at data, training, retrieval, prompt, threshold, or workflow layers; feedback comes from re-evaluation and deployment monitoring. This framework is deliberately broader than a fairness metric because complex systems cannot be governed by one number.

Beyond 2026, the most valuable research will likely make these control loops more reliable. We need culturally grounded evaluations that are not translations of English ontologies, causal methods that remain useful under partial identification, monitoring under privacy constraints, robust tests for agent trajectories, better calibration of LLM judges, methods for detecting synthetic-data homogenization, and formal interfaces between learned components and inspectable policy constraints. The endpoint is not an unbiased machine. It is a machine-mediated institution whose consequential dependencies are visible, bounded, contestable, and correctable.

Automated auditing is attractive because the test space of an open-ended model is effectively unbounded. An audit agent can generate demographic counterfactuals, paraphrase prompts, vary roles and conversation histories, search for tool-path divergence, and cluster thousands of outputs for human inspection. The danger is self-reference: if the same model family generates the tests, judges the outputs,

and proposes the remediation, correlated blind spots can create an illusion of coverage. A sound design deliberately introduces independence through different model families, deterministic checks, protected human-authored test sets, or reviewers who are not optimized by the same feedback loop.

Thinking of fairness as control also changes what an evaluation artifact should contain. The system needs a specification of harms and tolerances, sensors that measure relevant behavior, a diagnosis layer that proposes likely mechanisms, interventions at multiple architectural levels, and a re-test that determines whether the intervention created regressions elsewhere. Importantly, the state estimate should include uncertainty. A product should be able to say that evidence is insufficient for a language, subgroup, or new deployment region rather than forcing an unjustified pass/fail classification.

This framing suggests a practical research program for the next generation of assurance infrastructure. Evaluation suites should become versioned executable assets linked to system provenance. Agent traces should be queryable as evidence. Policy constraints should expose machine-checkable invariants where the institution can state them clearly. Human review should focus on unresolved normative and causal questions rather than manually repeating tests that software can reproduce. The result is not automated ethics; it is a better division of labor between measurement, explicit policy, learned behavior, and accountable judgment.

The long-term objective is therefore modest but demanding: not an “unbiased AI,” but a machine-mediated institution whose consequential dependencies are visible, whose harmful pathways are bounded, whose policies are contestable, and whose evidence can be renewed as the system changes.

The final position of the book is intentionally narrower than the slogan “unbiased AI” and stronger as an engineering claim. Fairness is a maintained body of evidence about bounded harms, explicit assumptions, inspectable policies, and a system that can be corrected when those assumptions stop holding.

CONCLUSION

The state of the art has advanced by repeatedly discovering that the object called “bias” was larger than the current measurement boundary.

Early work taught us to inspect datasets and subgroup error. Fairness theory showed that desirable criteria can conflict and that the conflict reflects different protected interests. Socio-technical analysis showed that valid metrics can still abstract away the institution that created the target. Foundation models then separated latent representation from visible behavior, and composed systems added retrieval, memory, tools, agents, and automated judges. Synthetic data and continuous deployment finally made time itself part of the system boundary.

The most defensible engineering method is therefore explicit and conditional. Define the harm. Identify the population. Explain the proxy and the causal assumptions. Measure more than one behavioral surface. Treat refusals and uncertainty as outcomes. Evaluate intersections and languages that matter for deployment. Audit retrieval exposure and agent paths. Validate automated judges. Re-run utility, calibration, robustness, privacy, and safety after mitigation. Preserve configuration and provenance so that the result can be reproduced after a model or product update.

There is no universal bias score and no universal fairness criterion. This is not an excuse for vagueness. It is a requirement for specificity. Science can reveal disparities, mechanisms, uncertainties, and trade-offs; institutions must still decide which differences are legitimate and which harms they are obligated to prevent. Good engineering keeps that normative ownership visible instead of hiding it inside a reward model or a dashboard.

The durable objective is not an AI system that can be declared unbiased. It is a system whose consequential dependencies can be observed, whose harmful pathways are constrained, whose assumptions are documented, whose decisions can be contested, and whose evidence is renewed as the world and the system change.

APPENDIX A. TECHNICAL GLOSSARY AND ORIENTATION LINKS

This appendix explains terms that a motivated high-school graduate or early undergraduate may encounter for the first time. Wikipedia links are provided for orientation and vocabulary, not as substitutes for the primary scientific literature cited in the bibliography.

A.1 Probability, statistics, and machine learning foundations

These terms describe how learning systems represent uncertainty, generalize from samples, and use measured variables. They are included so that readers can separate ordinary machine-learning vocabulary from fairness-specific claims in the main chapters.

[Calibration](#). A relationship between predicted probabilities and observed frequencies. A score of 0.8 is well calibrated when events assigned roughly that probability occur about eighty percent of the time.

[Base rate](#). The prevalence of an outcome in a population before considering a particular predictor. Different base rates are one reason fairness criteria can conflict.

[Distribution shift](#). A change between the data distribution used for training or evaluation and the distribution encountered later in deployment.

[Proxy variable](#). A measured variable used in place of a construct that is harder to observe directly. Proxies can be useful and still encode institutional inequality.

[Overfitting](#). Learning patterns that fit the training sample but do not generalize reliably to new data.

[Regularization](#). Methods that constrain a learning system to improve generalization or control complexity. Regularization is an example of deliberate inductive bias.

[Embedding](#). A numerical representation that maps discrete objects such as words, images, or users into a continuous vector space where relations can be learned.

[Classifier](#). A model that assigns an input to one or more categories. Many classical fairness metrics were developed for classifiers.

The linked pages are orientation aids for vocabulary and background. When the book makes a substantive claim about bias or fairness, the primary research in the bibliography rather than the glossary link is the evidential source.

A.2 Fairness, causality, and socio-technical reasoning

These concepts connect social concerns to measurable variables, decision criteria, and hypotheses about mechanism. Their technical meanings are narrower than everyday uses of words such as fairness or causality, so the explanations emphasize the distinctions needed to follow the engineering argument.

[Demographic parity](#). A fairness criterion requiring similar positive-decision rates across protected groups, regardless of the true outcome.

[Equalized odds](#). A criterion requiring equal false-positive and true-positive behavior across groups conditional on the actual outcome.

Counterfactual. A statement about what would have happened under an alternative condition. Counterfactual fairness uses causal models to reason about protected attributes.

Causal model. A model intended to represent directional relationships that support intervention or counterfactual reasoning, not only correlation.

Confounder. A variable that influences both a presumed cause and an outcome, potentially creating a misleading association.

Intersectionality. The idea that social categories can combine in ways that create experiences not captured by analyzing each category independently.

Construct validity. The degree to which a measurement actually captures the concept it claims to measure.

Socio-technical system. A system in which technical components and social institutions, people, incentives, rules, and workflows jointly determine outcomes.

No definition in this group should be read as the single canonical meaning of fairness. The criteria protect different relationships, while causal and socio-technical concepts clarify the assumptions needed to connect a numerical result to a real-world harm.

A.3 Language models, retrieval, and agents

These terms describe the components that surround contemporary foundation models. Understanding them is necessary because post-training, retrieval, reranking, tools, and agent logic can change behavior even when the base model is unchanged.

Large language model. A neural model trained on large text or multimodal corpora to predict, generate, and transform language conditioned on context.

Transformer. The neural architecture based on attention mechanisms that underlies most modern large language models.

Tokenization. The process of converting text into units used by a language model. Tokenization quality can differ across languages and scripts.

Instruction tuning. Fine-tuning a model on instruction-response examples so that it behaves more like an assistant than a raw next-token predictor.

RLHF. Reinforcement learning from human feedback, a family of methods that uses human preference data to shape model behavior.

RLAIF. Reinforcement learning from AI feedback, where model-generated or AI-evaluated preferences replace or supplement direct human feedback.

RAG. Retrieval-augmented generation, where external documents are retrieved and supplied to a generative model as context.

Reranker. A component that takes retrieved candidates and orders them again using a stronger relevance model or additional criteria.

Agent. An AI system that selects actions across multiple steps, often using tools, memory, retrieval, and feedback from the environment.

Sycophancy. A tendency of a language model to agree with or mirror a user's stated view even when evidence points elsewhere.

The architectural boundaries among these components vary across products. This is precisely why the main text treats the system boundary as something an audit must declare instead of assuming that the model alone is the relevant object.

A.4 Evaluation, robustness, privacy, and assurance

These concepts support the chapters on testing and continuous evidence. They provide vocabulary for asking whether a result generalizes, whether privacy constraints alter what can be measured, and whether an earlier fairness claim still applies after a system change.

[Benchmark](#). A standardized dataset or task used to compare systems. A benchmark is a measurement instrument and may have limited construct or deployment validity.

[Red teaming](#). Adversarial testing intended to discover failure modes by actively trying to make a system behave unsafely or incorrectly.

[Robustness](#). The ability of a system to maintain acceptable behavior under perturbations, noise, or changes in operating conditions.

[Uncertainty calibration](#). The degree to which a model's confidence reflects how often it is correct, including across subgroups or deployment conditions.

[Differential privacy](#). A formal privacy framework that bounds how much the output of an analysis can depend on any single individual record.

[Model card](#). Documentation describing intended use, evaluation, limitations, and relevant characteristics of a trained model.

[Dataset datasheet](#). Structured documentation of how a dataset was created, composed, maintained, and intended to be used.

[Provenance](#). Information about the origin, transformations, versions, and custody of data or model outputs.

[Audit trail](#). A record that makes important actions, changes, and evidence traceable after the fact.

[Neurosymbolic AI](#). Architectures that combine learned neural representations with symbolic, logical, or programmatic reasoning and constraints.

These terms become most important once a model leaves a static benchmark. They distinguish uncertainty about model behavior from uncertainty about changing populations, incomplete evidence, privacy-limited measurement, and the validity of past evaluations.

APPENDIX B. HOW TO READ BIAS AND FAIRNESS RESEARCH CRITICALLY

A technically sophisticated paper can still make an invalid generalization if its metric, benchmark, or deployment claim is broader than the experiment supports.

Begin by identifying the exact object being called bias. Is the paper measuring embedding association, token probability, generated language, a forced-choice decision, a rank, a refusal, an action trajectory, or the output of another model used as a judge? These are different behavioral surfaces. Evidence on one can motivate concern about another, but the transfer requires an argument.

Next inspect the construct. What social concern is the benchmark trying to represent, and how were the categories chosen? A benchmark built around English-language stereotypes may be methodologically sound for that setting and culturally weak elsewhere. Translation does not automatically create construct validity. The same caution applies to toxicity, sentiment, prestige, authority, and “professionalism,” all of which can be mediated by external classifiers with their own biases.

Then inspect experimental stability. Were prompts paraphrased? Was more than one sample generated? Were refusals separated from neutral answers? Were model and system-prompt versions reported? Could the benchmark have been present in training data? If an LLM judge was used, were position and style effects tested? A numerically precise result can still be unstable under small changes in the evaluation instrument.

Finally, compare the conclusion to the level of the experiment. Evidence that a method lowers metric Y on benchmark Z supports a narrow claim. It does not by itself establish that the model is fair in hiring, politically neutral, culturally unbiased, or safe in deployment. The strongest papers make transfer claims cautiously and expose uncertainty. The weakest convert local benchmark movement into a universal statement about fairness.

BIBLIOGRAPHY

The bibliography prioritizes primary research, influential surveys, and institutional sources used to ground the course. Wikipedia links in Appendix A are educational orientation aids rather than evidential sources.

- Haque, F., Fu, Z., Aditya, R., Xu, D., and Niu, X. (2026). Fair RAG: End-to-End Fairness Across Retrieval and Generation. Findings of ACL 2026. <https://aclanthology.org/2026.findings-acl.1358/>
- Wang, Y., Cai, Z., Bao, Y., Zhang, X., and Liu, Y. (2026). Observations and Remedies for Large Language Model Bias in Self-Consuming Performative Loop. ACL 2026. <https://doi.org/10.18653/v1/2026.acl-long.1561>
- Li, M., Chen, H., Wang, Y., Zhu, T., Zhang, W., Zhu, K., Wong, K.-F., and Wang, J. (2026). Understanding and Mitigating Bias Inheritance in LLM-based Data Augmentation on Downstream Tasks. ACL 2026. <https://aclanthology.org/2026.acl-long.1408/>
- Malik, S., Abdullah, H. M., Saha, S., and Sheth, A. (2026). Ask Me Again Differently: GRAS for Measuring Bias in Vision Language Models on Gender, Race, Age, and Skin Tone. Findings of EACL 2026. <https://doi.org/10.18653/v1/2026.findings-eacl.123>
- Islam, T. (2026). Who Gets Which Message? Auditing Demographic Bias in LLM-Generated Targeted Text. Findings of ACL 2026. <https://aclanthology.org/2026.findings-acl.452/>
- Proskurina, I., Gourru, A., and Velcin, J. (2026). The Fairness Collapse Phenomenon: Bias Amplification in Language Models Trained on Synthetic Data. arXiv:2608.04268. Preprint.
- Dwork, C., Hardt, M., Pitassi, T., Reingold, O., and Zemel, R. (2012). Fairness Through Awareness. Proceedings of the 3rd Innovations in Theoretical Computer Science Conference. <https://doi.org/10.1145/2090236.2090255>
- Hardt, M., Price, E., and Srebro, N. (2016). Equality of Opportunity in Supervised Learning. Advances in Neural Information Processing Systems 29.
- Kusner, M. J., Loftus, J., Russell, C., and Silva, R. (2017). Counterfactual Fairness. Advances in Neural Information Processing Systems 30.
- Caliskan, A., Bryson, J. J., and Narayanan, A. (2017). Semantics derived automatically from language corpora contain human-like biases. Science, 356(6334), 183-186. <https://doi.org/10.1126/science.aal4230>
- Bolukbasi, T., Chang, K.-W., Zou, J., Saligrama, V., and Kalai, A. (2016). Man is to Computer Programmer as Woman is to Homemaker? Debiasing Word Embeddings. Advances in Neural Information Processing Systems 29.
- Blodgett, S. L., Barocas, S., Daumé III, H., and Wallach, H. (2020). Language (Technology) is Power: A Critical Survey of "Bias" in NLP. ACL 2020. <https://aclanthology.org/2020.acl-main.485/>
- Bender, E. M., Gebru, T., McMillan-Major, A., and Shmitchell, S. (2021). On the Dangers of Stochastic Parrots: Can Language Models Be Too Big? FAccT 2021. <https://doi.org/10.1145/3442188.3445922>
- Mehrabi, N., Morstatter, F., Saxena, N., Lerman, K., and Galstyan, A. (2021). A Survey on Bias and Fairness in Machine Learning. ACM Computing Surveys, 54(6), Article 115. <https://doi.org/10.1145/3457607>
- Suresh, H., and Gutttag, J. V. (2021). A Framework for Understanding Sources of Harm throughout the Machine Learning Life Cycle. Equity and Access in Algorithms, Mechanisms, and Optimization. ACM. <https://doi.org/10.1145/3465416.3483305>
- Schwartz, R., Vassilev, A., Greene, K. K., Perine, L., Burt, A., and Hall, P. (2022). Towards a Standard for Identifying and Managing Bias in Artificial Intelligence. NIST Special Publication 1270. <https://doi.org/10.6028/NIST.SP.1270>
- Jacobs, A. Z., and Wallach, H. (2021). Measurement and Fairness. Proceedings of the 2021 ACM Conference on Fairness, Accountability, and Transparency. <https://doi.org/10.1145/3442188.3445901>
- Selbst, A. D., Boyd, D., Friedler, S. A., Venkatasubramanian, S., and Vertesi, J. (2019). Fairness and Abstraction in Sociotechnical Systems. FAT* 2019. <https://doi.org/10.1145/3287560.3287598>
- Kleinberg, J., Mullainathan, S., and Raghavan, M. (2016). Inherent Trade-Offs in the Fair Determination of Risk Scores. arXiv:1609.05807.
- Chouldechova, A. (2017). Fair Prediction with Disparate Impact: A Study of Bias in Recidivism Prediction Instruments. Big Data, 5(2), 153-163. <https://doi.org/10.1089/big.2016.0047>
- Obermeyer, Z., Powers, B., Vogeli, C., and Mullainathan, S. (2019). Dissecting racial bias in an algorithm used to manage the health of populations. Science, 366(6464), 447-453. <https://doi.org/10.1126/science.aax2342>
- Buolamwini, J., and Gebru, T. (2018). Gender Shades: Intersectional Accuracy Disparities in Commercial Gender Classification. Proceedings of Machine Learning Research, 81, 77-91. <https://proceedings.mlr.press/v81/buolamwini18a.html>
- Koenecke, A., Nam, A., Lake, E., et al. (2020). Racial disparities in automated speech recognition. Proceedings of the National Academy of Sciences, 117(14), 7684-7689. <https://doi.org/10.1073/pnas.1915768117>
- Gebru, T., Morgenstern, J., Vecchione, B., Vaughan, J. W., Wallach, H., Daume III, H., and Crawford, K. (2021). Datasheets for Datasets. Communications of the ACM, 64(12), 86-92.
- Bender, E. M., and Friedman, B. (2018). Data Statements for Natural Language Processing: Toward Mitigating System Bias and Enabling Better Science. Transactions of the Association for Computational Linguistics, 6, 587-604. https://doi.org/10.1162/tacl_a_00041
- Mitchell, M., Wu, S., Zaldivar, A., et al. (2019). Model Cards for Model Reporting. FAT* 2019. <https://doi.org/10.1145/3287560.3287596>

- D'Amour, A., Heller, K., Moldovan, D., et al. (2022). Underspecification Presents Challenges for Credibility in Modern Machine Learning. *Journal of Machine Learning Research*, 23.
- Nangia, N., Vania, C., Bhalerao, R., and Bowman, S. R. (2020). CrowS-Pairs: A Challenge Dataset for Measuring Social Biases in Masked Language Models. *EMNLP 2020*. <https://aclanthology.org/2020.emnlp-main.154/>
- Nadeem, M., Bethke, A., and Reddy, S. (2021). StereoSet: Measuring stereotypical bias in pretrained language models. *ACL-IJCNLP 2021*. <https://aclanthology.org/2021.acl-long.416/>
- Li, T., Khot, T., Khashabi, D., Sabharwal, A., and Srikumar, V. (2020). UNQOVERing Stereotyping Biases via Underspecified Questions. *Findings of EMNLP 2020*. <https://aclanthology.org/2020.findings-emnlp.311/>
- Dhamala, J., Sun, T., Kumar, V., et al. (2021). BOLD: Dataset and Metrics for Measuring Biases in Open-Ended Language Generation. *FAccT 2021*.
- Nozza, D., Bianchi, F., and Hovy, D. (2021). HONEST: Measuring Hurtful Sentence Completion in Language Models. *NAACL 2021*. <https://aclanthology.org/2021.naacl-main.191/>
- Parrish, A., Chen, A., Nangia, N., Padmakumar, V., Phang, J., Thompson, J., Htut, P. M., and Bowman, S. R. (2022). BBQ: A Hand-Built Bias Benchmark for Question Answering. *Findings of ACL 2022*. <https://aclanthology.org/2022.findings-acl.165/>
- Smith, E. M., Hall, M., Kambadur, M., Presani, E., and Williams, A. (2022). I'm sorry to hear that: Finding New Biases in Language Models with a Holistic Descriptor Dataset. *EMNLP 2022*. <https://aclanthology.org/2022.emnlp-main.625/>
- Costa-jussa, M. R., Andrews, P., Basta, C., et al. (2023). Multilingual Holistic Bias: Extending Descriptors and Patterns to Unveil Demographic Biases in Languages at Scale. *EMNLP 2023*.
- Liang, P., Bommasani, R., Lee, T., et al. (2022). Holistic Evaluation of Language Models. *arXiv:2211.09110*.
- Wang, B., Chen, W., Pei, H., et al. (2023). DecodingTrust: A Comprehensive Assessment of Trustworthiness in GPT Models. *arXiv:2306.11698*.
- Hofmann, V., Kalluri, P. R., Jurafsky, D., and King, S. (2024). AI generates covertly racist decisions about people based on their dialect. *Nature*, 633, 147-154. <https://doi.org/10.1038/s41586-024-07856-5>
- Zayed, A., Mordido, G., Shabanian, S., et al. (2024). Why Don't Prompt-Based Fairness Metrics Correlate? *ACL 2024*. <https://aclanthology.org/2024.acl-long.487/>
- Bang, Y., Chen, D., Lee, N., and Fung, P. (2024). Measuring Political Bias in Large Language Models: What Is Said and How It Is Said. *ACL 2024*. <https://aclanthology.org/2024.acl-long.600/>
- Sharma, M., Tong, M., Korbak, T., et al. (2023). Towards Understanding Sycophancy in Language Models. *arXiv:2310.13548*.
- Perez, E., Ringer, S., Lukosiute, K., et al. (2022). Discovering Language Model Behaviors with Model-Written Evaluations. *arXiv:2212.09251*.
- Denison, C., MacDiarmid, M., Barez, F., et al. (2024). Sycophancy to Subterfuge: Investigating Reward-Tampering in Language Models. *arXiv:2406.10162*.
- Zhao, Y., Wang, B., Wang, Y., Zhao, D., He, R., and Hou, Y. (2025). Explicit vs. Implicit: Investigating Social Bias in Large Language Models through Self-Reflection. *Findings of ACL 2025*. <https://aclanthology.org/2025.findings-acl.1/>
- Himelstein, R., LeVi, A., Youngmann, B., Nemcovsky, Y., and Mendelson, A. (2026). Silenced Biases: The Dark Side LLMs Learned to Refuse. *AAAI 2026*. <https://doi.org/10.1609/aaai.v40i44.41078>
- Asgari, A., Wu, H., Naziri, A., Kolahdoust, M., and Seyyed-Kalantari, L. (2026). Quantifying Metric and Model Agreement in Bias Evaluation of Large Language Models. *ACL 2026*. <https://doi.org/10.18653/v1/2026.acl-long.769>
- Nadeem, A., Dras, M., and Naseem, U. (2026). Fairness Evaluation and Inference Level Mitigation in LLMs. *Findings of ACL 2026*. <https://aclanthology.org/2026.findings-acl.1452/>
- Jung, D., et al. (2025). FLEX: A Benchmark for Evaluating Robustness of Fairness in Large Language Models. *Findings of NAACL 2025*. <https://aclanthology.org/2025.findings-naacl.199/>
- Guan, X., et al. (2025). SAGED: A Holistic Bias-Benchmarking Pipeline for Language Models with Customisable Fairness Calibration. *COLING 2025*. <https://aclanthology.org/2025.coling-main.202/>
- Yamamoto, T., et al. (2025). Bias Mitigation or Cultural Commonsense? Evaluating LLMs with a Japanese Dataset. *EMNLP 2025*. <https://aclanthology.org/2025.emnlp-main.874/>
- Faulborn, M., Sen, I., Pellert, M., Spitz, A., and Garcia, D. (2025). Only a Little to the Left: A Theory-grounded Measure of Political Bias in Large Language Models. *ACL 2025*. <https://aclanthology.org/2025.acl-long.1529/>
- Gubelmann, R., and Karray, G. (2025). Assessing Reliability and Political Bias In LLMs' Judgements of Formal and Material Inferences With Partisan Conclusions. *ACL 2025*. <https://aclanthology.org/2025.acl-long.1450/>
- Wu, X., Li, S., Wu, H.-T., Tao, Z., and Fang, Y. (2025). Does RAG Introduce Unfairness in LLMs? Evaluating Fairness in Retrieval-Augmented Generation Systems. *COLING 2025*. <https://aclanthology.org/2025.coling-main.669/>
- Hu, M., Wu, H., Zhu, R., et al. (2025). No Free Lunch: Retrieval-Augmented Generation Undermines Fairness in LLMs, Even for Vigilant Users. *Findings of EMNLP 2025*. <https://aclanthology.org/2025.findings-emnlp.984/>
- Li, Y., Guo, X., Gao, J., et al. (2025). LLMs Trust Humans More, That's a Problem! Unveiling and Mitigating the Authority Bias in Retrieval-Augmented Generation. *ACL 2025*. <https://aclanthology.org/2025.acl-long.1400/>

- Wang, P., Li, L., Chen, L., et al. (2024). Large Language Models are not Fair Evaluators. ACL 2024. <https://aclanthology.org/2024.acl-long.511/>
- Shi, L., Ma, C., Liang, W., Diao, X., Ma, W., and Vosoughi, S. (2025). Judging the Judges: A Systematic Study of Position Bias in LLM-as-a-Judge. IJCNLP-AACL 2025. <https://aclanthology.org/2025.ijcnlp-long.18/>
- Zhou, H., Huang, H., Long, Y., et al. (2024). Mitigating the Bias of Large Language Model Evaluation. CCL 2024. <https://aclanthology.org/2024.ccl-1.101/>
- Iso, H., et al. (2025). Evaluating Bias in LLMs for Job-Resume Matching: Gender, Race, and Education. NAACL Industry 2025. <https://aclanthology.org/2025.naacl-industry.55/>
- Seshadri, P., et al. (2025). Analyzing the Allocational Fairness of LLMs in Hiring Systems. IJCNLP-AACL 2025. <https://aclanthology.org/2025.ijcnlp-long.143/>
- Wu, X., Wang, Y., Wu, H.-T., Tao, Z., and Fang, Y. (2025). Evaluating Fairness in Large Vision-Language Models Across Diverse Demographic Attributes and Prompts. Findings of EMNLP 2025. <https://aclanthology.org/2025.findings-emnlp.1251/>
- Gustafson, L., et al. (2023). FACET: Fairness in Computer Vision Evaluation Benchmark. arXiv:2309.00035.
- Shumailov, I., Shumaylov, Z., Zhao, Y., Papernot, N., Anderson, R., and Gal, Y. (2024). AI models collapse when trained on recursively generated data. Nature, 631, 755-759. <https://doi.org/10.1038/s41586-024-07566-y>
- Deldjoo, Y., Jannach, D., Bellogin, A., Difonzo, A., and Zanzonelli, D. (2024). Fairness in recommender systems: research landscape and future directions. User Modeling and User-Adapted Interaction, 34, 59-108.
- Vassoy, B., and Langseth, H. (2024). Consumer-side fairness in recommender systems: a systematic survey of methods and evaluation. Artificial Intelligence Review, 57, 101.
- OpenAI. (2024). Evaluating fairness in ChatGPT. <https://openai.com/index/evaluating-fairness-in-chatgpt/>
- OpenAI. (2025). Defining and evaluating political bias in LLMs. <https://openai.com/index/defining-and-evaluating-political-bias-in-llms/>
- Anthropic and Collective Intelligence Project. (2023). Collective Constitutional AI: Aligning a Language Model with Public Input. <https://www.anthropic.com/research/collective-constitutional-ai-aligning-a-language-model-with-public-input>
- NIST. (2024, updated 2026). Artificial Intelligence Risk Management Framework: Generative Artificial Intelligence Profile. NIST AI 600-1. <https://doi.org/10.6028/NIST.AI.600-1>
- European Union. (2024, consolidated 27 July 2026). Regulation (EU) 2024/1689, Artificial Intelligence Act. <https://eur-lex.europa.eu/eli/reg/2024/1689/2026-07-27/eng>
- Council of Europe. (2024-2026). Framework Convention on Artificial Intelligence and Human Rights, Democracy and the Rule of Law, CETS No. 225. <https://www.coe.int/en/web/artificial-intelligence/the-framework-convention-on-artificial-intelligence>