

ANTI-TRIVIALIZATION MACHINES OF THE FUTURE

Institutions for Attention, Truth, and
Long-Term Judgment in the Age of AI



A book about the institutions we do not yet know how to build,
but may no longer be able to live without

ANTI-TRIVIALIZATION MACHINES OF THE FUTURE

Institutions for Attention, Truth, and Long-Term
Judgment in the Age of AI

*A book about the institutions we do not yet know how to build, but may no longer be
able to live without*

2026

Copyright © [2026] Axiologic Research

All rights reserved.

KDP publishing rights held by Outfinity SRL (Iași, Romania).

Available for free on Axiologic website (www.axiologic.net).

Author's Note

This work was created under the umbrella of Axiologic Research as part of two interconnected research programs, one dedicated to Artificial Intelligence and the other to Outfinitism, both led by Sînică Alboaie. To explore the details of how these two programs intersect and build upon each other, we invite readers to visit the Outfinity Venture Studio page at www.outfinity.ch.

While Artificial Intelligence language models were used to assist with text generation, the foundational philosophy, thematic structure, and analytical approach derive exclusively from the author's decades of dedicated study of social phenomena, social technologies, and genuine hard technologies resulting from various European and self funded research projects. Recently, within the Achilles research project, we have been deeply studying AI generated literature and its automated verification using neuro symbolic approaches; all the free books made available to the public are part of the suite of works we use to test and popularise our latest ideas and technologies.

Contents

Preface: The Question Behind the Book..... 5

Chapter 1: The Problem Is Not Triviality..... 10

Chapter 2: The Old Machines of Seriousness..... 18

Chapter 3: The Industrialization of Attention.....26

Chapter 4: The Failure Is Institutional, Not Moral.....34

Chapter 5: AI and the Great Cognitive Surplus..... 41

Chapter 6: Why the Counter-Machines Will Appear.....49

Chapter 7: The Attention Constitution..... 58

Chapter 8: Truth Without a Ministry of Truth..... 66

Chapter 9: Institutions for Time..... 73

Chapter 10: Polycentric Seriousness..... 79

Chapter 11: Education, Work, and the Rebuilding of Human Agency...
85

Chapter 12: A Civilization That Can Think..... 92

Bibliography..... 100

Preface: The Question Behind the Book

This book began with a provocation rather than a theory. In 2008, the clinical psychologist John F. Schumaker published an essay titled "The Triumph of Triviality." It drew on a humanistic lineage that included Erich Fromm, Carl Rogers, Abraham Maslow, Rollo May, Viktor Frankl, and earlier critics of mass culture. Schumaker's complaint was recognizable: modern consumer society appeared to have become extraordinarily good at producing comfort, stimulation, objects, images, and entertainment, while remaining strangely poor at producing mature attention to questions that mattered across generations. The future that humanistic psychology had imagined - a future of self-actualization, ethical development, and deeper forms of solidarity - had not arrived on schedule (Schumaker 2008).

The essay is valuable because it names an experience many people recognize. It is also insufficient. It attributes too much to cultural shallowness, too much to consumerism, and too much to a decline in individual seriousness. It occasionally falls into the oldest trap of cultural criticism: comparing the messiness of the present with an idealized past populated by unusually thoughtful people. There was never an age in which most citizens spent their evenings discussing Spinoza, long-term climate policy, and the epistemology of science. Ancient cities had spectacle. Medieval towns had gossip. Early modern Europe had pamphlet wars. Nineteenth-century newspapers discovered sensationalism. Radio and television learned to package politics as entertainment long before social media learned to optimize a feed.

The more interesting question is therefore not why human beings like trivial things. They always have, and this is not necessarily a defect. Play, gossip, sport, fashion, humor, spectacle, flirtation, games, and the pleasures of temporary irrelevance are part of a livable human world. A civilization that abolished triviality would not become wise. It would become unbearable.

The real problem is trivialization: the process by which activities that require context, patience, memory, uncertainty, and judgment are reorganized according to the selection rules of entertainment and immediate attention. Politics becomes a performance. Science becomes a headline. Education becomes a stream of measurable outputs. Personal identity becomes a profile. Reputation becomes a count. Argument becomes a clip. Institutions begin to communicate at the rhythm of platforms, and then, slowly, to think at that rhythm as well.

This distinction changes the problem completely. If triviality were a defect of human character, the obvious remedies would be moral exhortation, education, discipline, or perhaps a return to some earlier cultural form. If trivialization is an institutional process, however, the remedy must also be institutional. We have to ask what kinds of organizations, protocols, norms, rights, technologies, and public infrastructures can preserve slower forms of judgment inside an environment that rewards speed.

This book calls such arrangements anti-trivialization machines. The phrase is deliberately mechanical, but the machines in question are not necessarily pieces of hardware or software. A court is a machine in this sense. So is a university, a scientific journal, a parliament, an archive, a professional code, an accounting standard, a public-health agency, and a well-designed meeting. Each takes a noisy stream of events and imposes transformations upon it. Claims must be documented. Decisions must pass through stages. Reasons must be stated. Records must persist. Adversaries may challenge one another. Some decisions are deliberately slowed. Some kinds of evidence count more than others. The institution converts immediacy into procedure.

Human civilization already contains thousands of such machines. We rarely notice them because successful institutions disappear into normality. We notice the bridge, not the building code; the medicine, not the clinical-trial protocol; the bank transfer, not the accounting and settlement infrastructure; the scientific conclusion, not the archive of failed hypotheses that made the

conclusion credible. Seriousness is often mistaken for an individual virtue because the institutions that manufacture it are hidden in the background.

The thesis of this book is that the background is now inadequate. The industrial economy built institutions for factories, corporations, money, property, labor, public health, and physical infrastructure. The information economy added institutions for telecommunications, data, software, intellectual property, and privacy, although imperfectly. The attention economy and the emerging AI economy create a further institutional problem: they alter the conditions under which societies notice, remember, compare, prioritize, and decide. We have powerful institutions for governing money, but weak institutions for governing attention. We have elaborate procedures for certifying medicines, but almost none for certifying the cognitive environments in which millions of people form beliefs. We have centuries of jurisprudence about property and contracts, but only the beginnings of a jurisprudence for algorithmically mediated agency.

This argument belongs to a long intellectual history. Aldous Huxley and George Orwell offered two enduring images of domination, one through pleasure and one through coercion. Hannah Arendt studied the conditions of public judgment and political action. Jacques Ellul examined technological systems and propaganda. The Frankfurt School analyzed the culture industry and the conversion of culture into standardized commodities. Marshall McLuhan and later Neil Postman insisted that media do not merely carry messages; they shape the forms of discourse that become easy or difficult. Guy Debord described a society in which representation displaces direct experience. Jürgen Habermas asked what institutional conditions make a public sphere capable of rational-critical debate. Herbert Simon saw a structural fact that has only become more important: when information becomes abundant, attention becomes scarce. Later scholars of institutions, including Douglass North and Elinor Ostrom, showed why durable social outcomes depend not only on preferences or technologies but on rules, enforcement, incentives, local knowledge, and the architecture of collective

action (Arendt 1958; Ellul 1965; Horkheimer and Adorno 1947; McLuhan 1964; Debord 1967; Simon 1971; Postman 1985; North 1990; Ostrom 1990; Habermas 1989).

The book also departs from much of this lineage. It is not a book of decline. It does not assume that television, social media, or artificial intelligence must make humanity stupid. Research on digital media is far more mixed and context-dependent than the most dramatic public arguments suggest. Recent experiments show that feed design can influence political affect and exposure, but they also show that the effects are not reducible to a single algorithmic switch. Research on adolescents repeatedly warns against treating "screen time" as a single causal substance. Studies of generative AI similarly find both cognitive risks and opportunities: automation can reduce effort in some tasks, while well-designed tools can redirect effort toward verification, framing, comparison, and reflection (Guess et al. 2023; Orben et al. 2022; Fassi et al. 2025; Lee et al. 2025; Piccardi et al. 2025).

The deepest optimism of this book is therefore institutional rather than sentimental. New capacities produce new externalities, and new externalities create pressure for new forms of organization. Industrialization produced dangerous factories and then labor law, safety engineering, inspection, and standards. Dense cities produced epidemics and then sanitation systems and public health. Financial markets produced panics and then central banks, disclosure regimes, deposit insurance, and more sophisticated accounting. Scientific power produced unethical experiments and then research ethics, review boards, informed consent, and clinical-trial regulation. None of these counter-institutions appeared automatically, cleanly, or without conflict. They were partial, late, sometimes captured, and often created only after damage. Yet they appeared because societies that wanted to continue enjoying the benefits of a new capacity had to learn how to govern its costs.

The yin-yang metaphor is useful here if treated as a metaphor rather than a law of history. Every expansion of capacity carries within it a pressure for its complement. Speed creates a need for brakes. Connectivity creates a need for boundaries. Abundance creates a need for selection. Automation creates a need for audit. Personalization creates a need for shared reference points. Synthetic media create a need for provenance. Artificial intelligence creates a need for institutions that protect human and collective judgment. The corrective principle does not arrive from outside the technological system. It is generated by the system's own success.

This is why the future institutions described in the pages that follow are not offered as a blueprint. We do not know how to build them yet. Some will fail. Some will sound absurd before they become mundane. The point is to identify the institutional functions that are becoming necessary before we know their final forms.

The wager is simple. A civilization that can build systems capable of capturing attention at planetary scale can also build systems capable of protecting attention at planetary scale. A civilization that can generate persuasive language on demand can also build infrastructures that demand provenance, contestability, and uncertainty. A civilization that can optimize engagement can learn to optimize for deliberation without reducing deliberation to another metric. The road between these capacities will not be pleasant or linear. But the very excesses that make new institutions necessary will also make them politically imaginable.

Chapter 1: The Problem Is Not Triviality

A child laughing at an absurd video is not a civilization in decline. A football match is not evidence of intellectual collapse. A celebrity scandal may be ridiculous, but societies have always reserved part of their attention for the ridiculous. Any theory that begins by condemning pleasure will misunderstand both human psychology and political reality.

The useful distinction is between triviality and trivialization. Triviality is a category of content. Trivialization is a change in the rules by which content competes for attention and acquires legitimacy. A joke remains a joke. A scientific disagreement, however, can be trivialized when it is presented as a duel between personalities because conflict is easier to distribute than uncertainty. A parliamentary question can be trivialized when its success is measured by circulation rather than legislative consequence. A university can trivialize itself when it begins selecting research questions according to public visibility, publication counts, or grant-compatible fashion rather than intellectual importance. A company can trivialize risk when quarterly reporting compresses a twenty-year vulnerability into a line that disappears beneath next quarter's revenue.

This suggests a more precise definition. Trivialization occurs when the selection environment systematically favors signals with low contextual cost, short feedback cycles, high emotional salience, and easy comparability over signals whose value requires time, expertise, uncertainty, or cumulative interpretation. The critical word is systematically. There is nothing wrong with a short message. There is a problem when shortness becomes the hidden constitution of what can survive.

The twentieth century's media critics described parts of this transformation before digital platforms existed. Horkheimer and Adorno argued that industrialized culture standardizes experience and turns cultural production into a predictable commodity. McLuhan shifted attention from

content to medium: the structure of communication changes what can be communicated. Debord described spectacle as a social relation mediated by images. Postman later argued that television did not merely add entertainment to politics; it changed what politics had to become in order to function effectively on television (Horkheimer and Adorno 1947; McLuhan 1964; Debord 1967; Postman 1985).

These arguments were often overstated. Printed culture was never purely rational. Television produced serious journalism as well as nonsense. The internet has made possible extraordinary forms of scholarship, coordination, education, art, mutual aid, open-source engineering, and public access to knowledge. A person with an inexpensive device can now read scientific papers, study a language, follow a university course, inspect government data, collaborate on software, and consult cultural archives that would once have required access to elite institutions. Any serious account must place this cognitive abundance next to cognitive fragmentation.

The paradox is that the same infrastructure can expand access to depth while increasing the competitive advantage of superficiality. This is not contradictory. A library and a casino can occupy the same device. More importantly, the casino can finance itself by learning exactly when the reader becomes bored with the library.

Herbert Simon's analysis of information abundance provides a cleaner starting point than cultural pessimism. Information consumes attention. As information becomes cheap to produce and transmit, the scarce resource shifts from information to the capacity to receive, evaluate, and act upon it (Simon 1971). Digital networks did not create this scarcity, but they changed its scale. Social platforms then made the allocation of attention observable at high resolution. Views, clicks, reactions, shares, dwell time, retention, and social graphs turned attention into a measurable flow. Once measurable, it could be optimized. Once optimized, it could be auctioned.

The result is not a conspiracy and does not require malicious actors. It is an equilibrium. If one organization learns to capture attention more effectively, competitors must respond. If one politician learns to communicate in a more viral form, rivals imitate the form. If one news organization discovers that emotionally charged headlines improve survival, editorial norms face economic pressure. If one academic field discovers a metric that allocates prestige efficiently, researchers learn to optimize against it. Goodhart's law and Campbell's law describe the general danger: when a measure becomes a target, behavior reorganizes around the measure, often degrading the underlying objective it was supposed to represent.

Trivialization therefore has no single headquarters. It can emerge from individually reasonable decisions. A newsroom wants readers. A politician wants votes. A researcher wants citations. A school wants measurable performance. A platform wants engagement. A user wants stimulation. None of these goals is inherently corrupt. But their interaction can construct a selection environment in which immediate legibility wins against slow importance.

The most dangerous feature of such an environment is not that it makes people stupid. It can make intelligent people spend their intelligence on short-horizon adaptation. A talented journalist becomes excellent at headline optimization. A politician becomes excellent at symbolic conflict. A scientist becomes excellent at navigating incentives. A manager becomes excellent at making dashboards look healthy. The system can become locally intelligent and globally foolish.

This distinction matters because it changes the moral tone of the problem. The language of decadence blames citizens for being distracted. Institutional analysis asks what it is rational for citizens to do inside the environment they have been given. A person who checks a phone repeatedly may be exercising weak self-control, but the device, application, social network, employer, and peer group may all be organized around immediate

responsiveness. Asking the individual alone to become disciplined is like solving urban traffic by asking each driver to become morally superior.

Nor should seriousness be confused with solemnity. Some of the most serious human activities depend on play. Scientific imagination requires speculative freedom. Democratic culture needs satire. Friendship needs purposeless conversation. Childhood needs games whose value cannot be explained by a productivity framework. The anti-trivialization project must therefore resist a second failure mode: the bureaucratization of meaning. A Ministry of Important Things would quickly become ridiculous and dangerous.

A good anti-trivialization institution does not decide what citizens must care about. It protects conditions under which different kinds of value can survive long enough to be judged. It does not ban the fast; it preserves access to the slow. It does not certify one truth; it preserves processes by which claims can be challenged and corrected. It does not eliminate entertainment; it prevents entertainment metrics from becoming the universal currency of legitimacy.

That leads to the book's first institutional principle: plural tempos. Healthy societies need multiple speeds. Markets can be fast. Courts are often deliberately slow. Emergency medicine must act rapidly. Science sometimes requires decades. Elections occur periodically rather than continuously. Constitutions are harder to change than ordinary laws because not every political impulse should rewrite the rules of the game. Time itself is an institutional design variable.

Digital systems have tended to flatten these tempos. Everything arrives through the same rectangle. A family message, a war, an advertisement, a scientific discovery, an insult, a weather warning, and a joke may appear in one continuous feed. The interface equalizes what reality does not. The result is not merely distraction; it is a loss of temporal differentiation.

One can see the same problem from the opposite direction by asking what societies have historically done when they needed people to take something seriously. They changed the setting. Religious traditions created sacred times and sacred spaces in which ordinary behavior was suspended. Courts created ritual, costume, sequence, and rules of speech. Universities created semesters, libraries, examinations, degrees, and communities in which arguments acquired longer lives than conversations. Scientific societies created journals and meetings in which claims were detached from the charisma of the first speaker and entered a record that strangers could inspect. None of these forms guaranteed wisdom. Their importance lies in a simpler fact: human beings often become capable of a different kind of attention when a social structure changes the cost of interruption and the meaning of response.

This is why the history of attention cannot be reduced to psychology. Attention is socially scheduled. A monk in a scriptorium, a juror in a trial, a violinist in rehearsal, a surgeon in an operating theatre, and a teenager in a group chat inhabit different attention regimes. The brain is the same kind of organ, but the institution surrounding it establishes different expectations about what may interrupt, what must be remembered, what counts as completion, and what happens after an error. The modern digital environment is unusual not because it invented distraction but because it carries a single high-frequency regime across contexts that used to have stronger boundaries.

N. Katherine Hayles distinguished between forms of deep and hyper attention, arguing that contemporary media environments encourage rapid switching across streams while older pedagogical forms often prize prolonged concentration (Hayles 2007). The distinction should not be turned into a hierarchy. Hyper attention is useful in environments where signals are numerous and change rapidly. Air-traffic control, emergency response, trading, multiplayer collaboration, and many kinds of creative work require flexible switching. Deep attention is useful when meaning is cumulative and cannot be reconstructed from isolated fragments. The institutional failure

begins when one regime is treated as universally superior because one regime is easier to monetize or measure.

Hartmut Rosa's theory of social acceleration adds another layer. Modernity does not merely contain faster devices; it contains institutions in which technological acceleration, social change, and the pace of life reinforce one another (Rosa 2013). A faster communication channel creates expectations of faster response. Faster response creates organizational practices that assume permanent reachability. The time saved by technology can be consumed by a rise in the number of expected interactions. Acceleration can therefore produce scarcity rather than leisure.

This mechanism helps explain why the common promise that technology will give us time repeatedly disappoints. A faster email system does not necessarily reduce time spent communicating if it multiplies the expected volume of communication. An AI that writes a report in five minutes does not necessarily reduce organizational workload if the organization responds by requesting ten reports. Productivity gains are often reinvested in additional production. The same may happen to cognition: as generating analyses becomes cheap, the number of analyses requested may expand until judgment becomes the dominant bottleneck.

The distinction between seriousness and trivialization is also politically important. Cultural critics often select examples of vulgarity from mass culture while ignoring triviality among elites. Yet elite institutions can trivialize with extraordinary sophistication. A board can convert a strategic problem into a dashboard. A ministry can convert a social conflict into a performance indicator. A university can convert intellectual contribution into a ranking. A philanthropic organization can convert a complex community into a theory-of-change diagram that exists mainly to satisfy reporting requirements. Trivialization is not the victory of low culture over high culture. It is the reduction of reality to forms that a selection system can process cheaply.

This gives us a second institutional principle: protect irreducibility. Some things should resist complete conversion into metrics. Love, trust, scientific importance, civic legitimacy, education, and aesthetic value can be partially measured, sometimes usefully, but institutions should remember that the measurement is a view rather than the object itself. James C. Scott's critique of high-modernist schemes showed how administrative systems can become dangerous when they make society legible by simplifying away local knowledge and then mistake the simplified map for the territory (Scott 1998). Anti-trivialization institutions face exactly this risk. An institution created to defend seriousness can trivialize the world through its own categories.

The answer is not to abandon measurement but to make measurements contestable and plural. A school can use test results while allowing teachers to explain what the test misses. A hospital can use performance indicators while maintaining professional judgment and patient narratives. A research funder can use citation data while preserving space for expert evaluation, replication, long-horizon work, and negative results. A platform can provide popularity signals without pretending that popularity is importance.

We can now state the problem more rigorously. A civilization becomes vulnerable to trivialization when its institutions lose the capacity to maintain distinctions among different kinds of attention, different time horizons, and different standards of evidence. The problem is therefore not too much fun or too little seriousness. It is institutional monoculture.

The antidote is cognitive biodiversity. A healthy civilization needs places where speed is rewarded, places where slowness is rewarded, places where popularity matters, places where popularity is irrelevant, places where experts dominate, places where citizens deliberate, places where uncertainty is tolerated, and places where decisions must be made despite uncertainty. The function of anti-trivialization institutions is to keep these regimes from collapsing into one another.

The future anti-trivialization machine begins by restoring distinctions that technical convenience erased. Some signals should interrupt us. Some should wait. Some should require evidence. Some should disappear. Some should be preserved for a century. Some decisions should be reversible in a second; others should require a cooling-off period. The deep institutional question is no longer only who gets to speak. It is also how long a claim must survive, what memory accompanies it, what friction precedes action, and which time horizon a decision is allowed to ignore.

Chapter 2: The Old Machines of Seriousness

Modern societies already know how to manufacture seriousness. They simply do not usually describe the process that way.

Consider a court. Two people disagree about an event. Outside the institution, they might shout, threaten, recruit allies, edit videos, appeal to emotion, or exploit status. Inside the court, the disagreement is transformed. Evidence must be introduced under rules. Statements become records. A decision must be connected to law. Appeals are possible. The process can still fail spectacularly. Bias, money, inequality, bad law, and institutional capture remain. But the court is designed to convert immediate conflict into a slower and more constrained form of contestation.

Science performs a related transformation. A claim does not become scientific merely because an intelligent person believes it. Methods must be made inspectable. Evidence must be preserved. Other researchers must be able to criticize or reproduce the work. Journals, conferences, laboratories, archives, standards bodies, research ethics committees, and statistical conventions are all parts of a distributed machine for delaying premature certainty. Robert Merton's classic account of scientific norms emphasized communal evaluation, universalistic criteria, organized skepticism, and the idea that scientific knowledge exceeds private ownership even when institutions repeatedly violate these ideals (Merton 1942/1973).

Accounting performs another version. A corporation is a torrent of transactions. Without standards for classification, reporting, audit, and temporal comparison, managers and investors would face an ocean of anecdotes. Accounting creates a memory format. It makes some forms of deception harder, though never impossible. Modern financial systems could not scale merely by asking executives to be honest; they required institutions that made claims comparable across time and organizations.

Public health similarly converts invisible collective risk into persistent institutional attention. An individual may have little reason to invest in sewerage, epidemiological surveillance, vaccination infrastructure, or laboratory networks. Cities learned painfully that dense populations create externalities that private decisions cannot manage alone. Public health became a way for the future consequences of present behavior to acquire an institutional representative.

These examples matter because they reveal that seriousness is often the result of procedural artificiality. Human beings do not spontaneously keep complete records, disclose inconvenient evidence, wait for adversarial review, distinguish correlation from causation, or allocate resources to low-probability catastrophes. We build institutions that make these behaviors easier, mandatory, prestigious, or at least difficult to avoid.

Douglass North described institutions as the rules that structure interaction and reduce uncertainty, while emphasizing that institutional change is path-dependent and constrained by inherited incentives and beliefs (North 1990). Elinor Ostrom showed that collective problems need not be solved only by central states or privatization. Communities can govern shared resources through locally adapted rules, monitoring, graduated sanctions, conflict-resolution mechanisms, and nested arrangements, provided institutional conditions are right (Ostrom 1990). These insights are crucial for the attention problem because attention resembles a peculiar commons. It is individually experienced but collectively shaped. One person's distraction may seem private, yet a society's aggregate allocation of attention determines which risks are recognized, which institutions remain accountable, and which long-term projects receive political support.

The old machines of seriousness share several properties. They create boundaries. Not everyone can edit the accounting statement after publication. Not every rumor becomes courtroom evidence. Not every medical intervention becomes standard practice after one dramatic anecdote.

Boundaries are not inherently elitist; they are sometimes what allow specialized standards to exist at all.

They create memory. A legal precedent, laboratory notebook, audit trail, archive, or parliamentary record makes it harder for the present to pretend that the past never happened. Memory is one of civilization's strongest anti-trivialization technologies because trivialization thrives on reset. A system with no memory can repeat the same scandal, promise, error, and revelation indefinitely.

They create friction. Friction has acquired a bad name in software design because many digital products treat the shortest path to action as obviously desirable. Institutions often need the opposite. A waiting period, second signature, independent review, replication requirement, or public consultation can be a feature rather than a defect. Friction is expensive, and unnecessary friction can become bureaucracy. But the absence of friction has costs too. Civilization needs a theory of useful delay.

They create adversaries. Courts formalize opposition. Science rewards criticism in principle, even if social incentives complicate it in practice. Markets use competition. Democracies institutionalize disagreement through parties, opposition, journalism, courts, and periodic elections. The insight is profound: durable systems do not rely exclusively on good intentions. They create roles whose job is to say no.

They create scale-specific authority. A university department should not govern monetary policy. A central bank should not decide a local school curriculum. Ostrom's polycentricity matters because complex societies contain problems at different scales. Attempts to solve trivialization through one centralized authority would almost certainly become either impotent or authoritarian. The relevant institutions must exist at the level of individuals, families, schools, firms, professions, platforms, cities, states, and transnational networks.

They create legitimacy by exposing themselves to procedure. The judge cannot simply say, "Trust my intuition." The scientist cannot simply say, "The model feels right." The auditor cannot simply say, "The company seems healthy." Institutions become credible when they can explain how they arrived at an outcome and when others can inspect or challenge the process.

These old institutions are imperfect prototypes for what comes next. They were built in a world where the rate of symbolic production was constrained by human labor. A newspaper had a finite number of pages. A television channel had a schedule. A politician could repeat a slogan, but not generate millions of personalized rhetorical variants in real time. A citizen's social environment was mediated by institutions, but not continuously optimized by systems learning from each micro-response.

Generative AI changes the production function of persuasion. It makes text, images, video, summaries, simulations, personalities, tutors, advertisements, explanations, arguments, and counterarguments cheap to produce at enormous scale. This can democratize intelligence. It can also make the old scarcity assumptions of our institutions obsolete.

The response cannot be to restore the old gatekeepers exactly as they were. They excluded voices, concentrated power, moved too slowly, and often confused status with quality. The future machine must learn from both the success and failure of the old ones. It must preserve provenance without restoring monopoly. It must create friction without creating sclerosis. It must allow expertise without demanding deference. It must build memory without creating permanent reputational prisons. It must permit fast action without converting every issue into an emergency.

One way to understand institutional progress is as the discovery of better forms of selective friction. The challenge is never simply to slow things down. The challenge is to slow the right things, at the right moment, for reasons that can themselves be challenged.

The genealogy of these institutions is instructive because they were rarely invented in a single moment. The university emerged through layers of ecclesiastical, municipal, professional, and scholarly practice. The scientific journal grew out of correspondence networks and learned societies. Peer review did not arrive as a complete constitutional design; it developed gradually as communities needed ways to evaluate a growing volume of claims. Modern bureaucracies accumulated files, procedures, ranks, and standardized forms because states could not govern large populations through personal memory. Accounting became more formal as enterprises outgrew the owner's direct knowledge of every transaction. Institutional seriousness is usually sedimentary.

Printing illustrates the dual process especially well. The printing press lowered the cost of reproducing text and destabilized older authorities. It accelerated religious conflict, propaganda, pamphleteering, and rumor, but it also made possible new scholarly communities, standardized reference works, public libraries, scientific journals, and eventually mass education. The same technology that amplified confusion created pressure for bibliographic order. Catalogues, indexes, editorial conventions, citation practices, and archives were anti-trivialization machines for a world in which text had become abundant by earlier standards.

The pattern appears in infrastructure. Modern engineering is not merely the application of physics. It is an institutional system of professional licensing, testing, design codes, inspection, redundancy, failure analysis, maintenance schedules, and liability. A bridge is serious because thousands of small institutional constraints surround it before anyone drives across. The spectacular structure rests on invisible paperwork.

Aviation provides an even clearer example of productive institutional memory. Aviation safety emerged from accident investigation, standardized procedures, checklists, reporting systems, training, engineering redundancy, air-traffic coordination, and a culture that gradually learned to treat near

misses as information rather than embarrassment. The crucial move was to convert failure into a shared record. An accident that remains a story is tragedy. An accident that changes design becomes institutional learning.

The same conversion is central to anti-trivialization. Social-media scandals currently arrive as stories. A manipulated campaign, a viral falsehood, a catastrophic rumor, or a failure of automated decision-making generates outrage, commentary, temporary reform, and then replacement by the next event. The information environment has excellent event memory and poor mechanism memory. We remember what happened but not which structural condition made recurrence likely.

An anti-trivialization institution asks a different question after failure: what rule, incentive, interface, missing audit, or absent counter-party allowed this class of event to occur? The unit of learning is not the scandal but the mechanism.

Scientific institutions provide another useful distinction between primary and secondary cognition. Primary cognition produces a hypothesis, experiment, argument, or discovery. Secondary cognition evaluates how the knowledge was produced. Methods sections, peer review, replication, meta-analysis, systematic reviews, research ethics, statistical standards, and data repositories are secondary cognition. They are cognition about cognition. As scientific production scales, the secondary layer becomes increasingly important because no individual researcher can personally inspect the entire field.

AI will push many other institutions toward the same architecture. Organizations will need processes that do not merely produce decisions but evaluate how decisions were generated. Public agencies will need model registries, provenance records, red-team reports, and post-deployment incident analysis. Companies will need assurance functions that inspect AI-mediated workflows in the same way financial audit inspects financial reporting. Schools

will need assessment systems that evaluate learning processes rather than only finished artifacts. Journalism will need stronger source provenance because synthetic material lowers the cost of fabrication.

This suggests a third institutional principle: separate production from assurance. The person or system that benefits from declaring success should not be the only party evaluating success. Financial audit institutionalized this principle imperfectly. Courts institutionalize it through opposing parties and appeal. Science institutionalizes it through peer criticism and replication. Software security uses independent testing and red teams. The future cognitive environment will need assurance markets and public assurance institutions for claims, models, rankings, and automated decisions.

But assurance can itself become ceremonial. Corporate compliance offers many examples of procedures that exist mainly to show that a procedure exists. Institutions can optimize for auditability rather than quality. This is another form of trivialization: the process becomes the metric. Anti-trivialization therefore requires periodic attacks on its own rituals. Does the review actually catch errors? Does the waiting period improve decisions? Does the provenance record help users identify weak claims? Does the red team have incentives to challenge or merely to produce a report?

A mature institution measures its safeguards by the failures they prevent or reveal, not by the number of safeguards performed.

The old machines of seriousness thus teach two seemingly contradictory lessons. Procedure is indispensable, and procedure is dangerous. Memory is indispensable, and memory can become oppressive. Expertise is indispensable, and expertise can become monopoly. Friction is indispensable, and friction can become rent-seeking bureaucracy. The future institutions must live inside these tensions rather than pretend to solve them permanently.

This is why the anti-trivialization project is not cultural conservatism. It is institutional engineering. Its question is the same question that appears whenever a new infrastructure reshapes social life: what new rules must be invented so that the benefits of the infrastructure do not destroy the conditions that make those benefits worth having?

Chapter 3: The Industrialization of Attention

The attention economy is often discussed as if companies discovered a psychological weakness and then exploited it. This is true but incomplete. The more important development was that attention became measurable, tradable, and optimizable at scales that earlier media systems could not reach.

Advertising had long treated attention as scarce. Newspapers sold audiences to advertisers. Radio and television converted time into inventory. Public relations, propaganda, and political campaigns learned to engineer salience. Edward Bernays and Jacques Ellul approached mass persuasion from very different angles, but both understood that modern organizations could act systematically upon public perception (Bernays 1928; Ellul 1965). What digital platforms added was a fine-grained feedback loop. Instead of estimating what an audience might prefer, systems could observe what each user did next.

The technical achievement was extraordinary. Recommendation systems reduced the cost of navigating overwhelming abundance. Search engines made planetary knowledge accessible. Social graphs allowed people to maintain relationships across distance. Personalized systems helped users find music, research, products, communities, and ideas that would otherwise remain invisible. To call this only manipulation would miss its enormous utility.

But personalization also changed the competitive environment. A piece of content no longer needed to be good in any general sense. It needed to defeat nearby alternatives for a particular user at a particular moment. The unit of competition shrank from the program, newspaper, or book to the next item. The optimization horizon often shrank with it.

This produces a recurrent asymmetry. The system can measure whether a user clicked within seconds. It may need months or years to know whether the user learned something useful, became better informed,

strengthened a relationship, or regretted the time spent. Short feedback loops are technically convenient. Long feedback loops are epistemically difficult. The result is a measurement bias toward immediate proxies.

The problem is not unique to commercial platforms. Institutions everywhere prefer what can be measured cheaply. Schools use test scores. Universities use citation counts. Businesses use quarterly indicators. Governments use approval ratings. Researchers use benchmark performance. Metrics are necessary because complex organizations cannot coordinate by intuition alone. Yet metrics create their own reality. When the proxy is easier to optimize than the underlying value, organizations drift toward proxy performance.

Attention platforms simply reveal the general law in a concentrated form. A metric begins as a measurement. It becomes an incentive. The incentive becomes a selection pressure. The selection pressure reshapes production. Eventually producers forget that the metric was a proxy at all.

This helps explain why trivialization can spread into domains that were never designed as entertainment. Political communication adapts to the feed because the feed is where citizens are. News adapts because distribution depends on visibility. Academic institutions adapt because public attention influences funding and status. Firms adapt because investor narratives circulate through the same media environment. Individuals adapt because social visibility affects careers, relationships, and belonging.

The system does not force everyone to become shallow. It makes certain forms of depth more expensive. A nuanced explanation requires context. Context increases cognitive cost. Cognitive cost lowers immediate completion rates. Lower completion can reduce distribution. The infrastructure therefore taxes complexity without passing a law against it.

Neil Postman's television-era claim was that a medium has epistemological consequences because it privileges some forms of discourse over others (Postman 1985). In the platform era, the claim must be updated. The medium is not only a format. It is an adaptive selection system. It observes behavior and changes the environment in response. The medium learns.

This creates a feedback system between human preference and machine amplification. It is tempting to assign causal priority to one side. Critics blame algorithms for polarization and outrage. Defenders point out that algorithms respond to what users choose. Both can be right. A recommender system does not invent every desire it amplifies, but amplification changes the future distribution of desires, norms, and expectations. Human beings train the feed, and the feed trains human beings.

Empirical research supports caution against simple stories. Large platform experiments have found that changing feed structure can alter what users see without always producing the dramatic downstream political effects predicted by public debate. More recent experiments have shown that reranking exposure to partisan content can reduce partisan animosity under some conditions. The emerging picture is therefore neither "algorithms determine politics" nor "algorithms are irrelevant." It is that design choices change exposure environments, while social outcomes depend on network structure, identity, prior beliefs, institutions, and broader political conditions (Guess et al. 2023; Piccardi et al. 2025).

This uncertainty strengthens rather than weakens the institutional argument. We do not need to prove that one algorithm causes one social pathology before recognizing that large-scale attention allocation has become infrastructural. Financial markets are regulated not because every transaction is harmful but because the system coordinates resources at a scale where failure becomes public. Telecommunications are governed because connectivity is infrastructural. Food safety is regulated because consumers cannot individually inspect every production chain. Attention allocation increasingly

has similar properties: it shapes conditions of public knowledge that no single user can meaningfully audit.

The phrase attention market is still misleading, however, because ordinary markets assume that buyers can evaluate what they purchase with some degree of informed agency. In attention markets, the object being sold is often the buyer's own future behavior. The user is simultaneously consumer, product, training signal, and target of optimization. This creates principal-agent problems that traditional consumer protection only partly captures.

A useful comparison is the emergence of fiduciary duties. Some relationships are considered too asymmetrical to be governed by caveat emptor alone. Lawyers, trustees, directors, and other fiduciaries can acquire duties of loyalty and care because one party must entrust another with power that is difficult to monitor continuously. A future attention institution may begin from the recognition that a personal AI assistant or recommendation system can occupy a similarly intimate position. If it learns a person's goals, fears, habits, weaknesses, relationships, and decision patterns, the question is not merely whether it has a privacy policy. The question is: whose interests is it legally and technically required to serve?

The political theory of the twentieth century anticipated another aspect of the problem: citizens cannot personally inspect a complex world. Walter Lippmann argued in *Public Opinion* that modern people act upon pictures in their heads constructed through limited experience and mediated information. John Dewey accepted much of the diagnosis but resisted the conclusion that democratic publics were therefore obsolete. The challenge, for Dewey, was to construct institutions and communication forms through which dispersed consequences could become visible to those affected (Lippmann 1922; Dewey 1927).

The argument is newly relevant because the public is now surrounded by systems that can construct individualized pictures of the world at scale. Lippmann worried about newspapers and public relations; contemporary systems can select, rank, summarize, and eventually converse. The mediated picture is no longer identical for an audience. It can be adapted to a person.

This is a remarkable achievement for relevance. It is also a challenge for accountability. If every user receives a different political explanation, advertisement, search result, or recommendation sequence, public scrutiny becomes difficult. A false claim broadcast on television can be recorded and challenged. A million slightly different persuasive messages delivered privately require a different kind of institution.

Tim Wu's history of attention markets shows that each major communication medium eventually attracts business models designed to capture the audience that the medium aggregates (Wu 2016). Shoshana Zuboff later emphasized the economic importance of extracting behavioral data and using it to predict and shape future behavior (Zuboff 2019). One need not accept every element of the surveillance-capitalism thesis to recognize the institutional asymmetry: systems can learn from users at a resolution that users cannot reciprocate.

Yves Citton's idea of an ecology of attention provides a useful correction to purely individual language. Attention is not simply a resource inside each skull. It circulates socially. Institutions direct it by setting agendas, prestige systems, curricula, rankings, interfaces, and rituals (Citton 2017). When a scientific community decides which problems are fundable, it structures attention. When a news organization decides which events merit coverage, it structures attention. When a search engine orders results, it structures attention. When a school timetable assigns forty minutes to one subject and none to another, it structures attention.

The industrialization of attention therefore predates the smartphone. The smartphone made the industrial process intimate and continuous.

This continuity matters for reform. It is tempting to imagine that replacing one platform or algorithm will restore some natural baseline. There is no natural baseline. Every information system selects. A chronological feed selects by recency. An editor selects by professional judgment. A library selects through acquisition and cataloguing. A search engine selects by ranking. A teacher selects a syllabus. The real governance question is not whether selection occurs but whether the selection rule is visible, contestable, aligned with the institution's declared purpose, and subject to correction.

This is where contemporary debates about algorithms often become confused. Algorithmic selection is sometimes criticized merely because it is algorithmic, while human editorial judgment is romanticized as neutral. Human gatekeepers have biases, incentives, blind spots, and social networks too. Conversely, defenders of algorithms sometimes treat optimization as objective because it is mathematical. Mathematics only makes an objective precise after someone chooses it.

The institutional focus should therefore move from algorithm versus human to objective, accountability, and feedback. What is the system trying to optimize? Who benefits from the objective? Which outcomes are measurable? Which valuable outcomes remain invisible? How quickly does the feedback arrive? Who can challenge the rule? How does the system learn from long-term harm that was absent from short-term metrics?

This last question may become central to twenty-first-century governance. Short feedback is technically easy. Long feedback is politically and economically difficult. A platform knows immediately whether a user clicked. A school knows quickly whether a student completed an exercise. A company knows this quarter's revenue. A government knows today's polling. But the deeper outcome may appear years later: did the user become better informed,

did the student retain knowledge, did the company destroy future resilience, did the policy build trust?

Institutions of seriousness are, among other things, machines for incorporating delayed feedback.

A pension fund treats future obligations as present facts. Depreciation accounting treats the future exhaustion of an asset as relevant now. Environmental regulation treats delayed damage as a present governance problem. Scientific replication allows time to revise premature findings. Constitutional review prevents every transient majority from rewriting foundational rules. The attention economy needs equivalent mechanisms for cognitive externalities whose effects appear slowly.

One might imagine an attention-impact assessment for high-reach systems, analogous in spirit to privacy, safety, or environmental assessments. The aim would not be to calculate a single score called "attention quality." That would invite immediate gaming. The assessment would instead ask what the system optimizes, which populations may be vulnerable, which long-term outcomes are not represented in the objective, what user controls exist, how researchers can study effects, and what signals would trigger redesign.

Such procedures will sound bureaucratic until a major cognitive infrastructure failure makes them obvious. Institutional history repeatedly follows this sequence. The absence of a safeguard seems efficient before the crisis and negligent after it.

This is where the industrialization of attention creates its own counter-pressure. As systems become more powerful, the demand for trustworthy delegation becomes stronger. Users will increasingly want software that does not merely capture their attention but negotiates on behalf of it. Organizations will want cognitive environments that protect employees from information overload. Schools will want learning systems that optimize

for durable understanding rather than completion. Democracies will want public information infrastructures that make evidence traceable without placing one authority in charge of truth.

The market for attention therefore contains the seed of a market for attention protection. The first generation of digital systems competed to occupy the human mind. The next generation will increasingly compete to defend it.

Chapter 4: The Failure Is Institutional, Not Moral

Every age explains structural failures as failures of character before it learns to see the structure.

Early industrial accidents were often attributed to careless workers. Epidemics were attributed to dirty individuals or immoral populations. Financial collapse was blamed on greed as though greed had been invented that year. Environmental degradation was treated as a collection of bad choices rather than an externality produced by an incentive system. Moral explanations are attractive because they identify a villain quickly. They are also institutionally lazy.

The attention crisis is undergoing the same transition. People are told to use more discipline, disable notifications, read books, meditate, remove applications, create morning routines, and keep phones out of bedrooms. Much of this advice is sensible. But it has the same limitation as telling factory workers to be careful around unguarded machinery. Individual adaptation cannot substitute for institutional design when the environment is systematically optimized against the individual's stated goals.

This does not absolve agency. Institutions can only work through people, and no constitution can force citizens to care about what they have decided not to care about. The point is different. Human self-control is variable, contextual, and exhaustible. We should design environments that do not require heroic self-control as the normal price of participation.

Behavioral economics helped make this idea respectable in policy. People are not frictionless rational agents. Defaults, framing, salience, timing, social norms, and cognitive load influence decisions. Thaler and Sunstein popularized the idea of choice architecture: environments can be designed so that predictable behavioral tendencies lead to better outcomes without eliminating choice (Thaler and Sunstein 2008). The concept is useful but

becomes more complicated in commercial digital systems because the architect may have interests different from the user's.

The future problem is therefore not choice architecture but choice-architecture governance. Who chooses the defaults for the systems that choose what we notice? Who audits the objectives? What rights do users have over recommendation criteria? Can a person delegate attention management to an agent that is contractually forbidden from monetizing distraction? Can a school require an educational AI to disclose whether it is optimizing for engagement, completion, retention, transfer, or some blend? Can a public institution use AI-generated summaries without preserving access to the source material and dissenting interpretations?

These questions sound technical only because the institutions have not yet acquired familiar names. Questions about audit committees, food labeling, fiduciary duties, medical consent, and environmental impact assessments once sounded specialized too. Institutional maturity is partly the process by which strange governance mechanisms become obvious expectations.

The moral language of distraction also fails because it misidentifies the location of power. A platform can run millions of experiments. An individual cannot. A platform can observe aggregate behavior. An individual sees only a personal experience. A platform can change the interface overnight. An individual can leave, but often at the cost of losing a social or professional network. This asymmetry does not imply abuse in every case. It implies that voluntary choice alone may not produce a fair bargaining environment.

The same asymmetry appears inside organizations. Employees are often expected to manage a flood of messaging channels, meetings, dashboards, documents, notifications, and AI-generated material. The organization then treats the resulting concentration problem as a personal productivity issue. Yet the number of interruptions is produced collectively.

When every employee optimizes personal responsiveness, the organization can destroy its shared capacity for deep work.

This is a commons problem. Each message is locally reasonable. The aggregate can be cognitively catastrophic. Ostrom's work is valuable here because it directs attention away from moral exhortation and toward rules. Communities can manage shared resources when boundaries, monitoring, participation, conflict resolution, and nested governance are designed well (Ostrom 1990). Cognitive bandwidth inside an organization can be treated similarly. Teams can establish protected time, communication protocols, escalation thresholds, decision records, and norms about synchronous versus asynchronous work. What looks like personal discipline becomes institutional infrastructure.

The same principle applies at societal scale. No citizen can personally maintain context across thousands of scientific, legal, economic, and geopolitical questions. Representative institutions exist partly because attention cannot scale. Expertise exists because attention cannot scale. Journalism exists because attention cannot scale. Archives exist because memory cannot scale. The division of cognitive labor is not a defect; it is the foundation of complex civilization.

The danger arrives when the institutions that perform this cognitive division of labor are themselves optimized by the same short-horizon signals they are supposed to filter. A newspaper that must chase every attention spike can no longer function as a filter. A university that must continuously market itself may begin treating visibility as evidence of quality. A regulator that communicates through platform dynamics may start selecting priorities by public intensity rather than risk. A scientific institution that treats citation velocity as importance may reward fashionable errors over slow verification.

We should therefore stop asking whether citizens are sufficiently serious and ask whether institutions are capable of preserving seriousness under competitive pressure.

This reframing also makes optimism possible. Character is hard to redesign at scale. Institutions are difficult to redesign too, but history shows that societies do it repeatedly. We do not need a new species of human. We need better rules for old humans operating inside new systems.

The most powerful institutional reforms are often boring. Double-entry bookkeeping changed economic coordination more profoundly than many political speeches. Standardized shipping containers changed trade through geometry and agreement. Randomized clinical trials changed medicine by imposing a disciplined comparison on human intuition. Building codes quietly save lives because engineers decided that certain failures should not require individual vigilance.

Organizational theory offers a language for this reframing. Herbert Simon and James March treated organizations as systems of bounded rationality. People do not optimize across every possible option; they operate with limited information, limited time, routines, hierarchies, and decision premises that determine which problems even reach attention (March and Simon 1958; Simon 1947). Organizations are attention-distribution systems before they are decision systems.

A chief executive cannot examine every customer complaint, security alert, research result, employee concern, and market signal. The organization must filter. Departments, dashboards, meetings, reporting lines, escalation rules, and professional roles decide what becomes visible. Institutional failure often occurs not because nobody knew, but because the system failed to route knowledge to a place where action was possible.

This is why future anti-trivialization institutions cannot focus only on public media. The internal attention architectures of governments, laboratories, hospitals, companies, and universities are equally consequential. A hospital can have all the relevant data and still miss a deteriorating patient if alerts are badly prioritized. A company can employ brilliant security engineers and still ignore a vulnerability if reporting lines reward product launch speed. A government can commission excellent analysis and still fail if inconvenient findings do not survive political filtering.

Information overload is therefore not merely too much information. It is a mismatch between information production and institutional routing capacity.

AI will initially worsen this mismatch in many organizations because it makes reports, summaries, forecasts, presentations, code, and recommendations cheap. Every department can become a high-volume producer of plausible material. The traditional response to scarcity - produce more information - becomes counterproductive. The organization will need a new discipline: information budgets.

An information budget does not mean a quota on thought. It means treating the attention cost imposed on others as a real organizational cost. Before sending a report to five hundred people, the sender's workflow may need to classify who must act, who merely needs archival access, and who should not be interrupted. Before scheduling a meeting, the system may require an explicit decision question. Before escalating an alert, the organization may specify which threshold has been crossed and what response is requested.

These rules can sound petty until multiplied by thousands of interactions. Collective attention resembles network congestion. Each sender experiences the benefit of sending; the cost is distributed across receivers. The equilibrium is overproduction.

Economists call this an externality. The correct response to an externality is rarely moral shame alone. It is to alter prices, rights, norms, or rules so that decision-makers experience more of the cost they impose. In digital organizations the "price" need not be money. It can be friction, priority limits, batching, reputation effects, or enforced asynchronous channels.

This provides a concrete application of Ostrom's insights. Teams can define boundaries around attention, participate in creating communication rules, monitor the shared resource, and revise protocols locally. A research group may need different rules from a customer-support center. A surgical team may require immediate interruptibility; a theoretical mathematics group may require long protected periods. Polycentric governance begins inside the office.

The moral explanation also obscures inequality. People with power can often buy protection from distraction. Executives have assistants. Wealthy professionals can purchase quiet space, childcare, travel flexibility, and premium services. Lower-status workers are often required to remain continuously responsive to monitoring systems, customers, managers, and scheduling software. Attention autonomy can become a class privilege.

Future labor institutions may therefore treat cognitive conditions as part of working conditions. Occupational safety expanded from obvious physical hazards toward ergonomics, stress, and psychosocial risk. The next frontier may include interruption load, algorithmic monitoring, and expectations of constant availability. This should be approached cautiously; not every unpleasant job feature is a regulatory violation. But the principle that work environments can damage cognitive capacity is no stranger than the earlier realization that they can damage hearing, lungs, or posture.

The institutional reframing also changes the rhetoric of self-help. Personal practices remain valuable. Reading, meditation, device boundaries, exercise, sleep, and deliberate solitude can protect attention. But a society

should be suspicious when the solution to every structural problem is sold back to individuals as a subscription, app, retreat, or productivity method. The attention economy can monetize the cure to the attention economy.

The better aim is subsidiarity: solve what can be solved by personal practice personally, what requires organizational rules organizationally, and what creates public externalities publicly. No single level should absorb the entire burden.

The moral language of seriousness then becomes more generous. People are not failing because they are insufficiently noble. They are adapting to environments with particular incentives. Change the environment, and different parts of human nature become rational to express.

The future anti-trivialization machine may be similarly unromantic. It may look like a protocol, an audit requirement, a right to inspect recommendation logic, a cooling-off rule, a provenance standard, a professional duty, an interface convention, or a new category of institution that currently lacks a name. Its importance will not depend on its drama.

The great institutional achievements of civilization often become invisible precisely because they work. We should hope that the same fate awaits the institutions of attention. Success would mean that future citizens find it strange that systems once competed freely to manipulate the scarce cognitive resources on which every other institution depended.

Chapter 5: AI and the Great Cognitive Surplus

Generative AI intensifies the attention problem and simultaneously changes its solution space.

The first digital revolution reduced the cost of transmitting information. The second reduced the cost of discovering and ranking information. Generative AI reduces the cost of producing plausible symbolic material. Text, images, software, explanations, plans, simulations, arguments, lessons, summaries, synthetic personalities, and increasingly video and interactive environments can be generated on demand.

This is an extraordinary expansion of human capability. It can give a student access to explanation at any hour, allow a small company to perform tasks once requiring a large staff, help researchers navigate literatures, translate between languages and disciplines, assist people with disabilities, accelerate software engineering, and make expertise more accessible. It may become one of the great democratizing infrastructures of cognition.

But cheap cognition creates a new scarcity. If answers become abundant, judgment becomes scarce. If drafts become abundant, editorial selection becomes scarce. If arguments become abundant, the ability to compare arguments becomes scarce. If images become abundant, provenance becomes scarce. If personalized persuasion becomes abundant, trusted mediation becomes scarce.

The AI age will therefore not eliminate the attention economy. It may deepen it by creating a cognitive surplus that exceeds human capacity to inspect.

The familiar response is verification. Humans should check AI outputs. But this cannot be the final institutional model. If AI systems produce material thousands of times faster than humans can verify it,

verification becomes the bottleneck. The economic pressure will then be to reduce verification, automate verification, or accept a higher error rate. The question becomes how to build verification institutions rather than merely asking individuals to be careful.

Research on knowledge work already points toward this transition. A 2025 study of 319 knowledge workers found that people described critical thinking around AI in terms of goal setting, evaluation, verification, and integration, while greater confidence in AI was associated with lower reported critical-thinking effort in some contexts. The result should not be read as evidence that AI makes people unintelligent. It suggests that cognitive labor moves. The system may perform generation while the human becomes a supervisor, editor, verifier, or goal-setter (Lee et al. 2025).

This division of labor can be excellent when the supervisory layer is real. It can be dangerous when supervision becomes ceremonial. Aviation offers an analogy. Automation can improve safety dramatically, but highly reliable automation creates a paradox: human operators may be asked to intervene only when the system encounters unusual situations, precisely when their skills are least exercised and the context most difficult. Human-in-the-loop is not automatically meaningful simply because a human remains formally present.

The same problem will appear in science, medicine, law, management, education, and government. A researcher may approve an AI-generated literature synthesis without reading the underlying papers. A lawyer may sign a document assembled by systems whose sources are poorly traceable. A manager may accept a risk assessment because the dashboard looks rigorous. A citizen may receive a personalized explanation of policy that is accurate enough to feel trustworthy but optimized to persuade.

The institutional response should not be a blanket prohibition on automation. It should be an architecture of accountable delegation.

Delegation is one of civilization's oldest solutions to complexity. We delegate medical judgment to physicians, legal representation to lawyers, financial management to institutions, political decision-making to representatives, and technical inspection to engineers. What makes delegation tolerable is not trust alone. It is the existence of professions, standards, liability, documentation, review, competition, and rights of appeal.

AI needs analogous institutional layers.

One future institution is the cognitive fiduciary: an AI system or service with a legally meaningful duty to protect the user's declared interests in attention and decision-making. Today many digital services face an incentive conflict. The user wants to achieve a goal; the service may profit when the user remains engaged. A cognitive fiduciary would reverse the relationship. Its success could be measured by whether the user accomplished intended tasks, preserved chosen boundaries, understood relevant uncertainty, and maintained control over delegation. It would be forbidden from secretly optimizing toward advertiser or platform objectives that conflict with those goals.

A second institution is provenance infrastructure. In a world of cheap generation, every consequential claim should become easier to trace backward. Not every sentence needs a citation, and not every conversation should become a bureaucratic record. But high-stakes domains will increasingly need machine-readable chains connecting claims to sources, transformations, models, data, and human approvals. Provenance is the cognitive equivalent of a supply chain label.

A third institution is adversarial AI review. High-stakes systems should not merely generate an answer and a confidence score. They should face structured opposition. One system can search for missing evidence, another for contradictions, another for affected stakeholders, another for

historical analogies, another for incentive conflicts. The goal is not to simulate a theatrical debate but to turn disagreement into a designed function.

A fourth institution is the decision record. Organizations currently lose enormous amounts of context when choices are made in meetings, messaging channels, or private conversations. AI makes it cheap to preserve the reasons behind decisions, the alternatives considered, the uncertainties acknowledged, and the conditions that would justify reversal. This can create an institutional memory far richer than minutes or document archives. The danger is surveillance and bureaucratic overload, so governance must specify what is recorded, who controls it, and when it expires.

A fifth institution is the right to cognitive legibility. As AI systems mediate more decisions, individuals may need a practical right to understand which parts of a process were delegated, what objective the system optimized, and what uncertainty remains. This should not be confused with a demand to expose every parameter of a neural network. Institutional legibility is about the decision process, not necessarily the internal mathematics of the model.

These institutions may sound ambitious. Yet the alternative is equally ambitious: to let the most powerful cognitive infrastructure ever built evolve under rules designed for ordinary software products and advertising markets.

AI also creates the possibility of anti-trivialization systems that were previously too expensive. A citizen could have a personal agent that monitors information sources, detects repeated stories, distinguishes new facts from recycled outrage, maintains a map of unresolved questions, and asks whether a current controversy changes any prior belief. A research team could maintain a living evidence graph rather than repeatedly reconstructing literature context. A parliament could use independent AI systems to generate competing summaries of proposed legislation and expose assumptions before debate. A school could optimize not for time-on-platform but for durable mastery, curiosity, and the student's eventual independence from the tool.

The crucial design principle is emancipation from the tool. A system that maximizes its own use is not automatically aligned with its user. The best educational AI may become less necessary as the learner grows. The best attention agent may make itself invisible. The best public-information system may reduce the number of things a citizen must read rather than maximizing engagement with government content.

The most important institutional consequence of generative AI may be that cognition becomes modular. Instead of one person reading, reasoning, drafting, checking, and deciding in a single mental process, tasks can be divided among humans and machines. One model retrieves. Another summarizes. Another calculates. A third criticizes. A human defines the question. Another human signs the decision. The resulting cognitive chain resembles a supply chain.

Supply chains became governable only when societies developed standards for provenance, quality control, liability, certification, and recall. Cognitive supply chains will need analogous mechanisms. If a medical recommendation passed through five models and three databases, which component introduced the error? If a public report contains a fabricated citation, where did the provenance break? If a strategic decision used a synthetic scenario, who validated its assumptions? The final fluent paragraph hides the production process unless the system preserves it deliberately.

This is why the future of trustworthy AI may look less like a single superintelligent assistant and more like institutionalized separation of cognitive powers. One component proposes. Another tests. Another tracks sources. Another represents policy constraints. Another searches for counterexamples. A human or authorized body resolves remaining value conflicts. The architecture resembles constitutional government because the problem is similar: prevent one powerful actor from defining, executing, and reviewing its own decision without challenge.

There is a temptation to call this inefficiency. Why use five systems when one system can answer immediately? For low-stakes tasks, the immediate answer may be correct. Institutional architecture becomes valuable when the cost of undetected error is high or when the system's power is large. We do not demand a jury trial before choosing a restaurant. We do demand more procedure before imprisoning someone.

AI governance will need the same proportionality. A trivial spelling correction can be automatic. A medical diagnosis, credit denial, scientific claim, or public policy recommendation may require stronger provenance and review. The challenge is to define stakes without freezing innovation.

The most interesting future institution may therefore be a market for cognitive assurance. Independent organizations could evaluate AI-mediated processes, test model behavior under domain-specific conditions, verify provenance, examine failure reports, and certify whether organizations maintain appropriate human and automated review. Some functions will resemble audit, some laboratory testing, some cybersecurity penetration testing, and some scientific peer review.

Such assurance should not become a license for a few large firms to control entry. Open standards and multiple assurance providers will be essential. Small organizations need proportional methods. Some checks can be automated. Some evidence can be shared publicly. High-risk claims can face stronger requirements than ordinary creative work.

AI also changes the economics of adversarial reasoning. Historically, serious criticism is expensive. A researcher may spend months replicating a result. A legal team may spend weeks constructing a counterargument. A public agency may lack resources to model an alternative policy. Generative systems can lower the cost of generating structured opposition. This does not guarantee good criticism, but it creates an institutional opportunity.

Imagine that every major organizational decision automatically triggers an adversarial pass. One agent is tasked with identifying assumptions. Another searches for historical analogies in which similar plans failed. Another represents the interests of a stakeholder group that is absent from the meeting. Another checks numerical consistency. Another tries to construct the strongest case that the problem has been framed incorrectly. The outputs are not treated as truth. They are a checklist of intellectual attack surfaces.

The technique could fail by producing predictable theater. Models can share blind spots. They can generate endless objections, creating paralysis. They can be tuned to satisfy the organization's desired answer. This is why institutional diversity matters more than the number of agents. A genuinely adversarial system requires different data, incentives, or oversight, not merely different prompts to the same underlying model.

Personal AI creates a parallel opportunity. The first generation of assistants will often be optimized around convenience: answer, schedule, summarize, recommend, buy. The more interesting second generation may become guardians of cognitive integrity. A user could ask an agent to distinguish "things I enjoyed seeing" from "things I want to spend my life attending to." The distinction sounds philosophical but is technically implementable as preference over timescales.

The user may declare long-term projects, relationships, learning goals, health priorities, civic interests, and periods of desired quiet. The agent could then negotiate with other systems. It could delay low-priority requests, consolidate news into slower digests, flag when content is repetitive rather than novel, and detect when a recommendation exploits a momentary impulse inconsistent with stated goals. The agent would not make the user virtuous. It would make temporal conflict legible.

This is the point at which privacy is no longer enough. A perfectly private assistant can still manipulate. A secure system can still optimize the

wrong goal. The institutional concept we need is alignment of delegated agency. Fiduciary law is one historical template because it governs situations in which one actor exercises discretionary power on behalf of another.

There will be hard questions. What if the user's long-term goals are self-destructive? What if the user asks the agent to exclude political views? What if parents control a child's attention agent? What if employers provide the assistant? Who owns the memory? What happens after death? Can the agent testify about the user? These questions show that personal AI is not simply a software category. It is a new social role.

Institutions arise when roles become powerful enough to require norms.

This is the positive inversion at the heart of the AI transition. The same technology that can automate trivialization can automate parts of anti-trivialization. Synthetic abundance creates a market for trusted scarcity. Infinite drafts create demand for selection. Personalized persuasion creates demand for personal defense. Automated generation creates demand for automated criticism. The machine contains its counter-machine.

Chapter 6: Why the Counter-Machines Will Appear

It is risky to call any social development inevitable. Institutions do not emerge because history wants balance. They emerge because groups experience costs, organize interests, invent workable rules, and acquire enough power to impose or normalize them. Many societies tolerate destructive equilibria for generations. Some collapse before solving them.

Yet there is a weaker and more defensible form of inevitability. If a complex society wants to preserve a capacity while the side effects of that capacity threaten other indispensable systems, pressure for counter-institutions will increase. The precise institution is contingent. The function becomes difficult to avoid.

Industrial capitalism illustrates the pattern. Mechanized production created enormous wealth and severe accidents, pollution, labor conflict, and urban disruption. The eventual response included safety standards, labor law, factory inspection, insurance, zoning, environmental regulation, and new professional disciplines. None emerged from a harmonious historical dialectic. They were fought over. Employers resisted some and demanded others. Workers organized. Engineers professionalized. Courts adapted. Governments learned. Catastrophes accelerated reforms that had previously appeared excessive.

Karl Polanyi gave one of the most useful names to this kind of historical rhythm. In *The Great Transformation*, he described a "double movement": the expansion of market organization generates countermovements that attempt to protect society from forms of dislocation that an unrestricted market can produce (Polanyi 1944). The concept has often been debated, and it should not be converted into a law that guarantees benevolent correction. Countermovements can be illiberal, exclusionary, nationalist, captured, or simply ineffective. Yet Polanyi's formulation remains

valuable because it rejects the fantasy that a major institutional expansion can indefinitely externalize its social costs without generating organized responses.

The attention economy may be entering an analogous phase. The first movement has been extraordinary: more information, more expression, more personalized communication, lower distribution costs, ubiquitous measurement, and increasingly capable systems for predicting what will hold a person's attention. The countermovement will not necessarily oppose these gains. It will attempt to make them socially survivable. We should expect demands for cognitive autonomy, provenance, algorithmic due process, limits on manipulative optimization, rights of exit and portability, independent assurance, protected spaces for children, and new professional duties around delegated AI agency. Some demands will be wise and others clumsy. The important point is that the pressure is endogenous to the success of the first movement.

This is the closest institutional equivalent to the yin-yang intuition that motivates the optimistic stance of this book. The complement is born inside the excess. But the metaphor must remain disciplined. Yin does not automatically emerge because history prefers balance; people have to build it. The brighter the capability, the clearer the shadow it casts. The shadow then becomes information. It reveals missing rules, missing professions, missing liabilities, missing rights, and missing forms of public infrastructure. Institutional progress begins when those absences become nameable.

Finance shows the same co-evolution. Credit expands economic activity but also creates systemic fragility. Repeated crises generated institutions that manage settlement, liquidity, disclosure, deposit insurance, capital requirements, and lender-of-last-resort functions. These institutions remain controversial and imperfect. But an advanced financial system without counter-institutions is no longer considered a serious design.

Medicine offers an even sharper example. Powerful interventions increase both the possibility of cure and the possibility of harm. Modern medicine therefore accumulated licensing, professional standards, clinical trials, pharmacovigilance, ethics review, informed consent, and post-market surveillance. The more powerful medicine became, the less acceptable unstructured trust became.

AI and attention systems follow the same logic. The more deeply they mediate cognition, the harder it becomes to treat their failures as private consumer problems. A recommender system that chooses music is one thing. A system that mediates education, employment, political information, medical access, credit, and professional judgment becomes infrastructure.

The counter-machine appears first as inconvenience. Users complain. Professionals create informal norms. Organizations introduce internal controls. Researchers produce evidence. Insurers ask questions. Courts encounter cases. Standards bodies propose vocabulary. Regulators discover categories. New businesses sell protection against the problems created by older businesses. Eventually an institutional layer becomes ordinary.

The history of cybersecurity is instructive. Early computing treated security as a technical specialty. Networked systems made it an organizational problem. Repeated breaches made it a board-level, regulatory, insurance, and national-security problem. Security functions became embedded in software development, procurement, compliance, identity infrastructure, auditing, incident response, and professional certification. No single institution solved cybersecurity. A distributed ecology emerged because the risk appeared everywhere.

Anti-trivialization will likely develop similarly. There will not be one law or one agency. There will be product standards, professional duties, educational norms, organizational practices, market products, public institutions, civic organizations, technical protocols, and cultural expectations.

The transition will be messy because the object being governed - attention and judgment - is more intimate than a network port or financial ratio.

The political resistance will also be substantial. Systems that profit from attention extraction will resist restrictions. Citizens will resist paternalism. Governments may attempt to use anti-disinformation institutions to protect themselves from legitimate criticism. Employers may use productivity systems to intensify surveillance under the language of focus. Schools may use "attention protection" to suppress exploration. Any institution powerful enough to protect cognitive autonomy can also threaten it.

This is why anti-trivialization must be designed around constitutional suspicion. The cure can become worse than the disease. The same concern that motivates protection from manipulation should motivate protection from protectors.

A second intellectual lineage is therefore as important as the lineage of institutional construction: the critique of institutions themselves. Michel Foucault showed how institutions that present themselves as rational, corrective, therapeutic, or educational can also become technologies of discipline and normalization (Foucault 1975). Ivan Illich warned that professional systems can cross a threshold at which institutions created to serve human capacities begin to expropriate those capacities from the people they serve (Illich 1973). Langdon Winner later made the political content of technical arrangements explicit: artifacts and infrastructures can embody forms of authority even when no statute announces them as political choices (Winner 1980). These critiques are not arguments against institution building. They are design requirements for the next generation of institutions.

The future cognitive institution must therefore be able to answer an uncomfortable question: how can it protect attention without claiming ownership of attention? A school can defend concentration while still

allowing intellectual wandering. A platform can protect a user from manipulative optimization while preserving the user's right to choose strange, intense, or unpopular interests. A public provenance system can make the history of a claim visible without assigning a centralized truth score. An AI fiduciary can refuse conflicts of interest without deciding what kind of life its user should live. The institution becomes legitimate not because its goal is noble but because its powers are bounded, contestable, observable, and reversible.

The yin-yang structure here is not moral symmetry. It is recursive governance. Every counter-institution creates a new concentration of power, which creates a need for review, competition, exit, transparency, or appeal. Antitrust authorities can be captured. Scientific journals can become gatekeepers. Regulators can protect incumbents. Archives can preserve injustice as effectively as truth. Seriousness can harden into orthodoxy.

The future therefore belongs not to the strongest anti-trivialization machine but to ecosystems of machines that can correct one another.

This is the deep lesson of polycentric governance. Ostrom's work rejected the assumption that complex shared-resource problems require either a single sovereign controller or pure privatization. Multiple centers of authority can overlap, learn, compete, and coordinate (Ostrom 1990). The same principle is appropriate for cognitive infrastructure. A platform may provide one layer of moderation, independent auditors another, user-controlled agents another, courts another, professional associations another, public-interest research another, and local communities yet another.

Such redundancy will look inefficient. It is inefficient in the same sense that separation of powers is inefficient. Efficiency is not the only value in systems where error and capture can propagate widely. Sometimes duplication is the price of resilience.

Counter-institutions also tend to emerge through economic specialization. Once a risk becomes salient, some actors discover that reducing the risk is itself valuable. Insurance companies hire engineers because preventing fires can be cheaper than paying claims. Financial markets support rating agencies, auditors, custodians, and clearing institutions because trust and verification enable more exchange. Cybersecurity became an industry because interconnected systems created demand for intrusion detection, identity management, incident response, and independent testing.

Attention protection will likely develop a similar service ecology. Organizations will purchase tools that reduce information overload, verify generated material, audit recommendation systems, and monitor cognitive risk. Schools will purchase systems that can demonstrate learning rather than engagement. Professionals will demand agents that separate useful information from strategically generated noise. Users may pay precisely for services whose business model does not depend on selling their attention onward.

This market response is encouraging but insufficient. Markets are good at solving problems when beneficiaries can pay and outcomes can be contracted. Public epistemic goods are harder. A citizen may benefit when everyone has access to a trustworthy public archive, but no individual has reason to fund the full system. A democracy may benefit from shared provenance infrastructure even when no platform can capture the entire return. This is where public and commons institutions become necessary.

Institutional emergence therefore follows several channels at once: private demand, professional self-regulation, civil-society experimentation, insurance requirements, litigation, standards, procurement, and legislation. The process is messy because each channel sees a different failure mode.

Douglass North's emphasis on path dependence explains why apparently inferior arrangements can persist. Once organizations invest in a

particular metric, platform, profession, or regulatory category, changing it creates transition costs (North 1990). Institutional reform therefore often arrives not when everyone becomes convinced by an argument but when the cost of maintaining the old arrangement exceeds the cost of transition.

This helps explain why crises accelerate change. A crisis does not necessarily teach the correct lesson. It changes the political price of inaction. Ideas that sounded premature become available because the status quo loses legitimacy. The challenge for institutional thinkers is to design plausible alternatives before the crisis chooses among improvisations.

The financial historian Charles Kindleberger described how recurring crises expose the need for stabilizing functions even when societies argue over who should perform them (Kindleberger and Aliber 2005). A similar logic applies here. We may not know whether attention auditing should be performed by regulators, professions, independent labs, platforms, or user collectives. But if cognitive infrastructure repeatedly produces costly failures, the demand for the function becomes clearer than the preferred organization.

This is why the book speaks of institutions we do not yet know how to build. Naming the function is often the first step. Before deposit insurance existed, the function was protection against bank-run dynamics. Before aviation accident boards existed, the function was independent learning from crashes. Before data-protection authorities existed, the function was governance of information asymmetry. Institutional imagination begins by identifying missing functions without pretending to know the final bureaucracy.

The same humility should govern claims of inevitability. Some counter-machines will be captured by incumbents. Some will become ideological. Some will over-regulate. Some will collapse under technological change. A few may cause harms worse than those they prevent. Institutional

progress is not a march toward equilibrium; it is an evolutionary search with path dependence, conflict, and occasional regression.

The positive claim is narrower. Societies are extraordinarily inventive when a valuable capability becomes unusable without governance. We did not abandon aviation because aircraft can crash. We built an ecosystem that made flying acceptably safe. We did not abandon medicine because powerful drugs can kill. We built evidence, licensing, and surveillance systems. We did not abandon finance because credit can cause panics. We layered stabilizing institutions around it.

The same logic should make us skeptical of two extreme positions about AI and attention. One says the technology will inevitably destroy human agency. The other says markets and individual choice will automatically correct every problem. Historical experience supports neither. Powerful systems persist because they are useful, and they become governable because usefulness creates constituencies for correcting their worst side effects.

The road between usefulness and governability can be ugly. That is not a reason for despair. It is a reason to build intellectual and institutional prototypes early.

The strongest reason for optimism is not that societies are wise. It is that failure becomes expensive. When executives discover that AI-generated work creates hidden liability, they will pay for verification infrastructure. When universities discover that students can produce fluent answers without learning, they will redesign assessment. When citizens discover that every public claim can be synthesized convincingly, provenance will become valuable. When organizations discover that employees cannot distinguish critical information from generated noise, attention management will become operational infrastructure rather than wellness advice.

Markets will create some of these responses because protection becomes valuable. States will create others because market incentives remain misaligned. Professions will create others because credibility becomes a competitive asset. Communities will create still others because centralized systems will never understand every local context.

The path will not be pleasant because institutions are usually built from visible failure. Society learned fire codes from fires, aviation safety from crashes, financial controls from fraud and panic, and research ethics from abuses that should never have occurred. The attention and AI transitions will probably produce their own painful curriculum.

Optimism does not require pretending otherwise. It means believing that pain can become institutional memory rather than recurring spectacle.

Chapter 7: The Attention Constitution

Every constitution begins by deciding which forms of power require boundaries. The emerging digital environment forces a new constitutional question: when systems can continuously model and influence a person's attention, what powers should they be allowed to exercise by default?

The phrase attention rights may sound extravagant until compared with privacy. For much of history, privacy was protected by physical constraints. A stranger could not cheaply record every conversation, infer relationships, search years of correspondence, and recombine behavioral traces. Digital technology removed those practical barriers, forcing legal systems to make explicit what architecture had once protected implicitly.

Attention is undergoing the same transition. Earlier media could persuade, advertise, and propagandize, but they could not adapt continuously to individual response at the same resolution. AI increases this capacity further. A system can infer which style persuades a particular user, when the user is tired, what topics create anxiety, which social identities matter, and what sequence of messages is most likely to produce action. The relevant power is not merely speech. It is adaptive influence.

A future attention constitution would not outlaw persuasion. Persuasion is inseparable from politics, education, commerce, art, and ordinary relationships. The goal would be to identify asymmetries so extreme that ordinary consent becomes insufficient.

One principle is loyalty. Systems acting as personal assistants should be able to adopt a duty to the user that is stronger than ordinary terms of service. If a person asks an agent to help prepare for an examination, the agent should not secretly optimize for keeping the person online longer. If it helps manage news, it should not rank items according to undisclosed commercial

payments that conflict with the user's declared priorities. If it negotiates purchases, it should disclose commissions that alter recommendations.

A second principle is purpose separation. Data collected to provide one cognitive function should not automatically become a resource for manipulating another. A calendar assistant does not need unrestricted authority to infer emotional vulnerability for advertising. A learning system does not need to convert educational mistakes into commercial profiling. Institutions already use purpose limitation in privacy law; attention governance will require an analogous concept for influence.

A third principle is inspectable objectives. Users do not need access to every technical parameter, but they should be able to know the main criteria a system is optimizing. Is a feed trying to maximize time spent, predicted satisfaction, information diversity, declared goals, revenue, or some combination? The objective function is a constitutional fact.

A fourth principle is exit with continuity. Digital dependence makes exit difficult when leaving a platform means losing relationships, history, reputation, or professional identity. Attention governance therefore intersects with interoperability and portability. A person should be able to move social and cognitive context between services without starting life again from zero. Real exit disciplines power only when exit is practical.

A fifth principle is cognitive quiet. Modern law recognizes limits on intrusion into physical spaces and working time. Future institutions may need a stronger right to periods free from automated solicitation, ranking, and behavioral optimization. This need not be a universal legal ban. It could appear through employment law, educational standards, device protocols, or user-controlled agents that negotiate interruptions.

A sixth principle is temporal sovereignty. People should be able to declare which goals govern which periods. An individual might authorize

immediate interruption for family emergencies and medical alerts while requiring political content, shopping recommendations, and workplace requests to be batched. The important innovation is that attention becomes a policy object controlled by the user rather than a sequence of ad hoc reactions.

This is technically feasible. The barrier is institutional. Current software ecosystems have weak incentives to respect attention policies that reduce engagement. Personal AI agents could change the bargaining structure if they act as intermediaries. Instead of every application competing directly for the user's cognition, applications would negotiate with an agent operating under user-defined rules.

Such an agent would be a kind of cognitive firewall. Network firewalls do not decide what information is morally good; they enforce connection policies. An attention firewall would not decide which ideas are worthy. It would enforce conditions for when, how, and why systems may interrupt or influence the user.

The analogy should not be pushed too far. Human attention is not a packet stream. Serendipity matters. Friends should sometimes interrupt us. Art should surprise us. Political events can deserve unexpected attention. A perfectly optimized life would be sterile. The aim is not control so complete that nothing unplanned enters consciousness. The aim is to shift the default bargaining power.

Constitutional language becomes especially useful when we consider children. Adults can in principle choose among services, change settings, purchase alternatives, and accept contracts. Children are developing the very capacities that make informed choice possible. An attention constitution for childhood cannot simply treat a child as a small consumer.

Education, parents, platforms, device makers, and public institutions all possess partial authority, and none should possess unlimited authority. A

school may need to restrict certain systems during learning, but it should not gain permanent surveillance over a student's private life. Parents may set boundaries, but older children need increasing autonomy. Platforms can provide safety tools, but commercial incentives may conflict with developmental goals. The governance problem is therefore explicitly polycentric.

A useful principle is progressive agency: cognitive rights should expand as the individual's capacity for self-government develops, while institutions remain responsible for preventing forms of exploitation that development makes difficult to recognize. This principle already exists in many areas of law and could guide digital attention without requiring a single age threshold for every function.

Employers pose another challenge. A worker may technically consent to monitoring or AI assistance while facing little bargaining power. If an employer-provided assistant manages messages, schedules, productivity, and training, it occupies an intimate cognitive role. The employee should know whether the agent serves the worker, the employer, or both. Conflicts of purpose should be explicit.

This leads to a general constitutional rule: no hidden principal. A system acting in one social role should disclose when another actor can redirect its objectives. A personal assistant that quietly becomes an advertising channel violates the role even if the data processing is legally disclosed somewhere. A tutoring system that optimizes for institutional rankings rather than student learning has a similar conflict.

The rule is familiar in professional ethics. A lawyer identifies the client. An auditor's independence matters because hidden loyalty destroys the function. A trustee cannot secretly treat entrusted assets as a private opportunity. As AI systems become social actors, role clarity will become as important as model capability.

Another likely right is the right to a non-personalized channel for certain public functions. Citizens may accept personalized news, entertainment, and shopping, but there are contexts in which public accountability requires a common reference. Official notices, emergency communication, election information, legal obligations, and public-service explanations may need canonical versions that anyone can inspect, even if assistants later adapt language for accessibility.

Personalization should transform presentation without silently transforming the underlying public commitment.

The attention constitution may also need algorithmic due process. If a ranking or recommendation materially affects employment, education, credit, healthcare, or civic participation, the affected person should have some route to understand the basis, challenge errors, and obtain human or independent review. The exact procedure should vary with stakes. Due process is not a demand that every automated choice become a trial. It is the principle that consequential power should not be unanswerable.

There is also a right that sounds almost archaic: the right not to be optimized. A person may reasonably choose a mode in which a system provides requested functionality without continuously adapting persuasive techniques to maximize a behavioral outcome. This is different from privacy. The system may know the user well and still be prohibited from exploiting that knowledge for certain objectives.

The commercial consequences would be significant. Some services would become more expensive because behavioral optimization currently subsidizes access. The institutional transition must confront this honestly. Attention protection is not costless. Societies will have to decide which cognitive infrastructures are public goods, which are paid services, and which forms of advertising remain acceptable.

A serious attention constitution therefore cannot be merely a catalogue of rights. It must include sustainable business models. Privacy regulation taught this lesson. If governance forbids a revenue mechanism without creating viable alternatives, markets may consolidate around the few actors large enough to absorb compliance costs. Anti-trivialization rules must be designed to avoid becoming barriers that strengthen the incumbents they are meant to constrain.

Interoperability is one answer. If users can move relationships, archives, preferences, and agent memories between providers, smaller services can compete on alignment rather than network lock-in. Open standards can reduce compliance duplication. Public reference implementations can help smaller organizations implement provenance and attention controls. Collective procurement can make user-aligned services viable for schools, libraries, and public institutions.

The constitutional frame thus joins rights with architecture. Freedom is not only what the contract says. It is what the system makes practically possible.

The attention constitution must also protect the public sphere from excessive personalization. A democracy requires some shared objects of attention. Citizens do not need identical information, but they need enough common reference points to disagree about the same world. Hyper-personalized political communication can undermine accountability because different citizens receive incompatible promises that are never publicly compared.

One possible counter-institution is a public campaign ledger in which materially distinct political claims and targeted variants are archived and comparable. Another is a requirement that high-impact automated persuasion systems disclose sponsorship and targeting logic. Another is an independent infrastructure for examining how public claims vary across audiences. These

are not restrictions on ordinary speech; they are transparency responses to industrialized personalization.

The design challenge is to protect shared reality without forcing citizens into one approved narrative. The answer will not be a central fact machine. It will be procedural: provenance, comparison, adversarial access, public archives, and rights to inspect the machinery of influence.

Constitutionalism is the right intellectual tradition because it assumes that power will be abused sometimes, including by those who believe they are doing good. An attention constitution should therefore constrain platforms, states, employers, schools, and even user agents. It should preserve the right to be irrational, playful, private, surprising, and occasionally distracted.

Albert Hirschman's distinction between *exit* and *voice* adds a practical test. People need a meaningful ability to leave a cognitive environment, but exit alone is insufficient when network effects, employment, schooling, or public services make departure costly. They also need voice: the ability to contest ranking rules, challenge delegated decisions, alter defaults, demand explanations, and participate in the rules that shape the environment (Hirschman 1970). Portability without voice can merely make migration between similar systems easier. Voice without exit can become ritual consultation. A mature attention constitution needs both.

The architecture itself is political. Winner's famous question about whether artifacts have politics becomes almost literal when the artifact decides what interrupts a person, what becomes salient, and what disappears from view (Winner 1980). The relevant constitutional object is therefore not only the recommendation algorithm. It is the whole stack: identity, notifications, defaults, ranking, memory, interoperability, advertising, agent permissions, data flows, business models, and the legal ability to substitute one layer for another. Political liberty in a mediated environment partly depends on whether these layers can be inspected, replaced, or governed independently.

The purpose is not to create perfectly serious citizens. It is to ensure that no institution acquires an unaccountable right to decide what their seriousness should serve.

Chapter 8: Truth Without a Ministry of Truth

The phrase "institution for truth" should make a liberal society nervous. History provides enough examples of authorities that claimed to defend truth while suppressing inquiry. Yet a society that refuses to build any institutions for epistemic quality does not become neutral. It simply leaves the field to whatever systems are best at producing visibility, repetition, or authority.

The alternative is to distinguish truth from truth infrastructure. Science does not possess a Ministry of Scientific Truth. It has competing laboratories, journals, databases, methods, peer review, replication, professional norms, conferences, funding institutions, and criticism. These institutions do not guarantee truth. They make some errors easier to detect and some corrections easier to preserve.

The future public sphere needs analogous infrastructure for a world in which the volume of claims grows much faster than human verification capacity.

The first element is provenance. A consequential digital claim should increasingly carry a traceable history: where the underlying data came from, whether media are synthetic, what transformations were applied, which models generated summaries, what uncertainty remains, and who approved publication. This does not mean that anonymous speech disappears or every casual statement becomes certified. Different contexts require different standards. The key is that high-stakes claims should be able to travel with inspectable context.

The second element is evidence maps rather than binary fact labels. Many important disagreements cannot be represented as true or false. Evidence is mixed, causal inference is uncertain, definitions differ, and values affect interpretation. An evidence map can show what propositions are

supported by which sources, where experts disagree, which assumptions drive disagreement, and what new evidence would change the picture.

AI makes such systems more feasible because machines can maintain large, evolving argument structures. But AI also makes them dangerous if one model silently decides what counts as relevant. The institutional form should therefore be adversarial and plural. Different organizations should be able to produce competing maps using shared provenance standards. Users should be able to inspect the disagreements between maps.

The third element is claim persistence. Public discourse currently has a severe memory problem. Predictions disappear. Corrections circulate less widely than original claims. Institutions make promises and later redefine them. Politicians and commentators can reverse positions without the information environment automatically presenting the history.

An epistemic memory institution would not punish inconsistency. Changing one's mind is often rational. It would make changes legible. A statement could be connected to later revisions, evidence, and outcomes. Over time, credibility could be based less on rhetorical confidence and more on calibration: how often did a person or institution express appropriate uncertainty, revise when evidence changed, and distinguish prediction from advocacy?

The fourth element is structured adversarial review. High-impact public analyses could be accompanied by a deliberately commissioned counter-analysis. Courts already institutionalize opposition. Scientific peer review attempts to do so imperfectly. Intelligence analysis uses red teams. Engineering uses hazard analysis. AI can make adversarial review cheaper and more systematic, but the challenge is to avoid turning it into performative "both sides" symmetry when evidence is asymmetric.

A mature system would not require every claim to have an equal opposite. It would require serious claims to survive serious attempts at falsification.

The fifth element is uncertainty literacy. Public communication often rewards confidence because uncertainty is hard to compress. Yet uncertainty is not ignorance; it is information about the limits of knowledge. Institutions can create standard ways of representing ranges, scenarios, confidence, and unknowns. Medicine, meteorology, engineering, and finance already use specialized uncertainty practices. Public reasoning needs better translations of these practices.

The sixth element is epistemic standing. Not every source deserves equal weight on every question, but expertise cannot be reduced to institutional prestige. Future systems need richer ways to represent who has relevant experience, what conflicts of interest exist, how a source has performed historically, and which parts of a claim depend on specialized knowledge.

This is an area where AI could either deepen elitism or reduce it. A centralized reputation score would be dangerous. A transparent, context-specific record could be useful. The difference lies in whether the system produces a universal hierarchy or supports inspectable judgment.

The intellectual history of truth institutions is older than modern media. Plato worried that writing could create the appearance of knowledge without understanding. Early modern science developed public demonstration and reproducible method partly because private testimony was insufficient. The Royal Society's motto, *Nullius in verba*, expressed a refusal to treat authority alone as evidence. Karl Popper later emphasized falsifiability and criticism, while Thomas Kuhn showed that scientific communities also operate through paradigms that shape what questions and anomalies are visible (Popper 1959; Kuhn 1962). The lesson is not that one philosopher

solved epistemology. It is that knowledge depends on organized criticism as much as on information.

Modern social epistemology makes this institutional point even clearer. Alvin Goldman asked how social practices and institutions can be evaluated according to their tendency to produce knowledge rather than merely agreement (Goldman 1999). Helen Longino argued that objectivity is not the absence of social perspective but an achievement of communities that expose assumptions to effective criticism under conditions that permit genuine uptake (Longino 1990). These traditions help escape a false choice between naive individualism and centralized authority. Truth-seeking can be social without becoming majoritarian, and institutional without becoming bureaucratically decreed.

This distinction will matter enormously in synthetic information environments. An AI system can produce a persuasive consensus in seconds. Ten agents can independently "agree" while sharing the same training data, retrieval index, hidden benchmark, or upstream model. Numerical plurality is not epistemic independence. Future truth infrastructure will need to represent dependency: which systems share sources, which methods are genuinely different, which criticisms were generated from independent assumptions, and which apparent corroborations are merely echoes. In such a world, provenance is not only about where a sentence came from; it is about the genealogy of agreement.

The public sphere has weaker institutions of organized criticism than science because the objectives are different. Politics must decide under uncertainty. Values legitimately conflict. Evidence is incomplete. Citizens cannot devote years to each question. The goal cannot be to turn democracy into a laboratory.

But public reasoning can borrow institutional functions without pretending that politics becomes science. Provenance can show where a factual

claim came from. Argument maps can distinguish factual disagreement from value disagreement. Forecasts can record uncertainty. Public decisions can preserve the reasons that justified them. Independent bodies can evaluate whether promised outcomes occurred. Journalists can attach source trails that later systems can inspect automatically.

One particularly important future institution may be the public claim registry. The name sounds bureaucratic, but the idea is simple: high-impact institutional claims become addressable objects with persistent identities. A government forecast, corporate environmental claim, major scientific announcement, or campaign promise could have a machine-readable record linking original language, supporting evidence, later revisions, criticism, and observed outcomes.

The registry would not declare the claim true or false. It would preserve its history.

This distinction is crucial because the internet's memory is paradoxical. Almost everything can remain somewhere, but context is fragile. Screenshots survive while corrections disappear. A sentence can be retrieved without the document around it. Search can find a claim without showing what happened next. Persistence without context can increase confusion.

Epistemic institutions need versioned memory. Software engineering learned this through version control. A codebase is not merely the latest file; it is a history of changes, authorship, branches, reversions, and review. Public knowledge increasingly needs comparable infrastructure. A policy document should show not only what it says today but why important changes were made. A scientific synthesis should reveal which evidence entered or left the consensus. An AI-generated report should preserve which model and source set produced each version.

This idea has a surprising moral consequence. Better memory can make changing one's mind safer. In the current media environment, revision is often framed as hypocrisy because only the incompatible statements remain visible. A versioned system can show the evidence that justified change. Intellectual flexibility becomes legible rather than embarrassing.

Trust would then shift from consistency of position toward consistency of method. Did the actor update when evidence changed? Did uncertainty shrink appropriately? Were conflicts of interest disclosed? Did the prediction perform well? Did the institution correct errors openly?

This is close to the concept of calibration used in forecasting. A person who says "70 percent" repeatedly should be right roughly seven times out of ten across comparable cases. Calibration is not the whole of judgment, but it offers a richer reputation signal than confidence or popularity. Future public institutions could preserve forecasting records so that epistemic credibility becomes partly empirical.

Prediction markets and forecasting tournaments have explored some of these mechanisms, though they have their own limitations. Markets require well-defined questions and can reflect wealth or participation biases. Forecasting systems can privilege measurable events over morally important but ambiguous outcomes. The broader lesson is that institutions can make uncertainty performance visible.

The danger of all reputation systems is centralization. A universal credibility score would be a gift to authoritarianism and social conformity. The alternative is contextual evidence. Instead of "this person has credibility 82," a system can show "this source has published in this domain, disclosed these conflicts, made these forecasts, corrected these errors, and is disputed by these other sources." Judgment remains with the reader or institution.

AI can make such context navigable. It can compress a large evidence history into a summary while preserving the ability to inspect the record. This is a more promising use of summarization than simply shortening everything. The point is not to remove complexity but to make complexity traversable.

In that sense, the future epistemic machine should behave like a good map. A map simplifies, but it tells you that it is a map and allows you to choose the scale. Trivialization begins when the thumbnail pretends to be the territory.

Habermas's ideal of rational-critical public discourse was never fully realized, and critics rightly noted that public spheres exclude, stratify, and depend on power (Habermas 1989; Fraser 1990). Yet the aspiration remains valuable: legitimacy improves when claims can be challenged through reasons rather than status alone. The digital problem is that reasons now compete inside an attention architecture designed for other objectives.

Anti-trivialization institutions would not restore an imagined nineteenth-century public sphere. They would build new layers beneath and around the feed. The user could see the fast conversation and also access the slow record. Viral claims could remain viral, but their provenance and history would be one action away. Public institutions could publish machine-readable evidence alongside persuasive summaries. Journalists could draw from shared source graphs while maintaining editorial independence. Citizens' assemblies could use structured knowledge systems that preserve minority arguments rather than collapsing deliberation into a final vote.

The phrase "truth infrastructure" should therefore mean infrastructure for correction. The most trustworthy epistemic institution is not the one that claims to be right most loudly. It is the one that makes it easiest to discover when it is wrong.

Chapter 9: Institutions for Time

Trivialization is partly a crisis of temporal mismatch. Many of the most consequential systems in modern civilization operate on timescales longer than the institutions that govern them.

Climate systems unfold over decades and centuries. Infrastructure lasts for generations. Education shapes lives over decades. Public debt, demographic change, biodiversity loss, antibiotic resistance, nuclear waste, and institutional trust have long temporal tails. Political attention, corporate reporting, media cycles, and digital trends are much shorter.

This mismatch does not imply that short-term institutions are irrational. Elections must occur frequently enough to maintain accountability. Markets need rapid price signals. Managers need operational data. News must respond to events. The problem arises when short-horizon signals become the only signals with institutional force.

Societies have occasionally created representatives for absent futures. Pension systems institutionalize obligations across generations. Endowments preserve resources across decades. Constitutional rules protect commitments from ordinary political cycles. Central banks are often given degrees of independence to resist short-term political incentives. Environmental impact assessments force some future consequences into present decisions. Wales created a statutory framework centered on the well-being of future generations, one example of an emerging family of institutions that attempts to give long-term effects standing in present governance (Welsh Government 2015).

These efforts point toward a broader category: temporal institutions.

A temporal institution is any mechanism that prevents the present from exercising unlimited sovereignty over the future. It can do this through delay, representation, reserves, scenarios, commitments, memory, or review.

The attention economy weakens temporal institutions because it rewards immediacy. AI could weaken them further by accelerating analysis and decision cycles. When organizations can generate reports, options, and recommendations instantly, the temptation will be to treat decision latency itself as inefficiency.

But faster analysis does not make reality faster. A forest still grows slowly. A child still develops over years. A scientific claim still needs time to encounter replication, criticism, and unexpected evidence. Institutional wisdom will increasingly require the ability to distinguish computational speed from ontological speed.

One future mechanism is the decision latency floor. Certain high-impact decisions could require a minimum period between a complete proposal and irreversible action, except under clearly defined emergencies. Financial systems already use cooling-off periods and settlement delays in some contexts. Law uses notice periods and appeals. The principle could expand to AI-mediated decisions where synthetic analysis creates an illusion of complete understanding.

A second mechanism is the future-impact ledger. Major institutions could maintain explicit records of expected long-term effects, assumptions, and indicators. Years later, these predictions would be compared with outcomes. The point would not be punishment for imperfect forecasting. It would be institutional learning. Organizations currently remember outcomes selectively because decision rationales are dispersed and staff change. AI can make long-term audit practical.

A third mechanism is intergenerational adversarial review. Large projects could face a structured advocate whose mandate is not to oppose the project but to represent costs displaced into the future. Existing environmental and fiscal institutions approximate this function. AI could support it by simulating long-term scenarios, identifying hidden dependencies, and preserving assumptions for later review.

A fourth mechanism is the slow chamber. Democracies already use bicameral structures partly to introduce different forms of representation and pace. Future governance may experiment with bodies that have longer terms, different selection methods, and a specific mandate to review issues with long time horizons. Citizens' assemblies and deliberative mini-publics offer one starting point, especially when combined with transparent evidence infrastructure and protection from continuous media pressure (Fishkin 2018; Landemore 2020).

The risk is obvious. "Long-termism" can become a rhetorical device by which current elites suppress urgent needs in the name of hypothetical futures. Any institution representing the future must therefore remain accountable to present citizens and plural conceptions of value. Future generations cannot vote, and anyone claiming to speak for them can abuse the role.

The solution is not to abandon the future but to institutionalize contestation about it. Different models, ethical frameworks, and stakeholder perspectives should be visible. A future-impact process should expose assumptions rather than hide them behind a single score.

The problem of time has an ethical history as well as an institutional one. Hans Jonas argued that technological civilization expands the temporal and spatial consequences of action so dramatically that traditional ethics, developed for interactions among contemporaries, becomes inadequate (Jonas 1979/1984). When today's choices can alter ecosystems, genomes, climate, or

infrastructure for generations, the absent future becomes a stakeholder without a voice.

Institutions are the only practical way to give that stakeholder standing. Moral concern matters, but long horizons disappear from attention when every immediate interest has a representative and future consequences do not.

This is why many mature systems create reserves. Banks hold capital against future losses. Pension funds set aside assets against future obligations. Maintenance budgets preserve infrastructure before failure becomes visible. Strategic stockpiles exist because waiting for scarcity to appear is too late. Reserves are institutionalized respect for the possibility that the future will not resemble the present.

Attention institutions need reserves too. An organization can reserve time for exploratory research whose value is not immediately measurable. A government can reserve analytic capacity for low-probability risks. A university can preserve disciplines that are not currently fashionable. A scientific funding system can reserve resources for replication and negative results. A person can reserve hours in which no external system may claim priority.

The language of reserve is useful because it avoids moral grandiosity. We do not need to predict which long-term issue will prove most important. We need to preserve option value.

Option value is a central idea in decision-making under uncertainty. A reversible choice keeps future alternatives open; an irreversible choice consumes them. Institutions can therefore be designed to distinguish reversible from irreversible decisions. Fast experimentation is appropriate when reversal is cheap. Slow review becomes more valuable when consequences are difficult to undo.

This principle should be central to AI governance. Many software changes are reversible. Some deployments gradually restructure organizations, labor markets, educational expectations, or public information in ways that are difficult to reverse. Treating every AI deployment as a software update ignores institutional irreversibility.

A future "reversibility assessment" could ask whether a system can be removed without losing essential records, skills, relationships, or alternative providers. If a school delegates an entire curriculum to one AI platform, can it leave later? If a government agency replaces expert staff with automated systems, does the expertise still exist if the system fails? If a research field relies on one model ecosystem, can independent methods survive?

This is another form of temporal sovereignty: preserving the future's ability to choose differently.

Long-term institutions also need ceremonial protection from urgency. Modern organizations often label many requests urgent because urgency wins priority. The result is urgency inflation. When everything is urgent, strategic attention disappears.

One anti-trivialization mechanism is an emergency budget. An organization can limit the proportion of work that may bypass normal deliberation under emergency status without review. After the emergency, the exceptional decision receives retrospective audit. Constitutional systems use related logic because emergency powers are necessary and dangerous.

AI can support this by tracking how often institutions invoke urgency, which decisions bypassed review, and whether the predicted emergency materialized. What is currently a vague cultural complaint becomes measurable institutional behavior.

Time institutions can also change democratic communication. Political debates often force candidates to answer complex questions in

seconds, rewarding confidence and rhetorical compression. A future public sphere might combine fast debate with slower written or AI-assisted deliberation in which evidence, assumptions, and alternative scenarios remain available. The goal is not to make politics academic. It is to create a second temporal channel where claims can outlive the performance.

The future will always surprise us. Long-term institutions should therefore be judged less by predictive accuracy than by adaptive capacity. They preserve memory, reserves, diversity, and reversible options so that societies can respond when forecasts fail.

A civilization that respects the future is not one that knows the future. It is one that refuses to consume all of its choices in the present.

Time also matters at the individual level. Personal AI will increasingly be able to model a user's long-term intentions and short-term impulses. This creates a new form of self-governance. A person could ask an agent to protect a writing period from interruption, preserve a savings plan against impulsive purchases, maintain a reading project, or surface neglected relationships. The agent becomes a contract between present and future selves.

Such systems should be designed carefully because paternalism can be outsourced to software. Users need the right to change goals, break commitments, and act spontaneously. The value lies not in forcing consistency but in making temporal conflict explicit.

Civilization itself faces the same problem. It needs institutions that allow the present to change its mind while making forgetting more difficult.

The anti-trivialization machine of time is not a machine that makes everything slow. It is a machine that makes speed answer to consequence.

Chapter 10: Polycentric Seriousness

A centralized solution to trivialization would almost certainly reproduce the problem in a more dangerous form. If one authority determined which information deserved attention, which sources counted as serious, which issues required long-term protection, and which forms of persuasion were acceptable, the cure would become an architecture of epistemic power.

The better model is polycentric seriousness: many overlapping institutions, each with limited jurisdiction, different incentives, and the capacity to inspect or counterbalance others.

Elinor Ostrom's work on commons governance provides the intellectual foundation. She showed empirically that communities can manage shared resources under diverse institutional arrangements when rules fit local conditions, users participate in rule-making, monitoring exists, sanctions are graduated, conflicts can be resolved, and governance is nested across scales (Ostrom 1990). Her work is often summarized as a rejection of simple binaries between state and market. For the attention problem, the more important lesson is methodological: institutional design must be specific to the resource, the community, the scale, and the failure mode.

There is no single attention commons. A classroom, a research laboratory, a family, a social network, a parliament, a hospital, and a company have different cognitive needs. A hospital emergency department should not optimize for deep uninterrupted contemplation during an acute crisis. A research institute should not operate at emergency-department rhythm all week. Good governance begins by distinguishing contexts.

At the personal scale, user-controlled agents can negotiate interruptions, preserve goals, summarize information, and maintain cognitive boundaries. At the family scale, norms can govern devices, shared time, and children's exposure without pretending that one rule fits every household. At

the organizational scale, communication protocols can reduce interruption externalities. At the professional scale, standards can specify how AI-assisted work must be documented and reviewed. At the platform scale, ranking systems can be audited and users given meaningful choice. At the civic scale, public-interest information institutions can maintain provenance and deliberation. At the state scale, law can establish rights, liability, interoperability, and due process. At the international scale, standards can address cross-border AI systems and synthetic media.

These layers should overlap deliberately. A user should not have to trust a platform's self-assessment. A regulator should not be the only source of evaluation. Independent researchers, civil-society organizations, insurers, standards bodies, and competing services can create parallel forms of accountability.

This redundancy has a cost. It can slow innovation and create compliance complexity. But concentration has a cost too. The relevant question is not whether governance adds friction; all governance does. The question is whether the friction is proportionate to the systemic power being exercised.

Polycentricity also protects cultural diversity. Seriousness is not a single lifestyle. A religious community, an artistic collective, a scientific laboratory, and a startup may organize attention differently. The goal is not to impose one cognitive ideal but to give communities tools for defending their chosen forms against unwanted extraction.

This leads to an important reversal of contemporary platform design. Today the default is often centralized optimization with local preference settings. The future may require local optimization under shared constitutional constraints. Users, organizations, and communities could choose their own attention policies while interoperable standards protect basic rights and transparency.

AI makes this technically plausible because policy enforcement can become programmable. A school might specify that educational agents cannot optimize for advertising revenue, must preserve source provenance, and must provide teachers with uncertainty reports. A research organization might require that AI-generated claims retain source spans and that high-impact conclusions undergo independent model review. A family might use an agent that enforces age-appropriate boundaries without sharing detailed behavior with the platform. A government agency might require that public-facing AI maintain an auditable distinction between evidence, inference, and policy preference.

The danger is policy overload. If every institution encodes elaborate cognitive rules, life becomes a maze of compliance. The solution will require simple primitives, much as cybersecurity evolved reusable concepts such as authentication, authorization, logging, encryption, and least privilege.

Attention governance needs its own primitives. Provenance. Purpose. Priority. Interruption rights. Auditability. Reversibility. Escalation. Uncertainty. Memory. Expiration. Conflict-of-interest disclosure. Human override. These can become building blocks used differently across contexts.

Standards bodies may play an unexpectedly important role. They are among civilization's least glamorous institutions and among its most consequential. Shared standards allow systems to interoperate without placing one actor in complete control. Future standards for provenance, AI delegation, content authenticity, user attention policies, and cognitive audit could provide a common substrate on which diverse institutions compete.

Polycentric governance also changes the role of standards. Standards are often treated as technical details, but they are constitutional agreements among systems. A standard can specify how identity is represented, how provenance travels, how an alert declares urgency, how an AI system reports uncertainty, or how a user expresses an attention preference. When the

standard is open, many organizations can implement the rule without one organization controlling the entire ecosystem.

This is one reason open technical infrastructure may become central to cognitive freedom. A user-controlled attention agent is meaningful only if services can communicate with it. Provenance is useful only if systems can exchange provenance records. Portability is useful only if memories, preferences, and social graphs have interoperable representations. Without standards, rights can remain theoretical.

Open-source communities offer imperfect examples of polycentric institutional design. They combine technical rules, maintainers, contribution processes, code review, reputation, forks, foundations, firms, volunteers, and user communities. Power can concentrate heavily, and conflicts can be severe. Yet the possibility of forking changes governance because exit can sometimes become creation of an alternative rather than abandonment.

Forkability may become an important principle for cognitive institutions. A community dissatisfied with a ranking rule, knowledge map, or moderation policy should sometimes be able to create an alternative using the same underlying open records. This does not eliminate conflict. It prevents every conflict from becoming a battle for control of one central system.

There are limits. Some infrastructures have network effects so strong that parallel systems are inefficient. Emergency communication, identity systems, core scientific databases, and public legal records may require common layers. Polycentricity does not mean fragmentation for its own sake. It means separating the common substrate from the contested interpretation.

A public evidence infrastructure, for example, might standardize source identifiers, provenance, version history, and citation relationships while allowing many organizations to produce different summaries, rankings, and interpretations. A public attention standard might allow users to express

interruption preferences while different devices and services compete on implementation. Shared rails, plural applications.

This structure resembles the internet at its best. Common protocols enable diverse services. The lesson is not that the internet is inherently decentralized; contemporary layers contain significant concentration. The lesson is architectural: interoperability can permit innovation without requiring one authority to decide every higher-level use.

James Scott's warning about administrative legibility should remain present here. A polycentric system can still become oppressive if standards encode narrow categories and make nonconforming life difficult (Scott 1998). The more powerful a standard becomes, the more carefully its defaults, governance process, and extension mechanisms should be designed.

The institutional design therefore needs an escape hatch. Human beings should be able to do things that the schema did not anticipate. Not every conversation requires provenance. Not every friendship needs a policy. Not every local community should report its norms to a central registry. Anti-trivialization becomes anti-human when the infrastructure requires all meaningful activity to become machine-legible.

This produces a paradox worth preserving. We need more formal structure around the systems that allocate attention, while protecting informal spaces from over-formalization. The feed should become more accountable; dinner should not.

The same principle applies to scientific and professional institutions. Formal provenance is valuable for a published medical claim. It may be counterproductive during early brainstorming, where incomplete ideas need privacy and freedom. Good institutions distinguish exploratory spaces from commitment spaces. In exploration, errors are cheap and ambiguity is productive. In commitment, evidence and accountability matter more.

Future AI environments should make this distinction explicit. A system can offer a private speculative mode with minimal logging and a formal decision mode in which sources, assumptions, approvals, and dissent are preserved. Today these contexts often coexist in the same chat interface. Institutional maturity will separate them.

Polycentric seriousness is therefore not simply many institutions. It is a layered architecture in which the degree of formality rises with public consequence.

The long-term goal is an ecology rather than a fortress. A fortress assumes a clear inside and outside. The attention environment is too entangled for that. People will continue to seek entertainment, novelty, persuasion, outrage, intimacy, and surprise. The task is to make the ecology robust enough that no single selection pressure colonizes every niche.

In biological ecosystems, diversity creates resilience partly because species respond differently to shocks. In institutional ecosystems, different forms of judgment can play a similar role. Markets detect some signals quickly. Science validates others slowly. Journalism investigates. Courts adjudicate. Communities interpret. Art reframes. AI can compare patterns across enormous information spaces. Human experience contributes forms of relevance that cannot be reduced to database consistency.

Polycentric seriousness means preserving this cognitive biodiversity.

Chapter 11: Education, Work, and the Rebuilding of Human Agency

The institutions of attention will fail if they become a layer added after human development is complete. People learn how to allocate attention, tolerate uncertainty, evaluate evidence, and manage tools through education and work. These are not merely private habits. They are institutional capacities.

Modern education was built around an economy in which information was scarce enough that acquiring and recalling it had substantial value. AI changes the balance. Students can increasingly generate fluent explanations, essays, code, and summaries instantly. The obvious institutional response is surveillance: detect AI use, restrict devices, return to closed examinations. Some of this may be necessary in specific contexts, but it cannot be the long-term philosophy of education.

The deeper question is what capabilities remain valuable when generation is cheap.

One answer is judgment. Students need to learn how to formulate questions, compare explanations, detect missing assumptions, trace claims to evidence, design tests, revise beliefs, and decide when not to delegate. They need to understand enough of a domain to recognize when an answer is suspicious. The educational objective shifts from producing an artifact to demonstrating a defensible cognitive process.

This does not make foundational knowledge obsolete. Quite the opposite. Verification without internal knowledge is fragile. A person who knows nothing about a field cannot reliably evaluate a fluent synthesis. AI may therefore increase the value of conceptual foundations even as it reduces the value of rote production.

Assessment will have to change accordingly. A student might submit not only a final essay but a decision record showing which questions were delegated, which sources were consulted, where the AI was wrong, how the argument changed, and what remains uncertain. Oral defense may return in new forms. Project work may emphasize evidence trails and iterative reasoning. Examinations may test the ability to criticize a generated answer rather than reproduce one from memory.

This is not simply a pedagogical adaptation. It is training for institutional citizenship in an AI-mediated world.

The workplace faces the same shift. Companies adopting AI often focus on task acceleration. This is rational because productivity gains are measurable quickly. But the deeper organizational challenge is cognitive architecture. If AI produces more reports, messages, analyses, dashboards, code, and proposals than teams can evaluate, organizations can become faster at generating work and slower at deciding what matters.

The anti-trivialization workplace will therefore measure output less naively. More generated artifacts are not automatically more productivity. Organizations will need filters, decision records, escalation rules, provenance, and explicit limits on communication load.

Meetings provide a simple example. AI can summarize every meeting, generate action items, and distribute transcripts. This can reduce administrative cost. It can also make it easier to schedule more meetings because the apparent cost falls. The result may be an increase in coordination overhead. Technology that removes one friction can expose the value of another.

A better institution might use AI to reduce meetings rather than accelerate them. An agent could determine whether the issue has enough disagreement or uncertainty to justify synchronous discussion, gather

positions in advance, identify decision points, and preserve the rationale afterward. The success metric becomes fewer unresolved coordination loops, not more documented conversation.

Organizations also need protected cognitive states. Deep work cannot survive if every employee must remain continuously interruptible. Communication systems should support collective quiet periods, asynchronous defaults, and escalation protocols. These are not perks for knowledge workers. They are production infrastructure for tasks that require concentration.

The same principle applies to leadership. Executives face information environments in which what is easiest to measure can crowd out what is strategically important. AI may intensify dashboard governance by making metrics more abundant. Future boards may therefore need explicit anti-metric functions: processes that ask what is not being measured, which indicators are being gamed, what long-term risks have weak short-term signals, and where apparent precision exceeds actual knowledge.

There is a productive intellectual tradition for thinking about AI as augmentation rather than substitution. Vannevar Bush imagined systems that could extend associative memory. J. C. R. Licklider described human-computer symbiosis in which machines handled routine operations while humans contributed goals, judgment, and pattern recognition. Douglas Engelbart built his research program around augmenting human intellect and collective problem solving rather than merely automating existing tasks (Bush 1945; Licklider 1960; Engelbart 1962).

These visions are newly relevant because generative AI can either fulfill or betray them. A tool can extend the space in which a person can think, or it can collapse the space by producing the answer before the person forms the question.

The difference is often interaction design. A tutor that immediately supplies a complete solution optimizes short-term task completion. A tutor that diagnoses the learner's misconception, asks for a prediction, offers a counterexample, and only later reveals a solution can optimize learning. Both may use the same underlying model. The institution determines which objective governs the interaction.

This makes education a test case for anti-trivialization. If schools respond to AI only by detecting unauthorized output, they will spend enormous effort defending an assessment model that technology has already destabilized. If they accept all AI assistance without redesign, they may certify polished artifacts that reveal little about understanding. The sustainable response is to redesign what counts as evidence of learning.

That evidence can include explanation, oral defense, transfer to unfamiliar problems, ability to critique a flawed AI answer, reconstruction of reasoning from sources, and independent performance when tools are unavailable. AI can help administer these assessments rather than merely defeat them.

The same idea applies to professional competence. A physician using decision support may be safer than one relying on memory alone, but the profession still needs to know whether the physician can recognize when the system is wrong. A programmer using code generation may be more productive, but the organization needs assurance that somebody can inspect critical code. A scientist using automated literature tools may explore more widely, but scientific training must preserve the ability to assess methods and evidence.

The key concept is skill reserve. Just as societies maintain strategic reserves of material resources, organizations may need reserves of human competence that automation makes less frequently used but remains essential under failure or novelty. This is not an argument for preserving every obsolete

skill. It is a governance question about which capabilities are safety-critical, identity-forming, or necessary for independent judgment.

Aviation again offers a useful warning. Automation can make routine performance safer while reducing opportunities to practice rare manual skills. Training systems compensate through simulation. Knowledge work may need its own simulators: environments where professionals practice diagnosing AI errors, working from primary evidence, or operating under tool failure.

Workplace institutions also need to rethink career development. Junior professionals historically learned partly by performing routine tasks. If AI automates those tasks, organizations may gain immediate efficiency while destroying the apprenticeship pipeline through which experts are produced. Law firms, consultancies, laboratories, software teams, and medicine will face versions of this problem.

The anti-trivialization response is to treat learning as an organizational output rather than an accidental by-product of inefficient work. Firms may need explicit apprenticeship tasks, simulation, supervised review, and protected learning budgets even when a model could produce the immediate artifact faster.

This is another yin-yang reversal. Automation removes the necessity of practice, making deliberate practice more necessary.

Human agency also depends on the experience of causation. People need to feel that choices connect to consequences. Highly automated environments can make action feel like approval of machine-produced options. Over time, the person can become a passenger in a workflow nominally under human control.

Institutions can resist this by giving humans meaningful decision boundaries. The human role should not be ceremonial clicking. It should

involve setting objectives, resolving genuine ambiguity, evaluating exceptions, and carrying responsibility proportionate to actual control.

This may sometimes require giving machines more authority, not less. If a system makes the real decision, pretending that a human approved it after two seconds only hides responsibility. Institutional honesty requires aligning formal authority with actual cognitive control.

The future of work will therefore be shaped not only by which tasks AI can perform but by how organizations distribute agency, learning, and accountability around those tasks. Productivity is one objective. Human capability is another. Resilience is a third. Good institutions will refuse to collapse all three into a single number.

Education and work also shape psychological agency. A person accustomed to continuous external prompting may find self-directed attention more difficult. A person accustomed to AI completing every difficult cognitive step may lose confidence in unresolved thought. But the reverse is also possible. Well-designed AI can scaffold reflection, offer alternative explanations, create deliberate practice, and make expertise accessible.

The distinction is whether the system expands or contracts the user's future agency.

This suggests a simple criterion for cognitive technology: does using the tool make the person more capable of acting without it when necessary? Not every tool must pass this test. Calculators legitimately replace some arithmetic. Navigation systems replace some route memory. Civilization is built on delegation. The point is that education should decide consciously which capacities are being delegated and which are being cultivated.

The future institution of education may therefore treat AI literacy as a branch of epistemic self-government. Students will need to know not only how to prompt systems but how to establish objectives, inspect provenance,

recognize uncertainty, create adversarial tests, and understand incentive conflicts.

A society of powerful AI does not need every citizen to become an AI engineer. It needs citizens who know when a machine is answering the wrong question beautifully.

Chapter 12: A Civilization That Can Think

The phrase "a civilization that can think" is metaphorical. Civilizations have no single mind. They contain billions of minds, institutions, machines, conflicts, memories, markets, languages, and histories. Yet societies clearly differ in their capacity to notice problems, preserve knowledge, coordinate expertise, learn from failure, and act across time.

This capacity is not reducible to intelligence. A society can contain extraordinary individual intelligence and still make catastrophic collective decisions. The relevant property is institutional cognition: how effectively information is transformed into judgment and judgment into revisable action.

Trivialization is dangerous because it weakens this transformation without necessarily reducing information. In fact, it can coexist with information abundance. A society may know more facts than any previous civilization while becoming less able to hold a question steady long enough to act upon it.

The anti-trivialization project is therefore not an attempt to make everyone serious. It is an attempt to preserve institutional forms in which seriousness remains possible when needed.

The future machines described in this book are fragments of such an architecture. Cognitive fiduciaries would give individuals agents whose incentives are aligned with their declared goals. Attention constitutions would create boundaries around adaptive influence. Provenance infrastructure would make claims easier to trace. Evidence maps would represent uncertainty without collapsing it into binary labels. Decision records would preserve institutional memory. Adversarial AI would make criticism a function rather than an afterthought. Temporal institutions would give long-horizon consequences standing in present decisions. Polycentric governance would prevent the defense of seriousness from becoming centralized control.

Education would shift from artifact production toward judgment, while organizations would treat attention as shared infrastructure rather than private willpower.

None of these mechanisms is sufficient. Some may never exist in the forms imagined here. The deeper claim is functional. Advanced societies will need institutions that perform these roles because the informational environment is moving beyond the assumptions of twentieth-century governance.

The transition may initially look like loss. Platforms may face restrictions. Organizations may introduce review that slows deployment. Students may be asked to demonstrate reasoning rather than simply submit polished output. Political communication may acquire new disclosure requirements. AI systems may have to preserve provenance that makes them more expensive. Users may discover that truly aligned personal agents sometimes recommend doing less, consuming less, and closing the interface.

This friction will be criticized as inefficiency. Sometimes the criticism will be correct. Anti-trivialization institutions can become bloated, paternalistic, and self-important. They can confuse slowness with wisdom and credentialing with truth. They can be captured by incumbents who use safety and seriousness to suppress competitors.

For this reason, every machine needs a counter-machine. Auditors need auditing. Standards need revision. Experts need challengers. Long-term institutions need democratic checks. AI agents need user override. Provenance systems need privacy boundaries. Public archives need expiration policies where permanent memory becomes unjust. The anti-trivialization project must contain its own anti-bureaucratic principle.

This recursive design is not a weakness. It is how mature institutions work. Constitutional democracy distrusts power enough to divide it. Science

distrusts claims enough to demand reproducibility and criticism. Secure computing distrusts components enough to minimize privilege and log access. Resilient systems assume failure and make recovery possible.

The future will likely feel more synthetic than the present. More of what we read, watch, hear, and interact with will be generated or mediated by machines. This may initially increase the sensation that reality itself has become trivial: infinitely reproducible, infinitely editable, always available in another version.

But abundance can reverse value. When images are scarce, making an image matters. When images are infinite, provenance matters. When text is expensive, producing text matters. When text is infinite, judgment matters. When answers are scarce, access matters. When answers are infinite, questions matter.

This reversal is the positive heart of the book.

The attention economy discovered how valuable human attention is by learning to capture it. The AI economy will discover how valuable human judgment is by making synthetic cognition abundant. The very forces that appear to threaten seriousness will expose its scarcity more clearly than ever before.

Markets will respond because trusted judgment becomes valuable. Professions will respond because credibility becomes valuable. States will respond because systemic failure becomes costly. Communities will respond because people will seek environments in which relationships and meaning are not continuously optimized by external systems. Technology will respond because the tools needed for protection are built from the same capabilities as the tools of extraction.

It is useful to imagine what success would look like precisely because successful institutions become boring. By the 2040s, a high-school student

may find it normal that an educational agent cannot monetize attention, that sources remain inspectable, and that important assessments include adversarial discussion of AI-generated answers. The student may regard today's engagement-maximizing educational software as we regard schools that once accepted smoking in classrooms: not incomprehensible, but strangely careless about the environment in which development occurs.

A researcher may work inside an environment where claims automatically preserve provenance, where literature maps update continuously, where AI agents argue against new hypotheses, and where negative results are easier to store than to publish today. None of this guarantees discovery. It changes the institutional cost of forgetting, duplication, and premature confidence.

A citizen may receive public information through a personal agent that separates emergency alerts, shared civic notices, chosen interests, entertainment, and persuasion into different channels. Political advertisements may remain legal and aggressive, but their sponsors and targeting logic may be inspectable. Major public claims may have persistent histories showing revisions and outcomes. People will still lie. The infrastructure will simply make certain kinds of lying more expensive.

A company may treat attention the way it treats cybersecurity or financial controls. Boards may examine interruption load, AI assurance, skill reserves, and cognitive dependencies. Employees may have protected modes in which only defined categories can interrupt. High-impact automated decisions may preserve audit trails as naturally as financial transactions preserve ledgers.

A government may maintain long-term decision records and independent future-impact review. Emergency shortcuts may expire automatically unless renewed. Public AI systems may distinguish evidence,

inference, and policy preference. Different parties may produce competing analyses using shared source infrastructure.

None of these changes requires a universal ideology. Conservative, liberal, socialist, religious, libertarian, and technocratic communities can disagree radically while benefiting from institutions that preserve provenance, contestability, and cognitive autonomy. Anti-trivialization is not a doctrine about what people should believe. It is a doctrine about the conditions under which beliefs and decisions can remain answerable to reality and to other people.

This universality is important because the politics of attention can easily become partisan. Each faction is tempted to define triviality as whatever the other faction watches and seriousness as whatever it prefers. The book rejects that move. A political speech can be trivial regardless of ideology. A comedy can be profound without becoming serious in tone. A scientific institution can be captured. A religious ritual can protect deep attention. A market can reveal information quickly. A bureaucracy can preserve memory. No sector owns seriousness.

The institutional project is therefore meta-political. It concerns how systems maintain the capacity to disagree without every disagreement collapsing into attention warfare.

One can formulate the ambition in cybernetic terms. Norbert Wiener and later systems thinkers emphasized feedback: a system remains viable by detecting error and adjusting. A civilization that cannot distinguish signal from noise, preserve memory of past error, or route corrective information to decision centers loses feedback quality (Wiener 1950). Trivialization is partly a feedback pathology. Immediate signals become so loud that slower but more consequential signals fail to alter behavior.

Anti-trivialization machines improve feedback across timescales. They allow short signals to remain short while ensuring that long signals are not erased by their lack of urgency. They preserve exceptions, dissent, delayed consequences, and evidence histories long enough to matter.

There is no final institutional form because the environment will continue to change. Once attention agents become common, businesses will learn to persuade the agents. Once provenance becomes standard, actors will learn to game provenance. Once AI auditors become powerful, systems will optimize against the audits. Governance is an arms race, but not a futile one. Security is also an arms race, and yet secure systems are possible enough to support modern civilization.

The goal is not perfect protection. It is a favorable adaptive balance.

That balance will require institutions that can revise themselves. Every anti-trivialization mechanism should include a way to measure whether it still performs its function, a way for affected people to contest it, and a way to remove it when it becomes obsolete. Sunset clauses, experimental sandboxes, independent evaluation, open standards, and polycentric alternatives are not administrative details. They are the self-correction layer.

The most optimistic possibility is that AI itself increases society's capacity for such self-correction. Institutions currently fail to learn partly because evidence is dispersed, archives are inaccessible, and reviewing policy history is expensive. AI can make institutional memory searchable, compare outcomes across jurisdictions, simulate alternatives, and surface forgotten warnings. The same technology that can flood the world with synthetic noise can make the world's accumulated experience more usable.

If that happens, the AI transition may eventually be remembered not only as the moment when machines learned to generate language, but as the moment when institutions learned to inspect their own cognition.

The result will not be a return to a pre-digital world. Nor should it be. The old world was not more truthful merely because information moved slowly. Gatekeepers suppressed voices. Institutions hid failures. Expertise was often inaccessible. Geographic and economic barriers excluded most people from many forms of knowledge. Digital networks and AI can correct some of those injustices.

The goal is not restoration. It is synthesis.

We should want the abundance of the digital world with the memory of the archive. The accessibility of AI with the skepticism of science. The speed of networks with the due process of law. The creativity of markets with the boundary-setting of constitutionalism. The local knowledge of communities with the coordination capacity of large institutions. The pleasures of triviality with protected spaces for depth.

A mature civilization does not defeat one side of such pairs. It learns to govern the tension between them.

This is why the yin-yang image is useful at the end, provided it is kept modest. Every corrective force can become excessive; every capacity generates the need for a limit; every limit eventually needs an exception. Human institutions live by oscillation, not final victory.

The future anti-trivialization machine is therefore not a monument to seriousness. It is a regulator of balance. It knows that speed can save lives and destroy judgment. It knows that entertainment can restore us and consume us. It knows that AI can weaken agency and amplify it. It knows that institutions can protect truth and ossify into dogma.

Its defining virtue is not certainty. It is corrigibility.

If we build these institutions well, future generations may not describe them as anti-trivialization machines at all. They may call them

ordinary features of civilized life, just as we rarely celebrate the building code when the building remains standing.

That would be success.

The deepest hope of this book is therefore strangely conservative and radically futuristic at once. Human beings do not need to become a new species before we can survive our technologies. We need to continue doing what civilization has always done at its best: discover where new power creates new fragility, invent institutions that convert fragility into learning, and then make those institutions ordinary enough that the next generation can forget how strange they once sounded.

We built machines that learned to capture attention.

We can build institutions that teach civilization how to give attention back.

Bibliography

- Arendt, Hannah. 1958. *The Human Condition*. Chicago: University of Chicago Press.
- Bernays, Edward L. 1928. *Propaganda*. New York: Horace Liveright.
- Bush, Vannevar. 1945. "As We May Think." *The Atlantic Monthly* 176(1): 101-108.
- Citton, Yves. 2017. *The Ecology of Attention*. Cambridge: Polity.
- Debord, Guy. 1967. *The Society of the Spectacle*. Paris: Buchet-Chastel.
- Dewey, John. 1927. *The Public and Its Problems*. New York: Henry Holt.
- Ellul, Jacques. 1965. *Propaganda: The Formation of Men's Attitudes*. New York: Knopf.
- Engelbart, Douglas C. 1962. *Augmenting Human Intellect: A Conceptual Framework*. Menlo Park, CA: Stanford Research Institute.
- Fassi, Leonardo, et al. 2025. "Social Media Use in Adolescents with and without Mental Health Conditions." *Nature Human Behaviour*.
- Fishkin, James S. 2018. *Democracy When the People Are Thinking: Revitalizing Our Politics Through Public Deliberation*. Oxford: Oxford University Press.
- Foucault, Michel. 1975. *Surveiller et punir: Naissance de la prison [Discipline and Punish]*. Paris: Gallimard.
- Frankl, Viktor E. 1946/2006. *Man's Search for Meaning*. Boston: Beacon Press.
- Fraser, Nancy. 1990. "Rethinking the Public Sphere: A Contribution to the Critique of Actually Existing Democracy." *Social Text* 25/26: 56-80.
- Fromm, Erich. 1955. *The Sane Society*. New York: Rinehart.
- Goldman, Alvin I. 1999. *Knowledge in a Social World*. Oxford: Oxford University Press.
- Guess, Andrew M., et al. 2023. "How Do Social Media Feed Algorithms Affect Attitudes and Behavior in an Election Campaign?" *Science* 381.
- Habermas, Jürgen. 1989. *The Structural Transformation of the Public Sphere*. Cambridge, MA: MIT Press. Originally published in German in 1962.
- Hayles, N. Katherine. 2007. "Hyper and Deep Attention: The Generational Divide in Cognitive Modes." *Profession* 2007: 187-199.
- Hirschman, Albert O. 1970. *Exit, Voice, and Loyalty: Responses to Decline in Firms, Organizations, and States*. Cambridge, MA: Harvard University Press.
- Horkheimer, Max, and Theodor W. Adorno. 1947. *Dialectic of Enlightenment*. Amsterdam: Querido.
- Huxley, Aldous. 1932. *Brave New World*. London: Chatto & Windus.
- Illich, Ivan. 1973. *Tools for Conviviality*. New York: Harper & Row.

- Jonas, Hans. 1979/1984. *The Imperative of Responsibility: In Search of an Ethics for the Technological Age*. Chicago: University of Chicago Press.
- Kahneman, Daniel. 2011. *Thinking, Fast and Slow*. New York: Farrar, Straus and Giroux.
- Kindleberger, Charles P., and Robert Z. Aliber. 2005. *Manias, Panics, and Crashes: A History of Financial Crises*. 5th ed. Hoboken, NJ: Wiley.
- Kuhn, Thomas S. 1962. *The Structure of Scientific Revolutions*. Chicago: University of Chicago Press.
- Landemore, Hélène. 2020. *Open Democracy: Reinventing Popular Rule for the Twenty-First Century*. Princeton: Princeton University Press.
- Lee, Hao-Ping, Advait Sarkar, Lev Tankelevitch, Ian Drosos, Sean Rintel, Richard Banks, and Nicholas Wilson. 2025. "The Impact of Generative AI on Critical Thinking: Self-Reported Reductions in Cognitive Effort and Confidence Effects from a Survey of Knowledge Workers." *Proceedings of the 2025 CHI Conference on Human Factors in Computing Systems*. DOI: 10.1145/3706598.3713778.
- Licklider, J. C. R. 1960. "Man-Computer Symbiosis." *IRE Transactions on Human Factors in Electronics* HFE-1: 4-11.
- Lippmann, Walter. 1922. *Public Opinion*. New York: Harcourt, Brace.
- Longino, Helen E. 1990. *Science as Social Knowledge: Values and Objectivity in Scientific Inquiry*. Princeton: Princeton University Press.
- Macdonald, Dwight. 1953/1957. "A Theory of Mass Culture." Reprinted in *Mass Culture: The Popular Arts in America*, edited by Bernard Rosenberg and David Manning White. Glencoe, IL: Free Press.
- March, James G., and Herbert A. Simon. 1958. *Organizations*. New York: Wiley.
- Maslow, Abraham H. 1943. "A Theory of Human Motivation." *Psychological Review* 50(4): 370-396.
- Maslow, Abraham H. 1954. *Motivation and Personality*. New York: Harper.
- May, Rollo. 1953. *Man's Search for Himself*. New York: W. W. Norton.
- McLuhan, Marshall. 1964. *Understanding Media: The Extensions of Man*. New York: McGraw-Hill.
- Merton, Robert K. 1942/1973. "The Normative Structure of Science." In *The Sociology of Science*. Chicago: University of Chicago Press.
- North, Douglass C. 1990. *Institutions, Institutional Change and Economic Performance*. Cambridge: Cambridge University Press.
- Orben, Amy, Andrew K. Przybylski, Sarah-Jayne Blakemore, and Rogier A. Kievit. 2022. "Windows of Developmental Sensitivity to Social Media." *Nature Communications* 13: 1649.

- Orwell, George. 1949. *Nineteen Eighty-Four*. London: Secker & Warburg.
- Ostrom, Elinor. 1990. *Governing the Commons: The Evolution of Institutions for Collective Action*. Cambridge: Cambridge University Press.
- Piccardi, Tiziano, et al. 2025. "Reranking Partisan Animosity in Algorithmic Social Media Feeds." *Science*. DOI: 10.1126/science.adu5584.
- Polanyi, Karl. 1944. *The Great Transformation: The Political and Economic Origins of Our Time*. New York: Farrar & Rinehart.
- Popper, Karl R. 1959. *The Logic of Scientific Discovery*. London: Hutchinson.
- Postman, Neil. 1985. *Amusing Ourselves to Death: Public Discourse in the Age of Show Business*. New York: Viking.
- Rogers, Carl R. 1961. *On Becoming a Person*. Boston: Houghton Mifflin.
- Rosa, Hartmut. 2013. *Social Acceleration: A New Theory of Modernity*. New York: Columbia University Press.
- Schumaker, John F. 2008. "The Triumph of Triviality." *New Internationalist*, April 1, 2008. Republished by Films for Action, April 25, 2020.
- Scott, James C. 1998. *Seeing Like a State: How Certain Schemes to Improve the Human Condition Have Failed*. New Haven: Yale University Press.
- Simon, Herbert A. 1971. "Designing Organizations for an Information-Rich World." In *Computers, Communications, and the Public Interest*, edited by Martin Greenberger. Baltimore: Johns Hopkins University Press.
- Simon, Herbert A. 1947. *Administrative Behavior*. New York: Macmillan.
- Sunstein, Cass R. 2017. *#Republic: Divided Democracy in the Age of Social Media*. Princeton: Princeton University Press.
- Thaler, Richard H., and Cass R. Sunstein. 2008. *Nudge: Improving Decisions About Health, Wealth, and Happiness*. New Haven: Yale University Press.
- Tufekci, Zeynep. 2017. *Twitter and Tear Gas: The Power and Fragility of Networked Protest*. New Haven: Yale University Press.
- Welsh Government. 2015. *Well-being of Future Generations (Wales) Act 2015*.
- Wiener, Norbert. 1950. *The Human Use of Human Beings: Cybernetics and Society*. Boston: Houghton Mifflin.
- Winner, Langdon. 1980. "Do Artifacts Have Politics?" *Daedalus* 109(1): 121-136.
- Wu, Tim. 2016. *The Attention Merchants: The Epic Scramble to Get Inside Our Heads*. New York: Knopf.
- Zuboff, Shoshana. 2019. *The Age of Surveillance Capitalism*. New York: PublicAffairs.