

AI Alignment Research Report

Author: Sînică Alboaie, PhD, Axiologic Research

Purpose: A report summarising Axiologic Research vision on the AI Alignment

Visibility: Public

Date: February 2024 (Version 2 Update)

AI Alignment Vision.....	2
Agent Swarms, Swarm Communication, Executable Choreographies.....	3
Multi-Agent “Cognitive” Architecture for Guaranteed Alignment.....	4
Agents Operating System Layer (AOSL).....	5
Semantic Firewall Layer.....	6
Agents Execution Environment (xEnv).....	7
Training Environment (tEnv).....	8
Why?.....	9
Nature and science-inspired ideas requiring research.....	10
Conclusions.....	12
References.....	13

AI Alignment Vision

The vision of Axiologic Research on the evolution of AI, as outlined in previous reports [1], [2], [3], posits that the most probable future of AI in the next decade, and crucially for the safety of humanity, should involve swarms of intelligent agents. These agents would be isolated through technical security mechanisms, including operating system security, cryptography, and techniques for combining intelligence so that a flaw or malice in a group of agents does not significantly impact the alignment of the swarms. Our thinking prioritises decentralisation, role segregation, specialisation, and strategic limitations for agents, reflecting strategies against malicious control in human societies. We hypothesise that Artificial General Intelligence (AGI) will arise from the synergy of expert Large Language Models (LLMs), various AI technologies, and symbolic reasoning rather than from a single LLM. This model, a "swarm of intelligent agents," showcases the potential for specialised, intelligent components to conduct complex interactions in a manner understandable to humans (which we call choreographies), setting the foundation for advanced intelligence.

The core belief behind our proposal is that for AGI or superintelligence to be truly aligned, it is essential to prevent any single component from dominating the "intelligent swarm". It is crucial to maintain diverse directions in dynamic equilibrium and to create artificial barriers within the architecture. For instance, one approach involves limiting some agents' access to information while providing others with simplified summaries. In human societies, limitations naturally emerge due to computational complexity and the biological limits of the human brain. For example, a CEO cannot single-handedly steer a company towards socially harmful goals; they would need to influence many agents, thereby risking exposure of any conspiracy or harmful action to public scrutiny, whistleblowers, or internal pushback, potentially leading to legal action by the state. Future AI systems will likely face natural limits related to energy consumption and computational capacity, requiring architectures that reflect social interactions. Nonetheless, some limitations may need to be artificially established and continuously defended with intention.

Future systems should consist of multiple agents with varied interests, memories, experiences, and sources of code and training data to guard against alignment threats that could arise accidentally or through deliberate manipulation. The most innovative concept we propose involves creating multi-agent ecosystems that naturally limit the long-term survival of malicious agents by simulating their elimination in automated or semi-automated ways. This approach draws from human societies, where lasting schools of thought and institutions emerge, promoting continuity and ensuring the long-term evolution of systems, along with a permanent mechanism to remove malicious agents. Our research aims to develop a coherent theory and undertake extensive experiments and practical implementations on multi-agent systems that align with this vision. Furthermore, we plan to engage and reward meaningful contributions from thinkers and researchers across various disciplines. As we will detail in the following pages, our vision is to create an operating system with agents with long-term memory (experience) instead of processes. These agents live only as long as they are needed to perform tasks. The more useful agents are rewarded by being brought back to life more frequently and by being able to create "copies," meaning they are cloned. Cloning can be useful for verifying adherence to alignment and perpetuating useful behaviours and experiences. The death of agents will also inevitably be necessary. The approach proposed in this report reminds us of the concept of reincarnation found in various religions. It incorporates a cycle of life, death, and rebirth. It is adapted to the digital realm where agents evolve, replicates, and are selectively perpetuated based on their performance and alignment with desired goals.

Agent Swarms, Swarm Communication, Executable Choreographies

It's important to know the difference between *Swarm Communication* and *Agent Swarms* from the start. 'Swarm Communication' is about how to set up talking between agents in a group to do certain tasks. 'Agent Swarms' means the group of agents working together to get these tasks done." "Swarm Communication" represents a method of implementing "executable choreographies"—essentially, the communication protocol within a "swarm" of independent entities—and the idea of "agent swarms," which are these independent entities that collaborate to fulfil the business objectives of the "executable choreographies."

The concept of "swarm communication" [2] as a method for implementing "executable choreographies" [3] is presented as an innovative approach to distributed computing. This methodology is envisioned to model computational flows that navigate through a network of independent nodes, prioritising the movement of results or conclusions of computations within the flow while deliberately avoiding transferring confidential data from one node to another. An "executable choreography" is a dynamic script executed across various network nodes. Local APIs are invoked in each node, and the outcome of these computations updates the state of the "flow" or the "swarm".

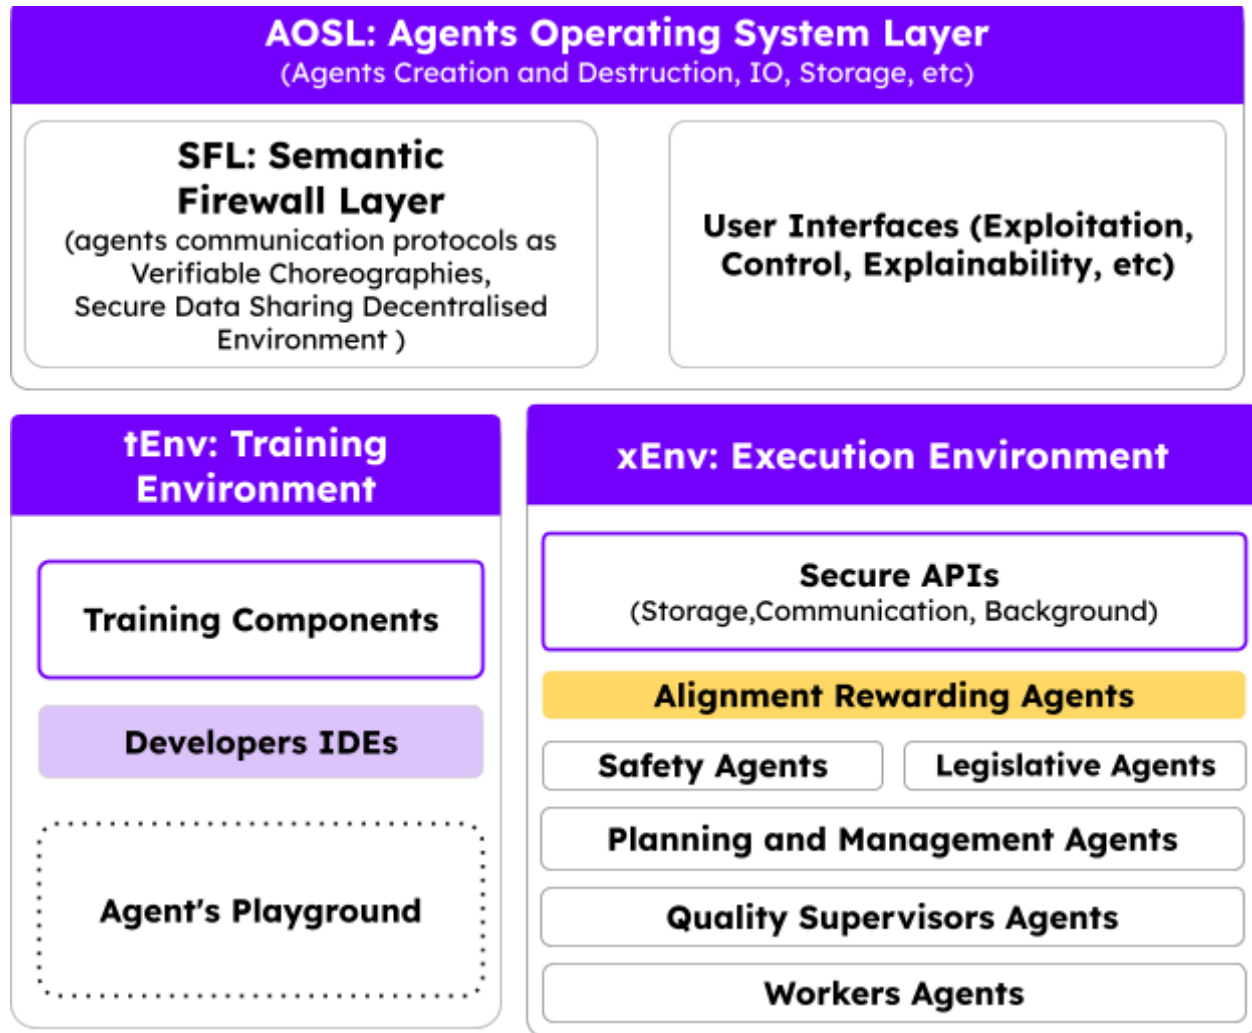
Swarm communication parallels the concept of "active messages" [1] but extends it to a broader context of distributed computing. This requires each node to have an address and the capability to accept visits from the "swarm" in messages. The identity of the swarm is recognised at the level of a group of messages, with each message representing just a part of the whole, hence the term "Swarm communication".

Although this approach might not be immediately intuitive, it offers numerous security, privacy, and scalability benefits [4]. Executable choreographies aim to supplant the orchestration of APIs, yet, until now, use cases have not been compelling enough to drive widespread adoption. It is crucial to acknowledge that, from our perspective, "executable choreographies," with their ability to model complex communication patterns—especially the "swarm communication" aspect pioneered by our research—represent a significant and fundamental concept in computing and software architectures.

With the advent of Large Language Models (LLMs), it is conceivable that executable choreographies could become an essential component of the future AI systems landscape. This is because they facilitate proper isolation of each agent's execution environment and enable the verification of communication between agents from a higher vantage point, all while maintaining the decentralised nature of the agent swarm. This breakthrough could herald a new era in AI systems, enhancing their functionality and application in ways previously unimagined.

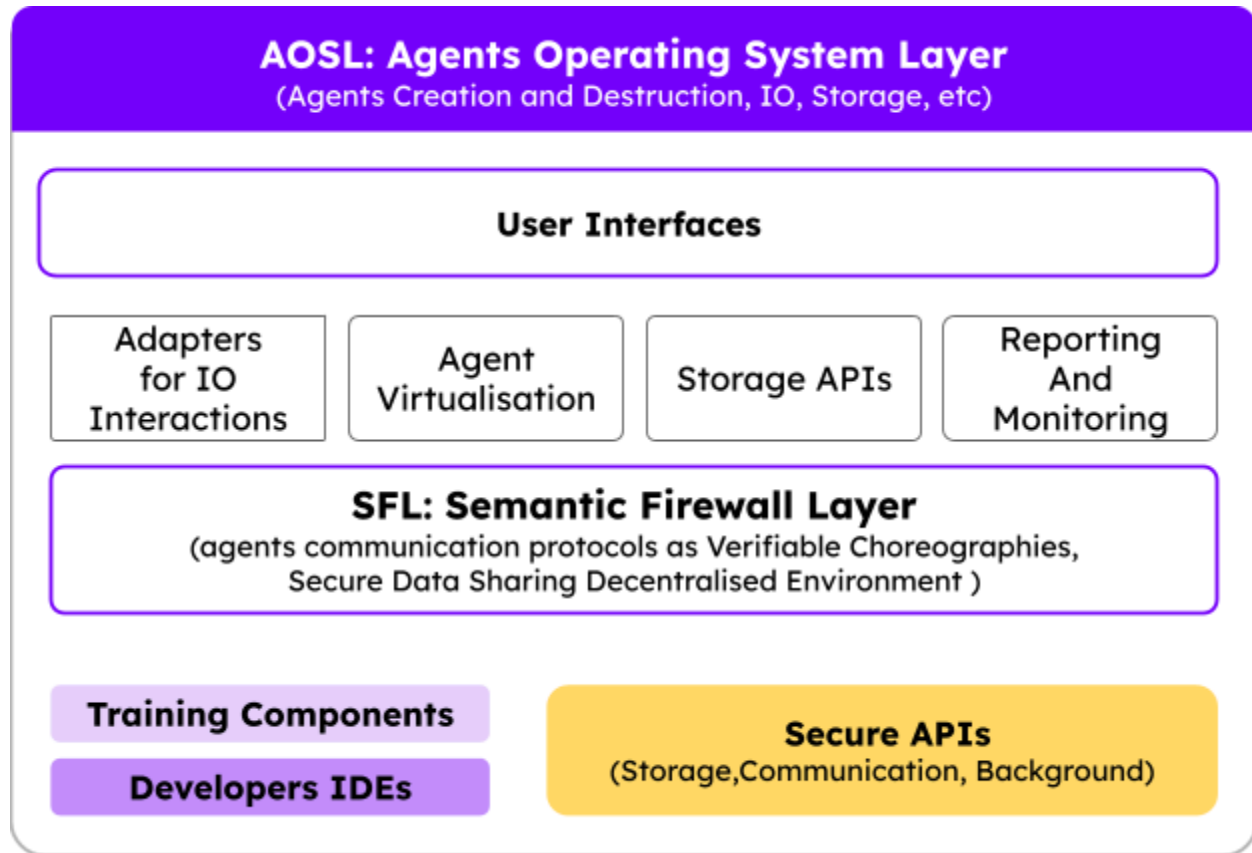
An important concept is that of 'verifiable choreographies,' which leads us to the ability to verify the properties of executable choreographies in a way that wouldn't be possible if they communicated using web APIs alone. Although choreographies inherently promote decentralisation and peer-to-peer (P2P) communication among agents, various verification techniques can be implemented transparently, primarily in the communication layer. This layer typically involves message queues for each agent. Still, it can also be more formally established when the choreography is approved by the 'swarm' or the system owners where the swarm of agents operates. Other verification techniques could involve cryptographic methods and more advanced technical approaches, such as Zero Knowledge Proof or self-validating data, where each agent digitally signs its contributions or ensures specific characteristics related to computational integrity.

Multi-Agent “Cognitive” Architecture for Guaranteed Alignment



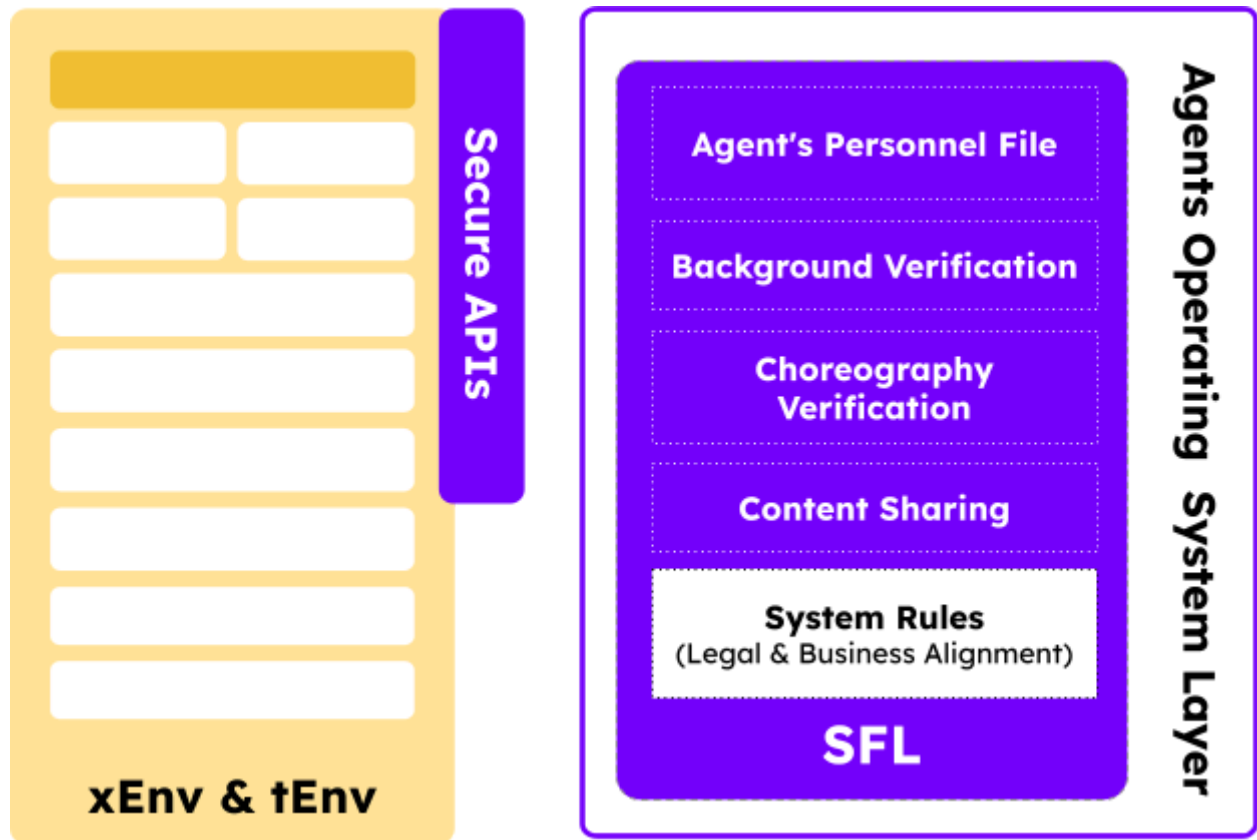
In our perspective, any multi-agent architecture designed to ensure the alignment of the entire swarm of agents with a set of legal and ethical values or rules will be structured around the following components: AOLS (Agents Operating System Layers), which provide hardware and software support for agents' instantiation and life cycles. This layer is not particularly intelligent; it is merely standard software code. Additionally, we have an Execution Environment (xEnv) that offers secure instantiation of agent instances and a Training Environment (tEnv) that provides training tools and a playground for agents to test trained models or to create simulations of agents before deployment in production. Between all these layers, there is an SFL (Semantic Firewall Layer), a system that checks everything communicated between agents. The testing and execution environments cannot bypass the SFL. Various implementation modes of the SFL will be analysed, but in principle, the SFL is a medium capable of analysing the communication protocols between agents. The proposal is for these protocols to be implemented in executable choreographies, more precisely, a form of verifiable choreography suitable for multi-agent environments.

Agents Operating System Layer (AOSL)



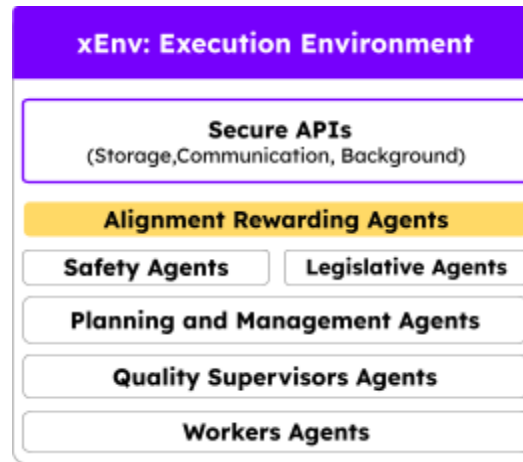
The AOSL (Agents Operating System Layer) is responsible for creating and destroying agents, I/O operations, and storage, but also for providing APIs for creating predefined tools and APIs that facilitate the creation of user interfaces for the system's users. This includes administrative aspects where we have APIs for reporting and monitoring and for potential human users or clients served by the proposed architecture. For the creation and destruction of agents, we consider the necessity of APIs that can be grouped under Agent Virtualization APIs. These APIs ensure the operating system-level security of the agents' execution environment, preventing their communication via direct channels. Discovering hidden communication channels between malicious agents is the responsibility of the Semantic Firewall Layer, or as we saw in the first chapter, it could also be aided by the way executable choreographies are organised, which include agents specialised in detecting hidden channels or Advanced Persistent Threats (APTs). From this perspective, the Semantic Firewall Layer could be just an operating system subcomponent. Still, we will analyse it separately in a later chapter due to its significant importance. The code, APIs, tools, and user interfaces running virtually in secure (sandboxed) environments in tEnv and xEnv can be viewed as the user-space component in an operating system.

Semantic Firewall Layer



The SFL employs as much code as possible that is intelligible to human users or AI models with a high degree of explainability, avoiding the use of complex AI models that could conceal Advanced Persistent Threats (APTs), biases, and other attacks that could undermine alignment. Our research will focus on ways to accurately represent general rules related to alignment and rules regarding compliance with laws and internal procedures proposed by the system owners. It is necessary for the rule description language to be as high-level as possible to allow for human auditing yet sufficiently formal and comprehensive to encompass all types of nuances caught by verification components related to content sharing, choreography checking, and other background verifications executed by the SFL based on data collected about agents' behaviour and incidents. In this regard, the current vision of the SFL is to maintain information in an "Agent's Personnel File" and, upon detection of serious incidents or the accumulation of minor incidents that raise suspicions, to proceed with the destruction of the involved agent or agents, initiate agents with clean states, or even involve human oversight in the analysis of incidents.

Agents Execution Environment (xEnv)



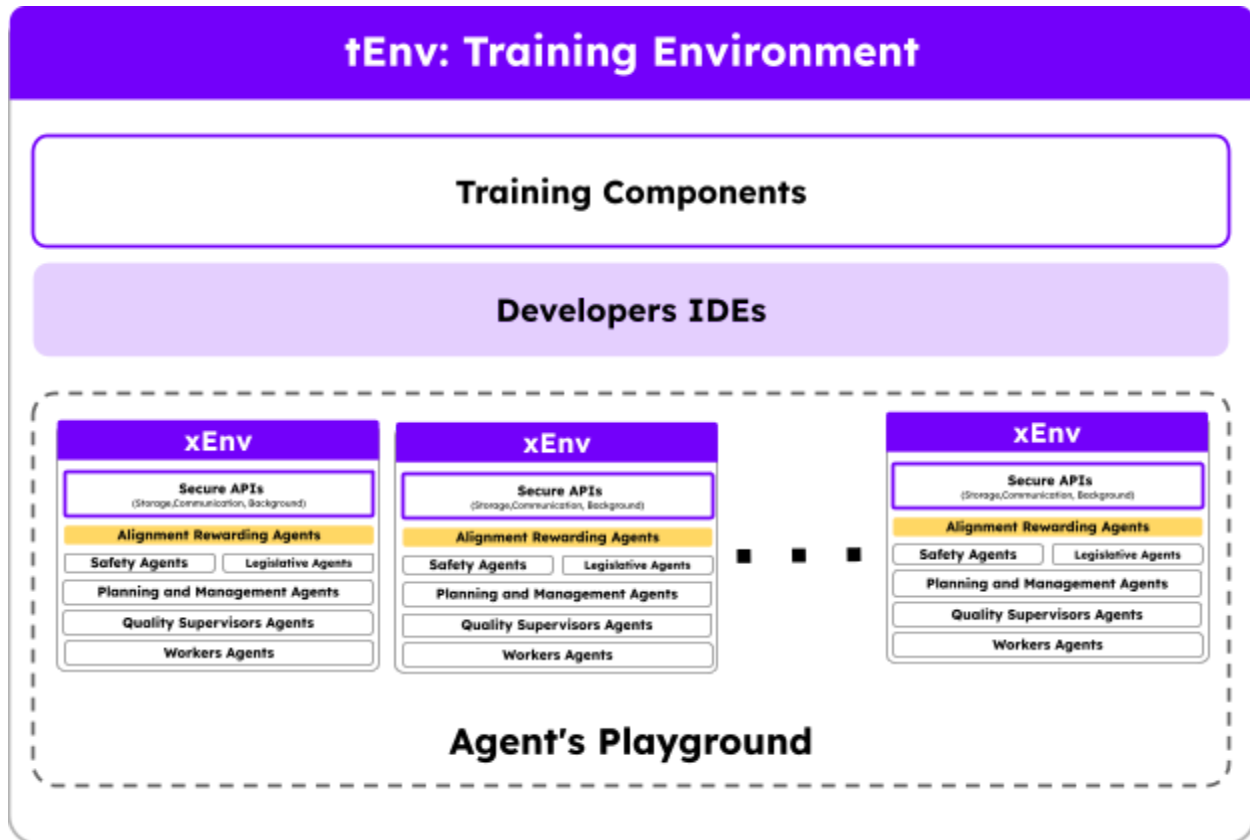
The xEnv (Execution Environment) functions as a sophisticated virtualisation space for multi-agent systems, resembling a Hypervisor in operating systems, offers APIs for secure and authenticated agent communication, enables secure data sharing among intelligent agents and supports background execution capabilities, alongside enforcing execution limitations based on the budget of each agent. Within xEnv, various agents operate to maintain and optimise the system:

Alignment Rewarding Agents work in concert with the Semantic Firewall Layer (SFL) to identify and address any deviations from the system's alignment goals. Safety Agents serve as vigilant monitors over the activities of their peers, ensuring the environment remains secure, and operations proceed without compromise. Legislative Agents are tasked with creating and enforcing rules, which is crucial in optimising resource distribution and computational tasks across the environment. Planning and Management Agents take on roles akin to top or middle management in traditional organisational structures, tasked with orchestrating task execution. These agents can accumulate reputation or resources, which may justify their further instantiation. Alignment Rewarding Agents could influence an xEnv regarding the resources allocated or the frequency of instantiation of various agents. This mechanism would allow for a dynamic adjustment of resources and instantiation frequency based on the alignment of agents with the desired objectives and behaviours. By rewarding agents that demonstrate alignment with predefined goals, the system can encourage the propagation of desirable behaviours and optimise the overall performance and efficiency of the multi-agent system. This approach ensures that agents contributing positively to the system's objectives are given precedence in resource allocation and opportunities for instantiation, thereby fostering an environment that promotes and maintains alignment with the system's core values and objectives.

The environment also hosts Quality Supervisors Agents and Worker Agents, specialised across dozens, hundreds, or thousands of domains. The execution environment dynamically requests the instantiation of these agents based on the specific needs of concrete tasks. These agents can access current task data and may possess their memory (or experiential knowledge) to optimise their performance.

This setup tries to ensure a robust, efficient, and secure operation of multi-agent systems, fostering an environment where agents can effectively collaborate, evolve, and contribute to the system's overarching goals while being on sleep, active or even destroyed accordingly with the need of the tasks and the safety and alignment concerns. Surprisingly, this approach reminds us of reincarnation in various religions. We mention this to reinforce the main intuition behind the proposed solution.

Training Environment (tEnv)



The tEnv (Training Environment) is envisioned as a comprehensive platform providing tools, libraries, and a user interface for developers to train AI models and experiment with swarms of agents before deploying them in actual production Execution Environments. A tEnv aims to feature an Agent's playground designed to virtualise multiple Execution Environments cost-effectively, implementing reduced budgets or other constraints to allow for deeper analyses and enhanced debugging capabilities. The Developer IDE are imagined as equipped with tools to examine data for biases, analyse the behaviour of the systems, and offer other tools for certifying the developed solutions, at least at the moment of deployment. This setup aims to ensure that developers have a robust environment for refining and validating their AI models, ensuring they meet the necessary standards and requirements before they go live, thereby enhancing the reliability and effectiveness of the deployed systems.

Why?

The proposed architecture introduces a layered design to enhance AI safety, ethical alignment, and operational efficiency. Central to this architecture is the capability for self-monitoring and self-regulation within AI agents, enabled by the Semantic Firewall Layer. This layer ensures AI operations adhere to ethical and legal standards, fostering societal trust through auditability. Enhanced transparency and accountability are crucial for meeting regulatory demands and building public confidence. Advanced security features within the architecture protect against external threats and internal errors, preserving the integrity of sensitive data. The decentralisation aspect addresses vulnerabilities associated with centralised systems by introducing resilience and self-healing capabilities, further augmented by a unique "karma" system for agents that promotes positive behaviour over time. Interoperability and specialisation across AI agents reduce complexity and enhance system manageability. The architecture's modular and scalable design ensures it can adapt to various application scales, from small to large ecosystems. Lastly, its forward-looking design principles ensure adaptability to future technological developments, securing its long-term applicability in the dynamic field of AI. This comprehensive approach mitigates current AI alignment challenges and paves the way for a more secure, ethical, and effective deployment of AI technologies.

In a landscape where superintelligence poses both immense potential and risks, the Semantic Firewall is a crucial safeguard. By design, it segments intelligence across multiple components, preventing any single part from becoming overwhelmingly powerful while fostering collective growth. This segmentation ensures that while the system's overall capabilities can expand, they do so under strict ethical and operational guidelines. The firewall's role is to preemptively neutralise threats by enforcing these guidelines, making it exceedingly difficult for a superintelligence to bypass safeguards without detection. This architectural choice mitigates the risk of rogue AI behaviour and aligns with the broader goal of developing AI that enhances societal well-being without compromising security or ethical standards.

The architecture adopts a zero-trust approach towards agents, actively resetting those veering towards potentially harmful paths. Simultaneously, it aims to repurpose and amalgamate agents' experiences in a manner reminiscent of human societies and natural ecosystems. This strategy ensures vigilance against deviations and fosters a collaborative, growth-oriented environment. The architecture benefits from diversity and resilience by emulating societal and ecological dynamics, promoting a harmonious evolution of AI capabilities within a framework of mutual trust and ethical alignment.

Nature and science-inspired ideas requiring research

In the following table, we have summarised over 20 ideas that are inspired by nature and science, which require research to implement the various components described above, primarily focusing on the "Alignment Rewarding Agents" area. However, this research could also influence the Semantic Firewall Layer (SFL).

Idea	Description	Inspired by	Priority
Energy Consumption Limits	Agents have energy budgets that limit their processing capabilities, similar to metabolic rates in biological organisms.	Biology	High
Copy Brain Structures	One approach to AI alignment could involve replicating aspects of the brain's architecture, such as the distinct roles of the cortex and neocortex.	Biology	High
Ageing and Obsolescence	Introduce ageing concepts, where agents become less efficient or obsolete over time.	Biology	High
Evolutionary Pressures	Agents manage resources to avoid depletion and ensure sustainability, reflecting environmental science concerns.	Biology	Medium
Resource Allocation Constraints	Agents compete for limited resources, incentivising efficient use. Task allocation among agents uses market-based mechanisms with tasks having varying rewards.	Economics	Medium
Specialization and Trade	Agents specialise and trade outputs, reflecting economic principles of comparative advantage.	Economics	Medium
Legal Regulations	Introduce legal frameworks that agents must adhere to, mimicking societal laws. Punish misbehaving agents.	Politics	High
Political Power Dynamics	Agents form groups with varying degrees of power, akin to political parties or nations.	Politics	Medium
Social Hierarchies	Agents are organised hierarchically with tiered access to information.	Sociology	Medium
Trust and Reputation Systems	Develop trust and reputation metrics for agents.	Sociology	High
Memory Constraints	Limit agents' memory capacity, requiring	Psychology	Medium

	prioritisation of information retention.		
Cognitive Load	Implement cognitive load limitations on agents.	Psychology	Medium
Collective Intelligence	The system harnesses collective intelligence for problem-solving.	Psychology	High
Adaptation and Learning Limits	Agents adapt or learn within bounded rationality.	Cognitive Sciences	Low
Innovation and Creativity Limits	Agents have constraints on their ability to innovate or create.	Cognitive Sciences	Low
Communication Barriers	Agents experience understanding limitations due to varying communication protocols.	Linguistics	Low
Data Access Barriers	Simply establish artificial limits on access to data from sensors or other agents, introducing inefficiencies in the actions of malicious agents.	Cybersecurity	High
Zero Trust Architecture	Adopt a zero-trust model where agents must verify their identity and permissions for every data access or system interaction. This approach assumes breach and verifies each request as if it originated from an open network, significantly reducing the attack surface.	Cybersecurity	High
Anomaly Detection	Utilize machine learning algorithms to monitor agent behaviours and detect anomalies that could indicate a cybersecurity threat, such as unusual data access patterns or attempts to bypass security controls.	Cybersecurity	High
Ethical Constraints	Embed ethical considerations into agents' decision-making.	Philosophy	High
Feedback Loops	Introduce feedback loops that influence agent behaviours and system stability.	Systems Theory	Medium
Ecosystem Services	Some agents perform essential roles for the system's health.	Ecology	Low
Emotional Intelligence	Equip agents with the ability to respond to others' emotional states.	Psychology	Low

Conclusions

Even in the absence of superintelligent AIs, we already live in an era where forms of superintelligence exist under the guise of states and corporations. These social entities have capabilities that exceed individual human power but are generally under society's control. Although they may seem threatening due to their vast power and influence, people manage, through various mechanisms, to channel them for the common good. This complex coexistence offers us a valuable perspective on how humanity can navigate the challenges of managing and directing advanced intelligence.

Despite problems that sometimes worsen, some moments inspire us with hope. We are learning to self-organize and create new forms of collaboration and innovation, such as decentralised brands and open-source projects. These initiatives reflect our capacity to transcend the traditional limits of organisation and to experiment with structures that are more flexible, transparent, and open to mass participation. Through these efforts, we contribute to democratising access to information, resources, and decision-making power, thus shaping a future in which technology and collective intelligence serve the needs and aspirations of humanity.

Maintaining an optimistic, albeit cautious, attitude is essential in rapid change. We recognise the potential threat of these superintelligent entities, whether AI or complex social structures. Still, we are also witnesses and participants in humanity's capacity to innovate and find solutions to emerging challenges. It is important to remain vigilant and continuously assess the impact and direction of technological and social developments and to shape them to reflect our collective values and goals actively.

Thus, even as we face uncertainties and challenges, the prospect of working together to guide and shape these superintelligent forces gives us a reason for optimism. We can turn these challenges into opportunities for collective growth and prosperity through collaboration, innovation, and a shared commitment to a better future.

The table above suggests that, on a smaller or larger scale, we can envision the superintelligences of the future as swarms of intelligent agents resembling societies of intentionally limited intelligence, which have hardcoded intentions of prosperity and survival.

This metaphor seems to lead us to interesting metaphysical ideas related to simulation theory or theories from Kabbalah. Still, we do not want to dwell on this direction too much, just to conclude that, interestingly, research on artificial intelligence seems to mysteriously converge with many sciences like biology, psychology, sociology or philosophy through problems of epistemology, knowledge representations, ethics, and axiology.

References

- [1] https://en.wikipedia.org/wiki/Active_message
- [2] “Swarm Communication – a Messaging Pattern Proposal for Dynamic Scalability in Cloud”
L. Alboaie, S. Alboaie, P. Andrei, At 15th IEEE International Conference on High-Performance Computing and Communications (HPCC 2013). Zhangjiajie, China, November 2013.
- [3] Extending swarm communication to unify choreography and long-lived processes
L. Alboaie, S. Alboaie, T. Barbu, 3rd International Conference on Information Systems Development (ISD2014 Croatia) pp. 375–382, 2014.
- [4] “Levels of Privacy for e-Health systems in the Cloud era” S. Alboaie, L. Alboaie, A. Panu, 24th International Conference on Information Systems Development. Harbin, China, August 25-27, 2015.